

Welcome to the course!

LINEAR CLASSIFIERS IN PYTHON



Michael (Mike) Gelbart

Instructor, The University of British
Columbia

Assumed knowledge

In this course we'll assume you have some prior exposure to:

- Python, at the level of *Intermediate Python for Data Science*
- scikit-learn, at the level of *Supervised Learning with scikit-learn*
- supervised learning, at the level of *Supervised Learning with scikit-learn*

Fitting and predicting

```
import sklearn.datasets

newsgroups = sklearn.datasets.fetch_20newsgroups_vectorized()

X, y = newsgroups.data, newsgroups.target
```

```
X.shape
```

```
(11314, 130107)
```

```
y.shape
```

```
(11314,)
```

Fitting and predicting (cont.)

```
from sklearn.neighbors import KNeighborsClassifier
```

```
knn = KNeighborsClassifier(n_neighbors=1)
```

```
knn.fit(X,y)
```

```
y_pred = knn.predict(X)
```

Model evaluation

```
knn.score(X, y)
```

```
0.99991
```

```
from sklearn.model_selection import train_test_split  
  
X_train, X_test, y_train, y_test = train_test_split(X, y)
```

```
knn.fit(X_train, y_train)
```

```
knn.score(X_test, y_test)
```

```
0.66242
```

Let's practice!

LINEAR CLASSIFIERS IN PYTHON

Applying logistic regression and SVM

LINEAR CLASSIFIERS IN PYTHON



Michael (Mike) Gelbart

Instructor, The University of British
Columbia

Using LogisticRegression

```
from sklearn.linear_model import LogisticRegression
```

```
lr = LogisticRegression()  
lr.fit(X_train, y_train)  
lr.predict(X_test)  
lr.score(X_test, y_test)
```


LogisticRegression example

```
import sklearn.datasets
wine = sklearn.datasets.load_wine()

from sklearn.linear_model import LogisticRegression
lr = LogisticRegression()
lr.fit(wine.data, wine.target)

lr.score(wine.data, wine.target)
```

```
0.972
```

```
lr.predict_proba(wine.data[:1])
```

```
array([[ 9.951e-01,  4.357e-03,  5.339e-04]])
```

Using LinearSVC

LinearSVC works the same way:

```
import sklearn.datasets

wine = sklearn.datasets.load_wine()
from sklearn.svm import LinearSVC

svm = LinearSVC()

svm.fit(wine.data, wine.target)
svm.score(wine.data, wine.target)
```

0.893

Using SVC

```
import sklearn.datasets
wine = sklearn.datasets.load_wine()

from sklearn.svm import SVC
svm = SVC() # default hyperparameters
svm.fit(wine.data, wine.target);

svm.score(wine.data, wine.target)
```

1.

Model complexity review:

- **Underfitting:** model is too simple, low training accuracy
- **Overfitting:** model is too complex, low test accuracy

Let's practice!

LINEAR CLASSIFIERS IN PYTHON

Linear decision boundaries


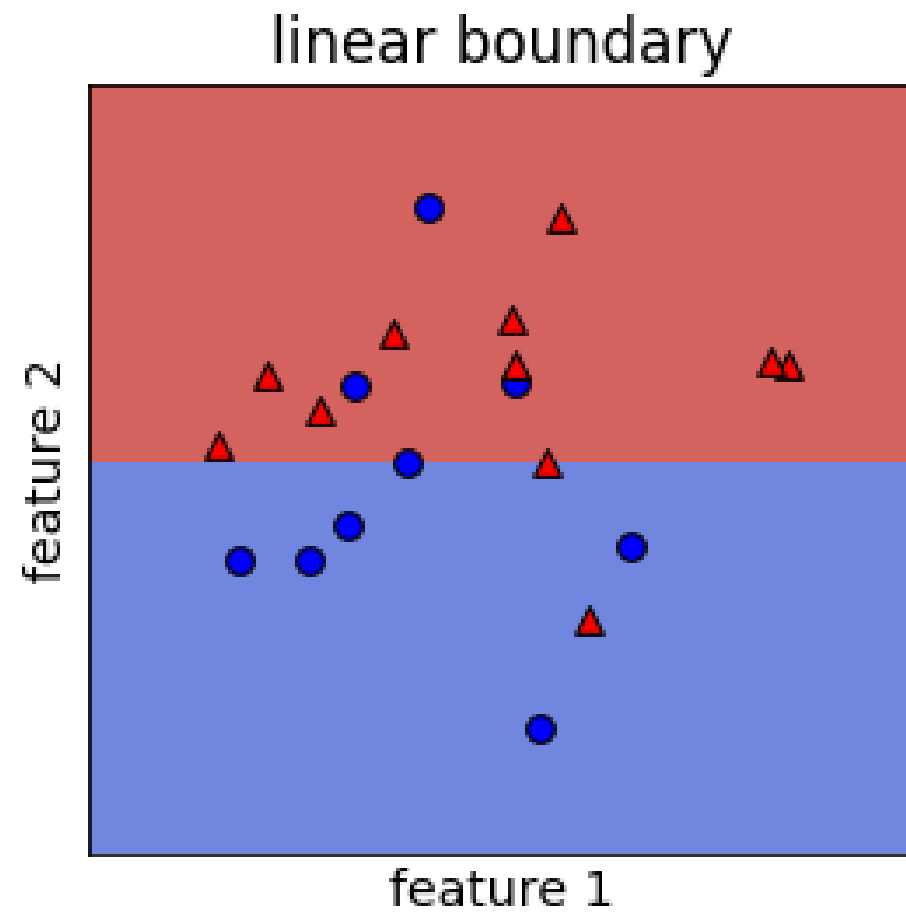
LINEAR CLASSIFIERS IN PYTHON



Michael (Mike) Gelbart

Instructor, The University of British
Columbia

Linear decision boundaries




Definitions

Vocabulary:


- **classification**: learning to predict categories
- **decision boundary**: the surface separating different predicted classes
- **linear classifier**: a classifier that learns linear decision boundaries
 - e.g., logistic regression, linear SVM
- **linearly separable**: a data set can be perfectly explained by a linear classifier

Linearly separable data

not linearly separable



linearly separable



Let's practice!

LINEAR CLASSIFIERS IN PYTHON

Linear classifiers: prediction equations

LINEAR CLASSIFIERS IN PYTHON



Michael (Mike) Gelbart

Instructor, The University of British
Columbia

Dot Products

```
x = np.arange(3)
x
```

```
array([0, 1, 2])
```

```
y = np.arange(3,6)
y
```

```
array([3, 4, 5])
```

```
x*y
```

```
array([0, 4, 10])
```

```
np.sum(x*y)
```

```
14
```

```
x@y
```

```
14
```

- `x@y` is called the dot product of `x` and `y`, and is written $x \cdot y$.

Linear classifier prediction

- raw model output = coefficients · features + intercept
- Linear classifier prediction: compute raw model output, check the sign
 - if positive, predict one class
 - if negative, predict the other class
- This is the same for logistic regression and linear SVM
 - `fit` is different but `predict` is the same

How LogisticRegression makes predictions

raw model output = coefficients · features + intercept

```
lr = LogisticRegression()
```

```
lr.fit(X,y)
```

```
lr.predict(X)[10]
```

```
0
```

```
lr.predict(X)[20]
```

```
1
```

How LogisticRegression makes predictions (cont.)


```
lr.coef_ @ X[10] + lr.intercept_ # raw model output
```

```
array([-33.78572166])
```


```
lr.coef_ @ X[20] + lr.intercept_ # raw model output
```

```
array([ 0.08050621])
```


The raw model output



The raw model output



The raw model output

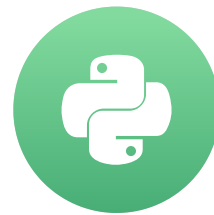


Let's practice!

LINEAR CLASSIFIERS IN PYTHON

What is a loss function?

LINEAR CLASSIFIERS IN PYTHON



Michael Gelbart

Instructor, The University of British
Columbia

Least squares: the squared loss

- scikit-learn's `LinearRegression` minimizes a loss:

$$\sum_{i=1}^n (\text{true } i\text{th target value} - \text{predicted } i\text{th target value})^2$$

- Minimization is with respect to coefficients or parameters of the model.
- Note that in scikit-learn `model.score()` isn't necessarily the loss function.

Classification errors: the 0-1 loss

- Squared loss not appropriate for classification problems (more on this later).
- A natural loss for classification problem is the number of errors.
- This is the **0-1 loss**: it's 0 for a correct prediction and 1 for an incorrect prediction.
- But this loss is hard to minimize!

Minimizing a loss

```
from scipy.optimize import minimize
```

```
minimize(np.square, 0).x
```

```
array([0.])
```

```
minimize(np.square, 2).x
```

```
array([-1.88846401e-08])
```

Let's practice!

LINEAR CLASSIFIERS IN PYTHON

Loss function diagrams


LINEAR CLASSIFIERS IN PYTHON




Michael (Mike) Gelbart

Instructor, The University of British
Columbia


The raw model output




0-1 loss diagram




Linear regression loss diagram




Logistic loss diagram



Hinge loss diagram



Hinge loss diagram



Let's practice!

LINEAR CLASSIFIERS IN PYTHON

Logistic regression and regularization


LINEAR CLASSIFIERS IN PYTHON




Michael (Mike) Gelbart

Instructor, The University of British
Columbia

Regularized logistic regression



Regularized logistic regression



How does regularization affect training accuracy?

```
lr_weak_reg = LogisticRegression(C=100)
lr_strong_reg = LogisticRegression(C=0.01)

lr_weak_reg.fit(X_train, y_train)
lr_strong_reg.fit(X_train, y_train)

lr_weak_reg.score(X_train, y_train)
lr_strong_reg.score(X_train, y_train)
```

```
1.0
0.92
```

regularized loss = original loss + large coefficient penalty

- more regularization: lower training accuracy

How does regularization affect test accuracy?

```
lr_weak_reg.score(X_test, y_test)
```

```
0.86
```

```
lr_strong_reg.score(X_test, y_test)
```

```
0.88
```

regularized loss = original loss + large coefficient penalty

- more regularization: lower training accuracy
- more regularization: (almost always) higher test accuracy



L1 vs. L2 regularization

- Lasso = linear regression with L1 regularization
- Ridge = linear regression with L2 regularization
- For other models like logistic regression we just say L1, L2, etc.

```
lr_L1 = LogisticRegression(penalty='l1')  
lr_L2 = LogisticRegression() # penalty='l2' by default  
  
lr_L1.fit(X_train, y_train)  
lr_L2.fit(X_train, y_train)
```

```
plt.plot(lr_L1.coef_.flatten())  
plt.plot(lr_L2.coef_.flatten())
```

L2 vs. L1 regularization



Let's practice!

LINEAR CLASSIFIERS IN PYTHON

Logistic regression and probabilities

LINEAR CLASSIFIERS IN PYTHON




Michael (Mike) Gelbart

Instructor, The University of British
Columbia

Logistic regression probabilities


Without regularization
($C = 10^8$):

- model coefficients: `[[1.55 1.57]]`
- model intercept: `[-0.64]`



Logistic regression probabilities


Without regularization
($C = 10^8$):



- model coefficients:
`[[1.55 1.57]]`
- model intercept:
`[-0.64]`


Logistic regression probabilities

Without regularization
($C = 10^8$):



- model coefficients:
`[[1.55 1.57]]`
- model intercept:
`[-0.64]`


With regularization ($C = 1$):



- model coefficients:
`[[0.45 0.64]]`
- model intercept:
`[-0.26]`

How are these probabilities computed?

- logistic regression predictions: sign of raw model output
- logistic regression probabilities: "squashed" raw model output



Let's practice!

LINEAR CLASSIFIERS IN PYTHON

Multi-class logistic regression

LINEAR CLASSIFIERS IN PYTHON



Michael (Mike) Gelbart

Instructor, The University of British
Columbia

Combining binary classifiers with one-vs-rest

```
lr0.fit(X, y==0)
```

```
lr1.fit(X, y==1)
```

```
lr2.fit(X, y==2)
```

```
# get raw model output  
lr0.decision_function(X)[0]
```

```
6.124
```

```
lr1.decision_function(X)[0]
```

```
-5.429
```

```
lr2.decision_function(X)[0]
```

```
-7.532
```

```
lr.fit(X, y)  
lr.predict(X)[0]
```

```
0
```

One-vs-rest:

- fit a binary classifier for each class
- predict with all, take largest output
- pro: simple, modular
- con: not directly optimizing accuracy
- common for SVMs as well
- can produce probabilities

"Multinomial" or "softmax":

- fit a single classifier for all classes
- prediction directly outputs best class
- con: more complicated, new code
- pro: tackle the problem directly
- possible for SVMs, but less common
- can produce probabilities

Model coefficients for multi-class

```
# one-vs-rest by default
lr_ovr = LogisticRegression()

lr_ovr.fit(X,y)

lr_ovr.coef_.shape
```

```
(3,13)
```

```
lr_ovr.intercept_.shape
```

```
(3,)
```

```
lr_mn = LogisticRegression(
    multi_class="multinomial",
    solver="lbfgs")
lr_mn.fit(X,y)

lr_mn.coef_.shape
```

```
(3,13)
```

```
lr_mn.intercept_.shape
```

```
(3,)
```

Let's practice!

LINEAR CLASSIFIERS IN PYTHON

Support Vectors

LINEAR CLASSIFIERS IN PYTHON




Michael (Mike) Gelbart


Instructor, The University of British
Columbia

What is an SVM?

- Linear classifiers (so far)
- Trained using the hinge loss and L2 regularization




- Support vector: a training example **not** in the flat part of the loss diagram
- Support vector: an example that is incorrectly classified **or** close to the boundary
- If an example is not a support vector, removing it has no effect on the model
- Having a small number of support vectors makes kernel SVMs really fast




Max-margin viewpoint

- The SVM maximizes the "margin" for linearly separable datasets
- Margin: distance from the boundary to the closest points



Max-margin viewpoint

- The SVM maximizes the "margin" for linearly separable datasets
- Margin: distance from the boundary to the closest points



Let's practice!

LINEAR CLASSIFIERS IN PYTHON

Kernel SVMs


LINEAR CLASSIFIERS IN PYTHON




Michael (Mike) Gelbart

Instructor, The University of British
Columbia



Transforming your features



Transforming your features





Transforming your features





transformed feature =
 $(\text{original feature})^2$

Transforming your features



transformed feature =
 $(\text{original feature})^2$

Transforming your features




transformed feature =
 $(\text{original feature})^2$

Kernel SVMs

```
from sklearn.svm import SVC
```


```
svm = SVC(gamma=1) # default is kernel="rbf"
```



Kernel SVMs

```
from sklearn.svm import SVC
```

```
svm = SVC(gamma=0.01) # default is kernel="rbf"
```




- smaller `gamma` leads to smoother boundaries

Kernel SVMs

```
from sklearn.svm import SVC
```

```
svm = SVC(gamma=2)    # default is kernel="rbf"
```



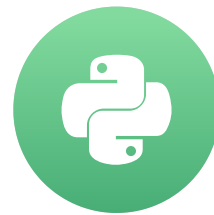
- larger `gamma` leads to more complex boundaries

Let's practice!

LINEAR CLASSIFIERS IN PYTHON

Comparing logistic regression and SVM

LINEAR CLASSIFIERS IN PYTHON



Michael (Mike) Gelbart

Instructor, The University of British
Columbia

Logistic regression:

- Is a linear classifier
- Can use with kernels, but slow
- Outputs meaningful probabilities
- Can be extended to multi-class
- All data points affect fit
- L2 or L1 regularization

Support vector machine (SVM):

- Is a linear classifier
- Can use with kernels, and fast
- Does not naturally output probabilities
- Can be extended to multi-class
- Only "support vectors" affect fit
- Conventionally just L2 regularization

Use in scikit-learn

Logistic regression in sklearn:

- `linear_model.LogisticRegression`

Key hyperparameters in sklearn:

- `C` (inverse regularization strength)
- `penalty` (type of regularization)
- `multi_class` (type of multi-class)

SVM in sklearn:

- `svm.LinearSVC` and `svm.SVC`

Use in scikit-learn (cont.)

Key hyperparameters in sklearn:

- `C` (inverse regularization strength)
- `kernel` (type of kernel)
- `gamma` (inverse RBF smoothness)

SGDClassifier

`SGDClassifier` : scales well to large datasets

```
from sklearn.linear_model import SGDClassifier
```

```
logreg = SGDClassifier(loss='log')
```

```
linsvm = SGDClassifier(loss='hinge')
```

- `SGDClassifier` hyperparameter `alpha` is like `1/C`

Let's practice!

LINEAR CLASSIFIERS IN PYTHON

Conclusion

LINEAR CLASSIFIERS IN PYTHON



Michael (Mike) Gelbart

Instructor, The University of British
Columbia

How does this course fit into data science?

- Data science
- → Machine learning
- →→ Supervised learning
- →→→ Classification
- →→→→ Linear classifiers (this course)

Congratulations & thanks!

LINEAR CLASSIFIERS IN PYTHON