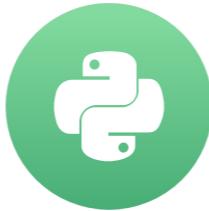


# What is Keras?

INTRODUCTION TO DEEP LEARNING WITH KERAS



**Miguel Esteban**  
Data Scientist & Founder

# Theano vs Keras

```
import theano
import theano.tensor as T
from theano.ifelse import ifelse
import numpy as np
from random import random

# Define variables
x = T.matrix('x')
w1 = theano.shared(np.array([random(),random()]))
w2 = theano.shared(np.array([random(),random()]))
w3 = theano.shared(np.array([random(),random()]))

a2 = 1/(1+T.exp(-T.dot(x,w2)-b1))
x2 = T.stack([a1,a2],axis=1)
a3 = 1/(1+T.exp(-T.dot(x2,w3)-b2))

a_hat = T.vector('a_hat') #Actual output
cost = -(a_hat*T.log(a3) + (1-a_hat)*T.log(1-a3)).sum()
dw1,dw2,dw3,db1,db2 = T.grad(cost,[w1,w2,w3,b1,b2])

[w1, w1-learning_rate*dw1],
[w2, w2-learning_rate*dw2],
[w3, w3-learning_rate*dw3],
[b1, b1-learning_rate*db1],
[b2, b2-learning_rate*db2]

]

# You can (finally) train your model
cost = []
for iteration in range(30000):
    pred, cost_iter = train(inputs, outputs)
    cost.append(cost_iter)
```

```
from keras.layers import Dense
from keras.models import Sequential

# Define model and add layers
model = Sequential()
model.add(Dense(2,input_shape=(2,),activation='sigmoid'))
model.add(Dense(1,activation='sigmoid'))

model.compile(optimizer='adam', loss='categorical_crossentropy')

# Train model
model.fit(inputs,outputs)
```

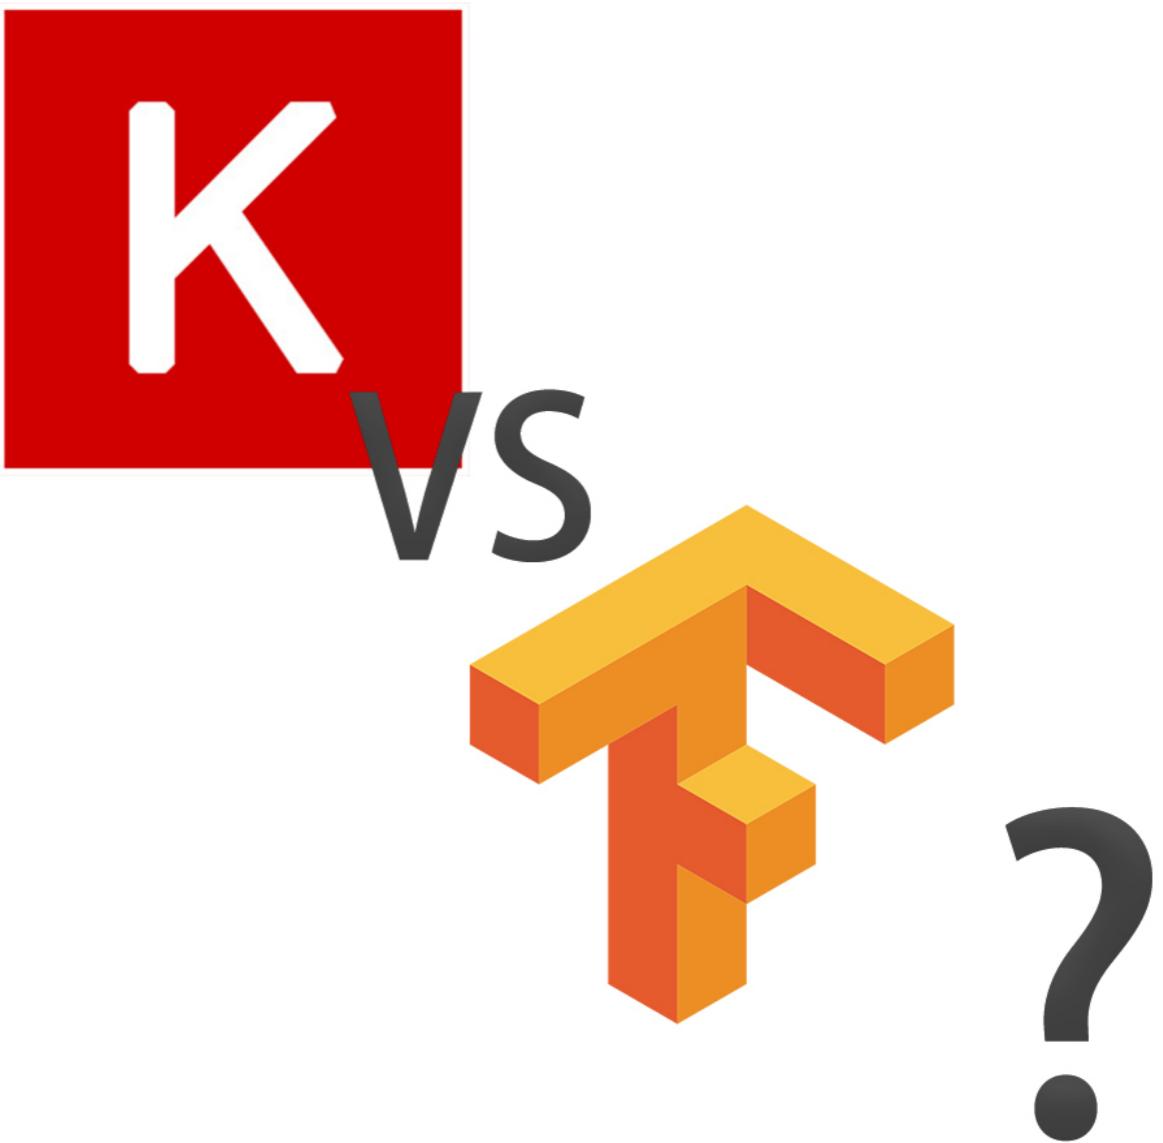
# Keras

- Deep Learning Framework
- Enables fast experimentation
- Runs on top of other frameworks
- Written by François Chollet



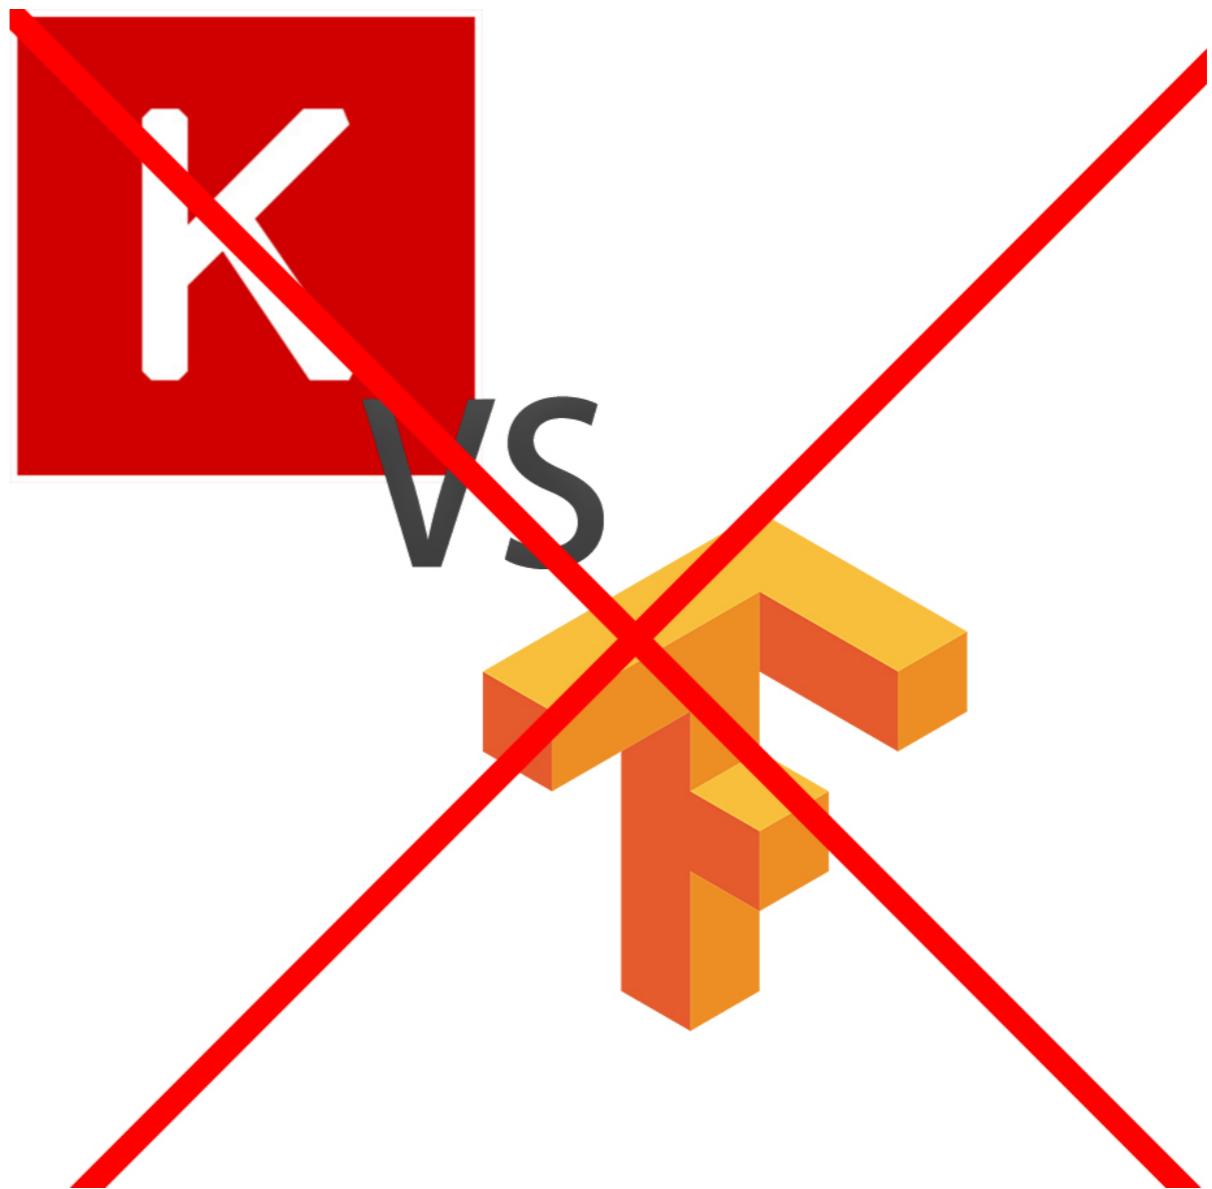
# Why use Keras?

- Fast industry-ready models
- For beginners and experts
- Less code
- Build any architecture
- Deploy models in multiple platforms

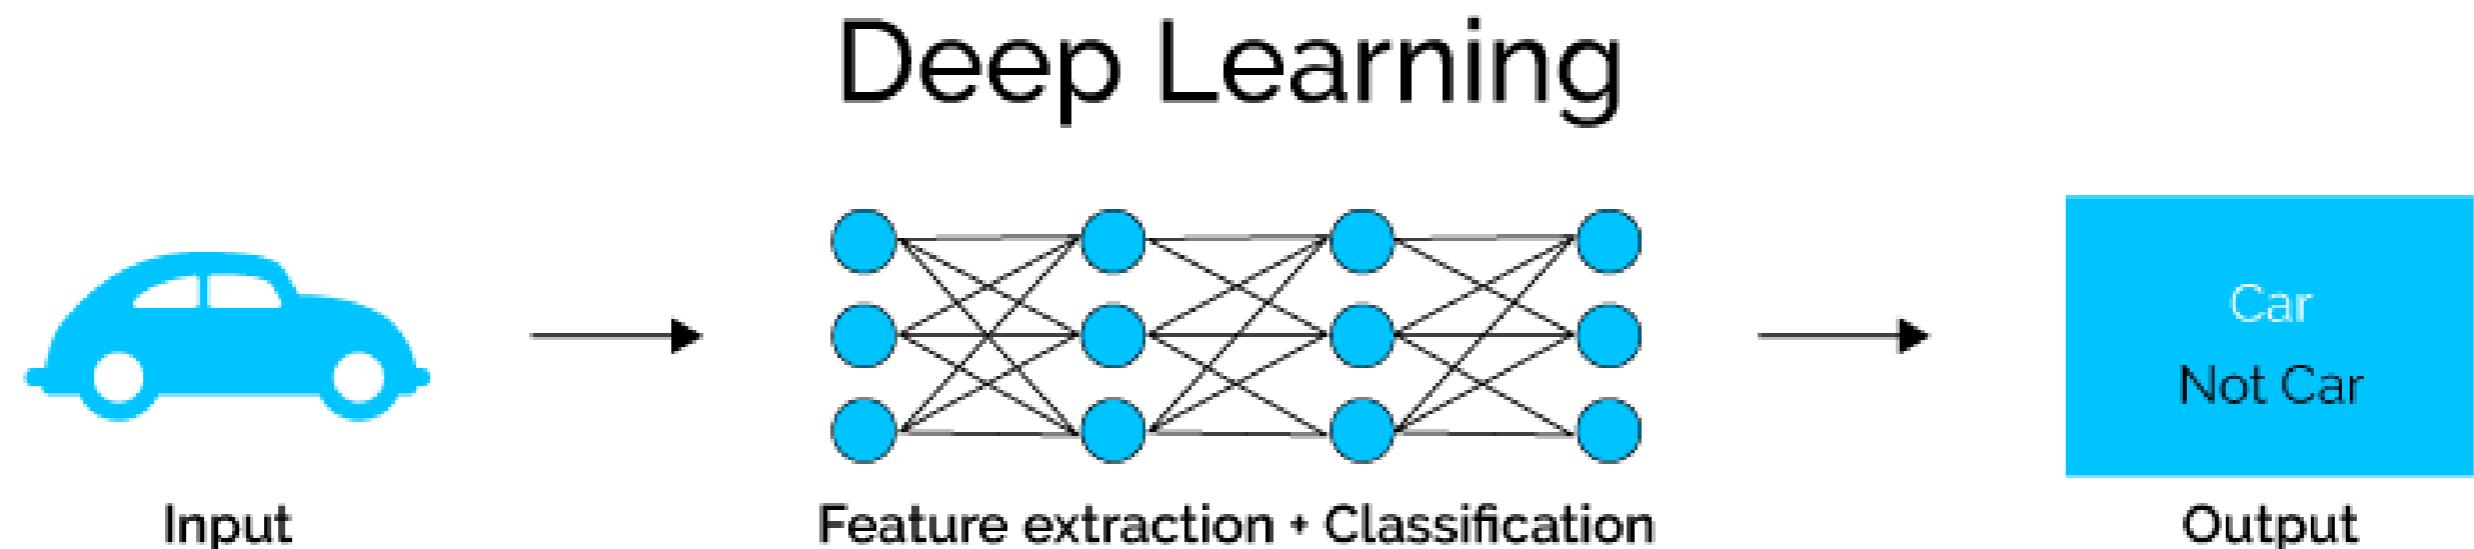
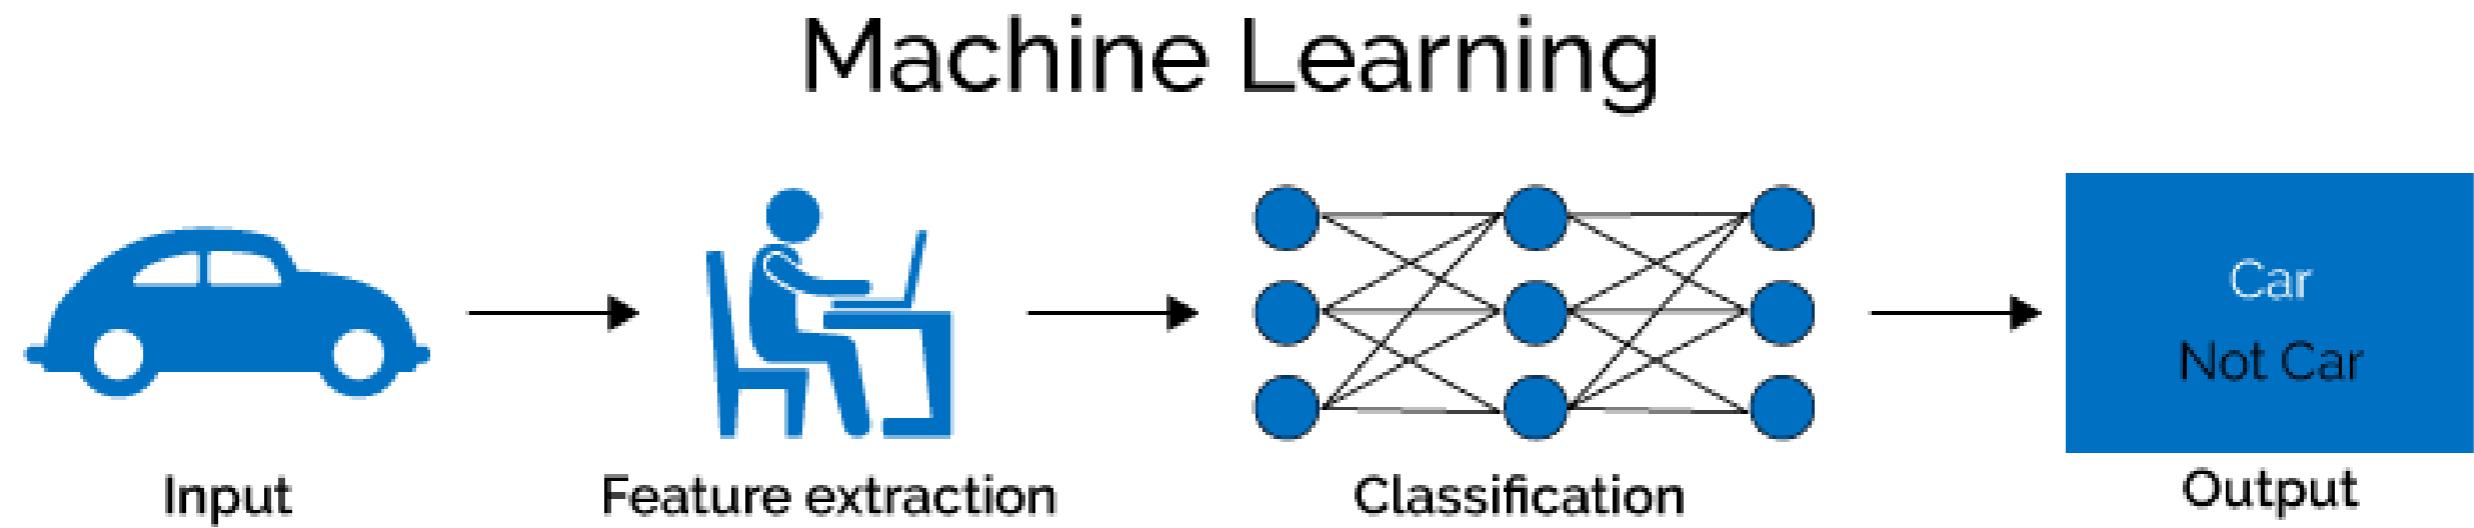


# Keras + TensorFlow

- TensorFlow's high level framework of choice
- Keras is complementary to TensorFlow
- You can use TensorFlow for low level features

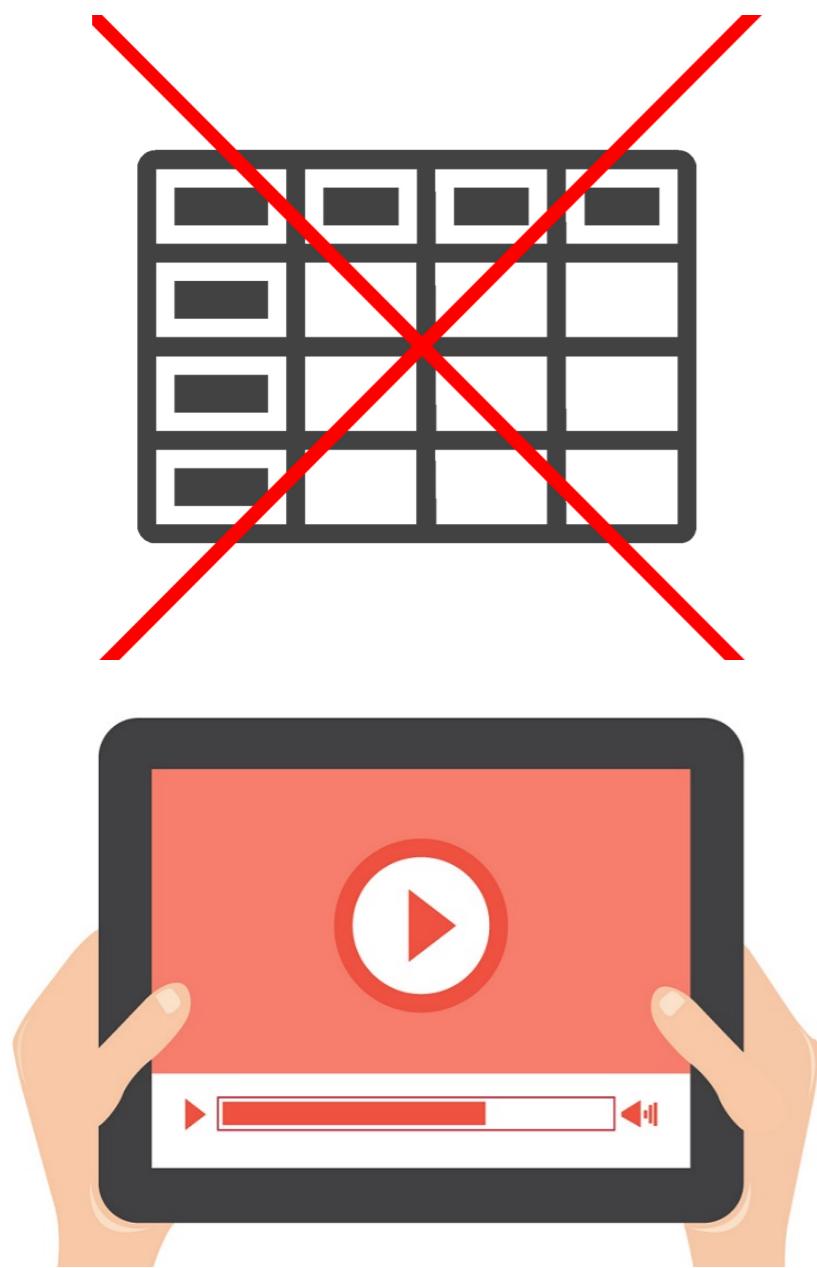


# Feature Engineering



<sup>1</sup> Towards Data Science

# Unstructured data

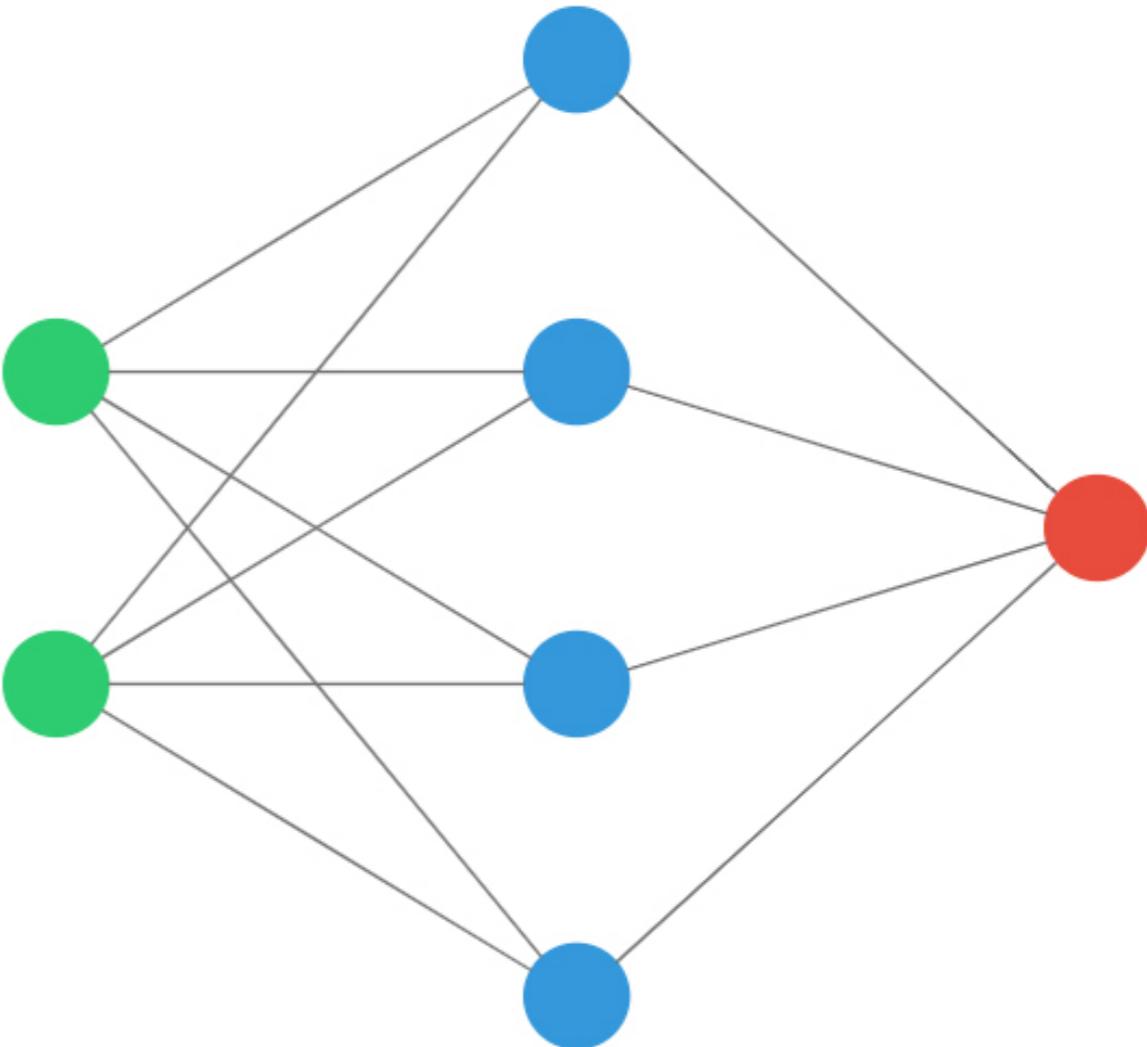


# So, when to use neural networks?

- Dealing with unstructured data
- Don't need easily interpretable results
- You can benefit from a known architecture

**Example:** Classify images of cats and dogs

- **Images -> Unstructured data**
- You don't care about why the network knows it's a cat or a dog
- You can benefit from convolutional neural networks



# Let's practice!

INTRODUCTION TO DEEP LEARNING WITH KERAS

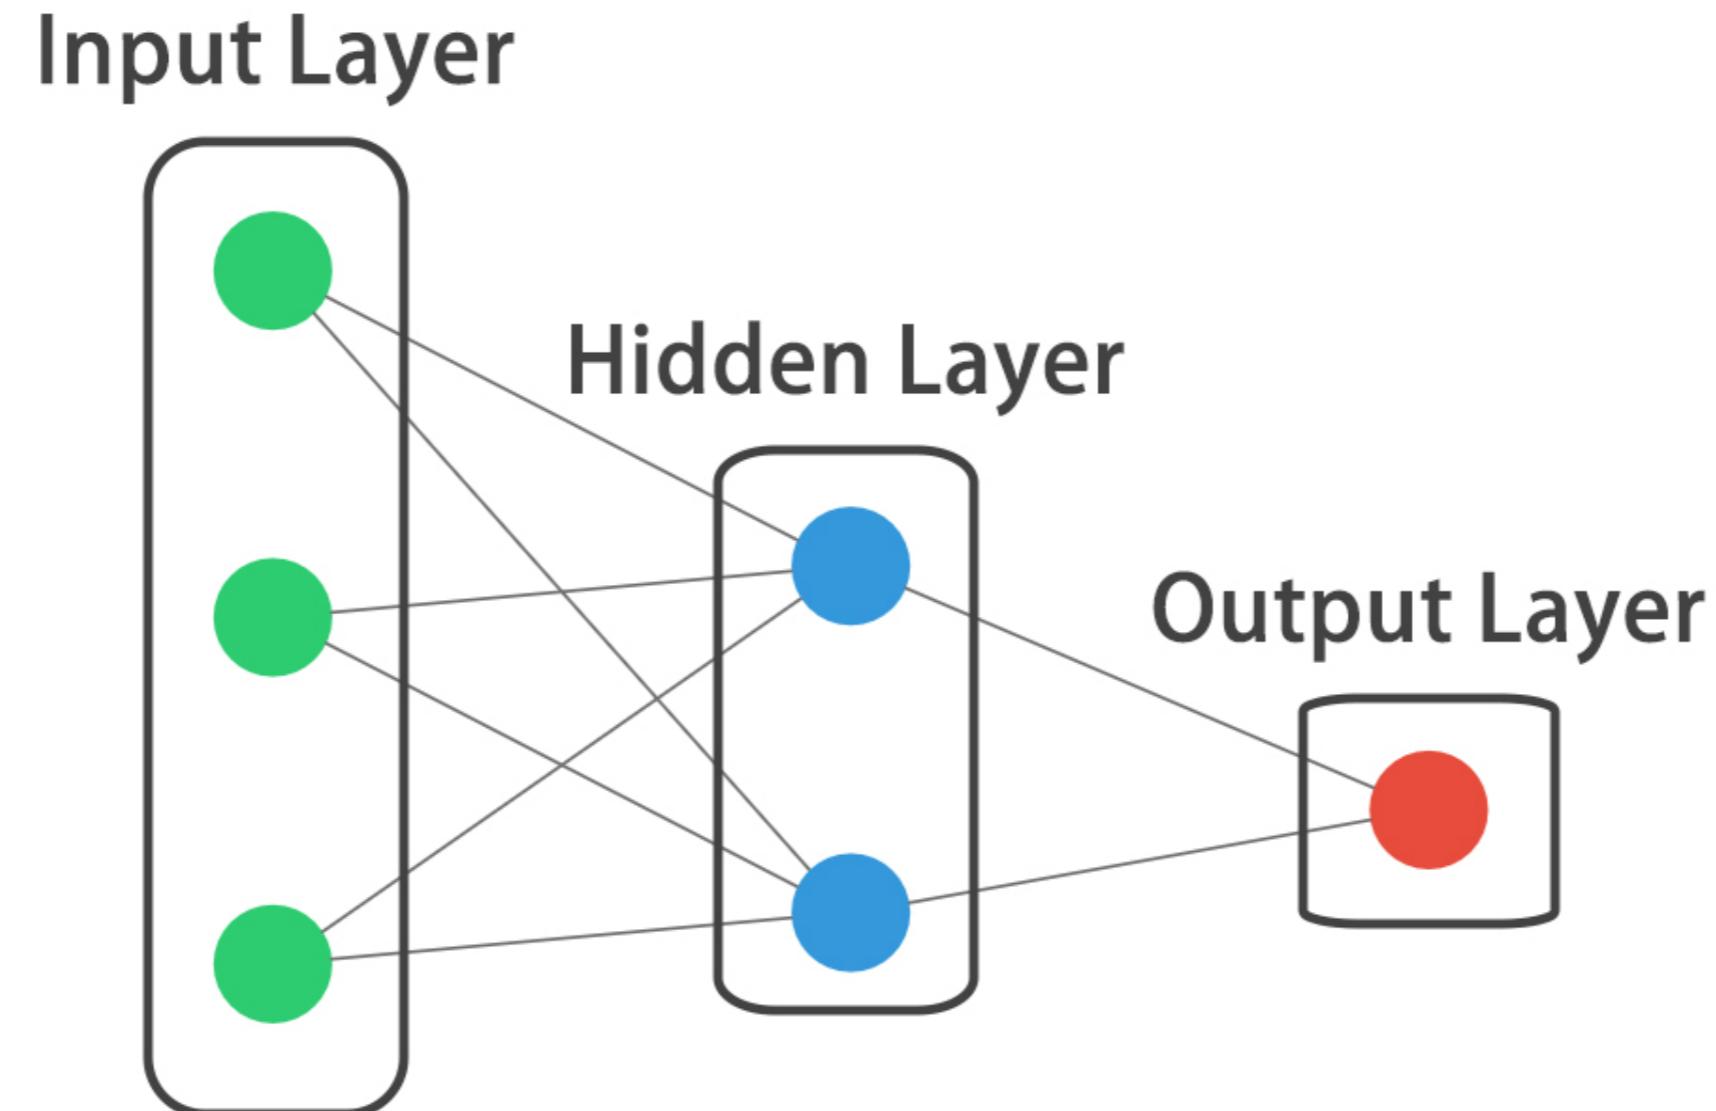
# Your first neural network

INTRODUCTION TO DEEP LEARNING WITH KERAS

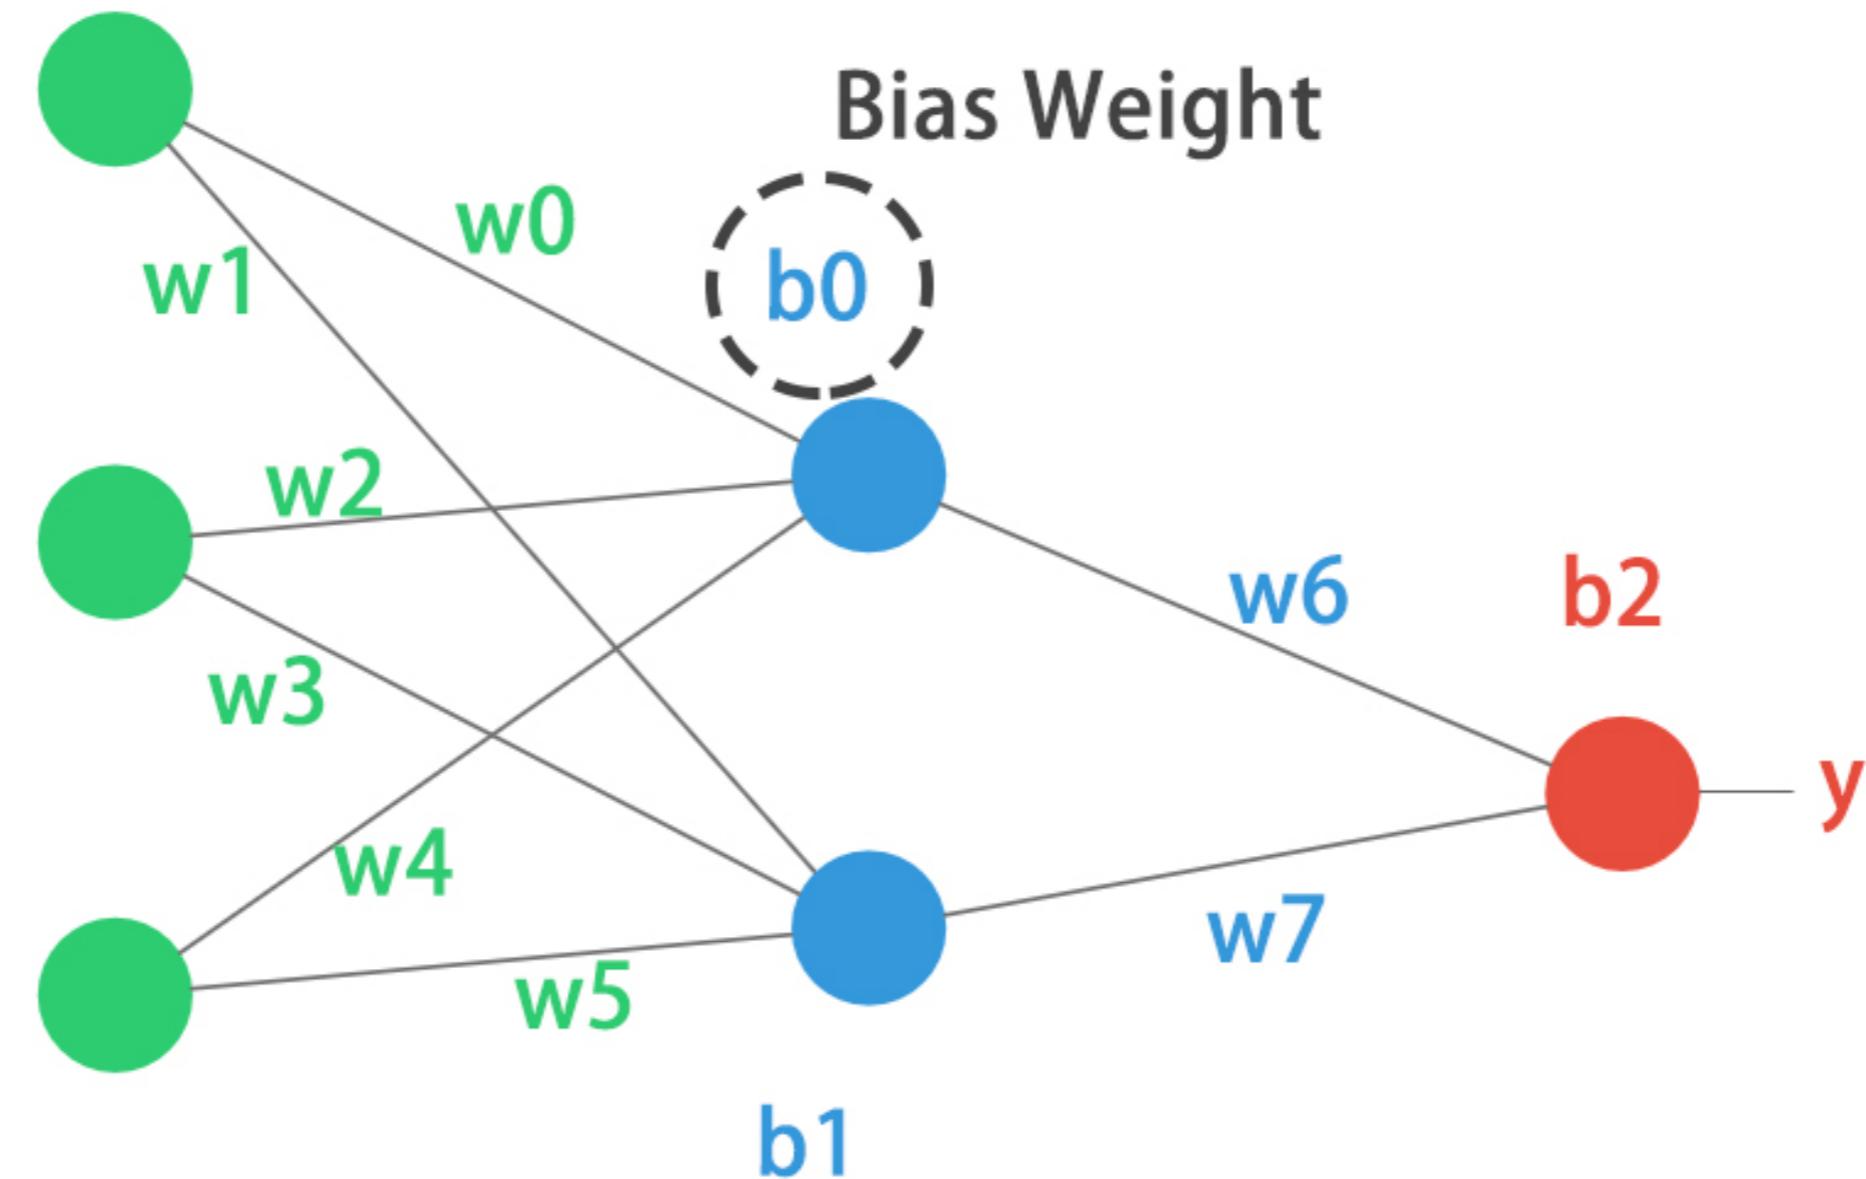


**Miguel Esteban**  
Data Scientist & Founder

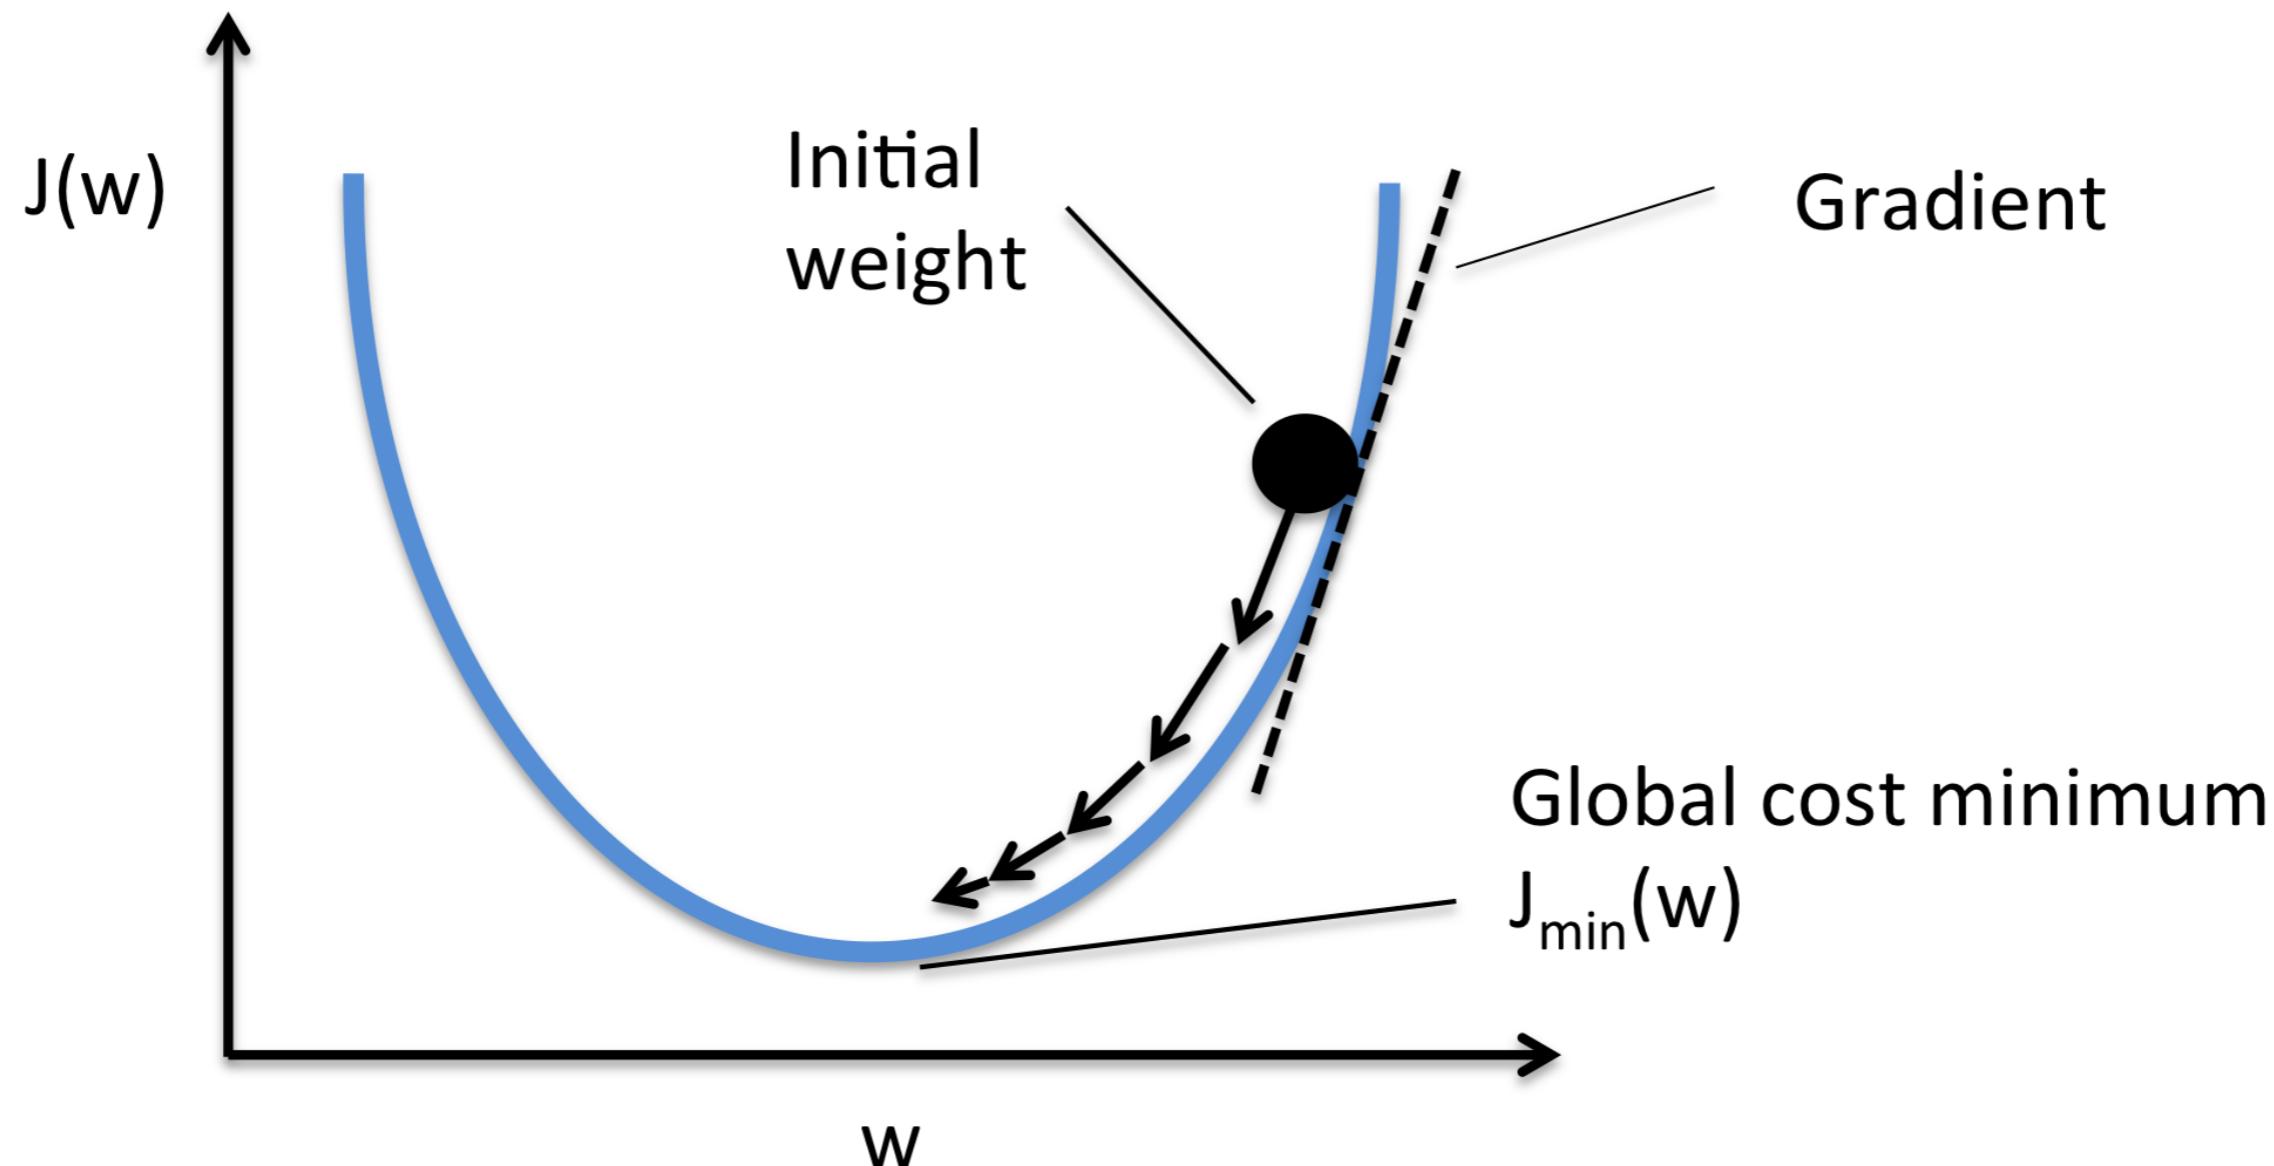
# A neural network?



# Parameters



# Gradient descent

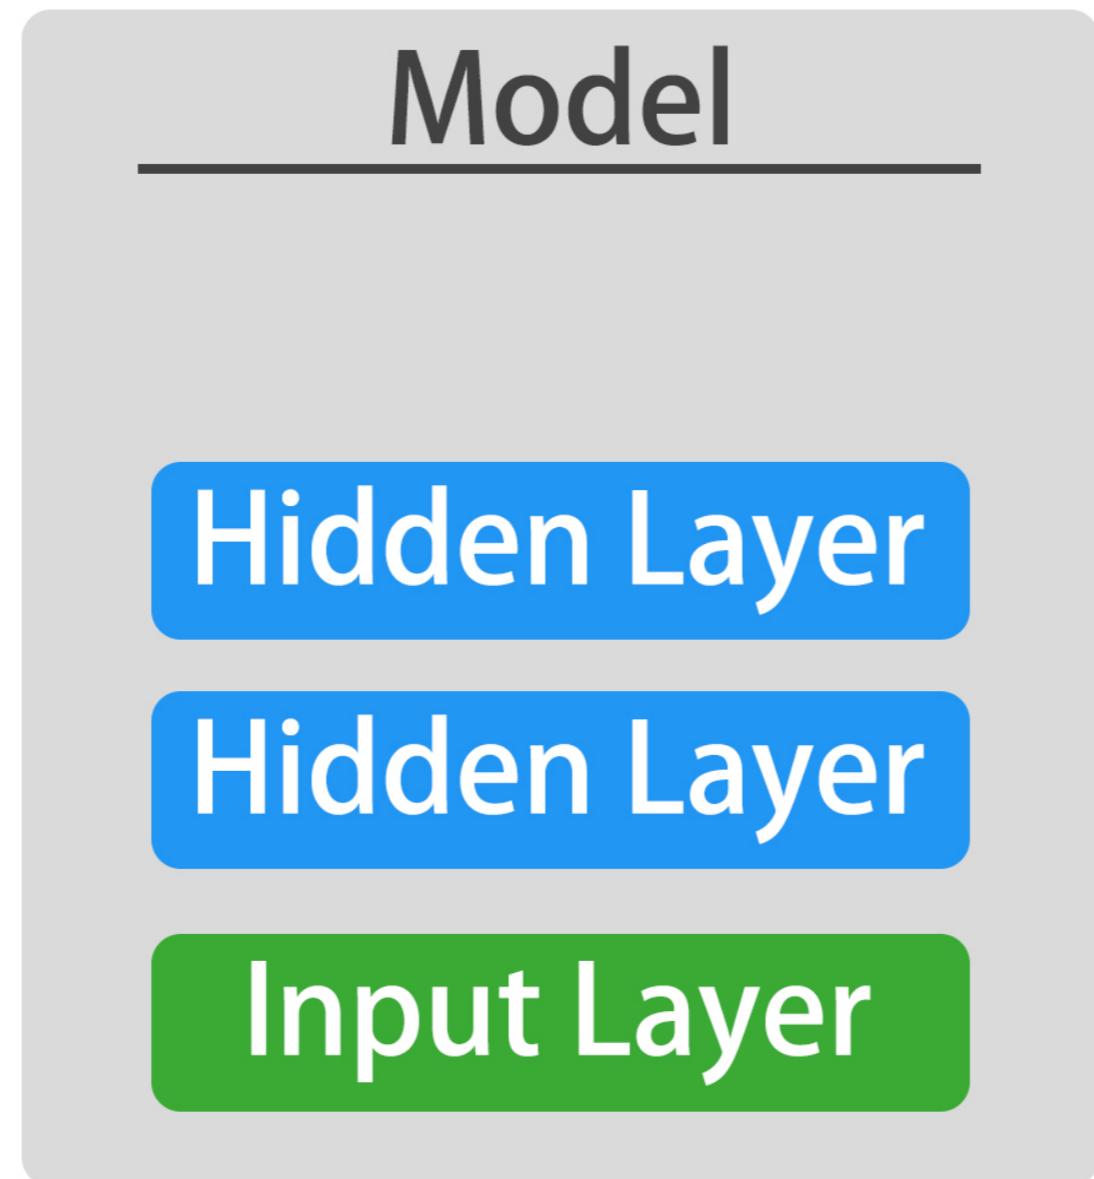


# The sequential API

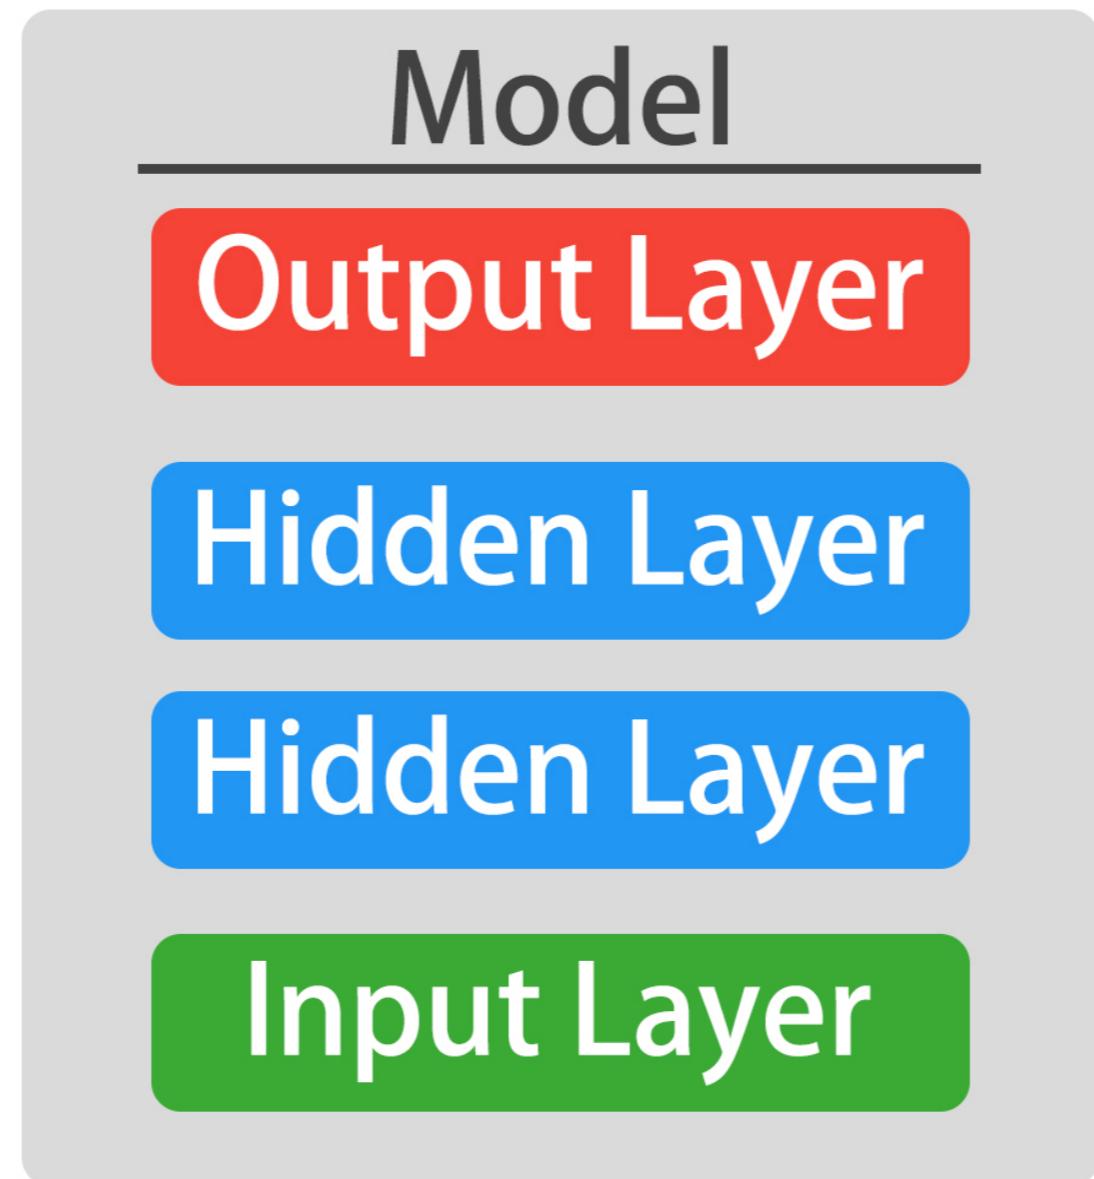
Model

Input Layer

# The sequential API

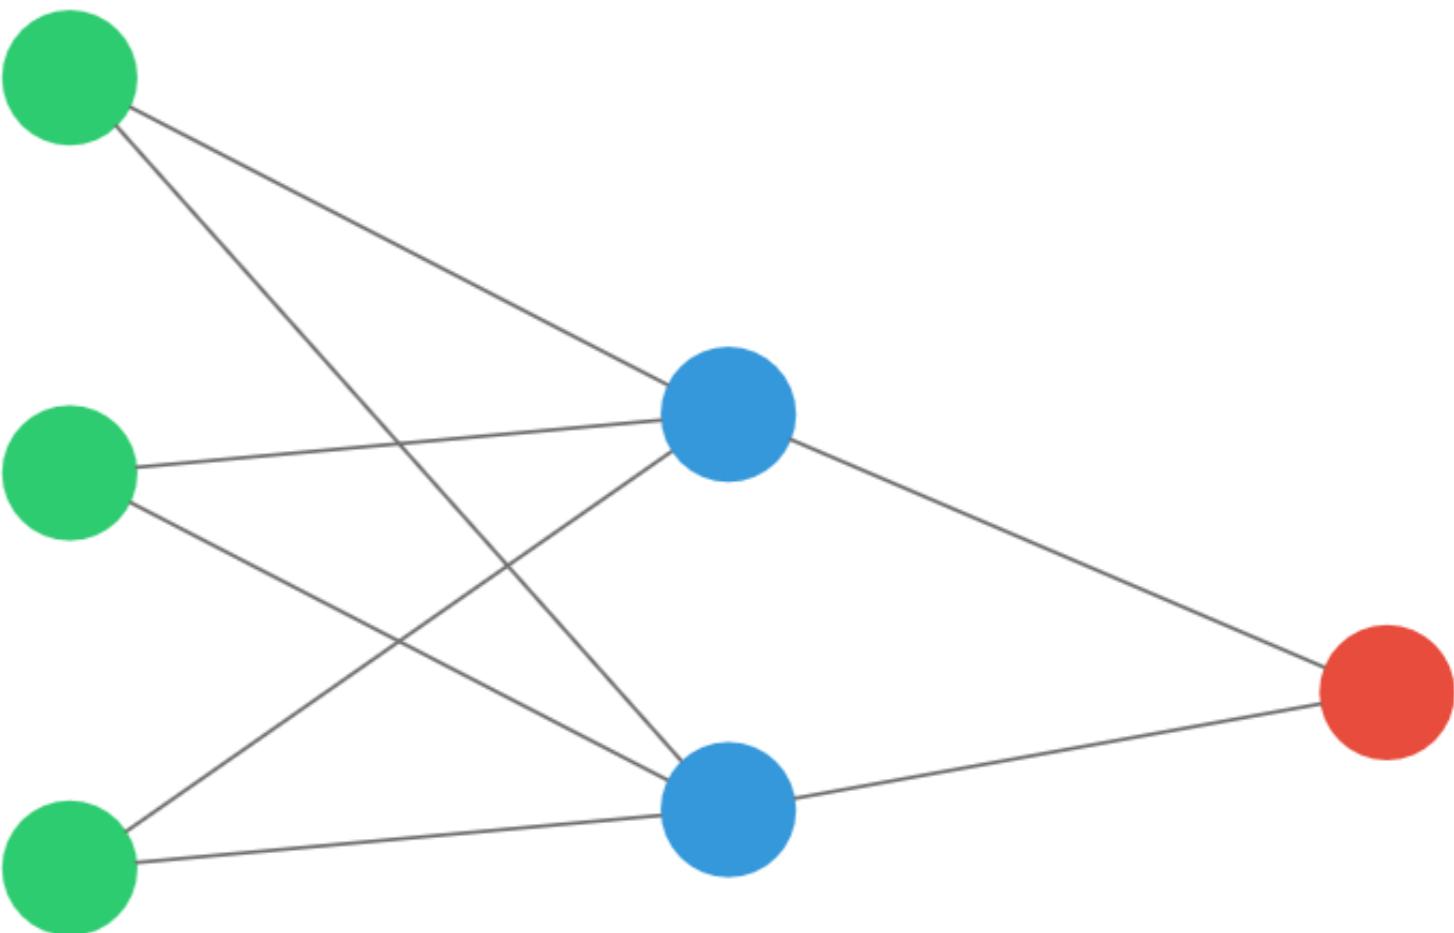


# The sequential API



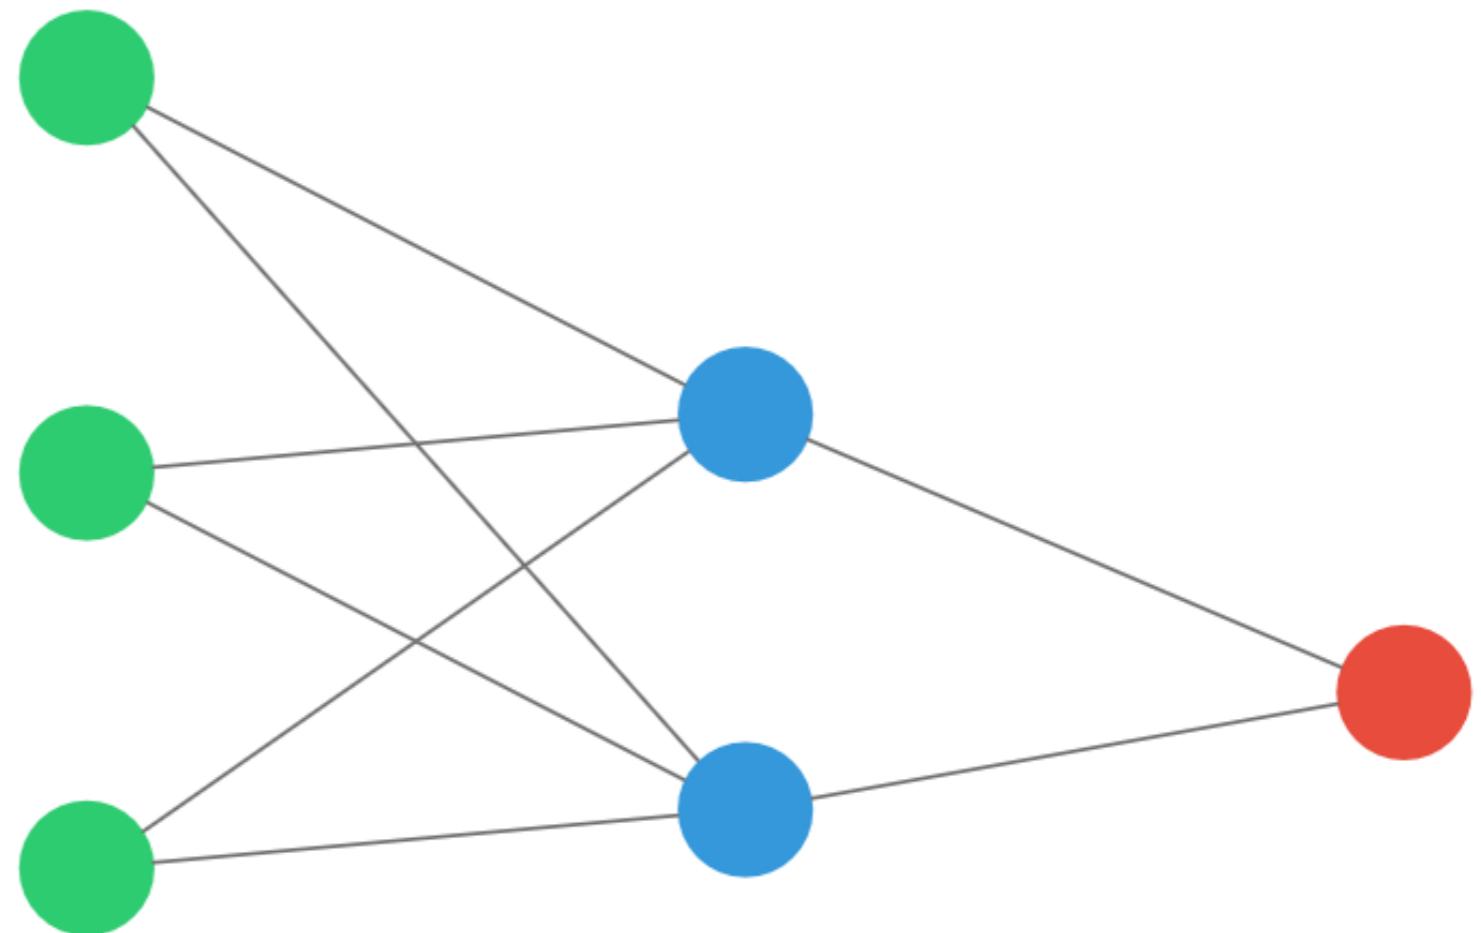
# Defining a neural network

```
from keras.models import Sequential  
  
from keras.layers import Dense  
  
# Create a new sequential model  
model = Sequential()  
  
# Add an input and dense layer  
model.add(Dense(2, input_shape=(3,)))  
  
# Add a final 1 neuron layer  
model.add(Dense(1))
```



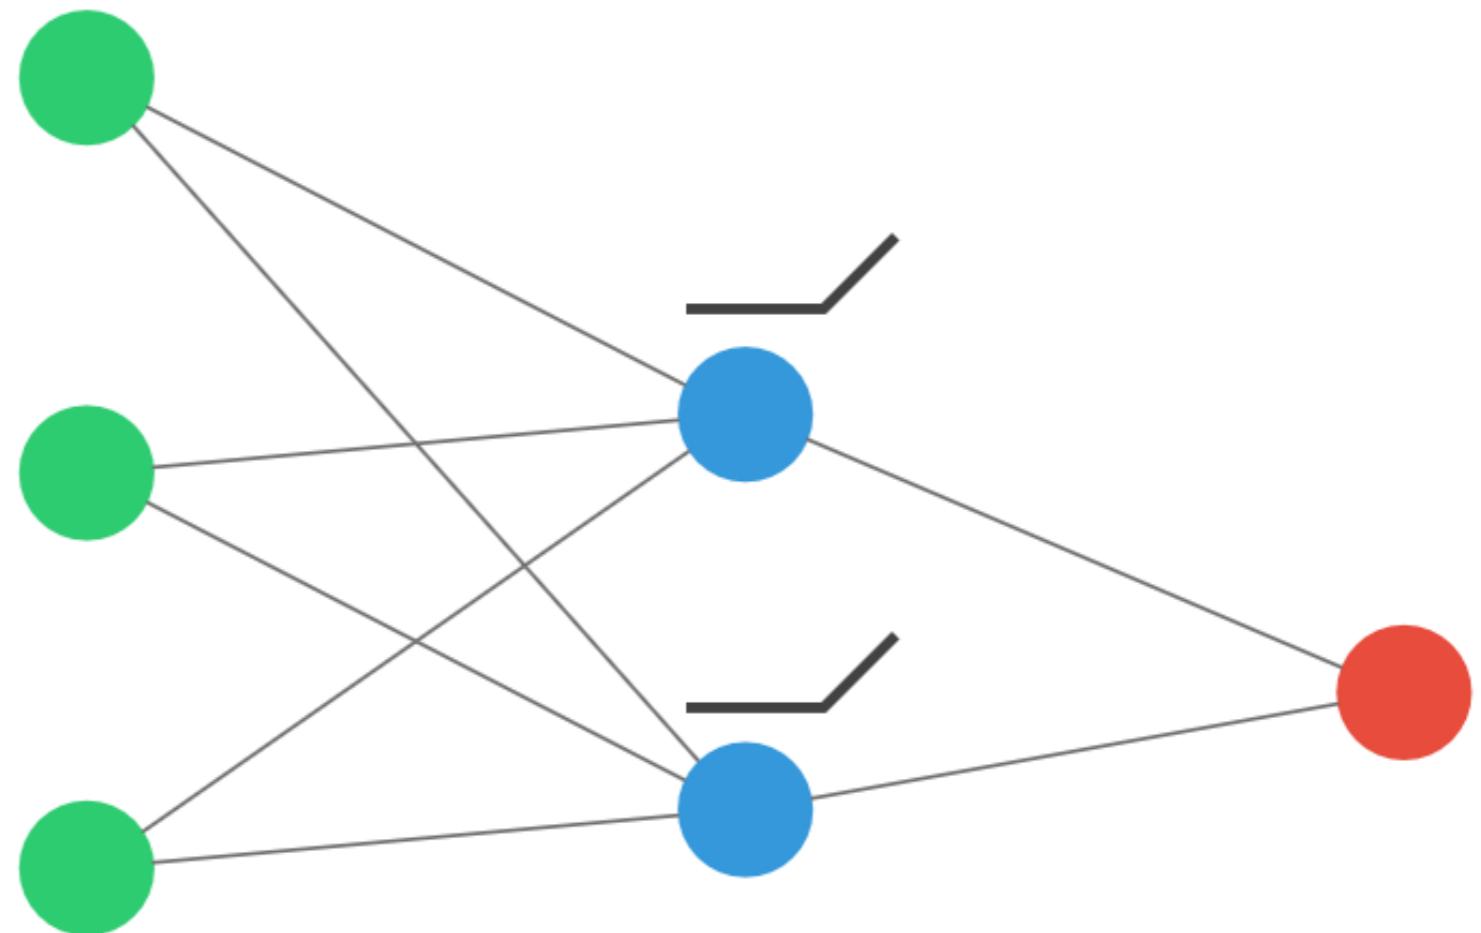
# Adding activations

```
from keras.models import Sequential  
  
from keras.layers import Dense  
  
# Create a new sequential model  
model = Sequential()  
  
# Add an input and dense layer  
model.add(Dense(2, input_shape=(3,)))  
  
# Add a final 1 neuron layer  
model.add(Dense(1))
```



# Adding activations

```
from keras.models import Sequential  
  
from keras.layers import Dense  
  
# Create a new sequential model  
model = Sequential()  
  
# Add an input and dense layer  
model.add(Dense(2, input_shape=(3,),  
               activation="relu"))  
  
# Add a final 1 neuron layer  
model.add(Dense(1))
```

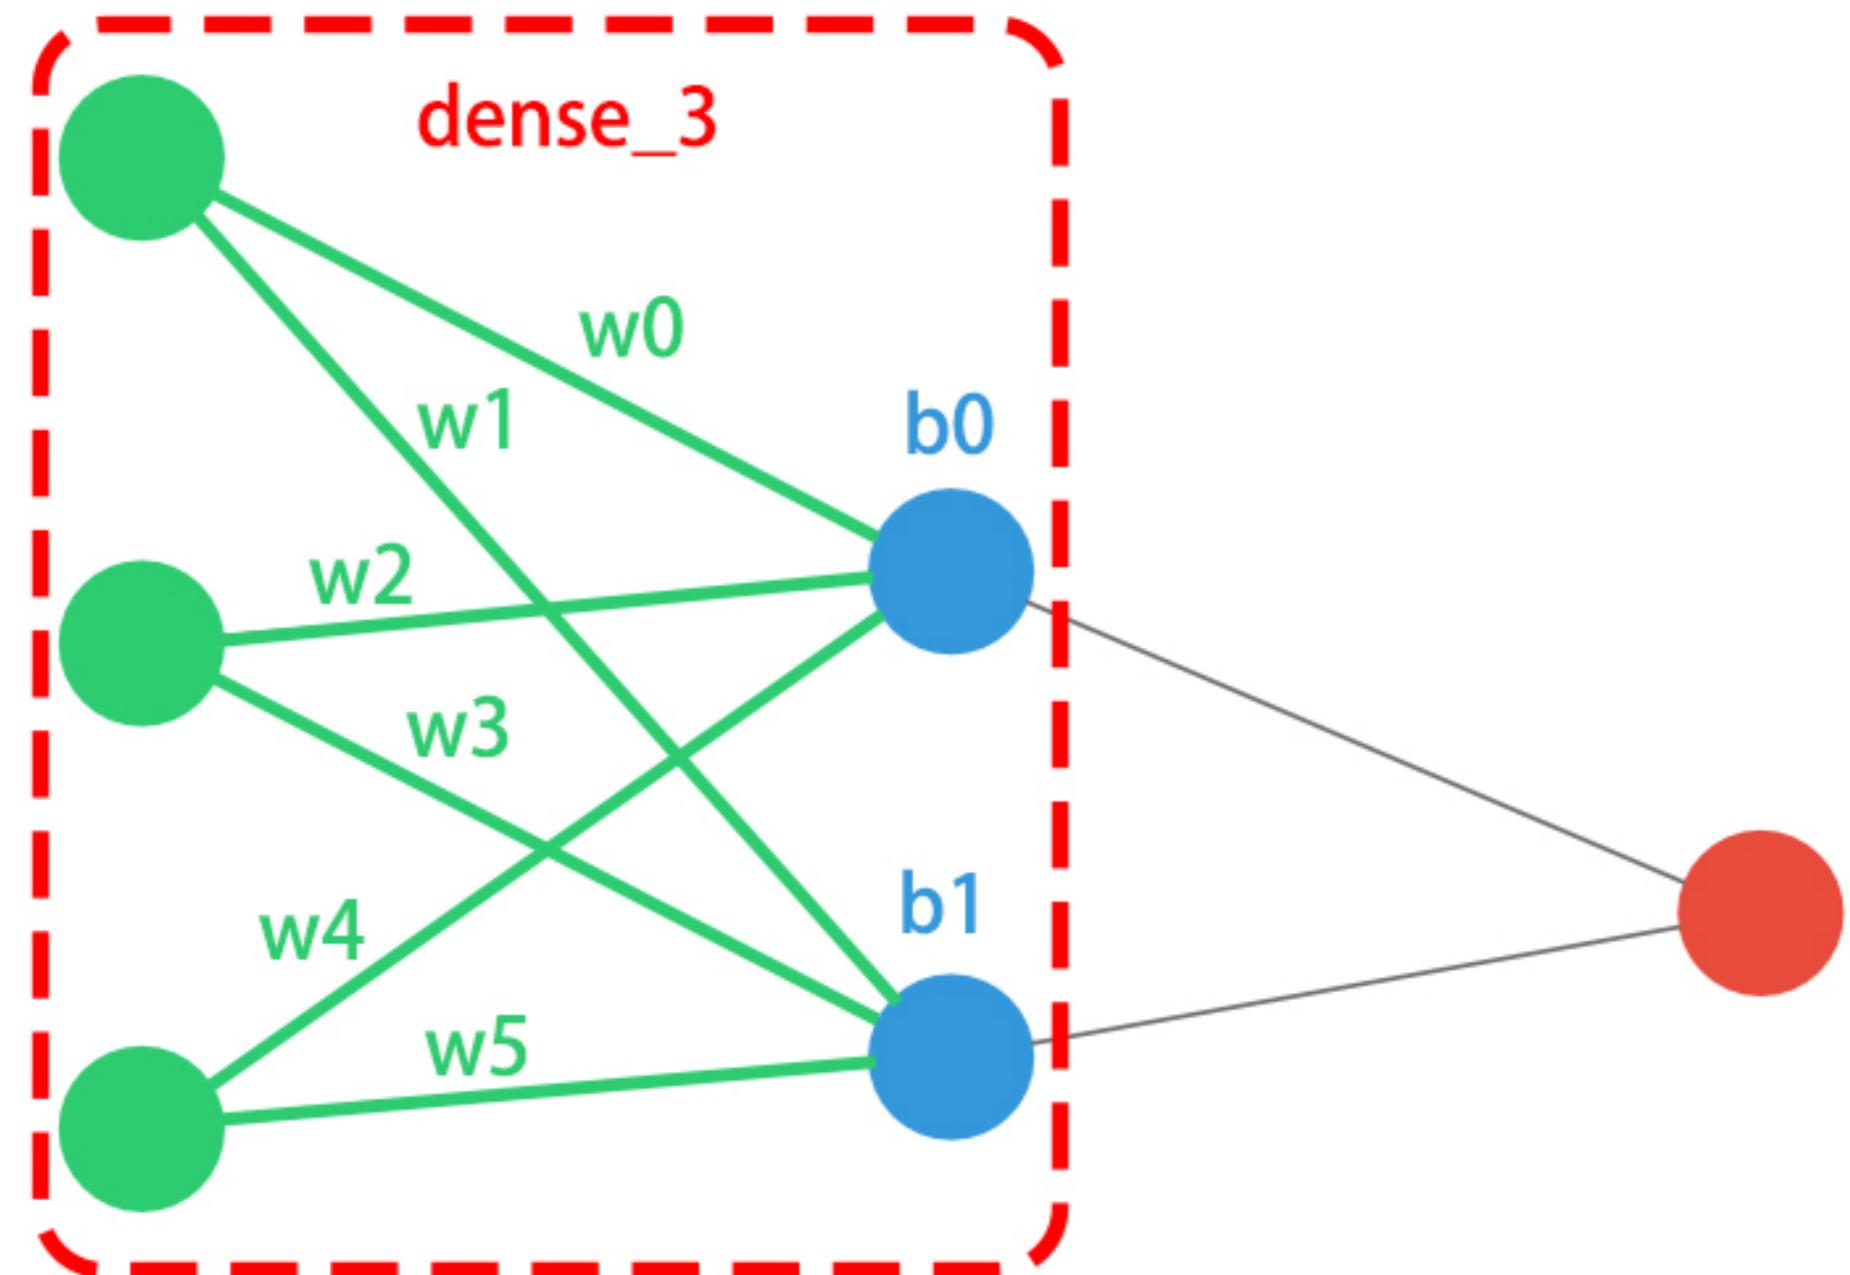


# Summarize your model!

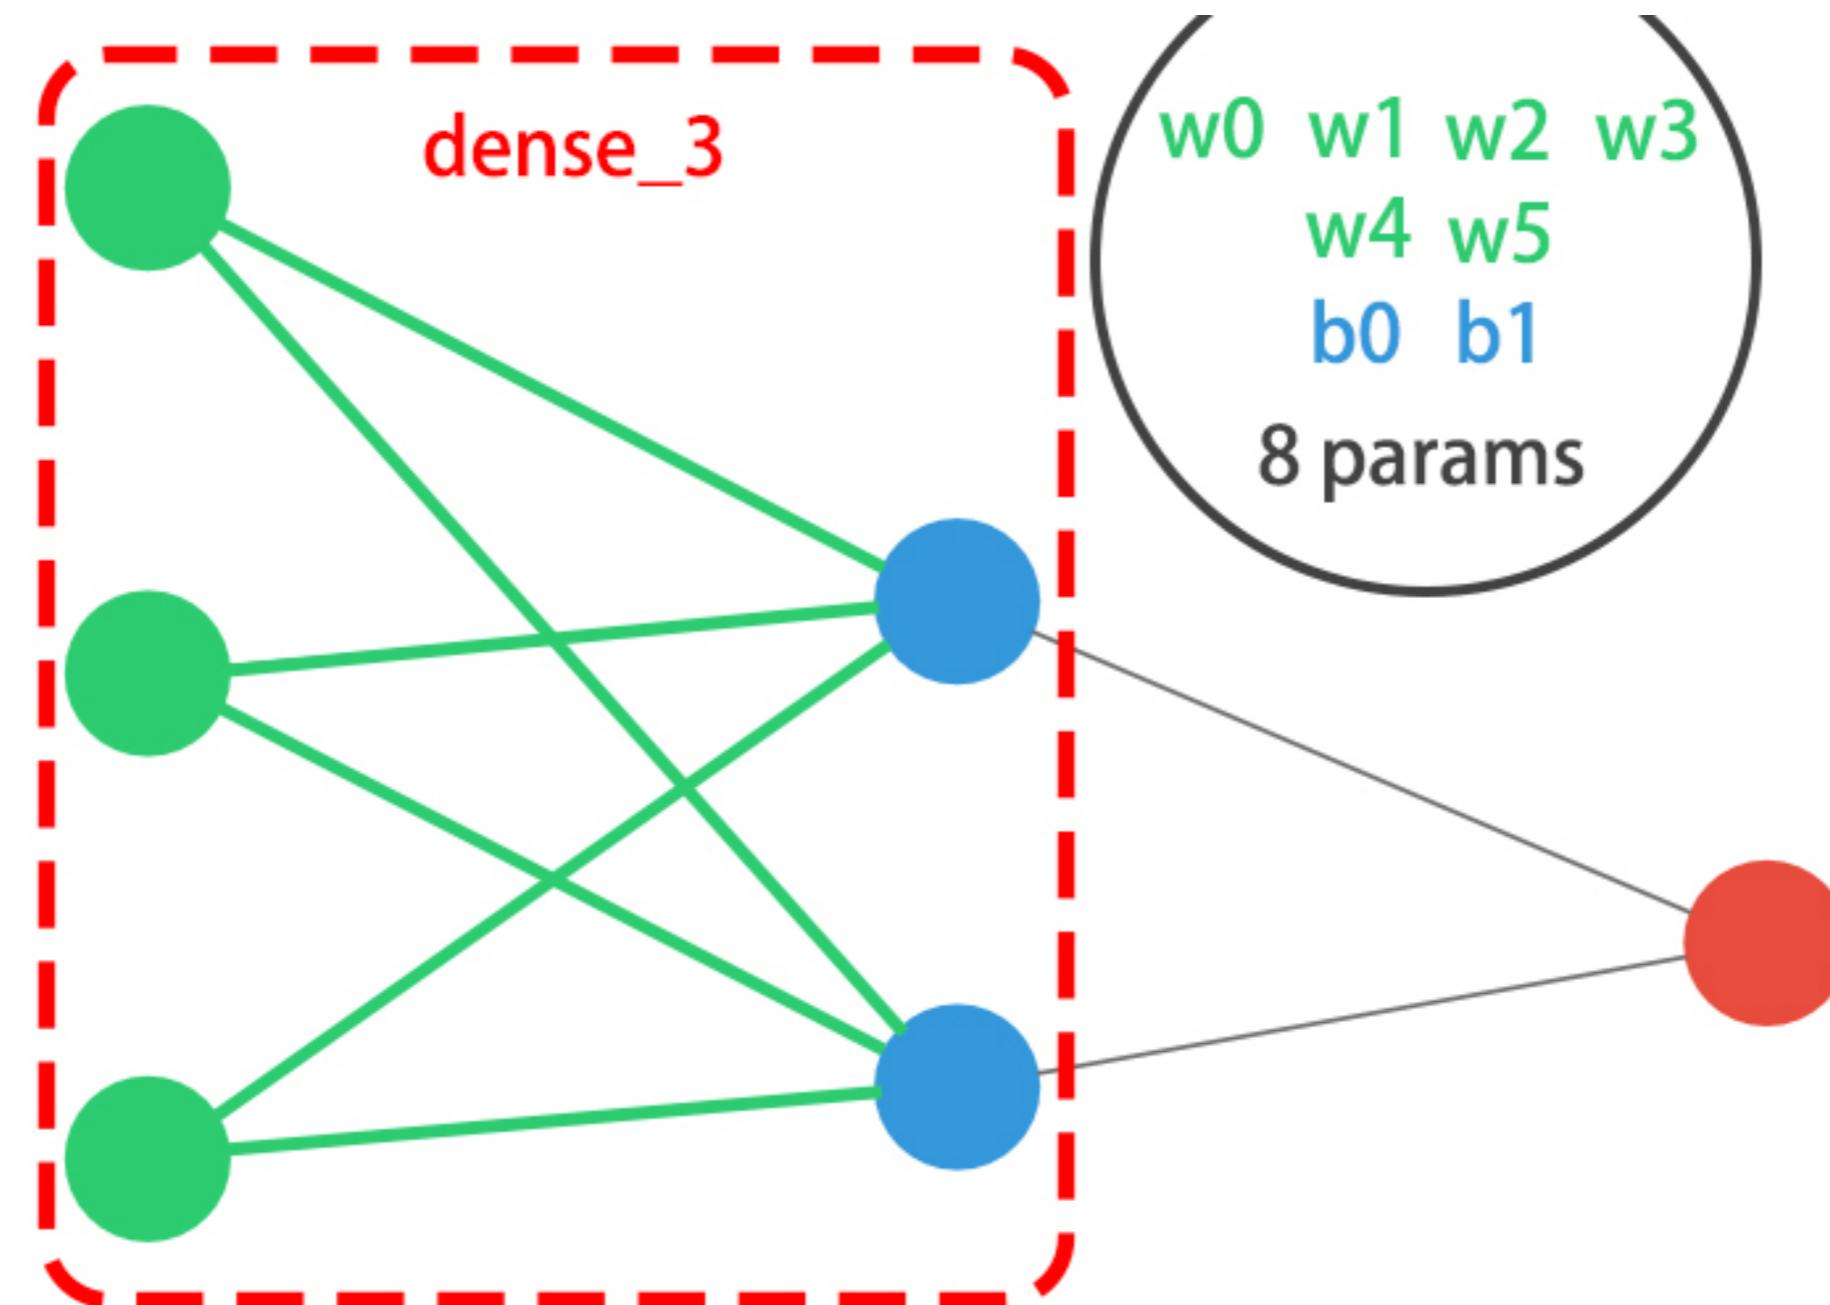
```
model.summary()
```

Layer (type)	Output Shape	Param #
<hr/>		
dense_3 (Dense)	(None, 2)	8
<hr/>		
dense_4 (Dense)	(None, 1)	3
<hr/>		
Total params:	11	
Trainable params:	11	
Non-trainable params:	0	

# Visualize parameters



# Visualize parameters



# Summarize your model!

```
model.summary()
```

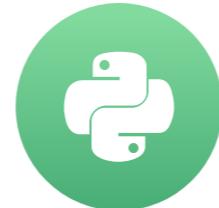
Layer (type)	Output Shape	Param #
=====		
dense_3 (Dense)	(None, 2)	--> 8 <--
-----		
dense_4 (Dense)	(None, 1)	3
=====		
Total params: 11		
Trainable params: 11		
Non-trainable params: 0		

# Let's code!

## INTRODUCTION TO DEEP LEARNING WITH KERAS

# Surviving a meteor strike

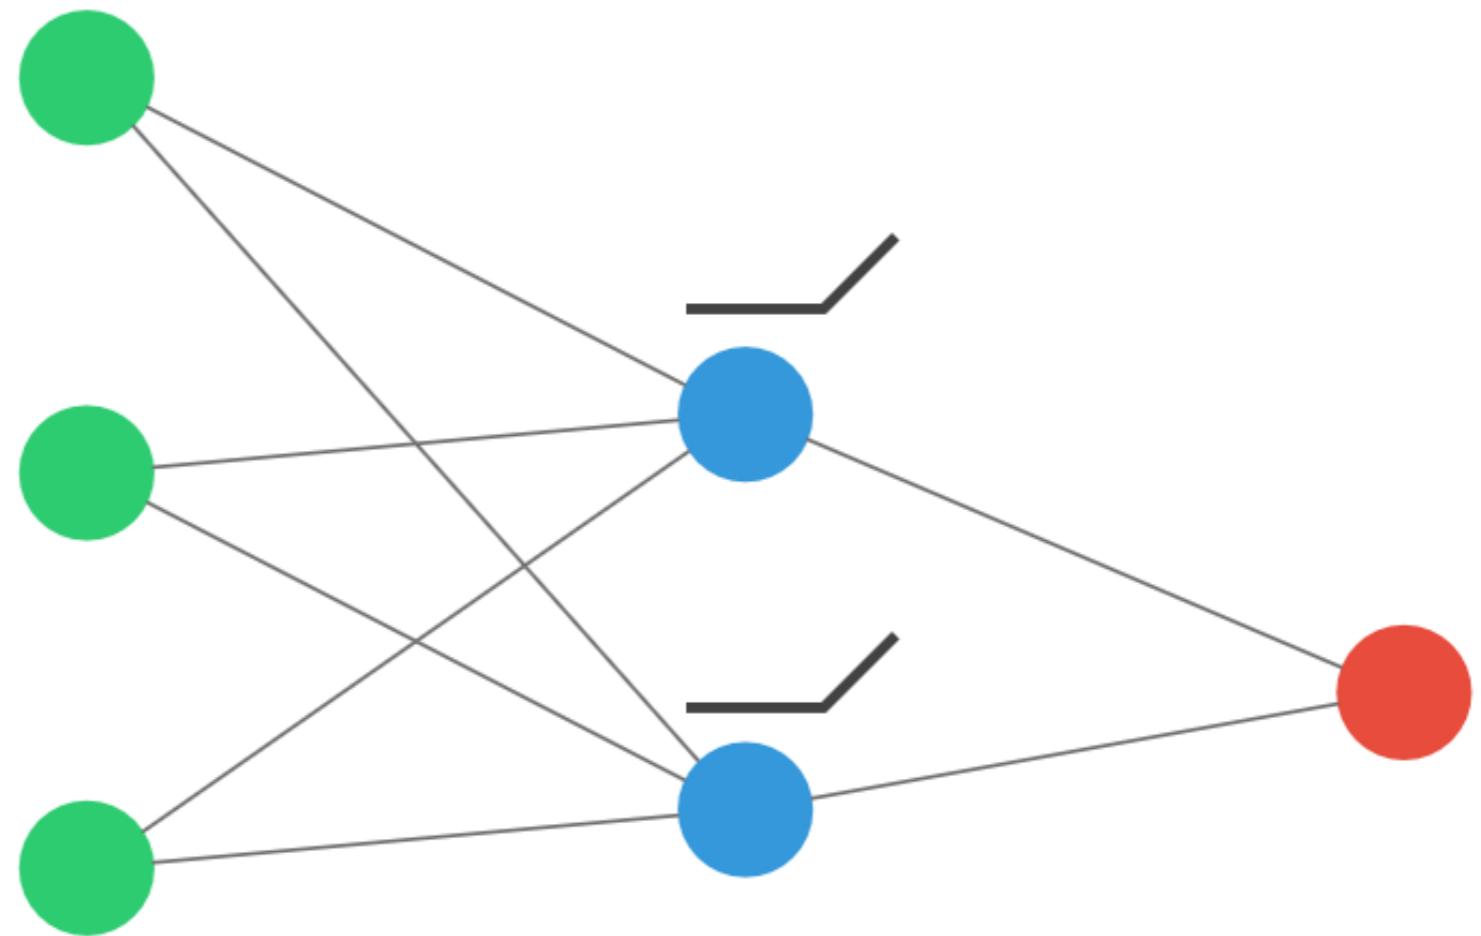
INTRODUCTION TO DEEP LEARNING WITH KERAS



**Miguel Esteban**  
Data Scientist & Founder

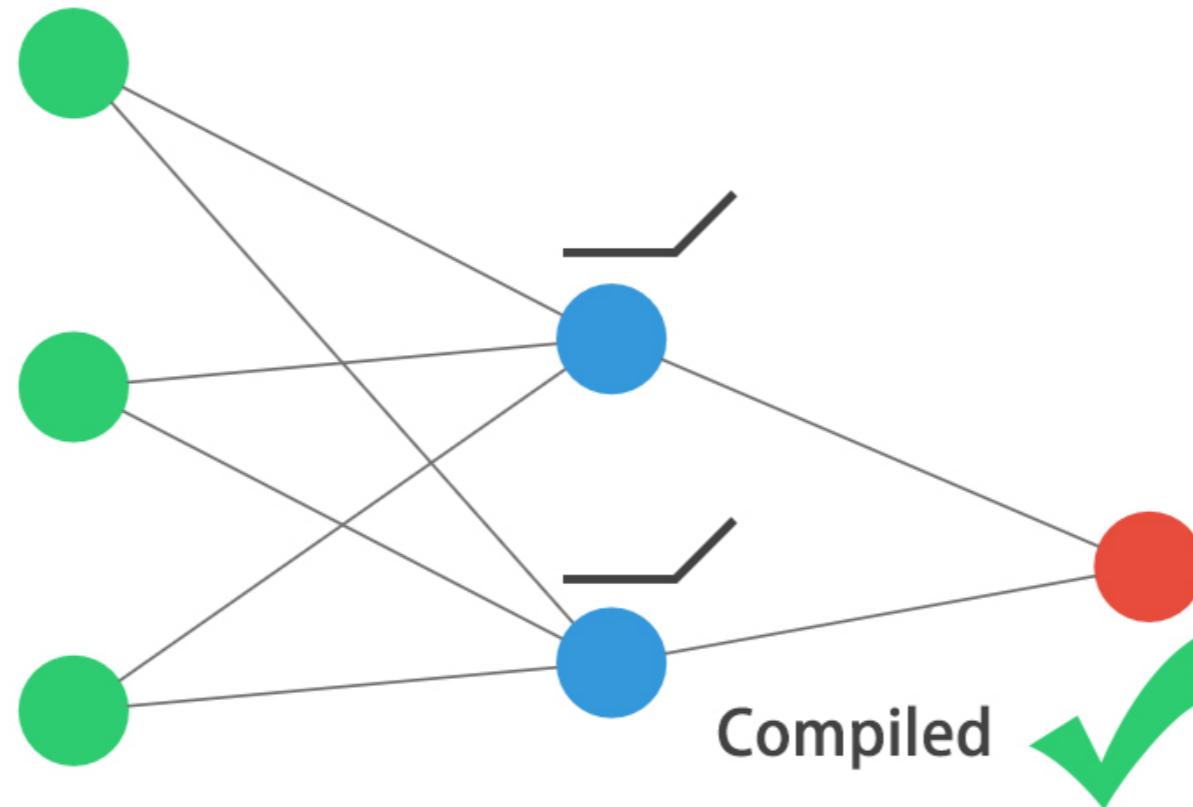
# Recap

```
from keras.models import Sequential  
from keras.layers import Dense  
  
# Create a new sequential model  
model = Sequential()  
  
# Add an input and dense layer  
model.add(Dense(2, input_shape=(3,),  
               activation="relu"))  
  
# Add a final 1 neuron layer  
model.add(Dense(1))  
<
```



# Compiling

```
# Compiling your previously built model  
model.compile(optimizer="adam", loss="mse")
```



# Training

```
# Train your model  
model.fit(X_train, y_train, epochs=5)
```

```
Epoch 1/5  
1000/1000 [=====] - 0s 242us/step - loss: 0.4090  
Epoch 2/5  
1000/1000 [=====] - 0s 34us/step - loss: 0.3602  
Epoch 3/5  
1000/1000 [=====] - 0s 37us/step - loss: 0.3223  
Epoch 4/5  
1000/1000 [=====] - 0s 34us/step - loss: 0.2958  
Epoch 5/5  
1000/1000 [=====] - 0s 33us/step - loss: 0.2795
```

# Predicting

```
# Predict on new data  
preds = model.predict(X_test)  
  
# Look at the predictions  
print(preds)
```

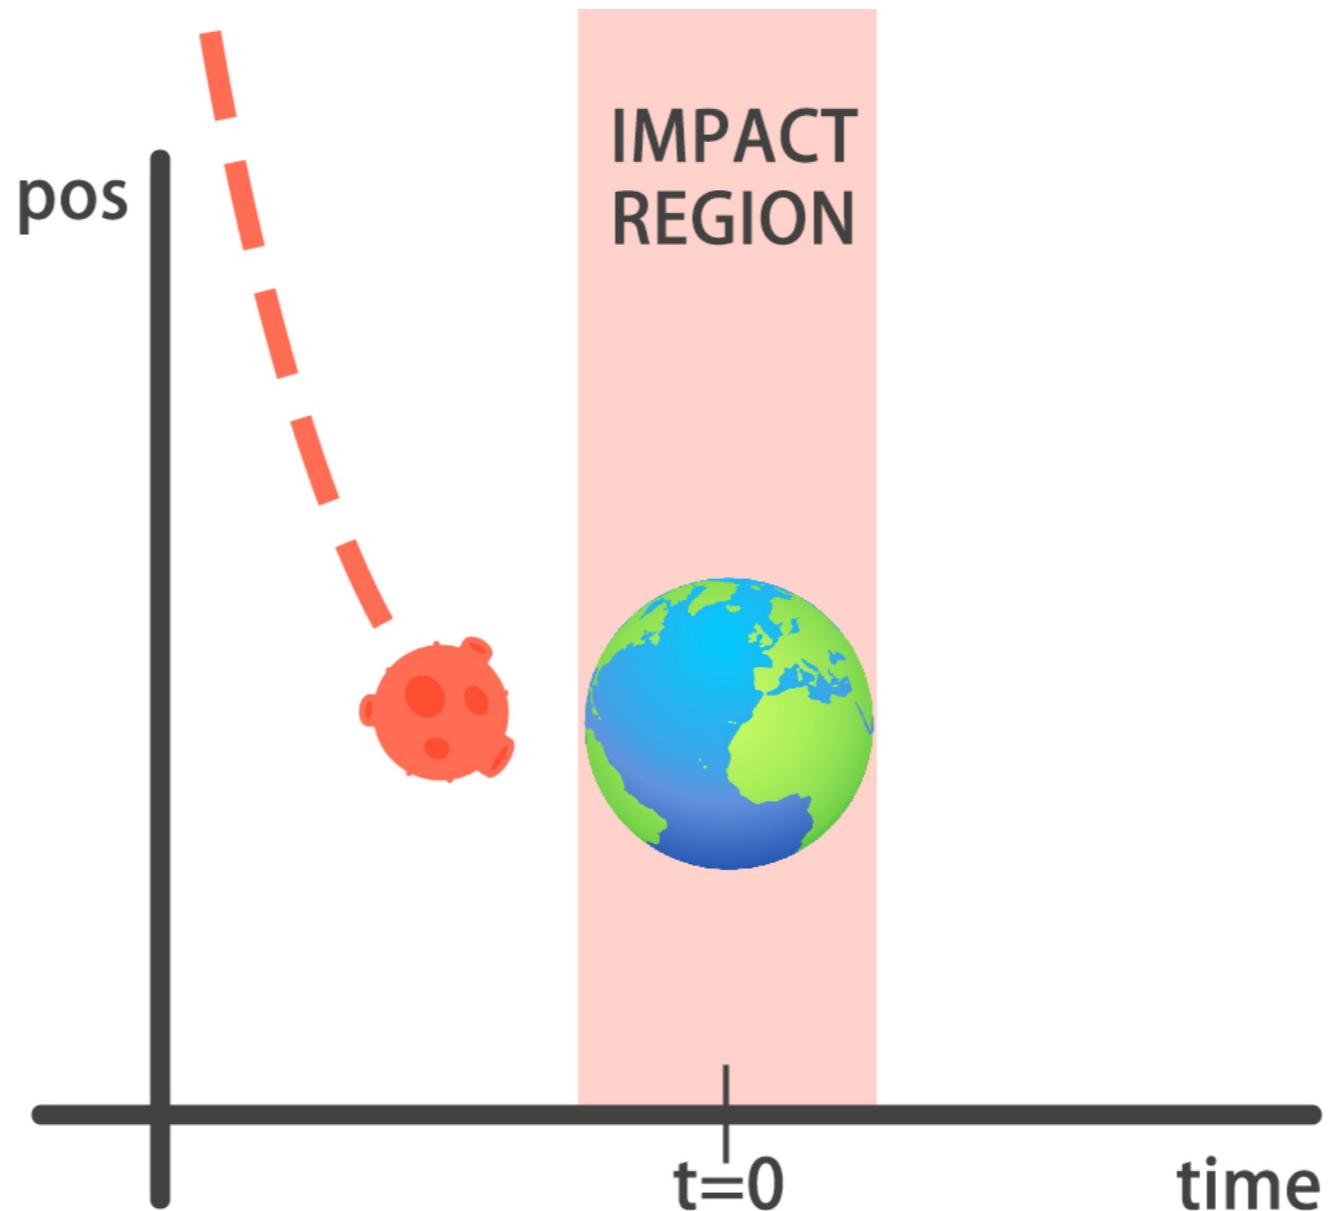
```
array([[0.6131608 ],  
       [0.5175948 ],  
       [0.60209155],  
       ... ,  
       [0.55633   ],  
       [0.5305591 ],  
       [0.50682044]])
```

# Evaluating

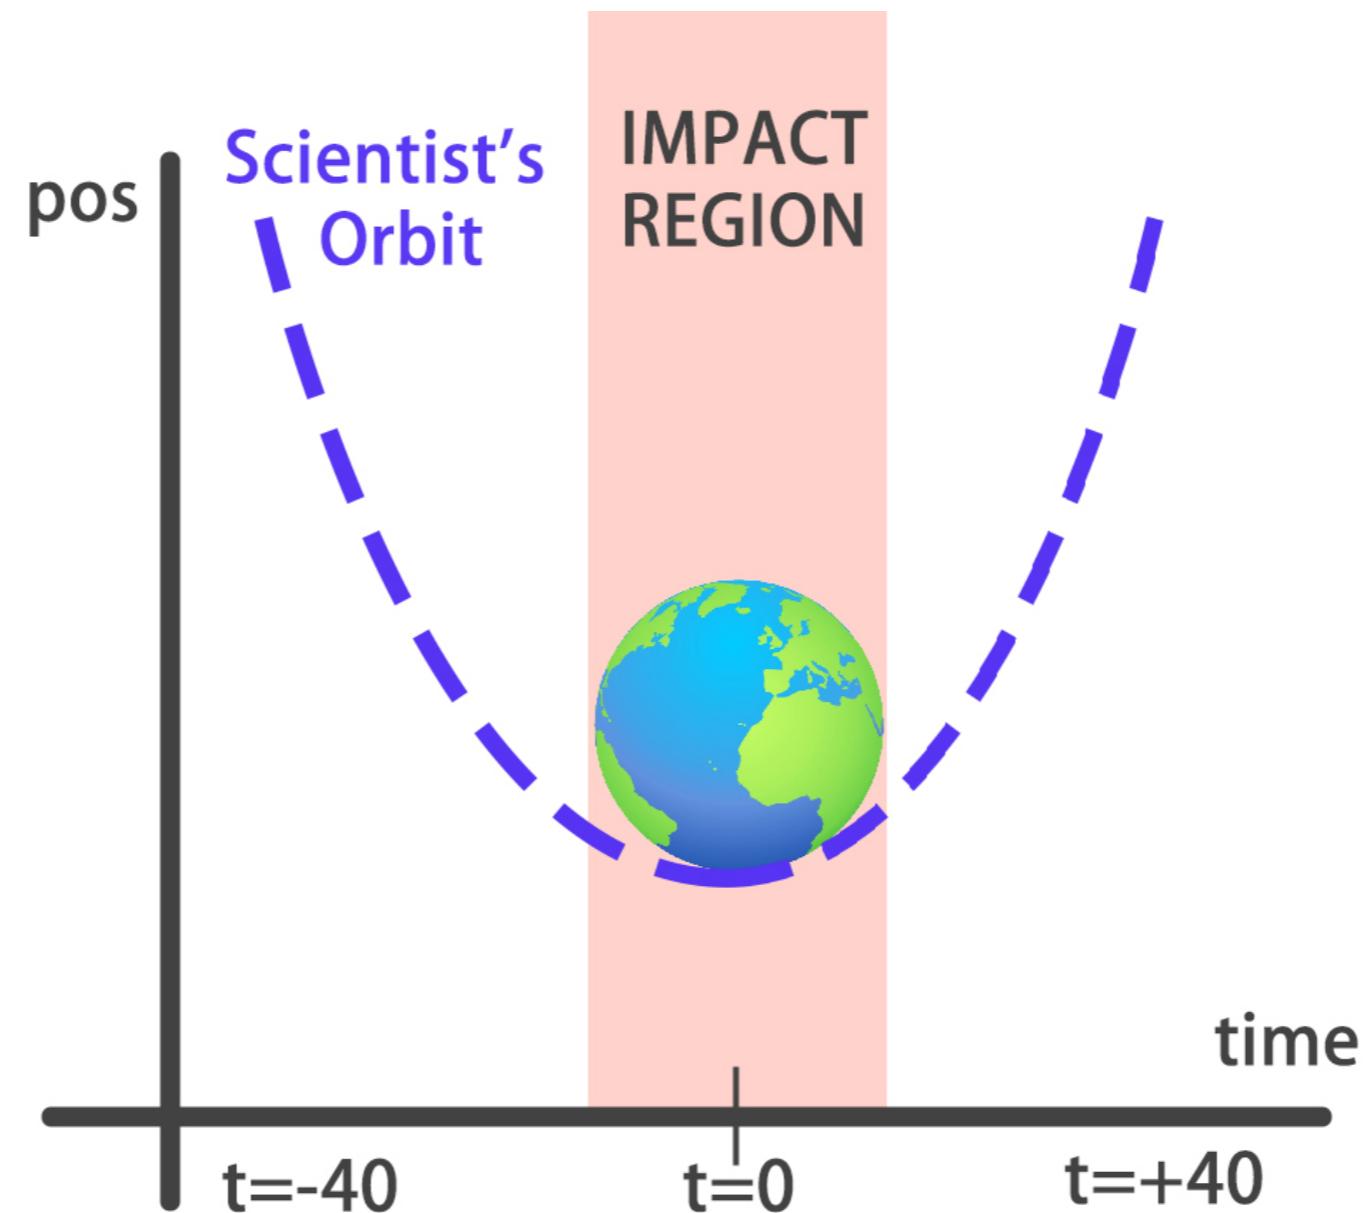
```
# Evaluate your results  
model.evaluate(X_test, y_test)
```

```
1000/1000 [=====] - 0s 53us/step  
0.25
```

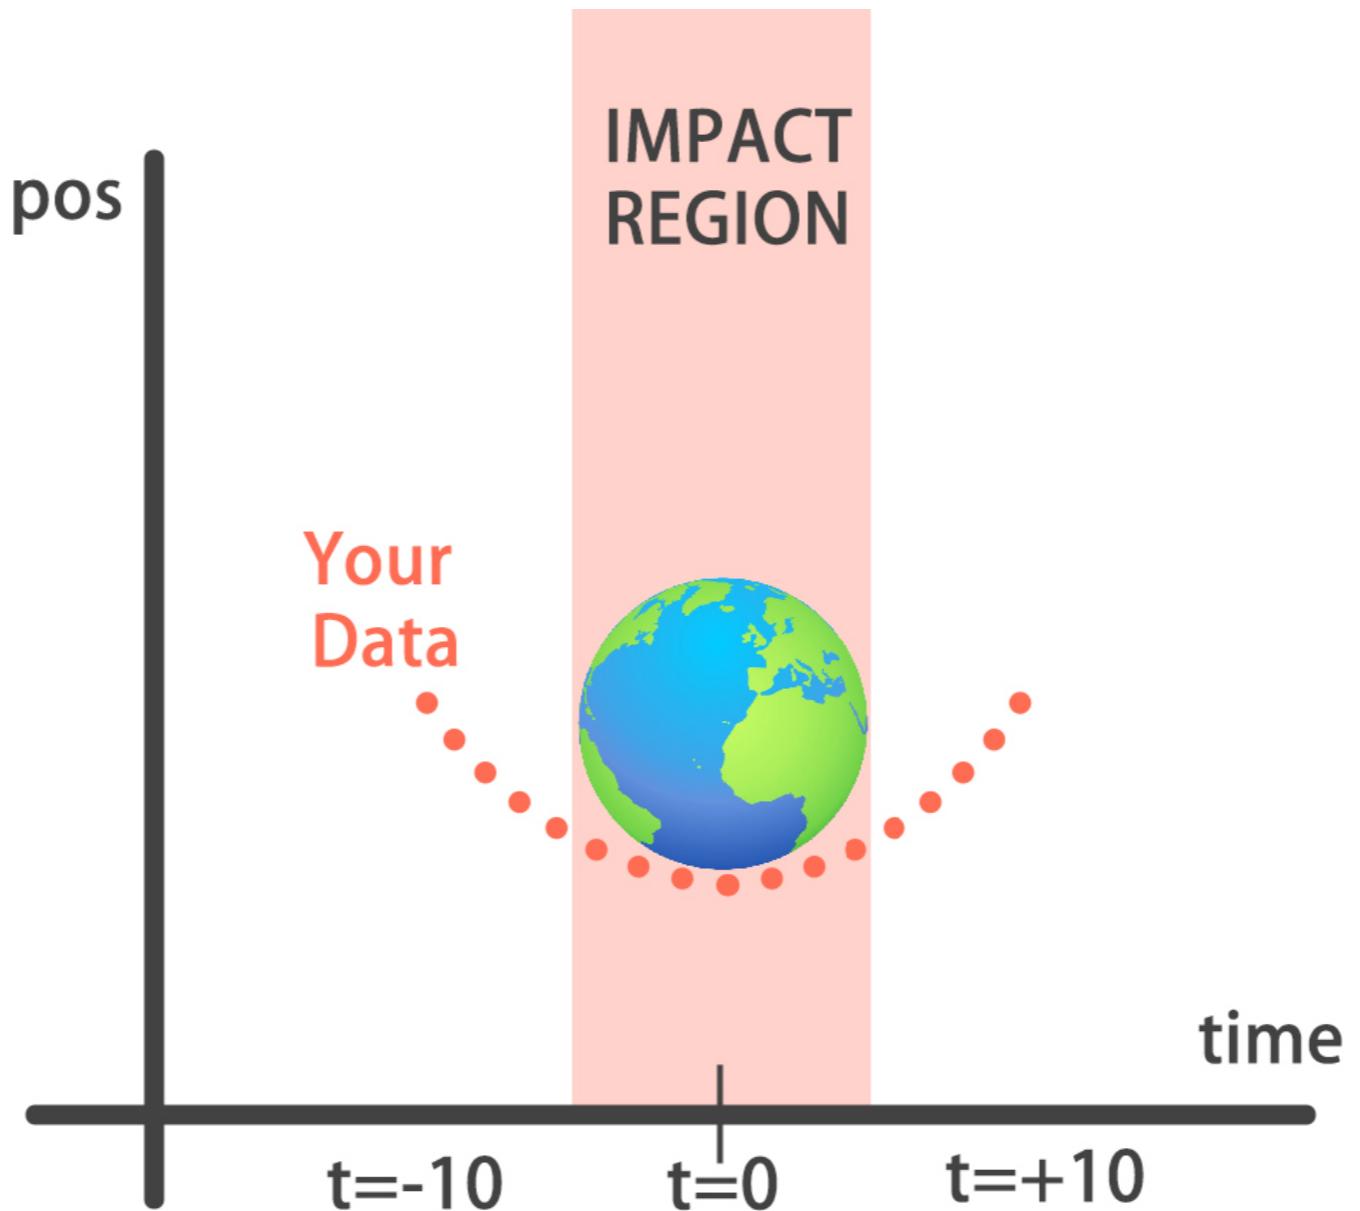
# The problem at hand



# Scientific prediction



# Your task

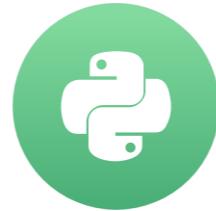


# Let's save the earth!

INTRODUCTION TO DEEP LEARNING WITH KERAS

# Binary classification

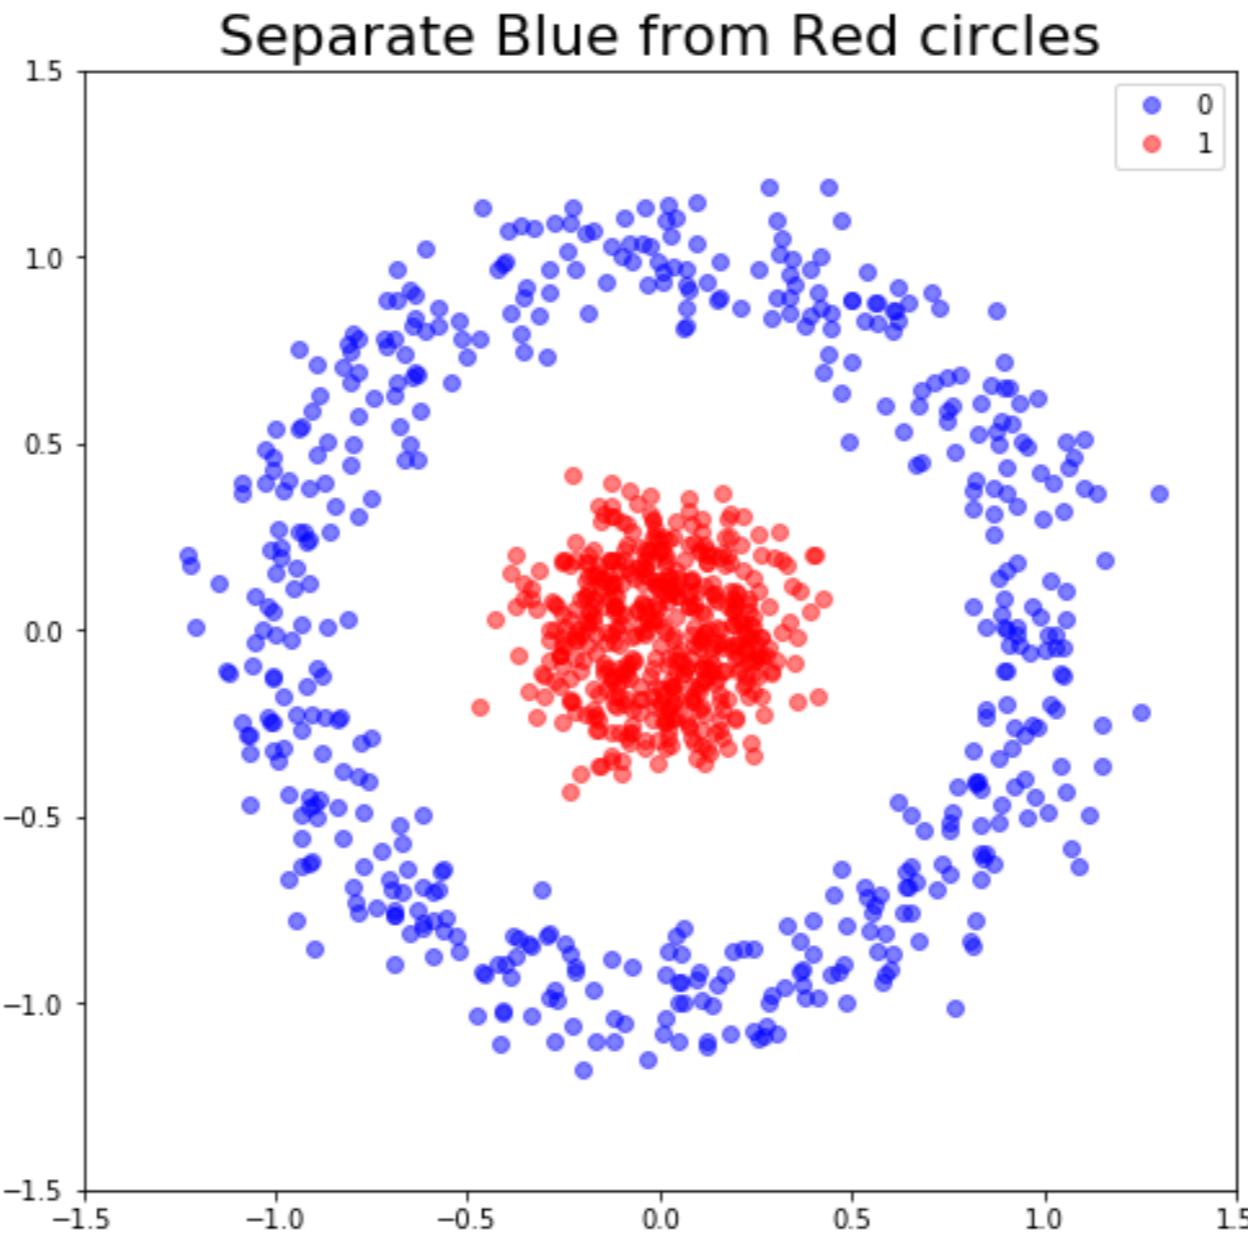
INTRODUCTION TO DEEP LEARNING WITH KERAS



Miguel Esteban

Data Scientist & Founder

# When to use binary classification?



# Our dataset

coordinates

[0.242, 0.038]

[0.044, -0.057]

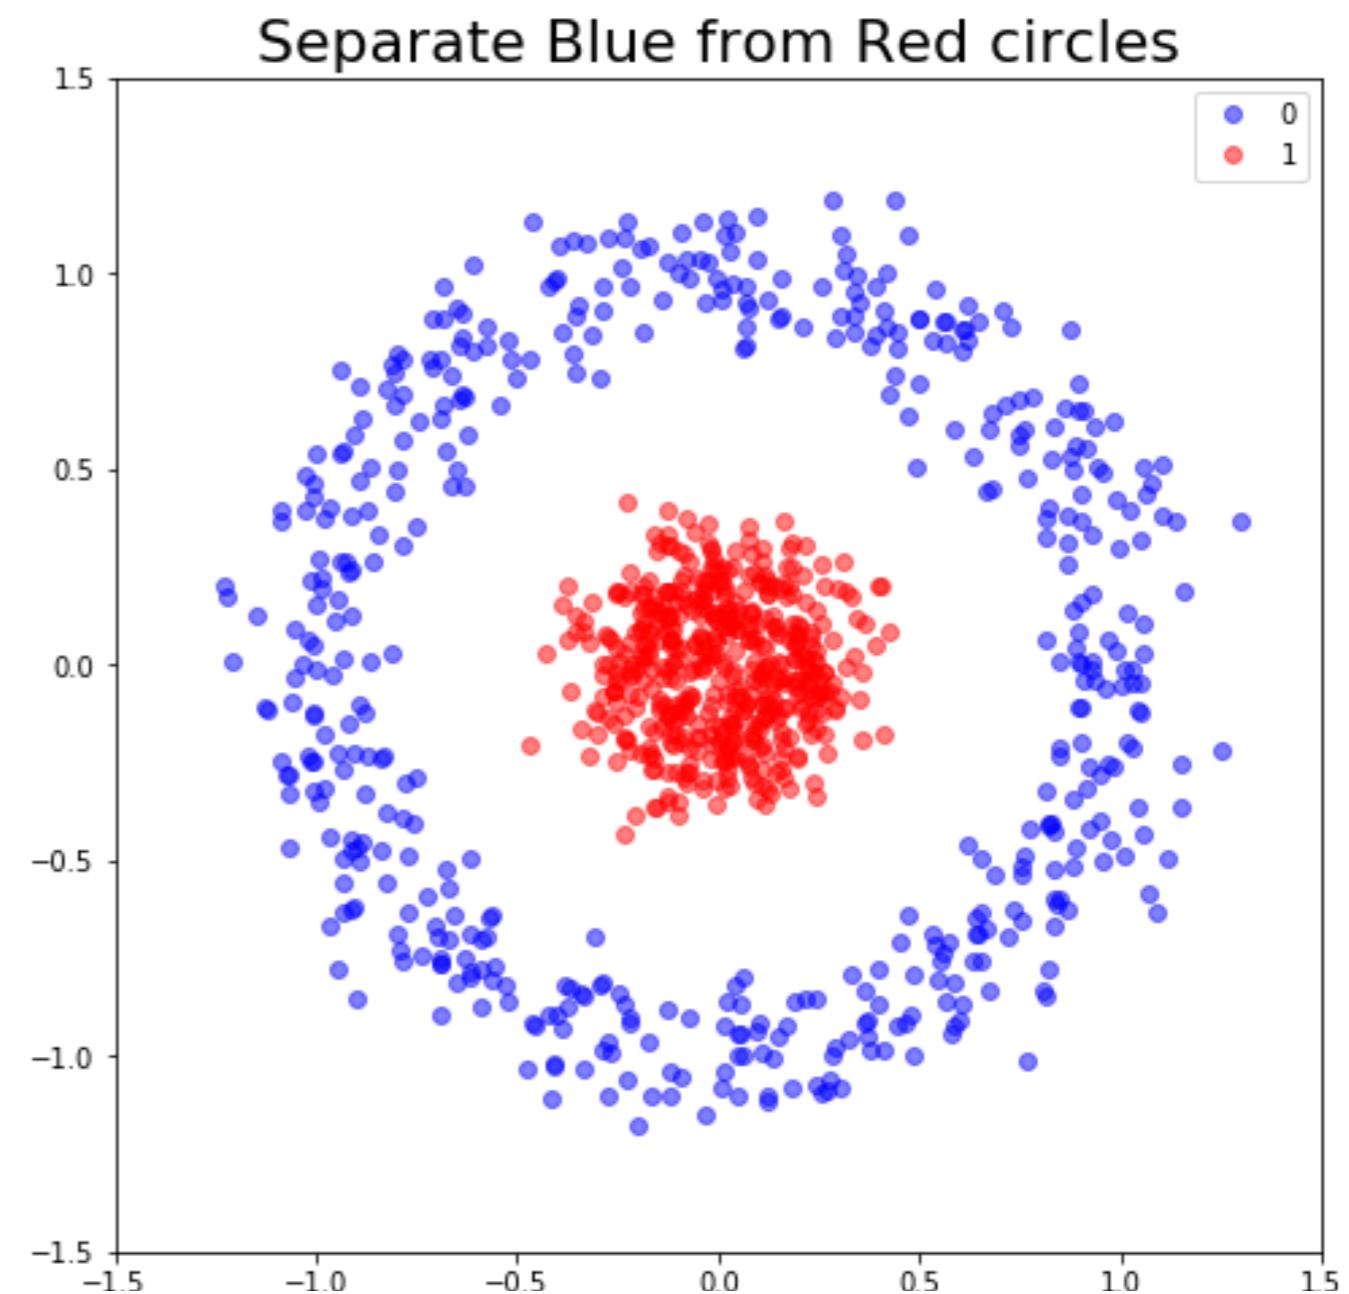
[-0.787, -0.076]

labels

1

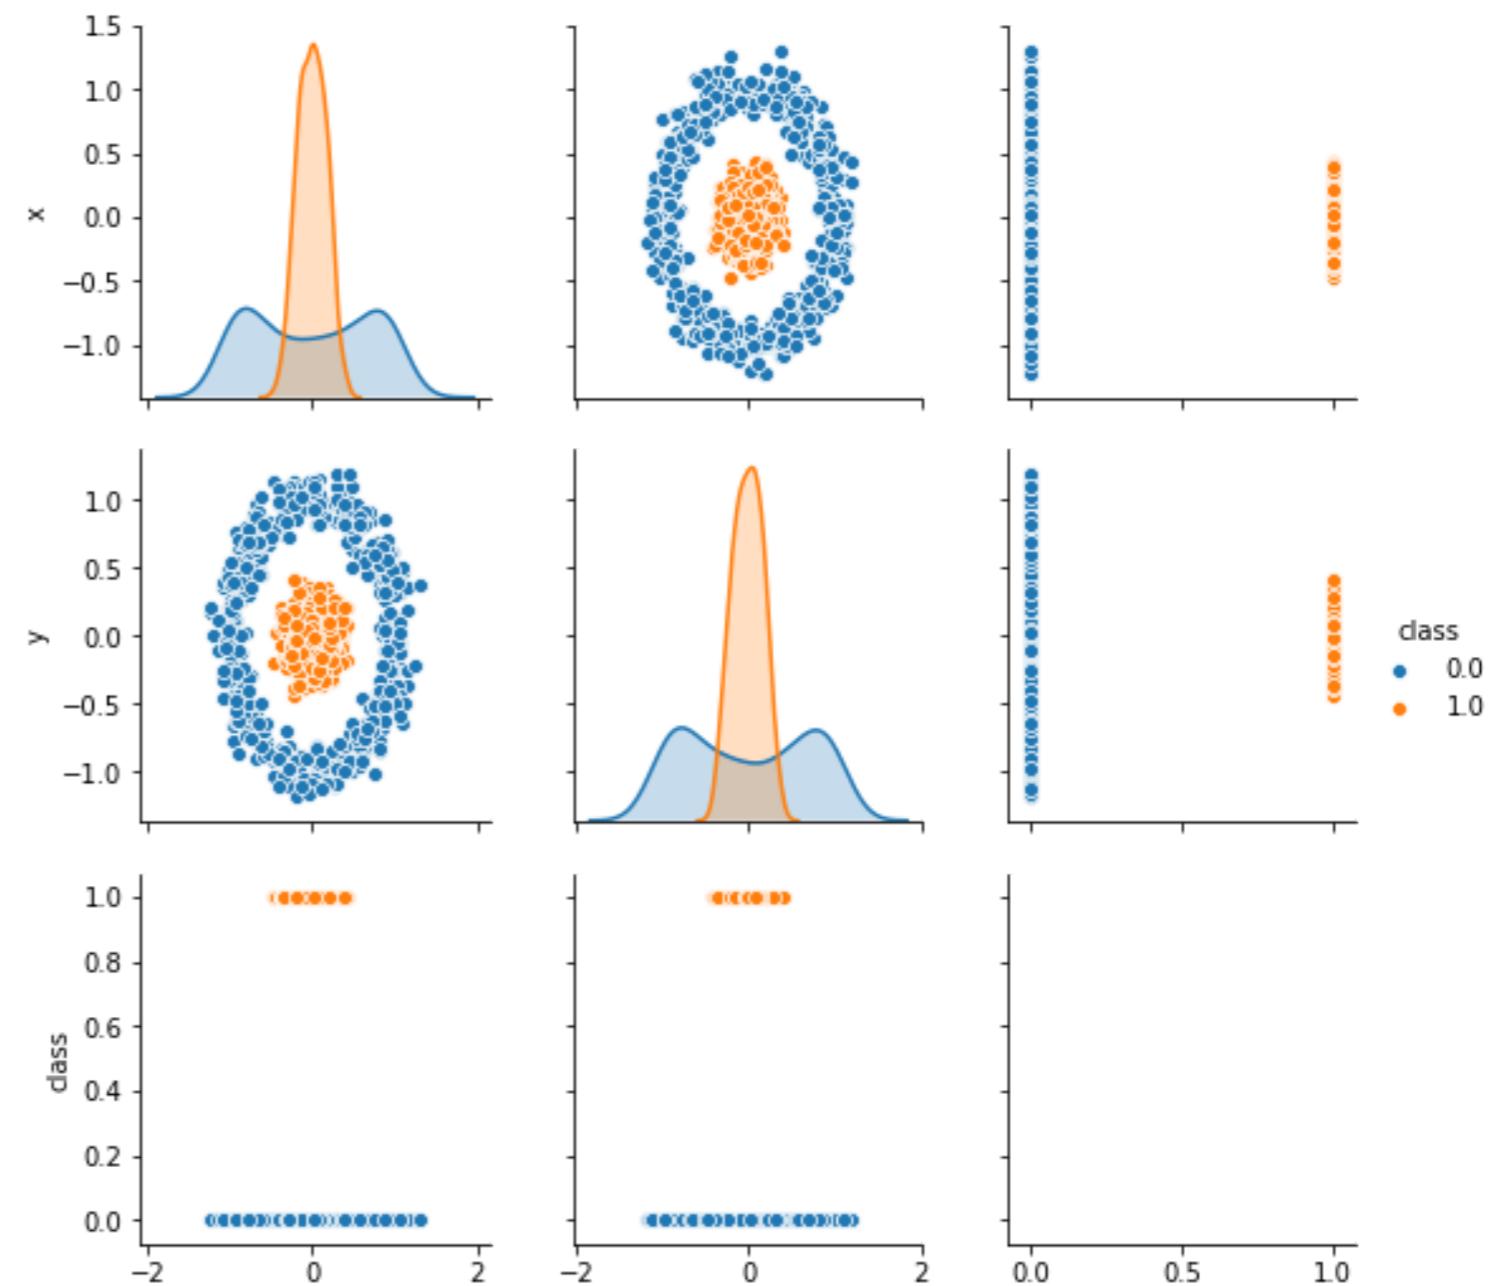
1

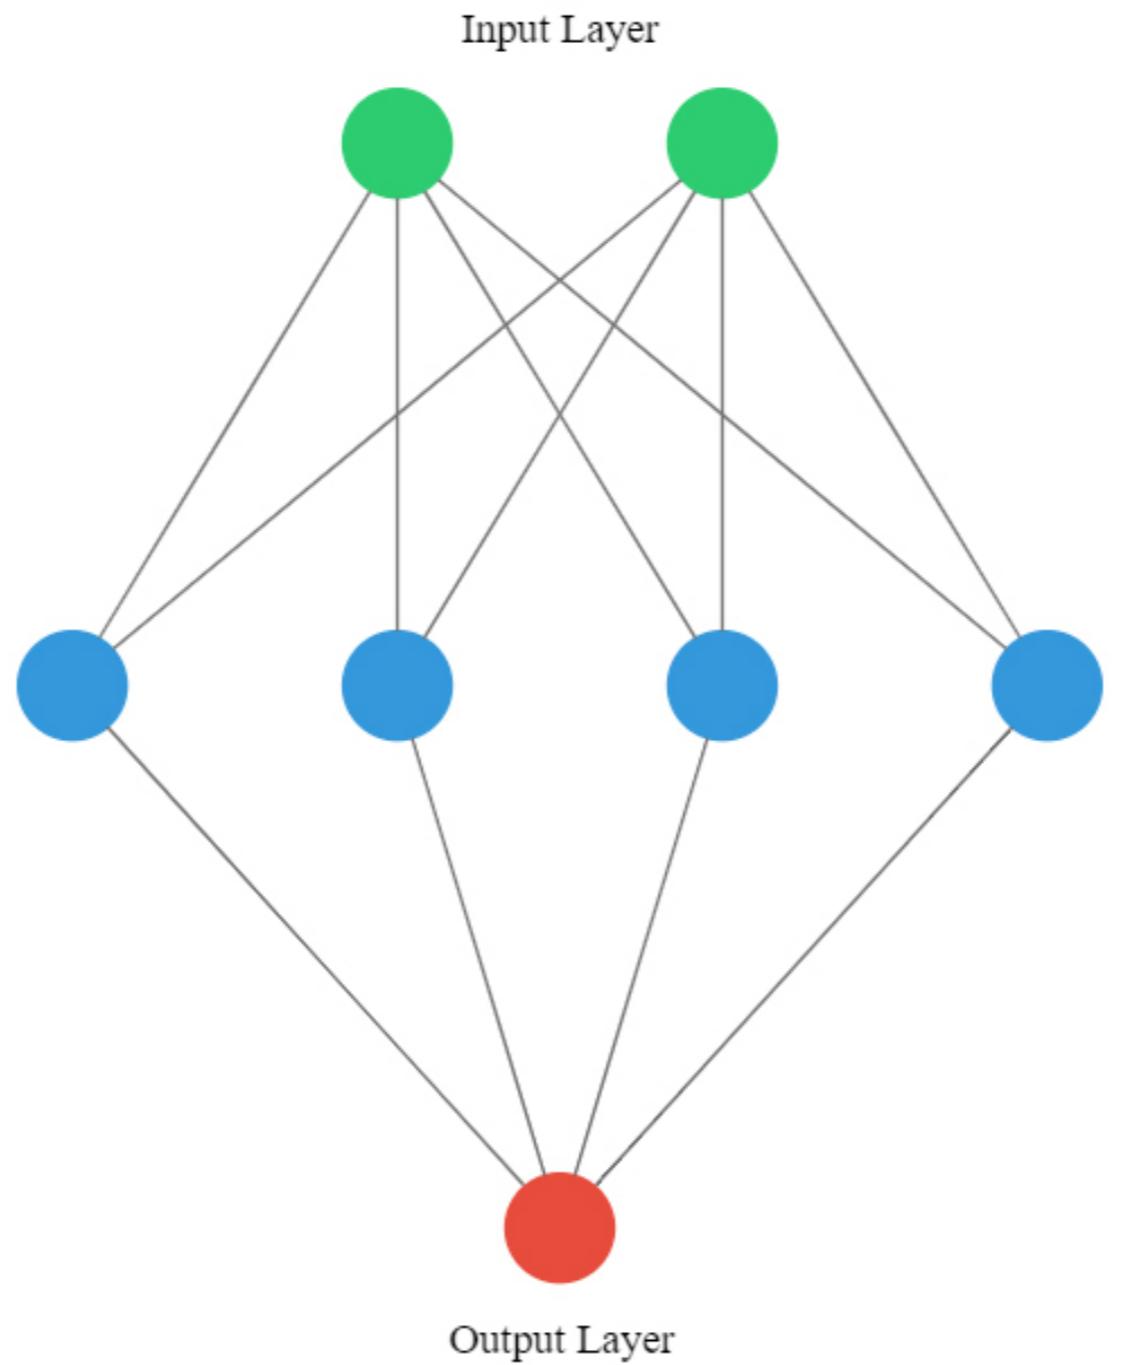
0

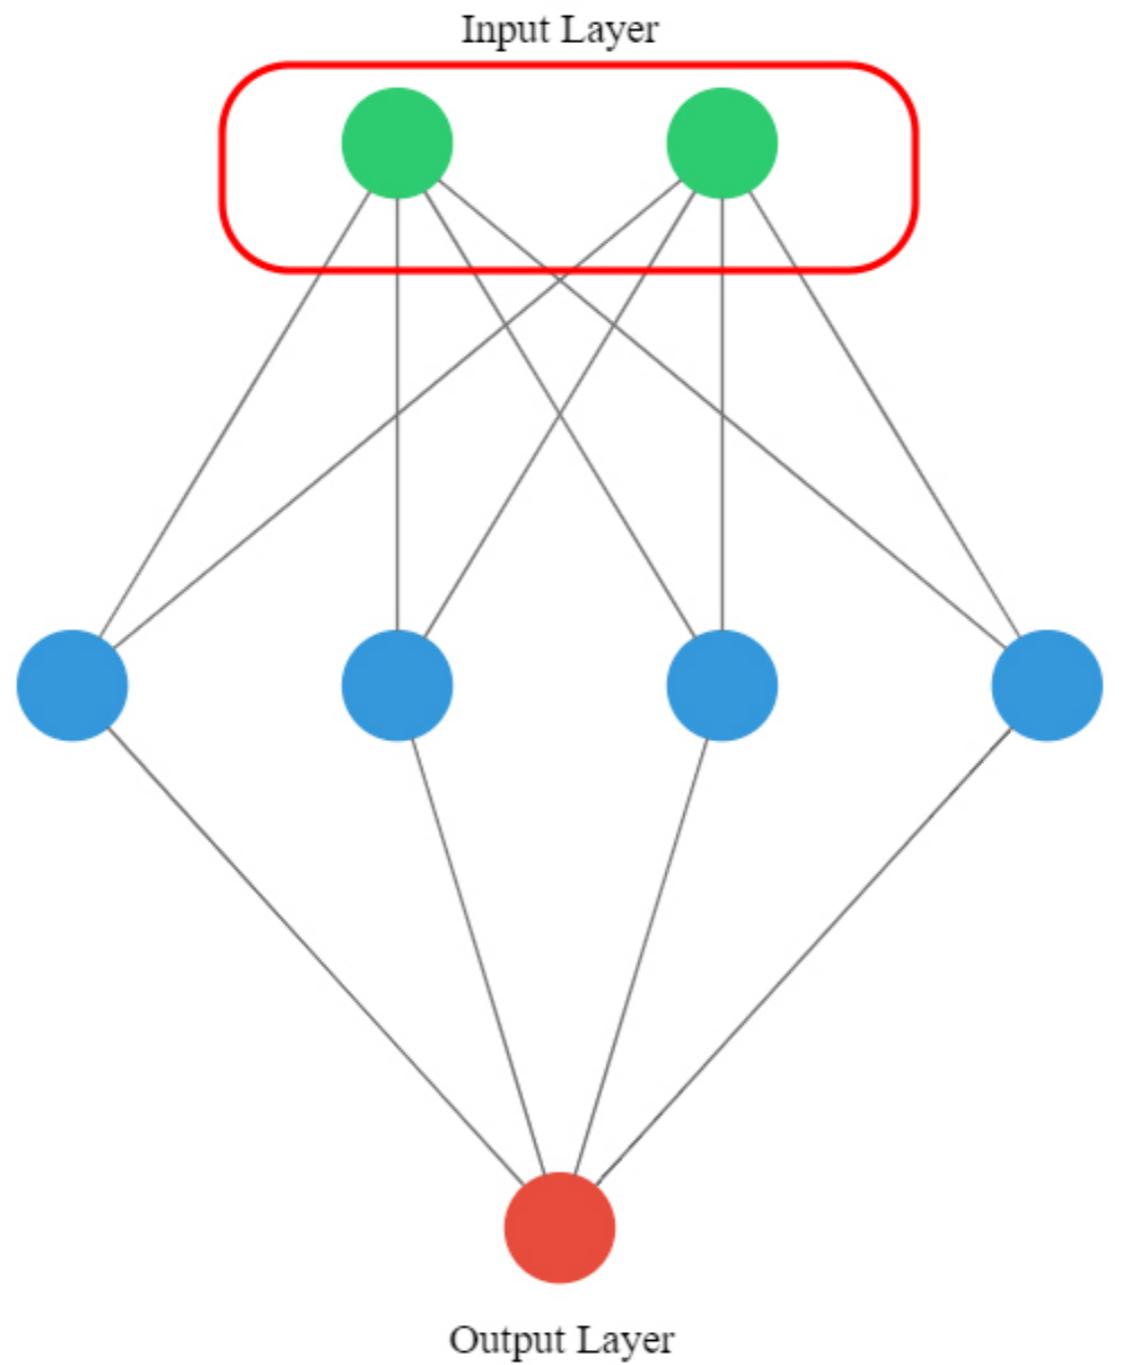


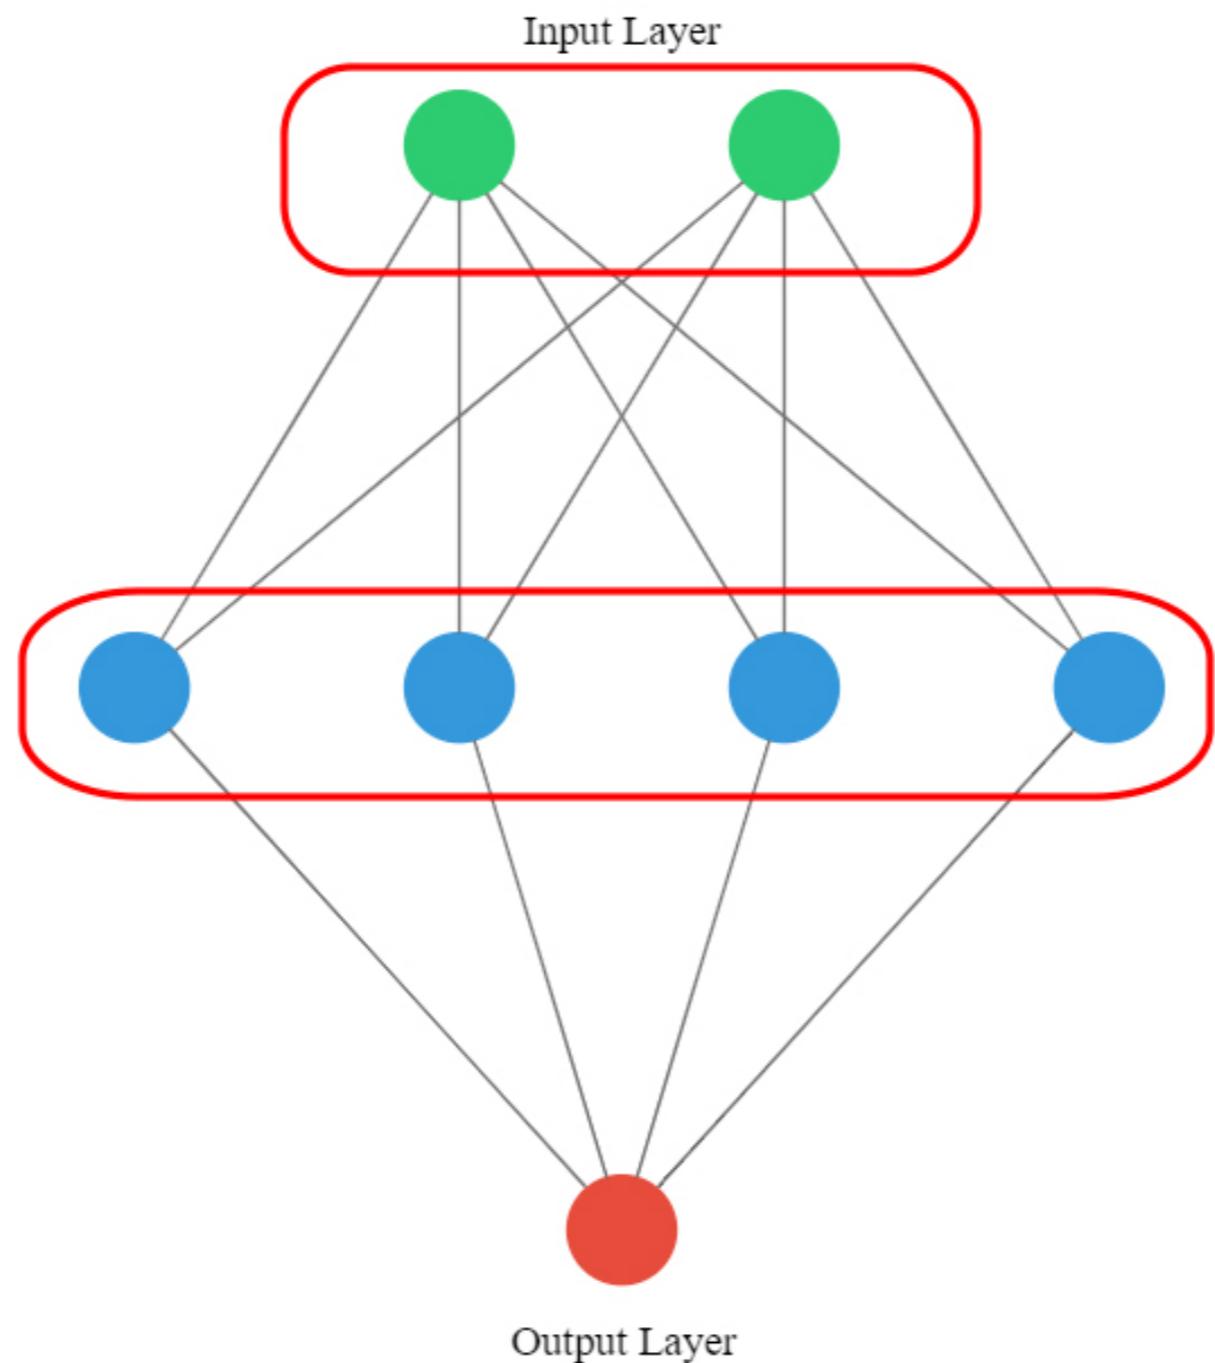
# Pairplots

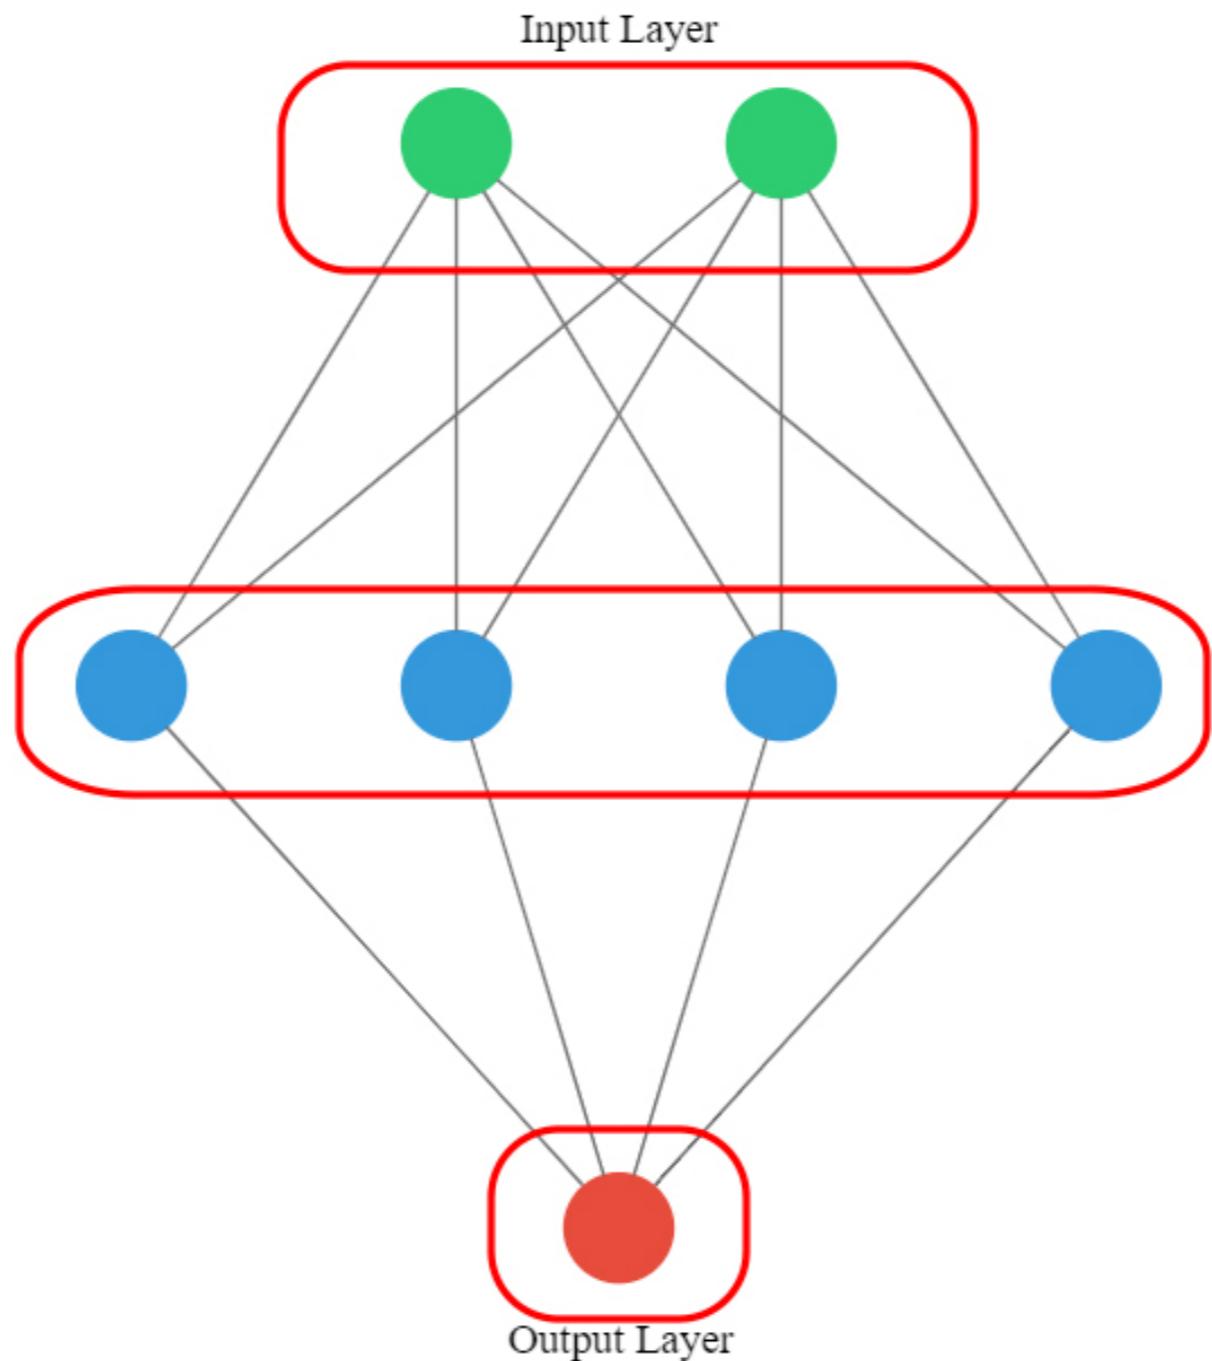
```
import seaborn as sns  
  
# Plot a pairplot  
sns.pairplot(circles, hue="target")
```

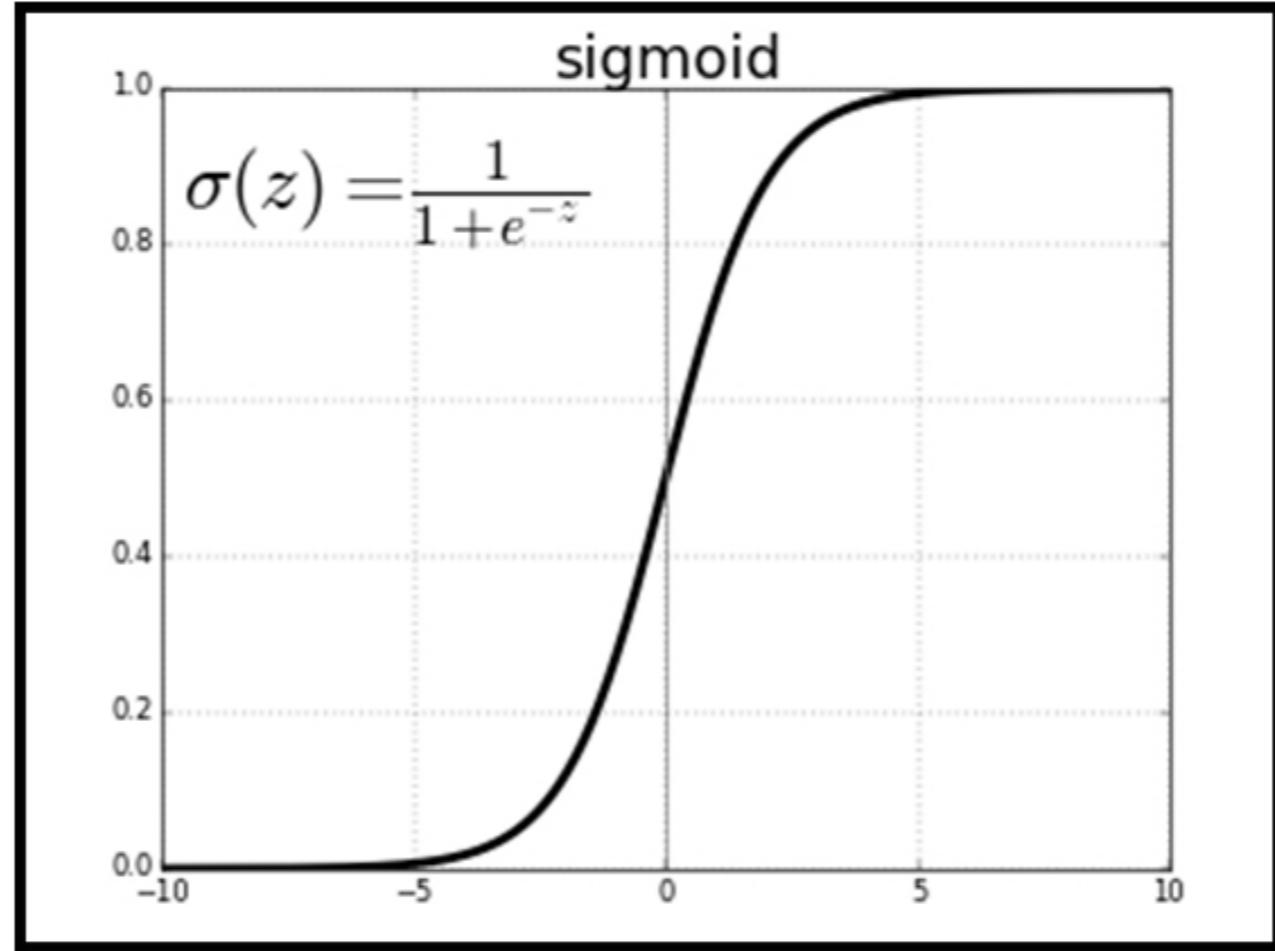












Output Layer

# The sigmoid function

neuron  
output

3 →



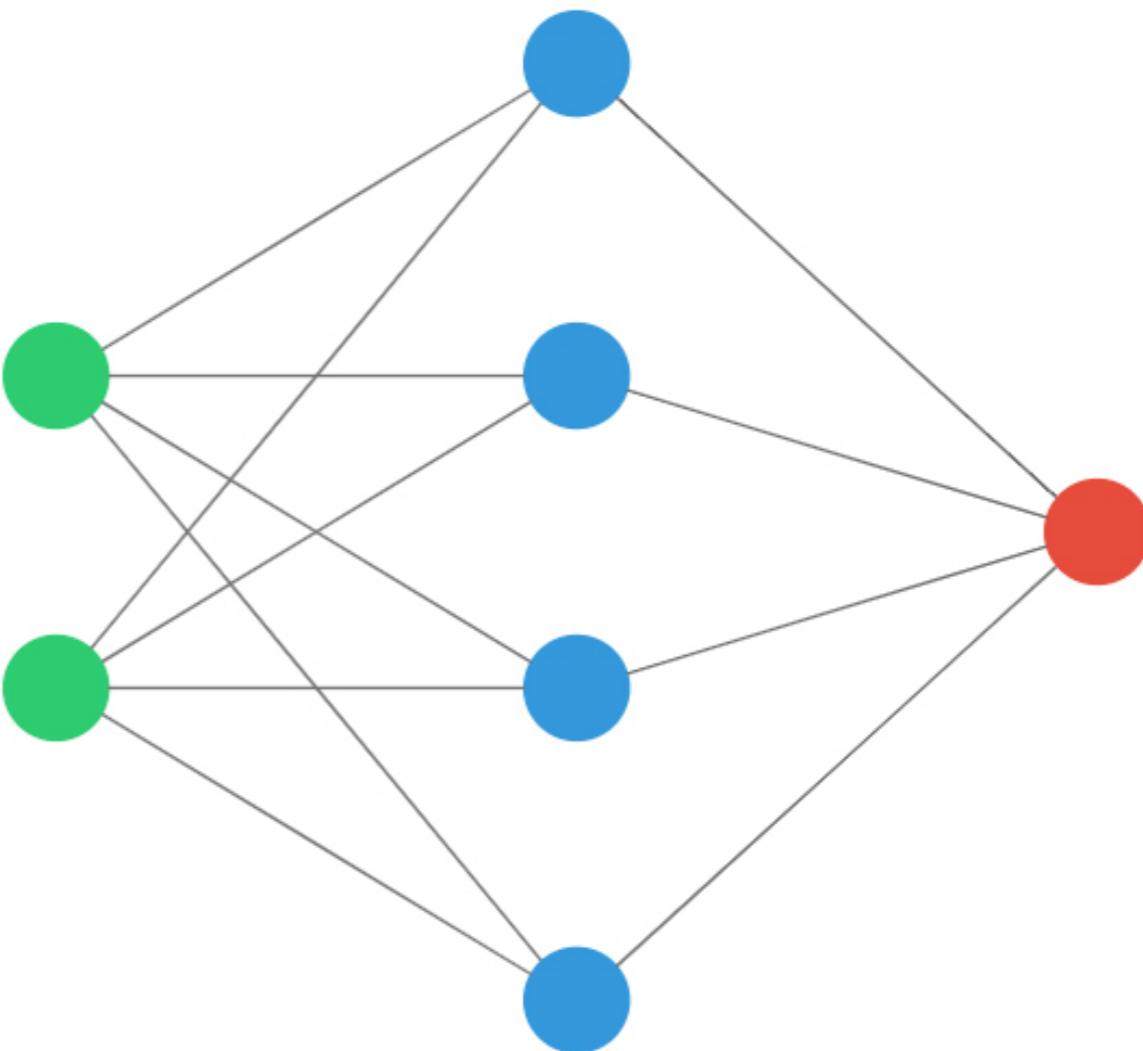
transformed  
output

0.95 → 1

rounded  
output

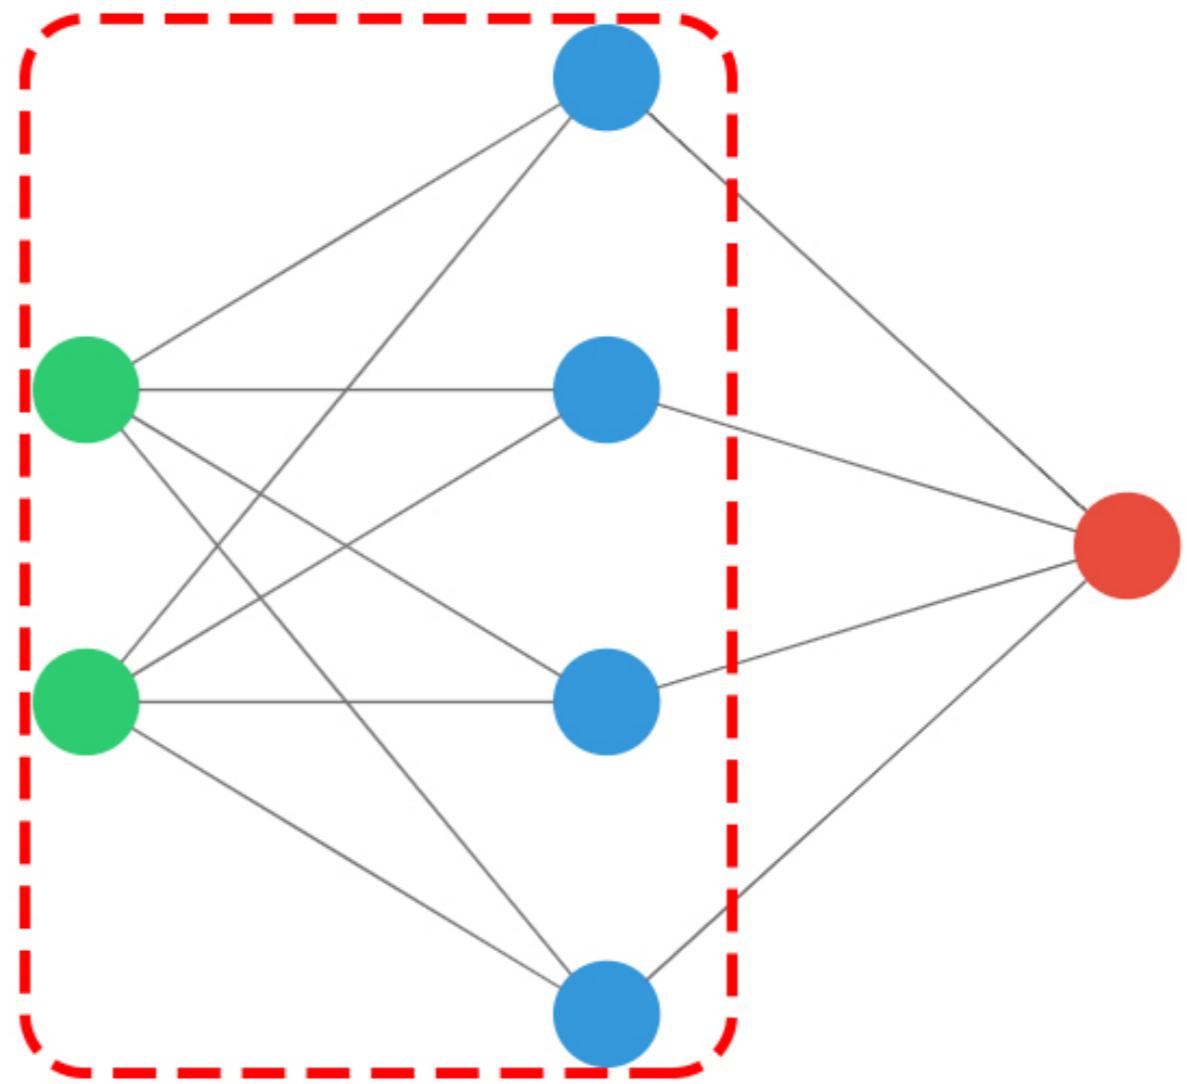
# Let's build it

```
from keras.models import Sequential  
from keras.layers import Dense  
  
# Instantiate a sequential model  
model = Sequential()
```



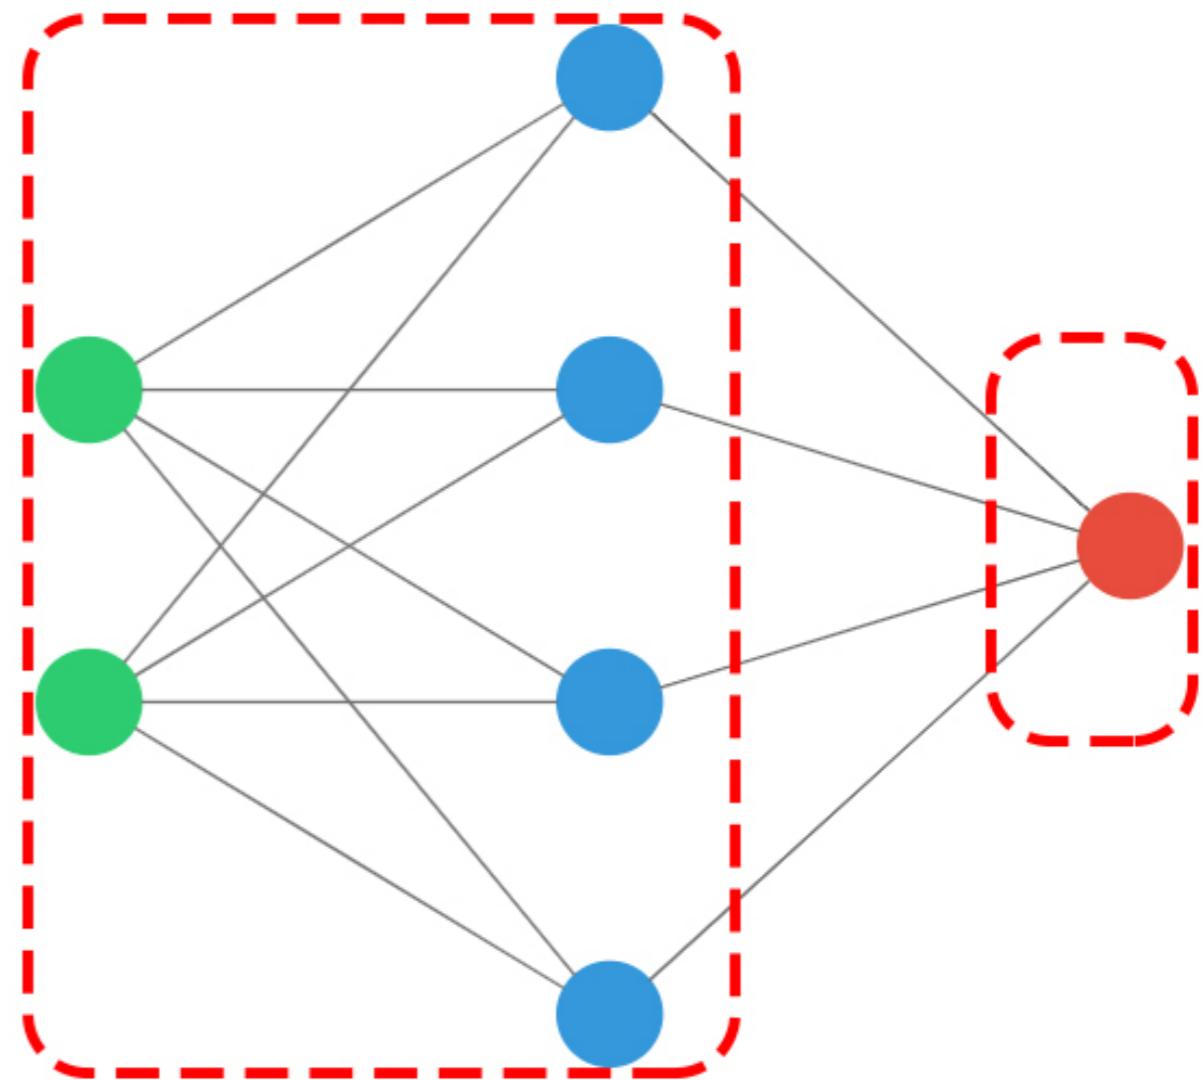
# Let's build it

```
from keras.models import Sequential  
from keras.layers import Dense  
  
# Instantiate a sequential model  
model = Sequential()  
  
# Add input and hidden layer  
model.add(Dense(4, input_shape=(2,),  
               activation='tanh'))
```



# Let's build it

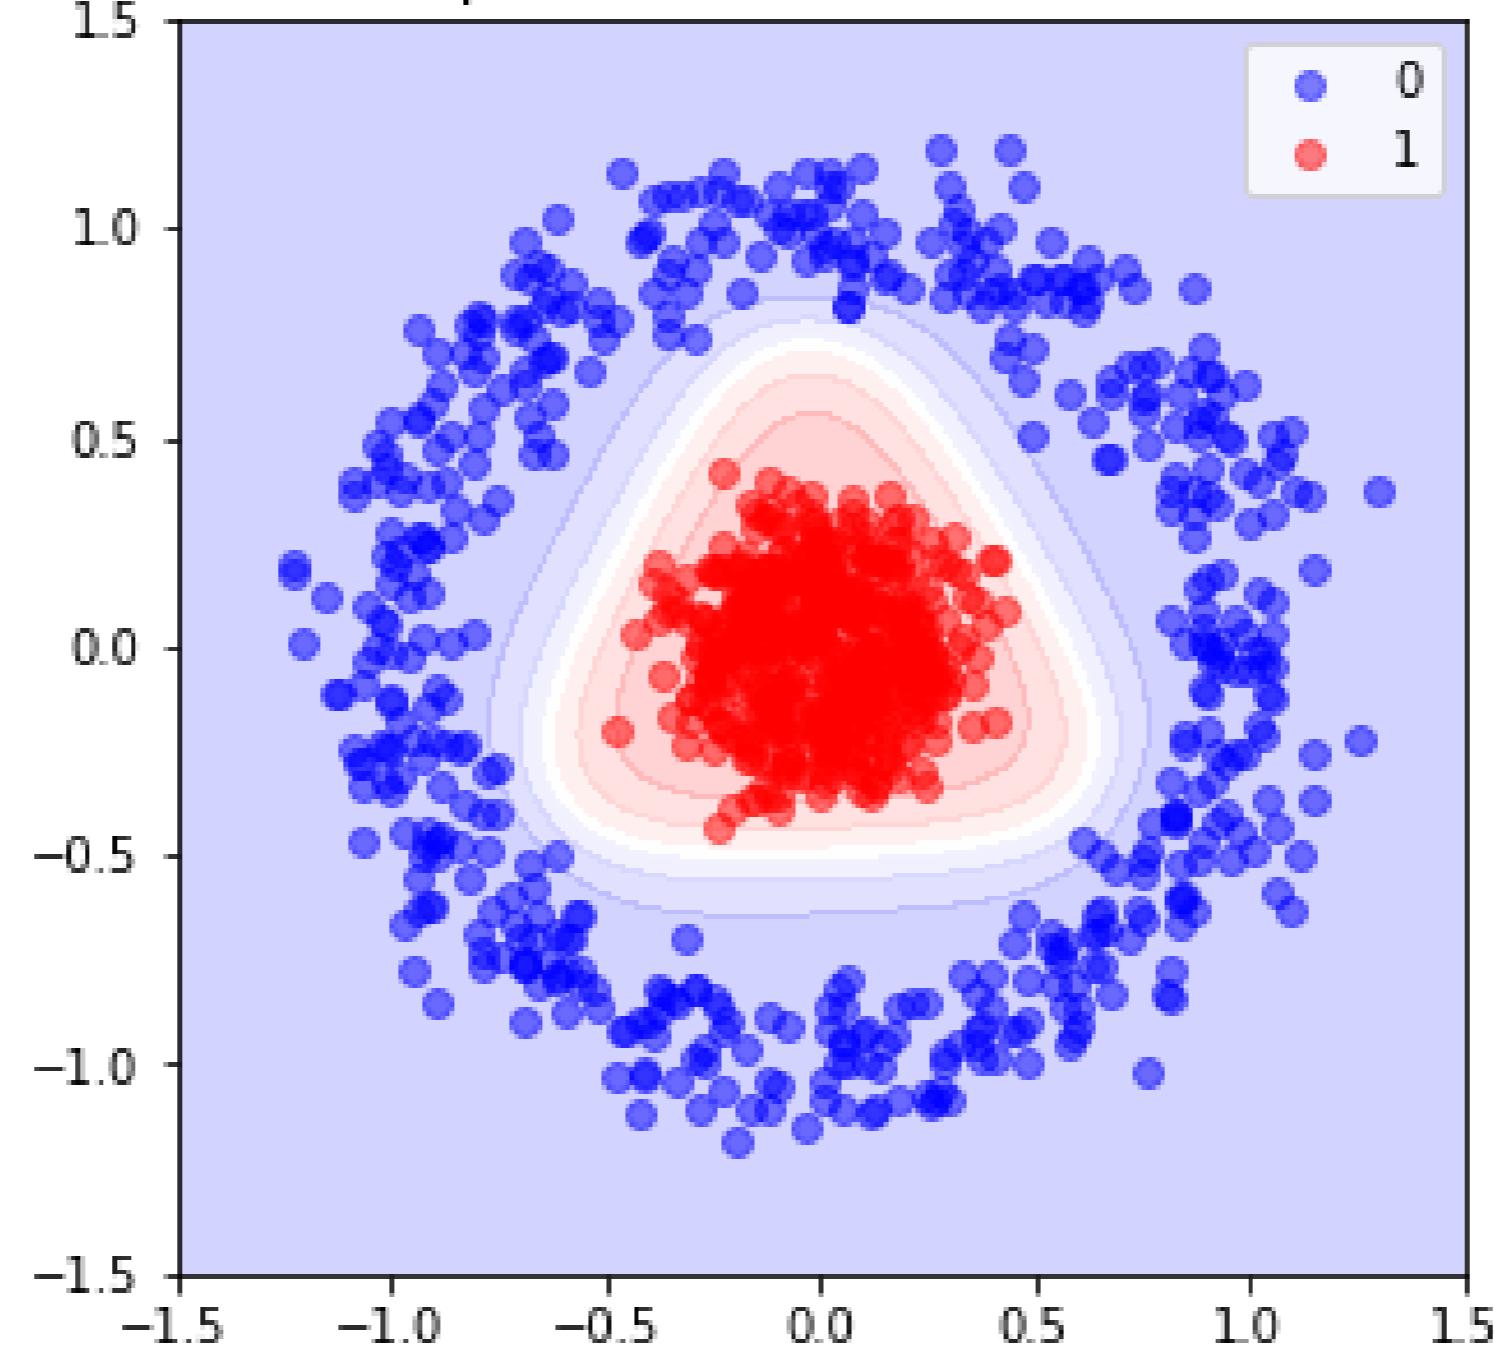
```
from keras.models import Sequential  
from keras.layers import Dense  
  
# Instantiate a sequential model  
model = Sequential()  
  
# Add input and hidden layer  
model.add(Dense(4, input_shape=(2,),  
               activation='tanh'))  
  
# Add output layer, use sigmoid  
model.add(Dense(1,  
               activation='sigmoid'))
```



# Compiling, training, predicting

```
# Compile model  
model.compile(optimizer='sgd',  
              loss='binary_crossentropy')  
  
# Train model  
model.train(coordinates, labels, epochs=20)  
  
# Predict with trained model  
preds = model.predict(coordinates)
```

## Separate Blue from Red circles

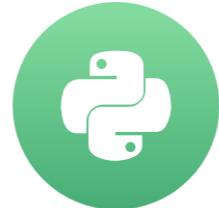


# Let's practice!

INTRODUCTION TO DEEP LEARNING WITH KERAS

# Multi-class classification

INTRODUCTION TO DEEP LEARNING WITH KERAS



**Miguel Esteban**  
Data Scientist & Founder

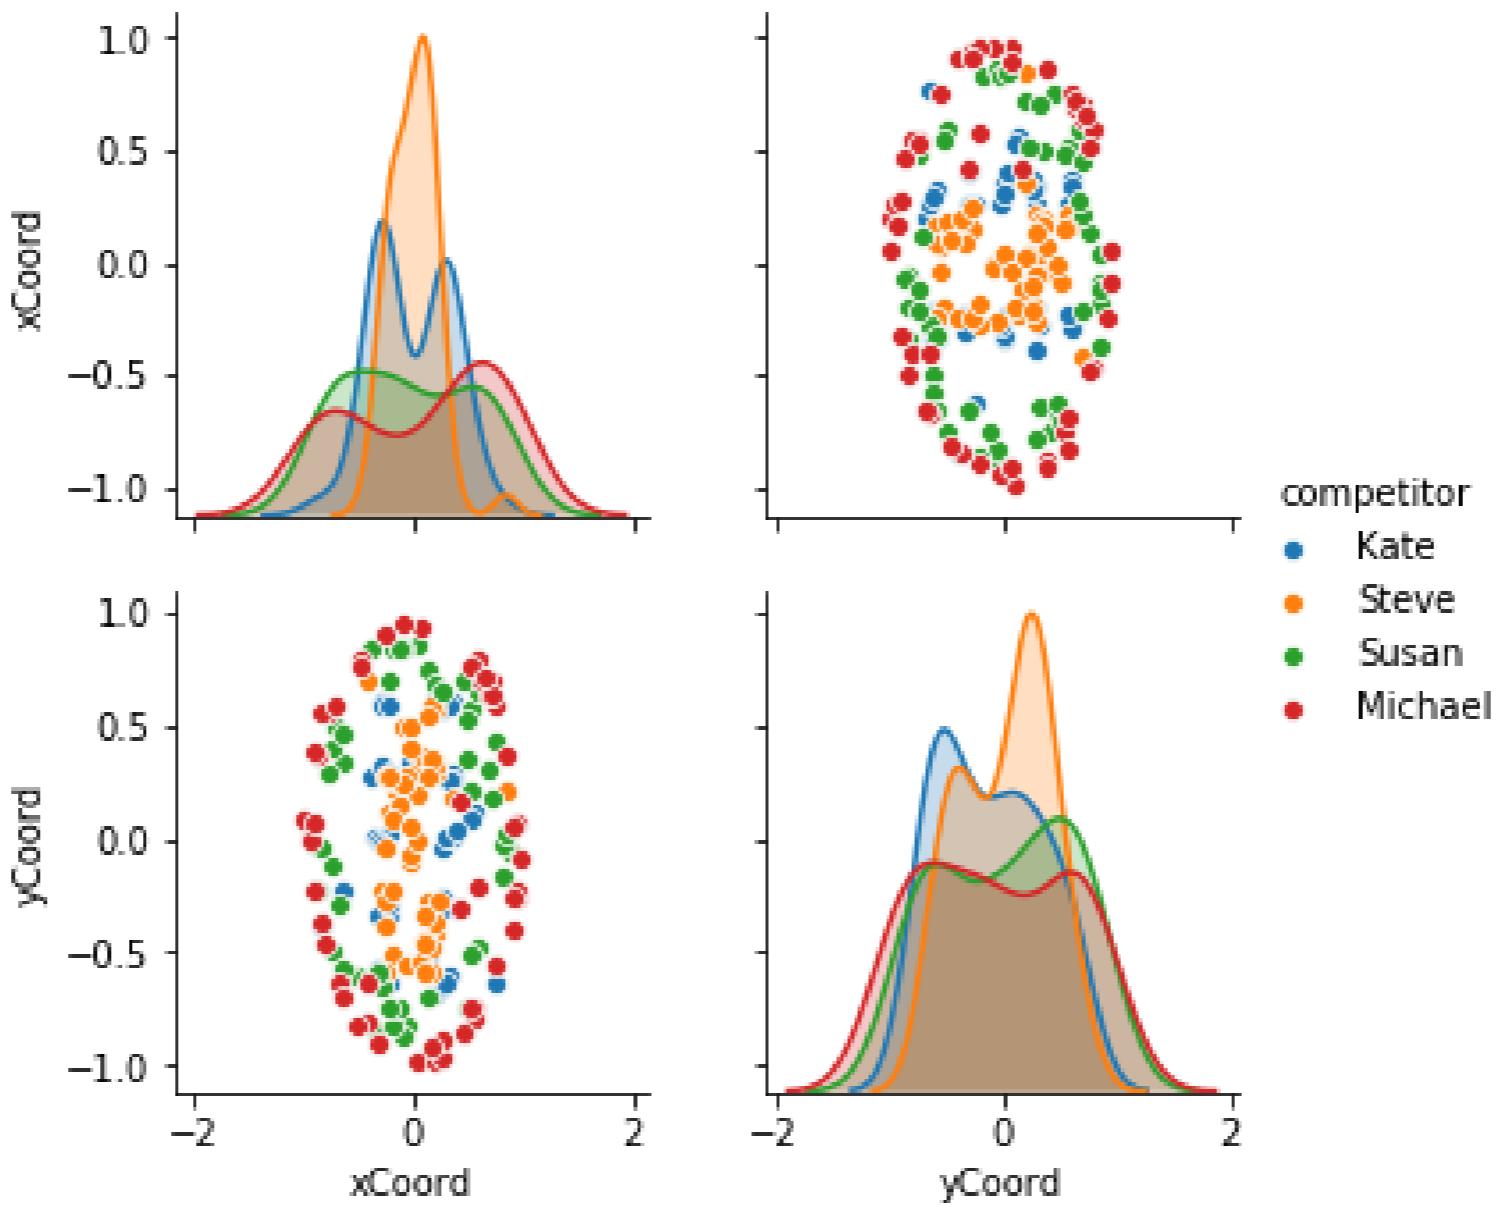
# Throwing darts



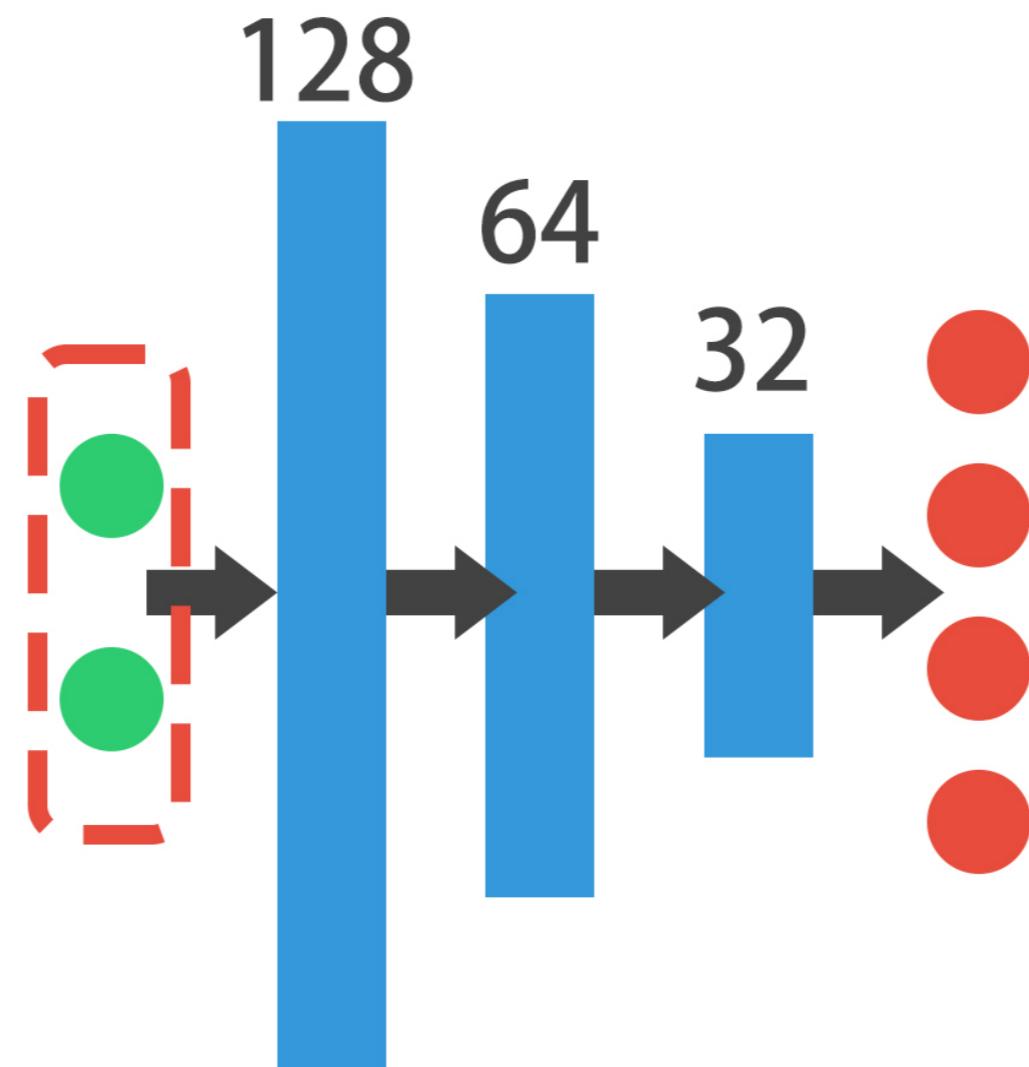
# The dataset

<b>xCoord</b>	<b>yCoord</b>	<b>competitor</b>
-0.037673	0.057402	Steve
-0.331021	-0.585035	Susan
-0.123567	0.839730	Susan
-0.086160	0.959787	Michael
-0.902632	0.078753	Michael

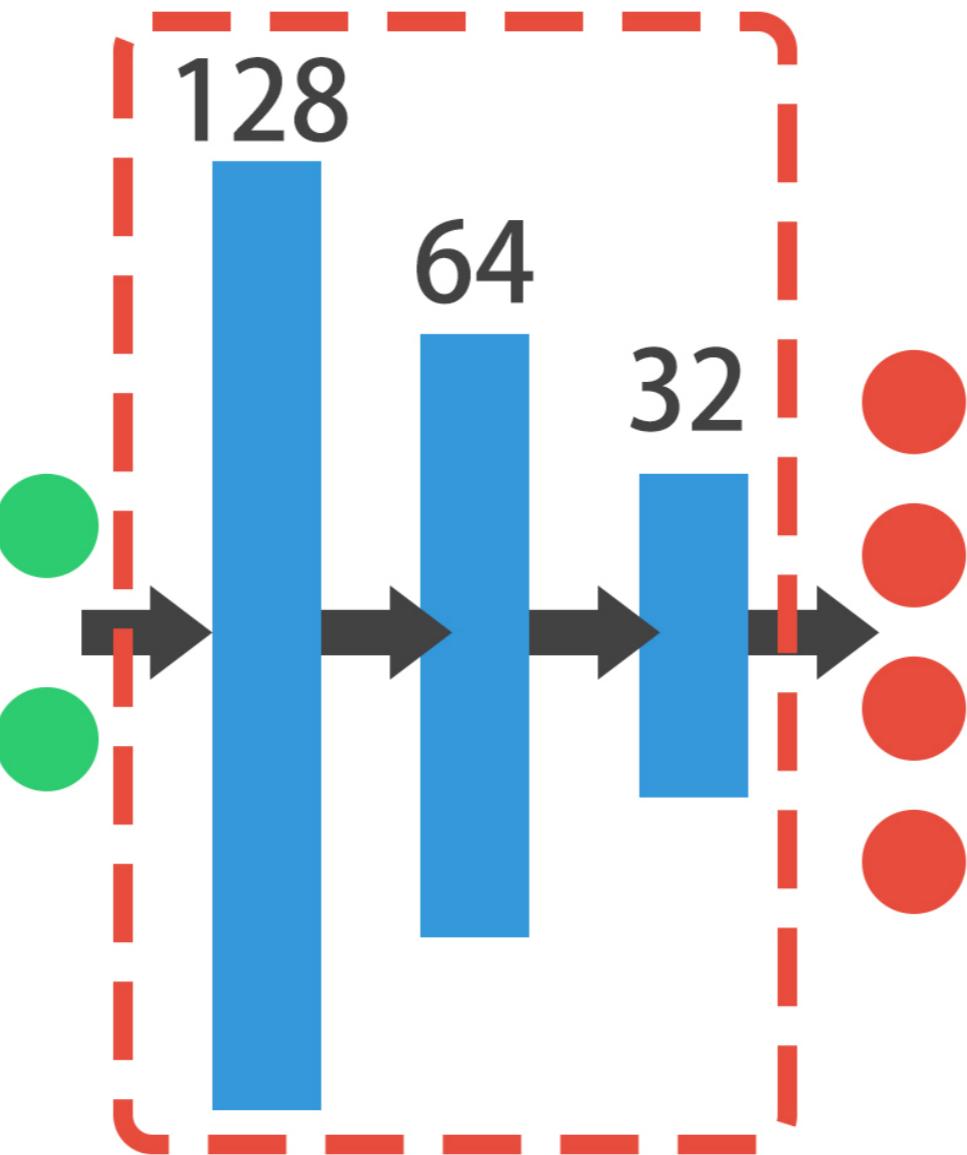
# The dataset



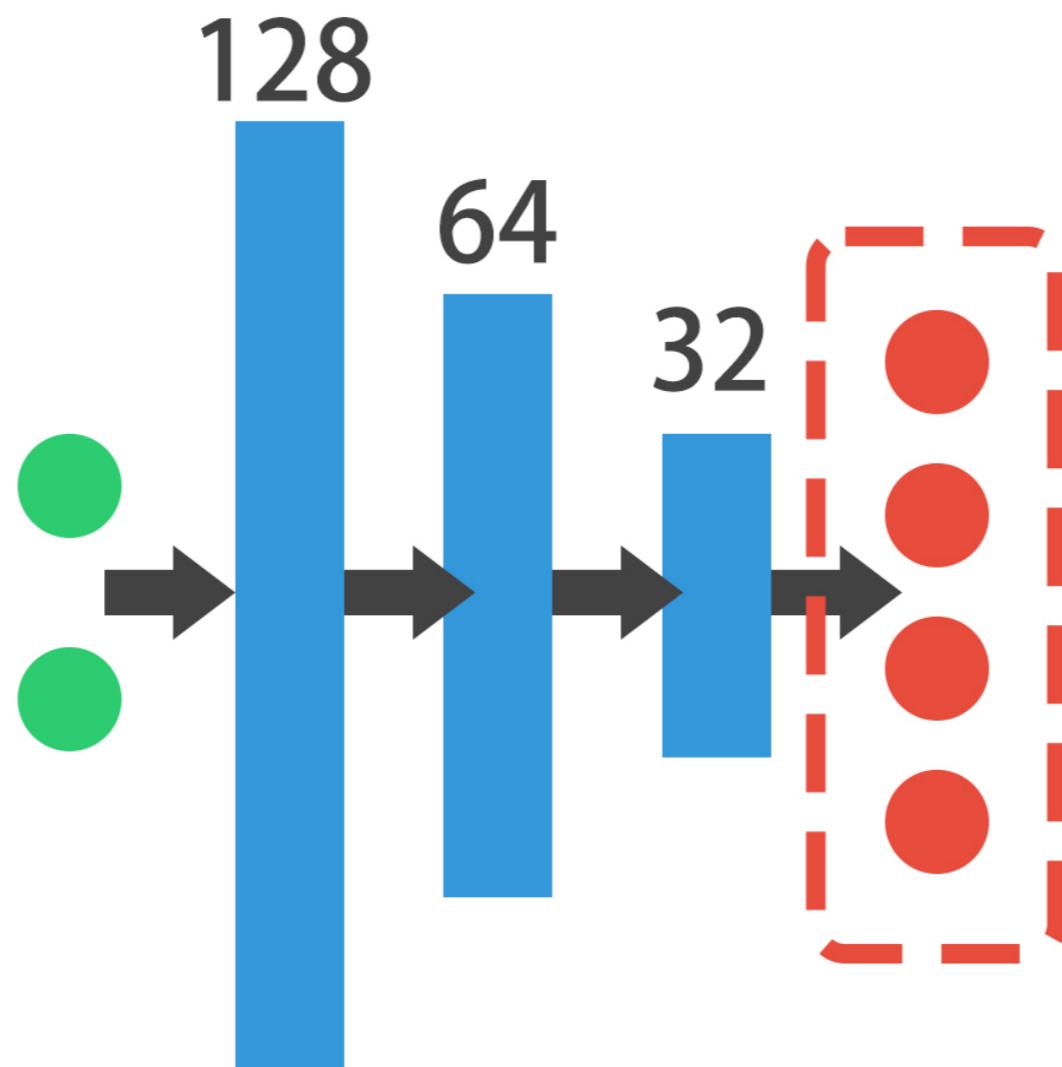
# The architecture



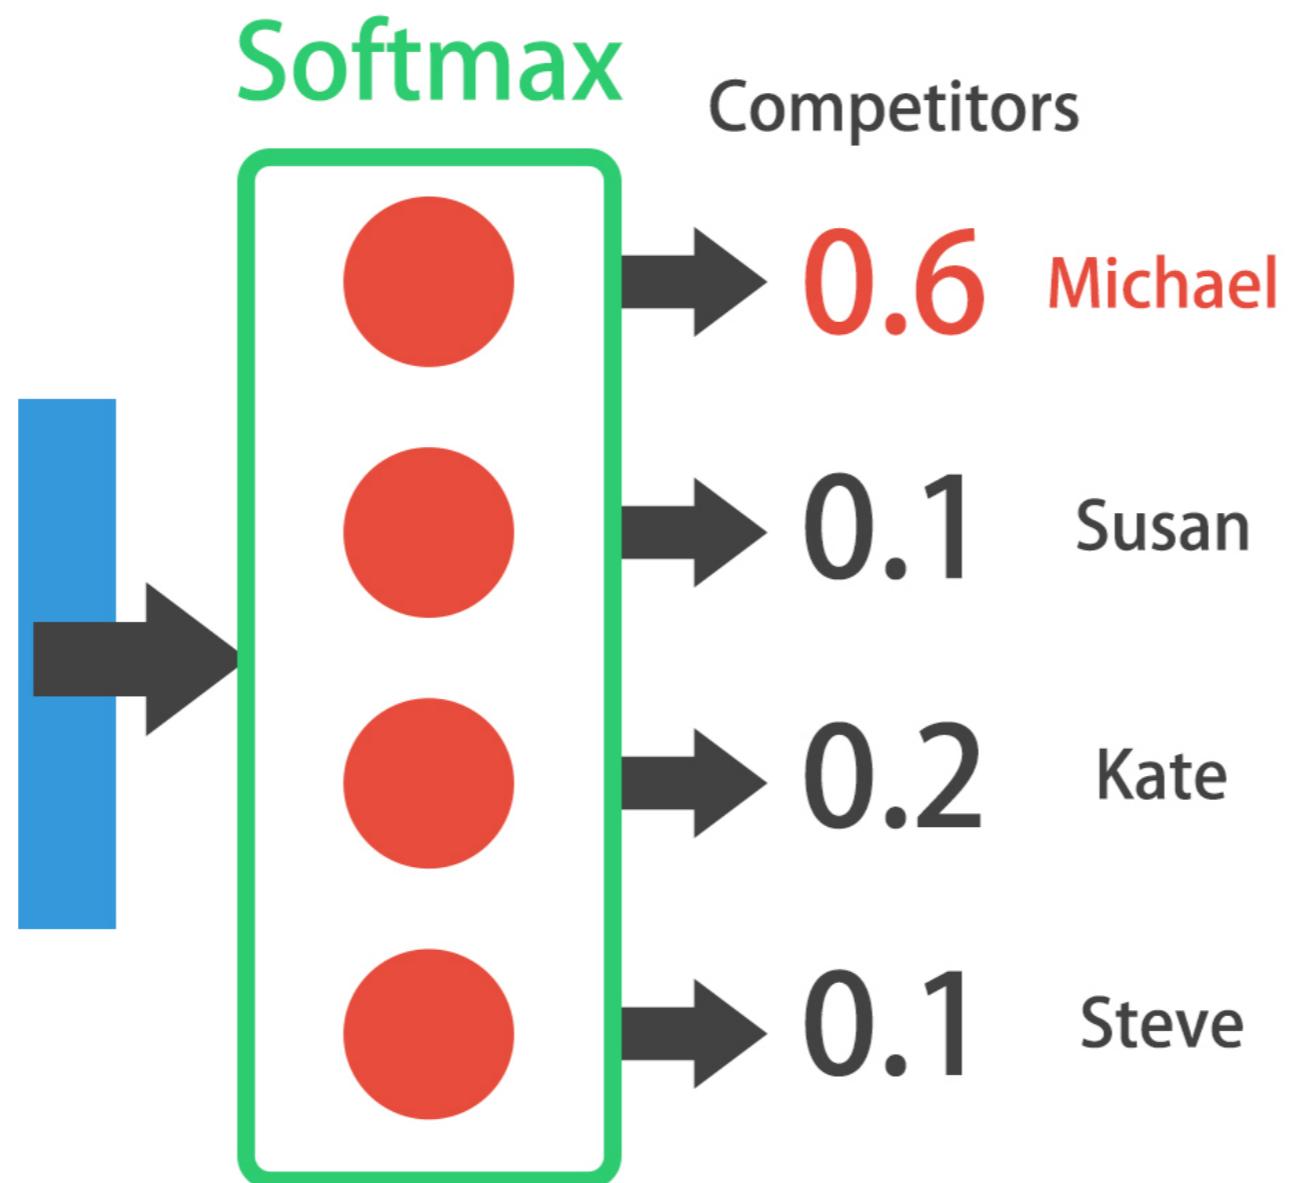
# The architecture



# The architecture



# The output layer

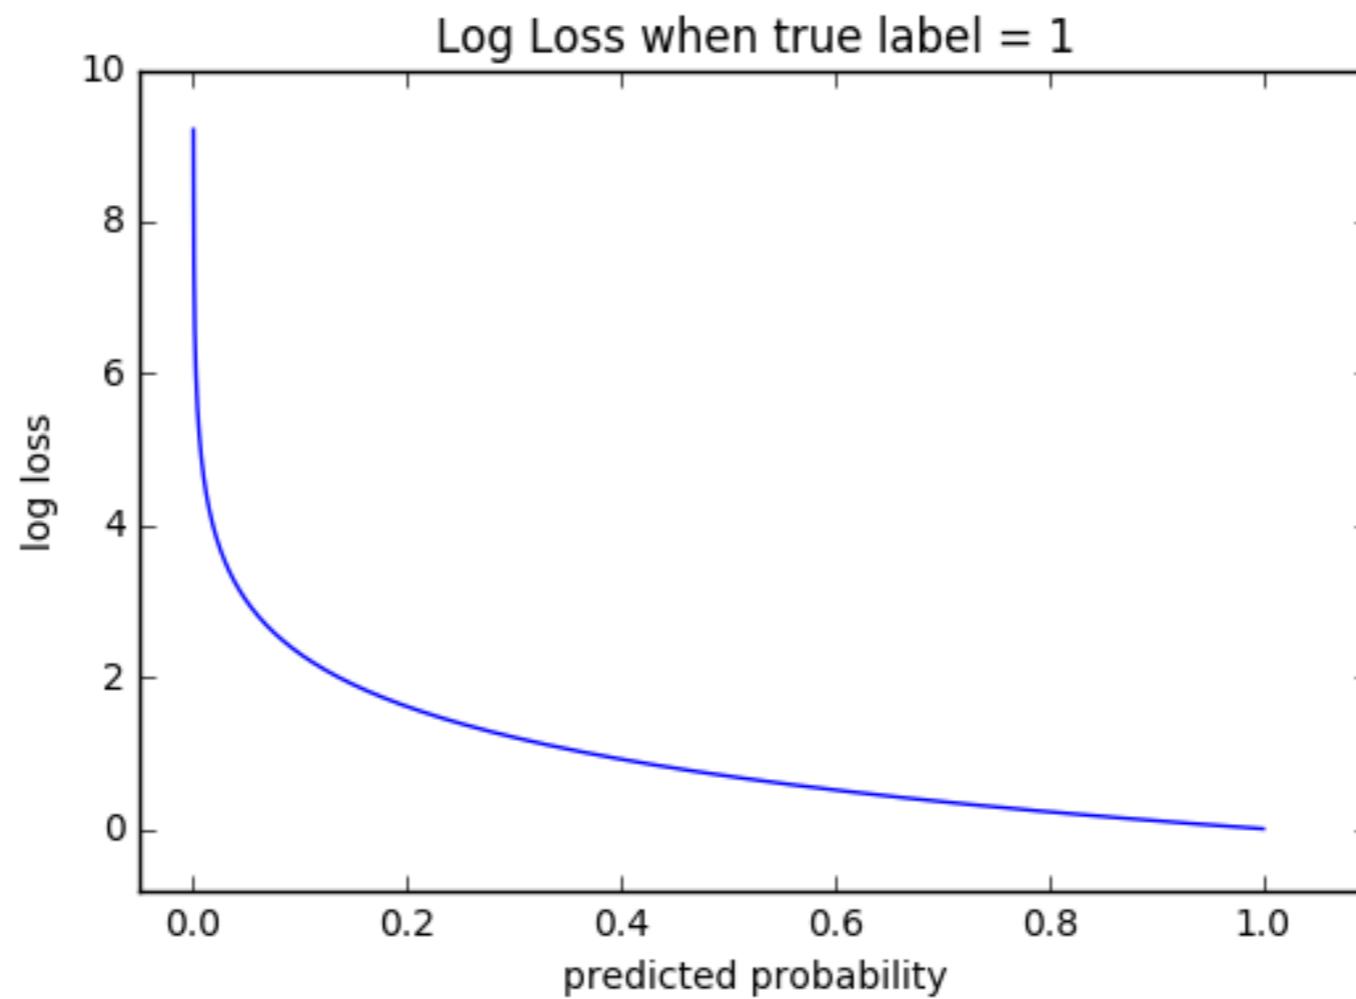


# Multi-class model

```
# Instantiate a sequential model  
# ...  
# Add an input and hidden layer  
# ...  
# Add more hidden layers  
# ...  
# Add your output layer  
model.add(Dense(4, activation='softmax'))
```

# Categorical cross-entropy

```
model.compile(optimizer='adam', loss='categorical_crossentropy')
```



# Preparing a dataset

```
import pandas as pd
from keras.utils import to_categorical

# Load dataset
df = pd.read_csv('data.csv')

# Turn response variable into labeled codes
df.response = pd.Categorical(df.response)
df.response = df.response.cat.codes

# Turn response variable into one-hot response vector
y = to_categorical(df.response)
```

# One-hot encoding

Label Encoding

Food Name	Categorical #	Calories
Apple	1	95
Chicken	2	231
Broccoli	3	50



One Hot Encoding

Apple	Chicken	Broccoli	Calories
1	0	0	95
0	1	0	231
0	0	1	50

# Let's practice!

INTRODUCTION TO DEEP LEARNING WITH KERAS

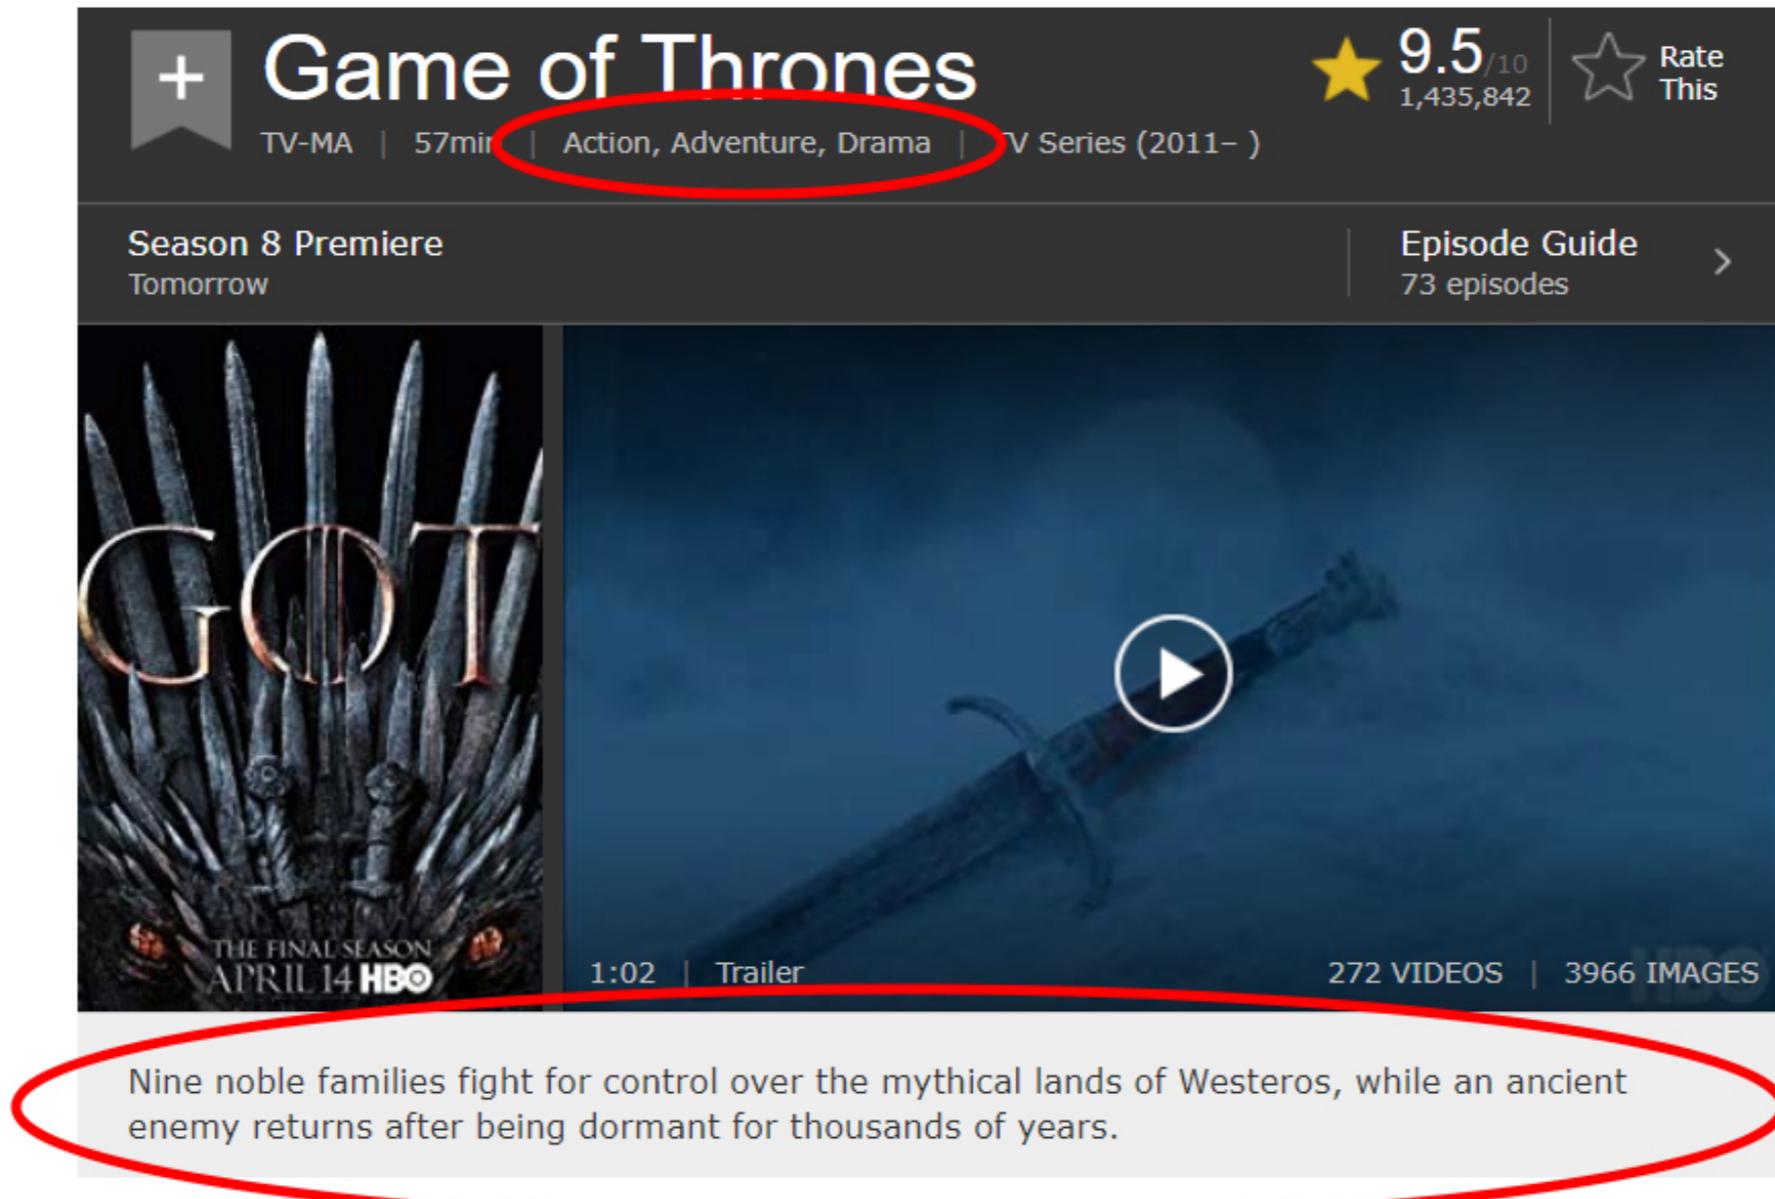
# Multi-label classification

INTRODUCTION TO DEEP LEARNING WITH KERAS

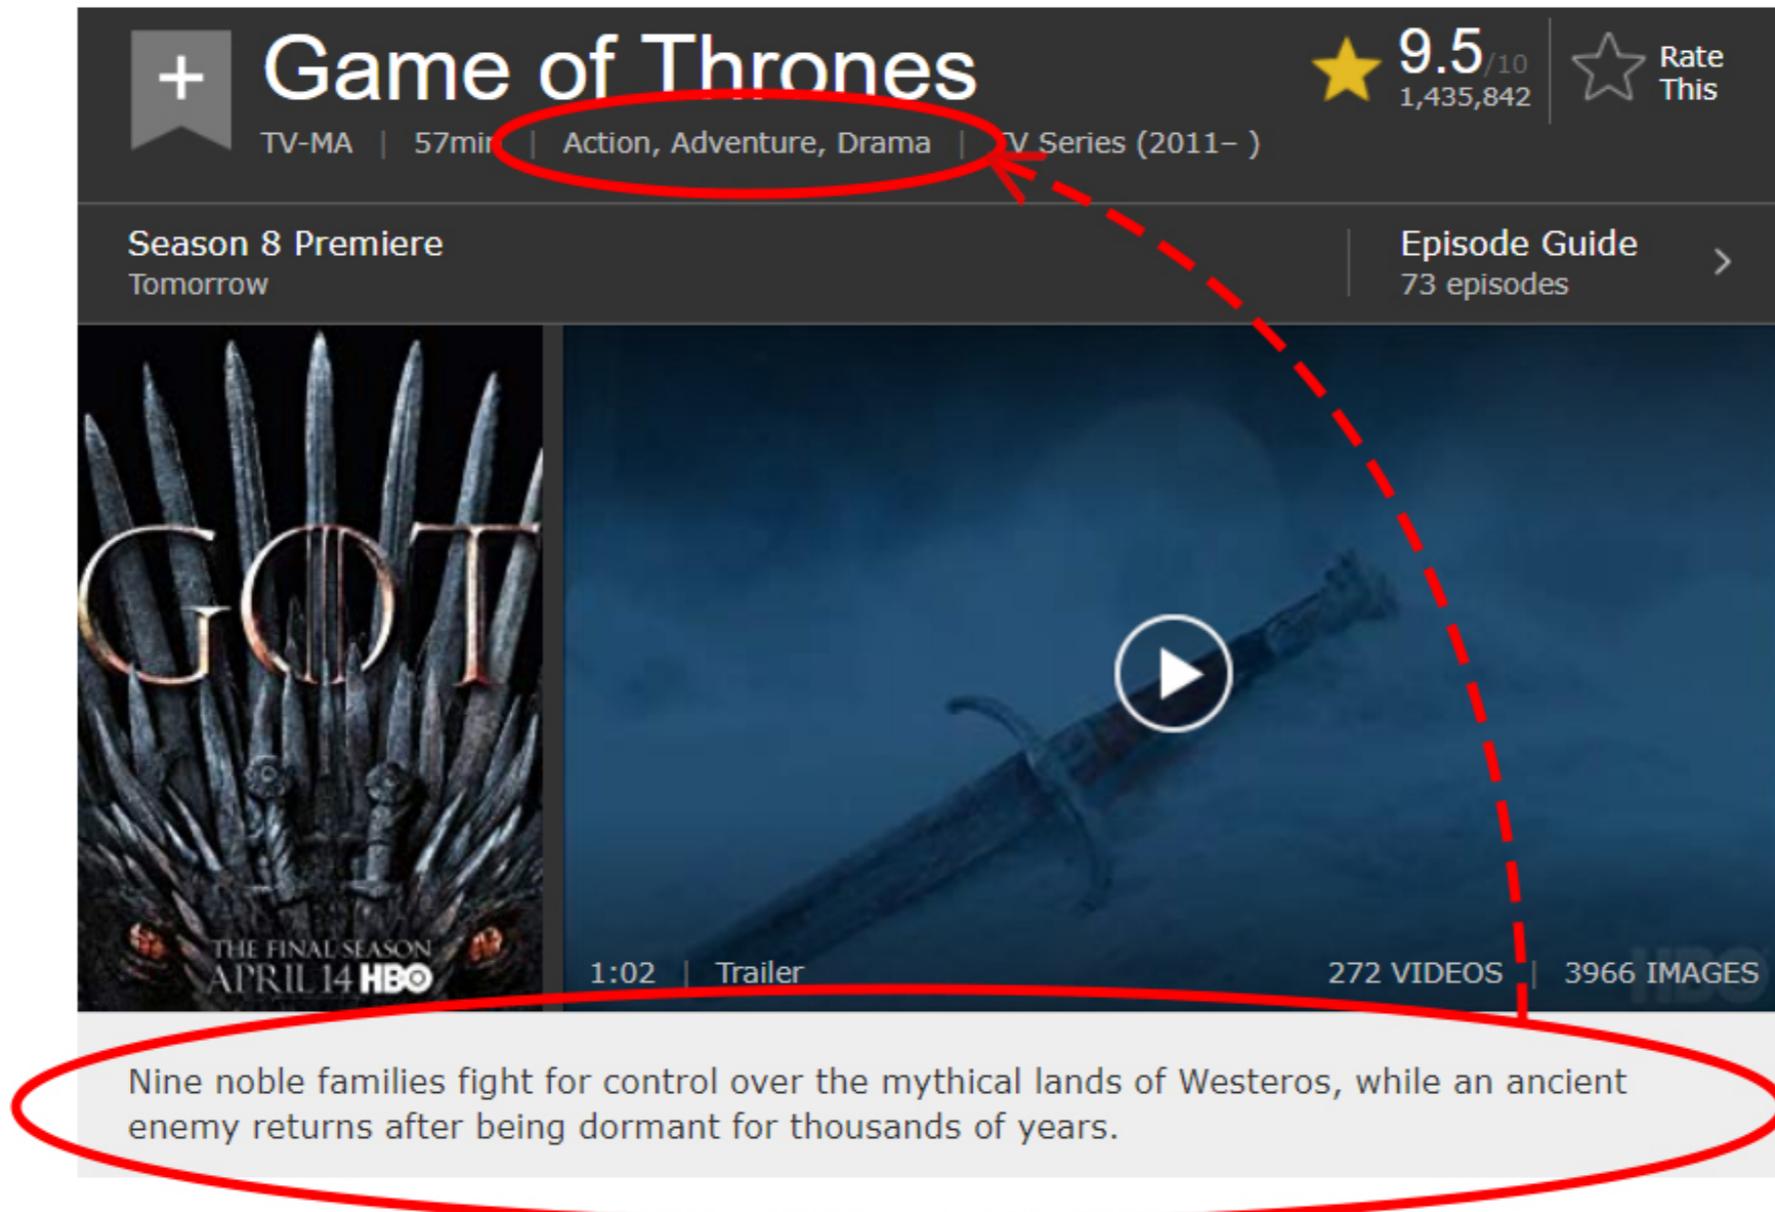


**Miguel Esteban**  
Data Scientist & Founder

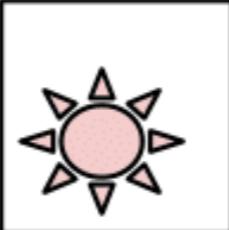
# Real world examples



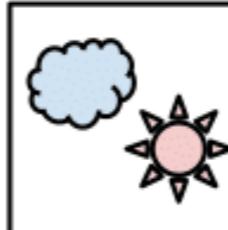
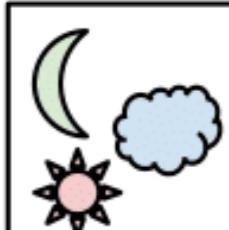
# Real world examples



## Multi-Class

C = 3	Samples		
			
	Labels (t)		
	[0 0 1]	[1 0 0]	[0 1 0]

	Samples		
			
	Labels (t)		
	[1 0 1]	[0 1 0]	[1 1 1]

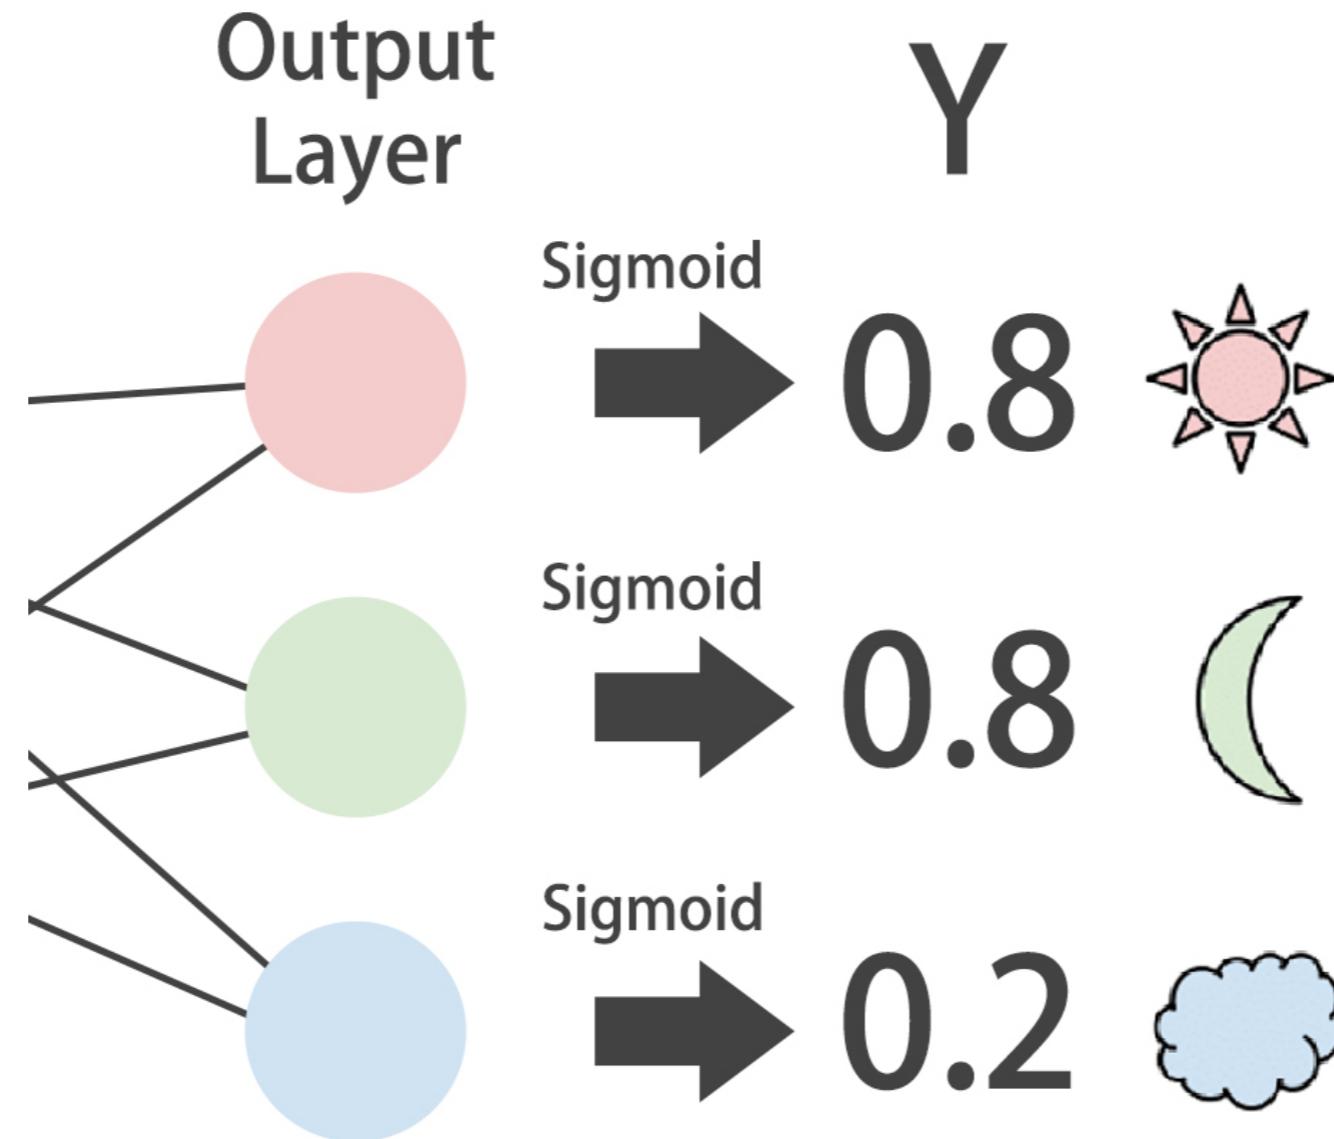
<sup>1</sup> [https://gombru.github.io/2018/05/23/cross\\_entropy\\_loss/](https://gombru.github.io/2018/05/23/cross_entropy_loss/)

## Multi-Label

# The architecture

```
from keras.models import Sequential  
from keras.layers import Dense  
  
# Instantiate model  
model = Sequential()  
  
# Add input and hidden layers  
model.add(Dense(2, input_shape=(1,)))  
  
# Add an output layer for the 3 classes and sigmoid activation  
model.add(Dense(3, activation='sigmoid'))
```

# Sigmoid outputs

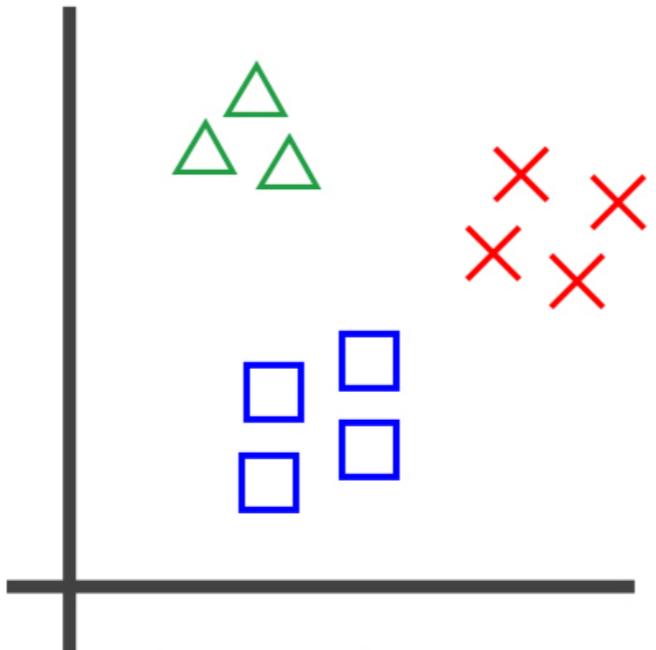


```
# Compile the model with binary crossentropy  
model.compile(optimizer='adam', loss='binary_crossentropy')
```

```
# Train your model, recall validation_split  
model.fit(X_train, y_train,  
          epochs=100,  
          validation_split=0.2)
```

```
Train on 1260 samples, validate on 280 samples  
Epoch 1/100  
1260/1260 [=====] - 0s 285us/step  
- loss: 0.7035 - acc: 0.6690 - val_loss: 0.5178 - val_acc: 0.7714  
...
```

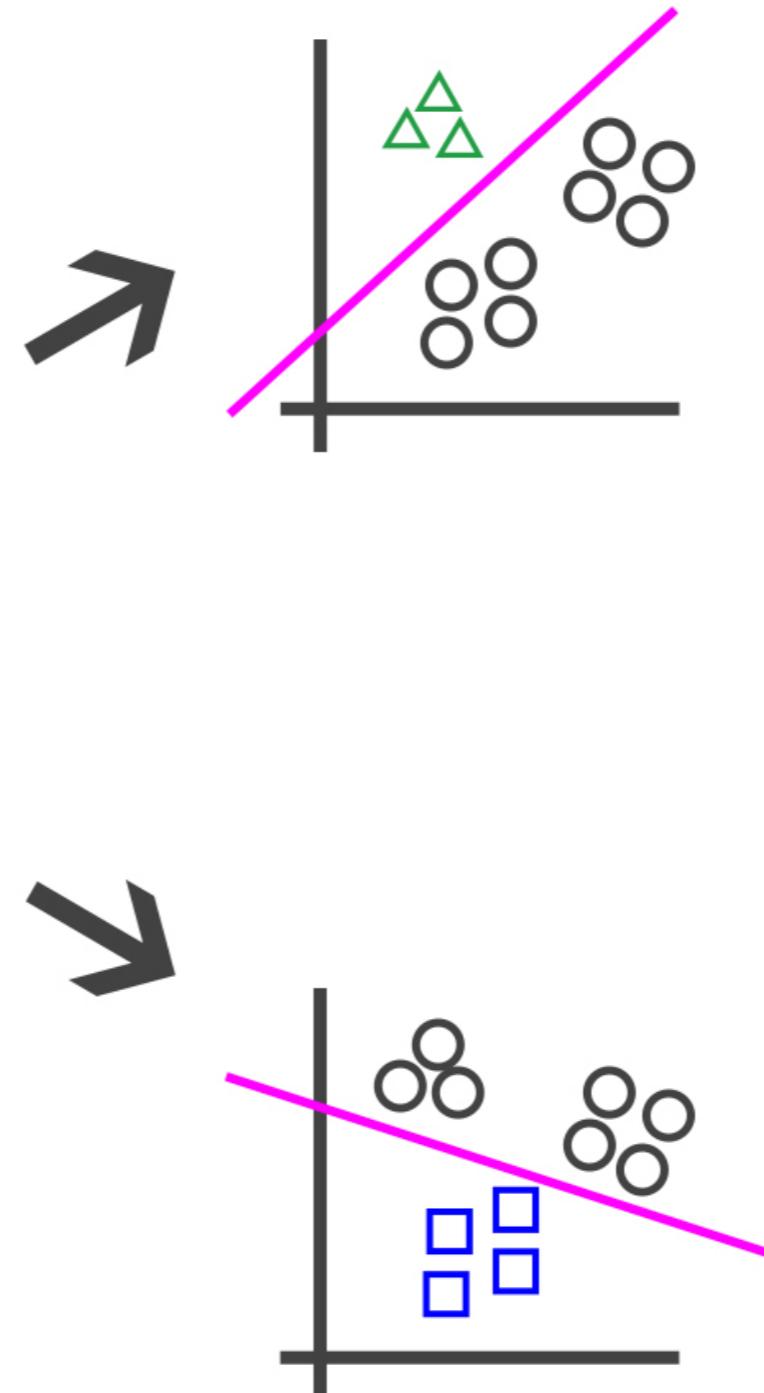
## One-vs-rest



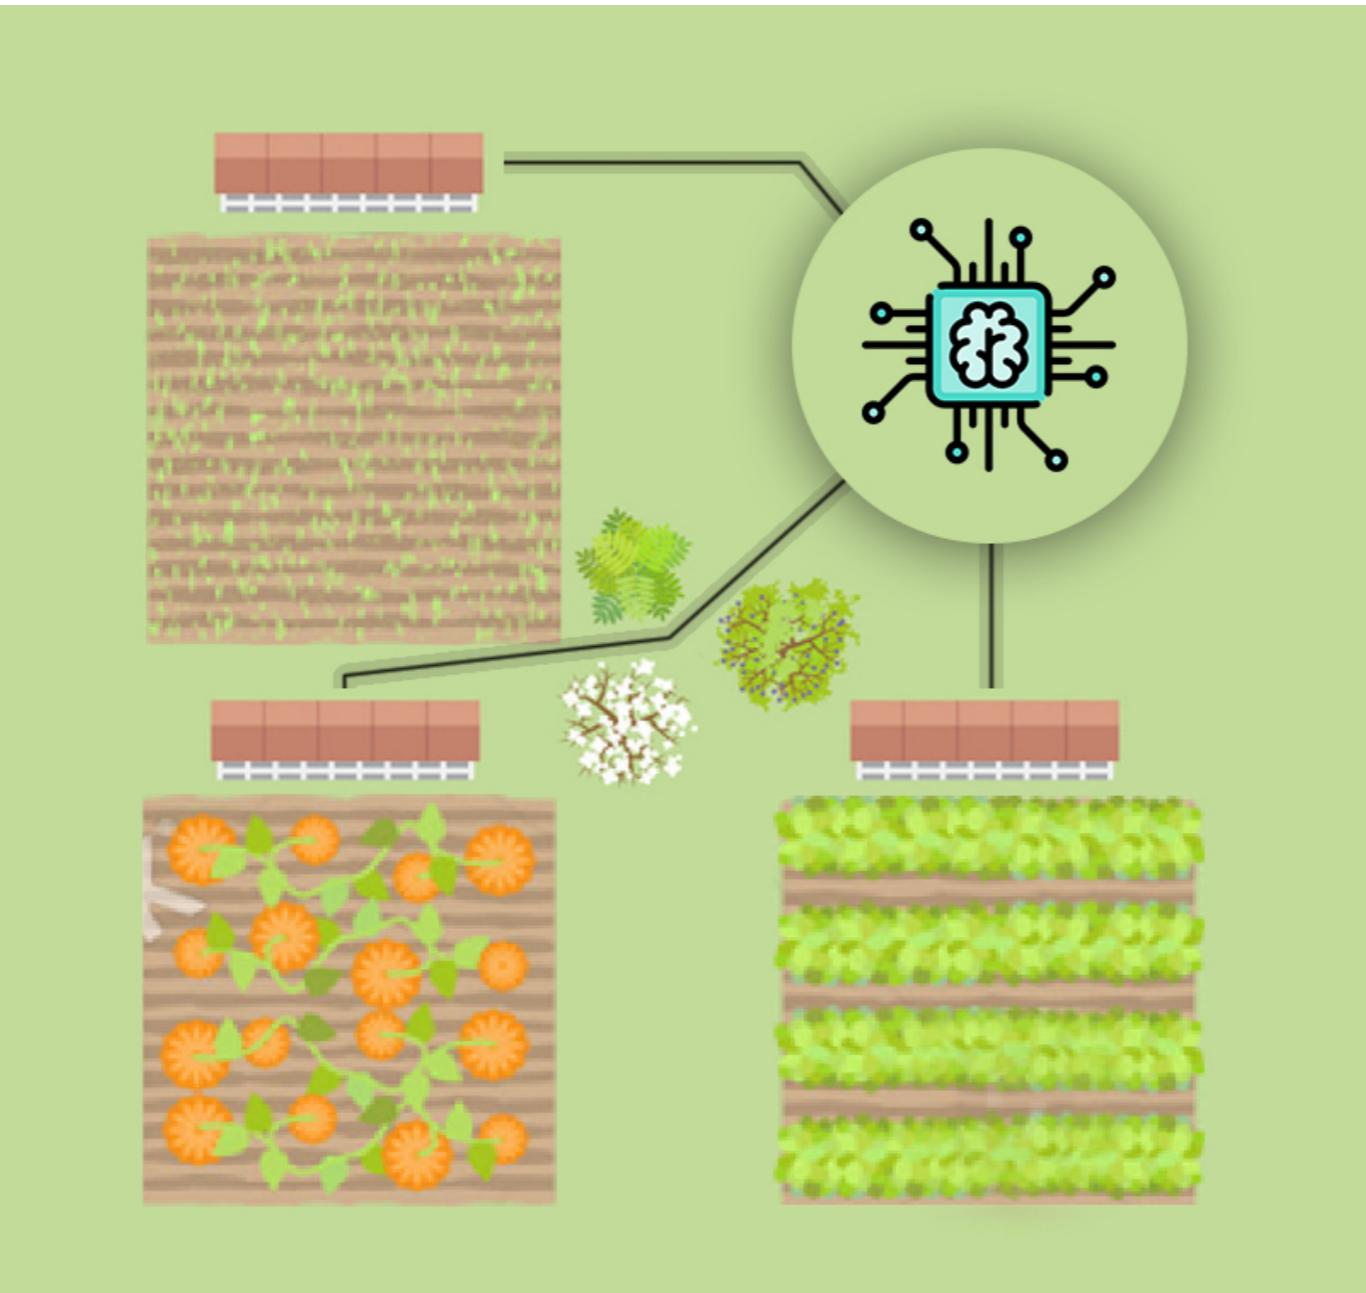
class 1:  $\triangle$

class 2:  $\square$

class 3:  $\times$



# An irrigation machine



# An irrigation machine

## Sensor measurements

	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19
0	4.0	3.0	4.0	4.0	1.0	1.0	4.0	0.0	3.0	3.0	2.0	2.0	3.0	1.0	4.0	2.0	2.0	2.0	1.0	4.0
1	1.0	1.0	6.0	3.0	2.0	3.0	4.0	5.0	1.0	3.0	3.0	2.0	3.0	2.0	2.0	4.0	0.0	1.0	2.0	4.0
2	1.0	5.0	7.0	6.0	4.0	0.0	0.0	6.0	0.0	1.0	1.0	3.0	4.0	2.0	1.0	0.0	4.0	1.0	1.0	3.0
3	0.0	1.0	3.0	3.0	7.0	1.0	2.0	2.0	0.0	4.0	3.0	2.0	4.0	2.0	0.0	2.0	3.0	4.0	1.0	2.0
4	1.0	5.0	2.0	2.0	1.0	0.0	3.0	3.0	1.0	1.0	5.0	1.0	1.0	2.0	2.0	4.0	3.0	3.0	0.0	5.0

## Parcels to water

	0	1	2
0	0	0	0
1	1	1	1
2	1	1	0
3	1	1	1
4	0	0	0

# Let's practice!

INTRODUCTION TO DEEP LEARNING WITH KERAS

# Keras callbacks

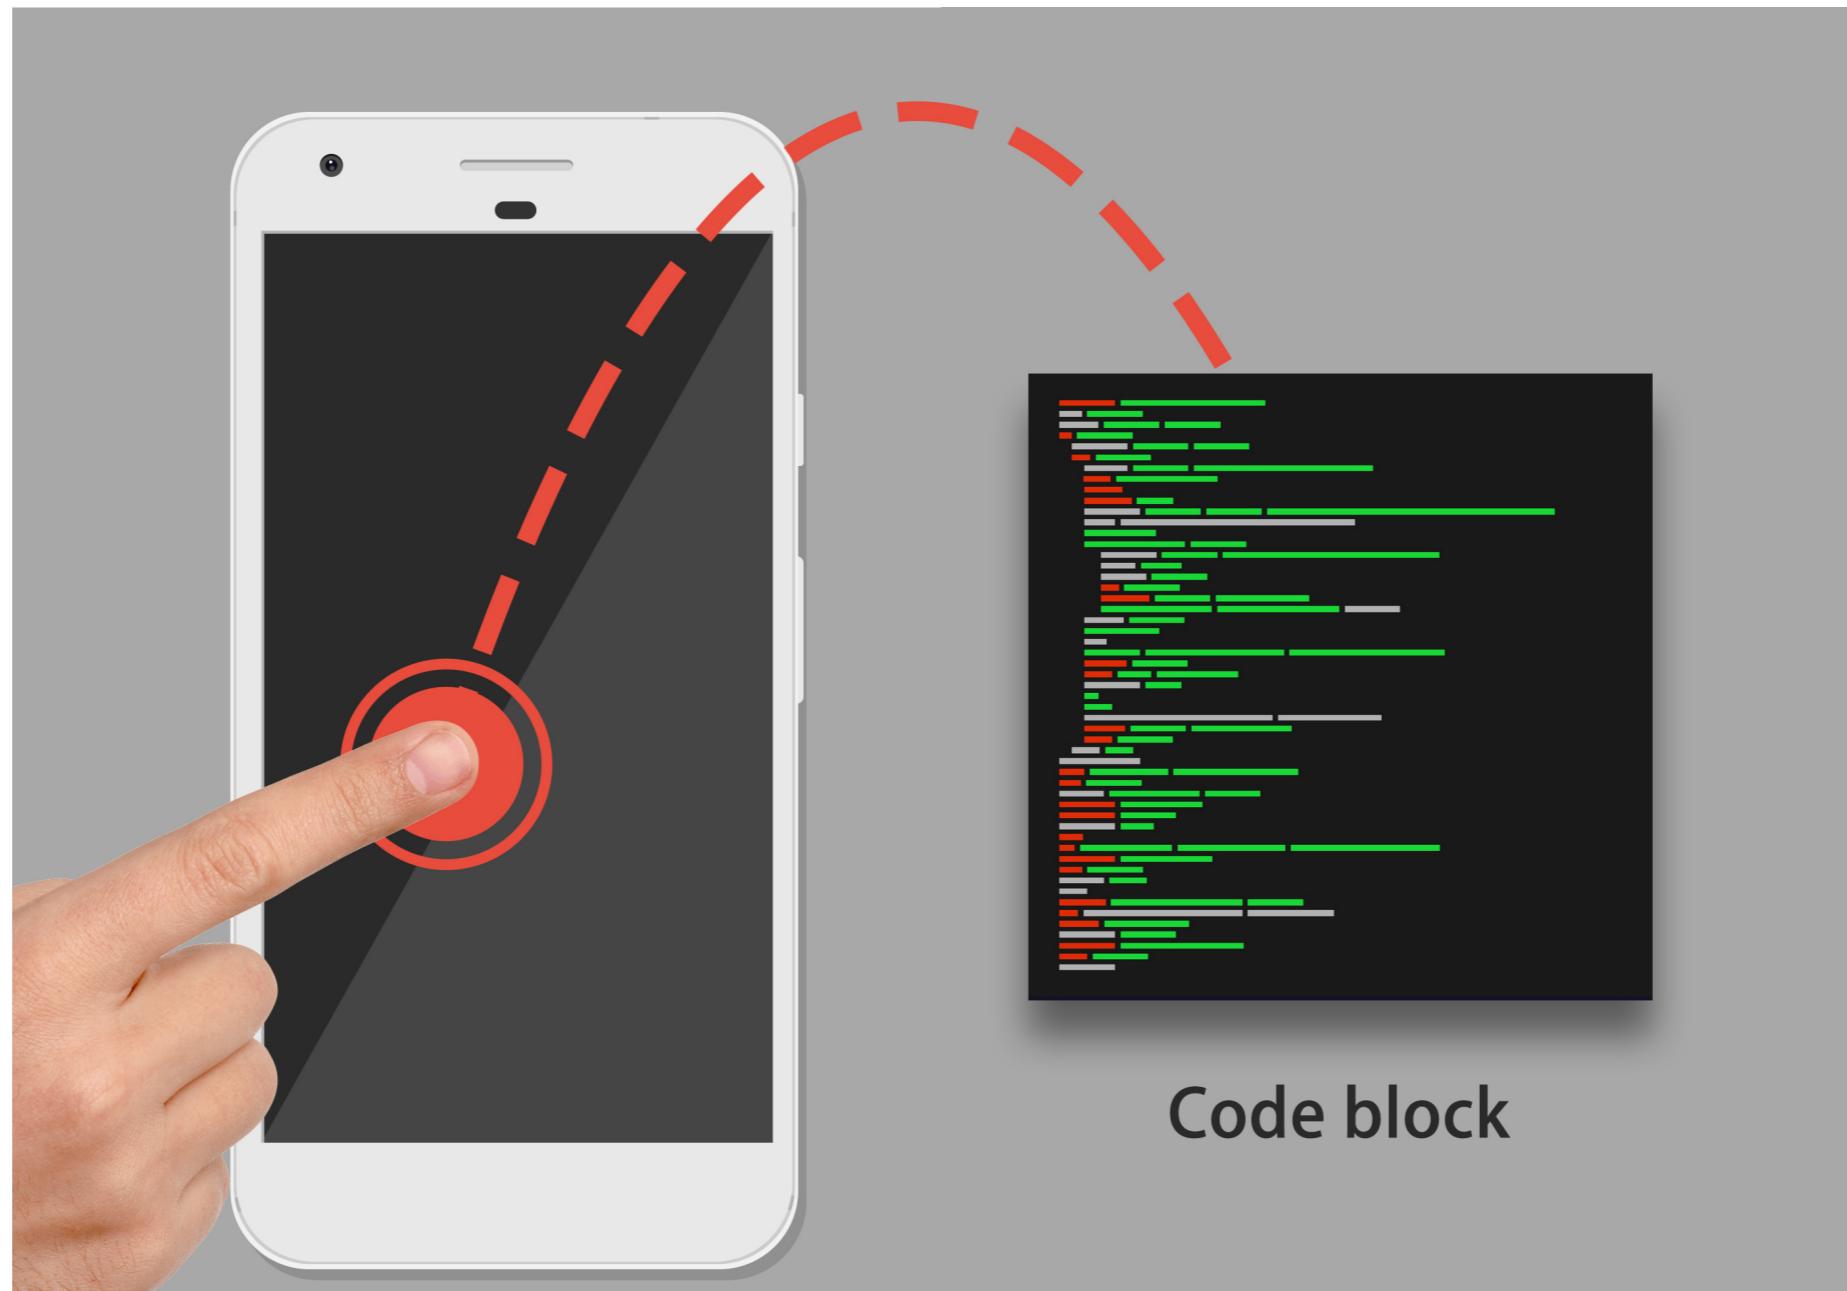
INTRODUCTION TO DEEP LEARNING WITH KERAS



Miguel Esteban

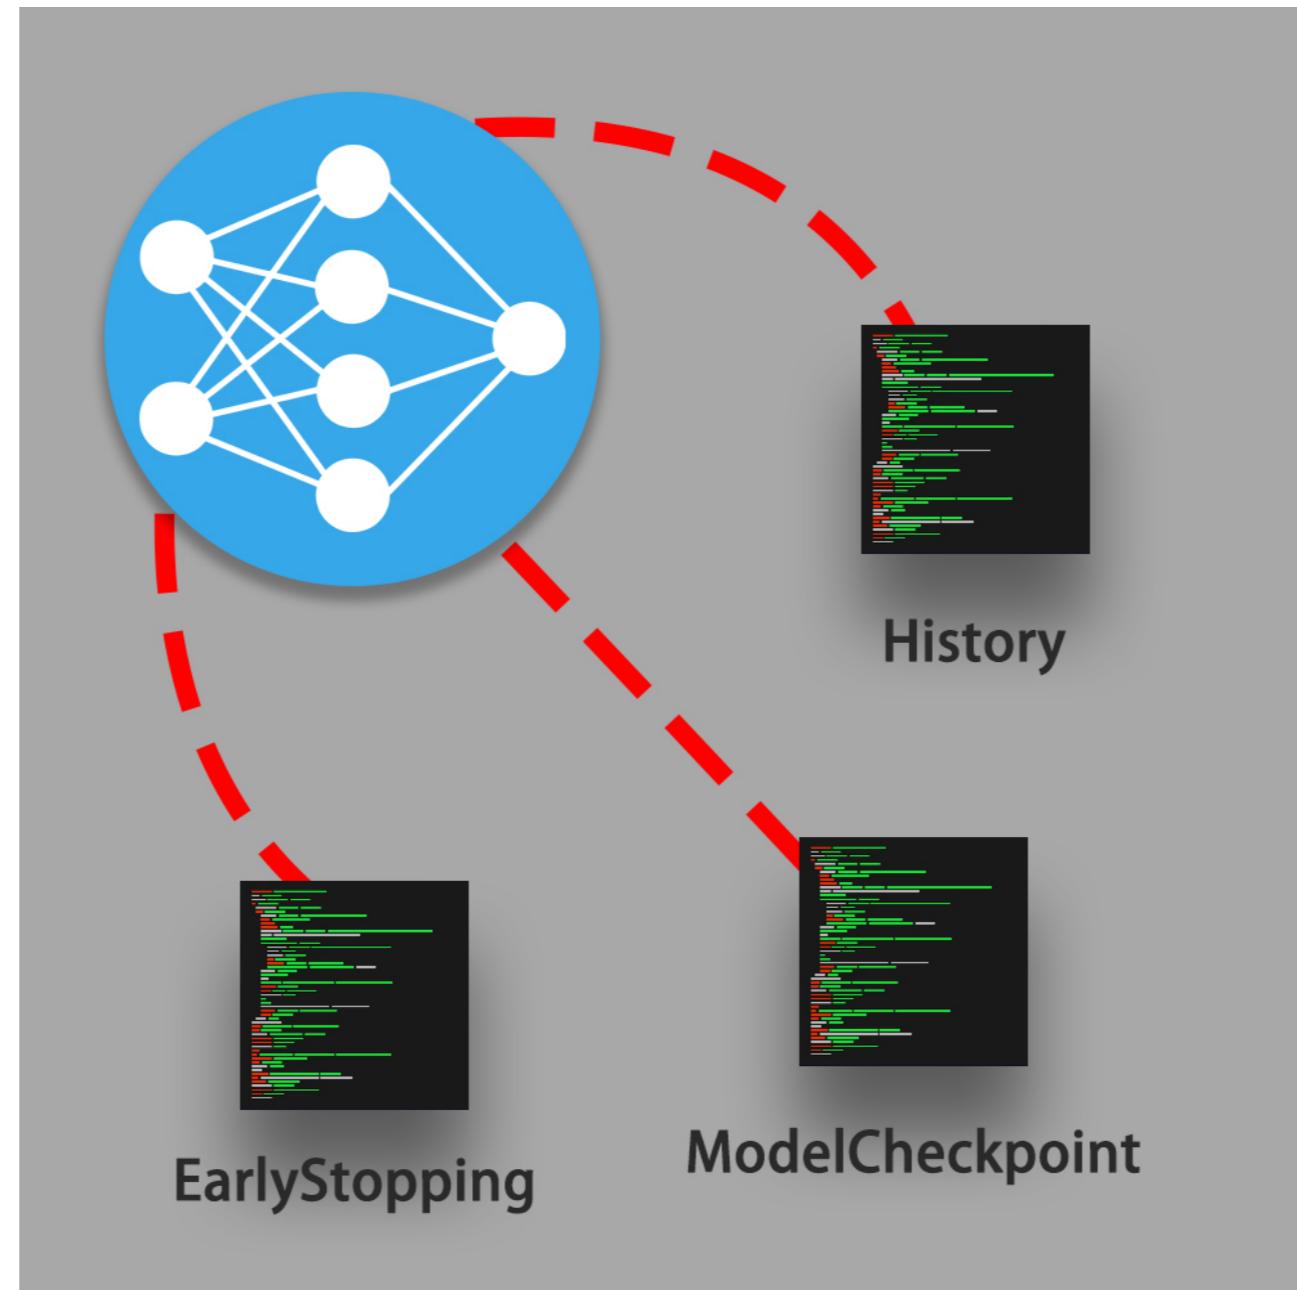
Data Scientist & Founder

# What is a callback?



Code block

# Callbacks in Keras



# A callback you've been missing

```
# Training a model and saving its history  
history = model.fit(X_train, y_train,  
                      epochs=100,  
                      metrics=[ 'accuracy' ])  
  
print(history.history[ 'loss' ])
```

```
[0.6753975939750672, ..., 0.3155936544282096]
```

```
print(history.history[ 'acc' ])
```

```
[0.7030952412741525, ..., 0.8604761900220599]
```

# A callback you've been missing

```
# Training a model and saving its history
history = model.fit(X_train, y_train,
                     epochs=100,
                     validation_data=(X_test, y_test),
                     metrics=[ 'accuracy' ])

print(history.history[ 'val_loss' ])
```

```
[0.7753975939750672, ..., 0.4155936544282096]
```

```
print(history.history[ 'val_acc' ])
```

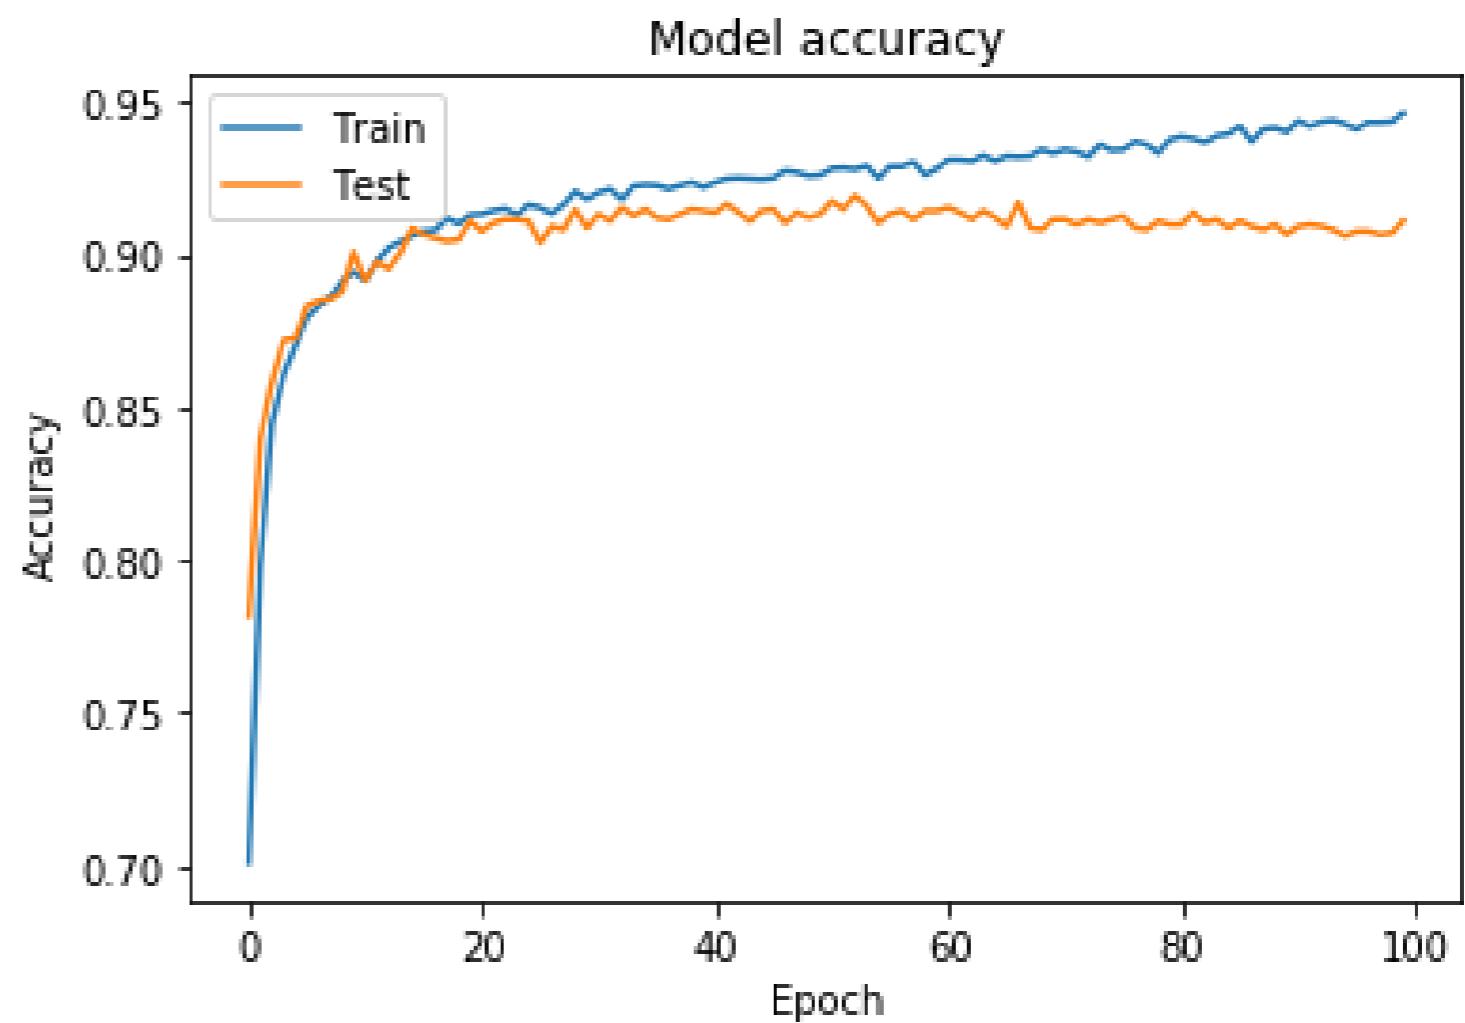
```
[0.6030952412741525, ..., 0.7604761900220599]
```

# History plots

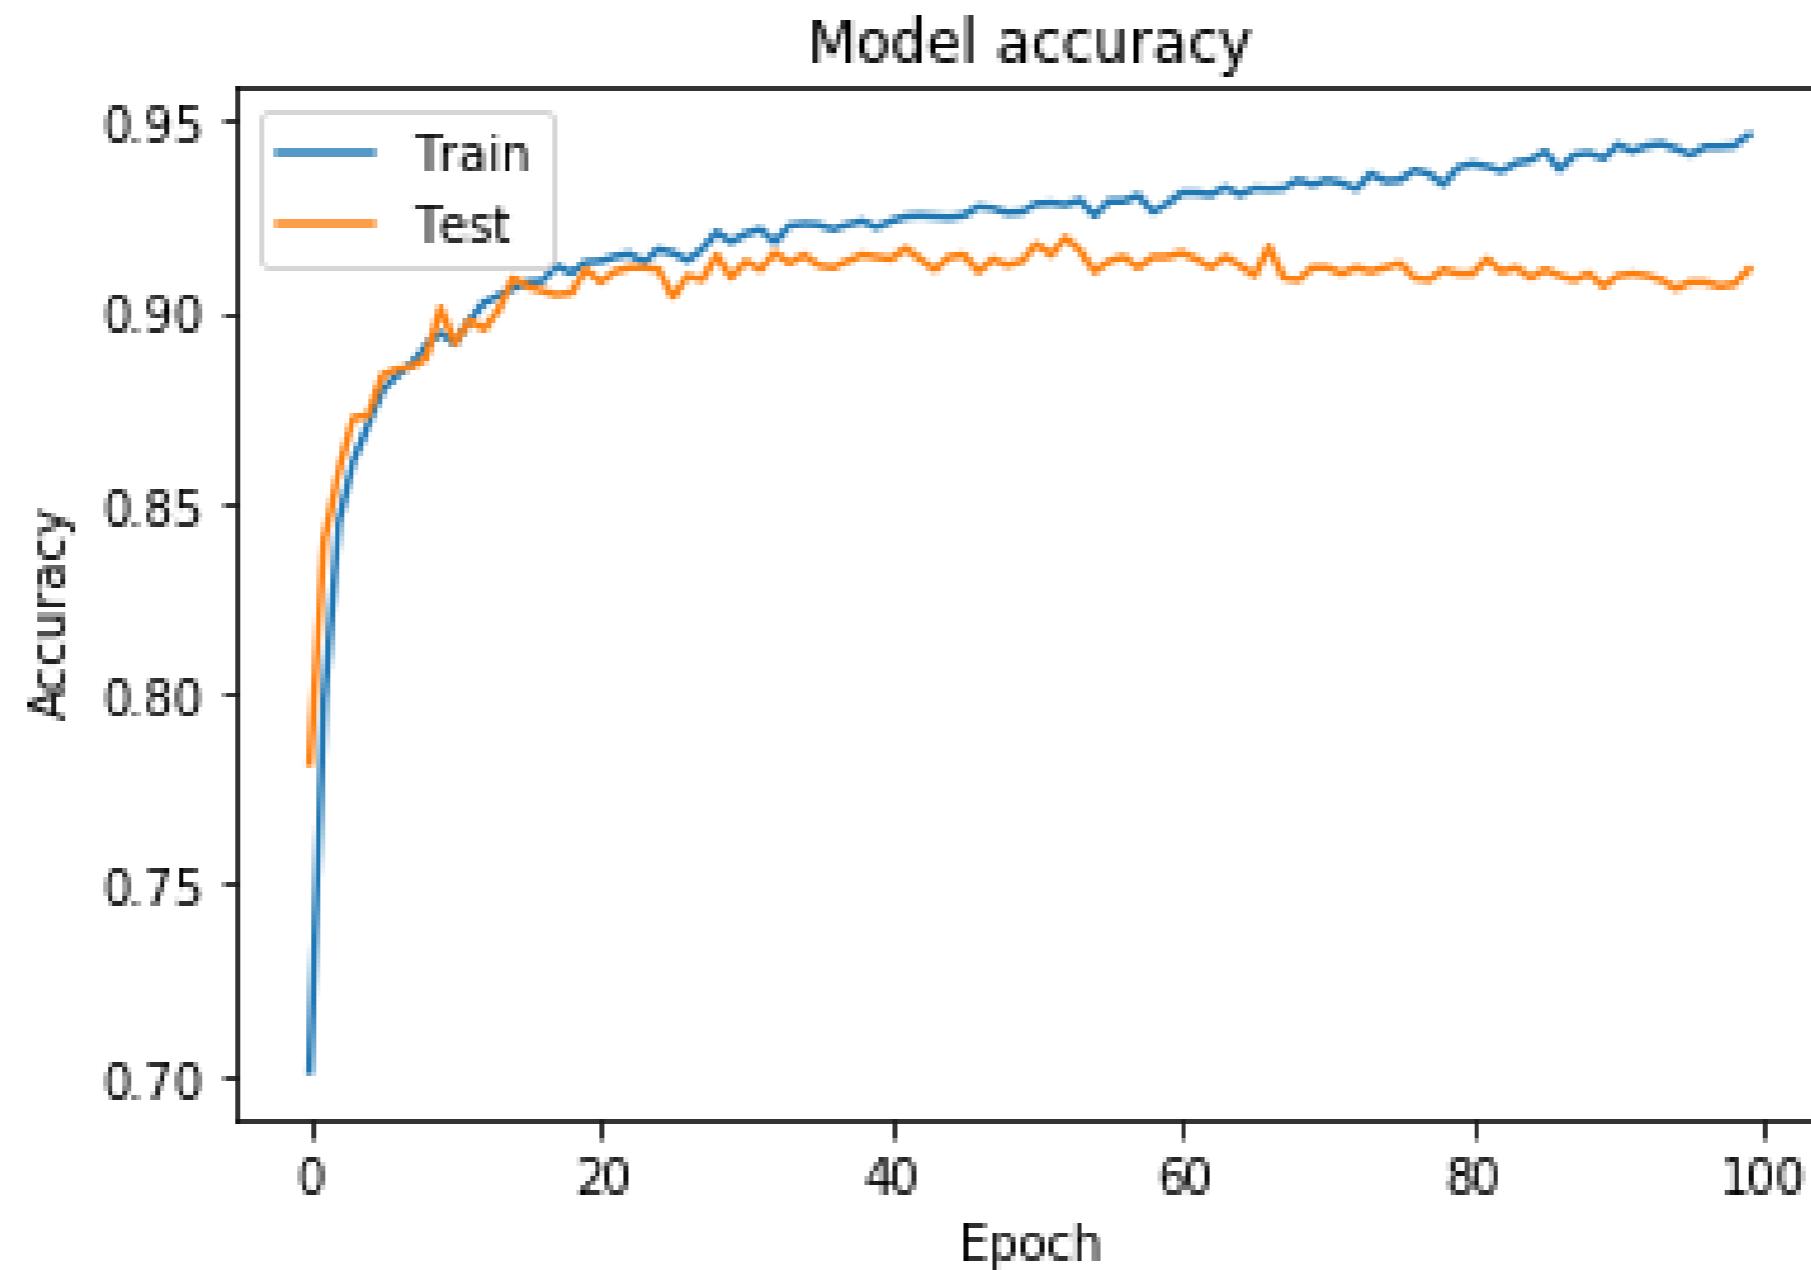
```
# Plot train vs test accuracy per epoch
plt.figure()

# Use the history metrics
plt.plot(history.history['acc'])
plt.plot(history.history['val_acc'])

# Make it pretty
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'])
plt.show()
```



# History plots



# Early stopping

```
# Import early stopping from keras callbacks
from keras.callbacks import EarlyStopping

# Instantiate an early stopping callback
early_stopping = EarlyStopping(monitor='val_loss', patience=5)

# Train your model with the callback
model.fit(X_train, y_train, epochs=100,
           validation_data=(X_test, y_test),
           callbacks = [early_stopping])
```

# Model checkpoint

```
# Import model checkpoint from keras callbacks
from keras.callbacks import ModelCheckpoint

# Instantiate a model checkpoint callback
model_save = ModelCheckpoint('best_model.hdf5',
                             save_best_only=True)

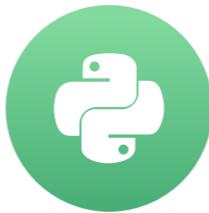
# Train your model with the callback
model.fit(X_train, y_train, epochs=100,
           validation_data=(X_test, y_test),
           callbacks = [model_save])
```

# Let's practice!

INTRODUCTION TO DEEP LEARNING WITH KERAS

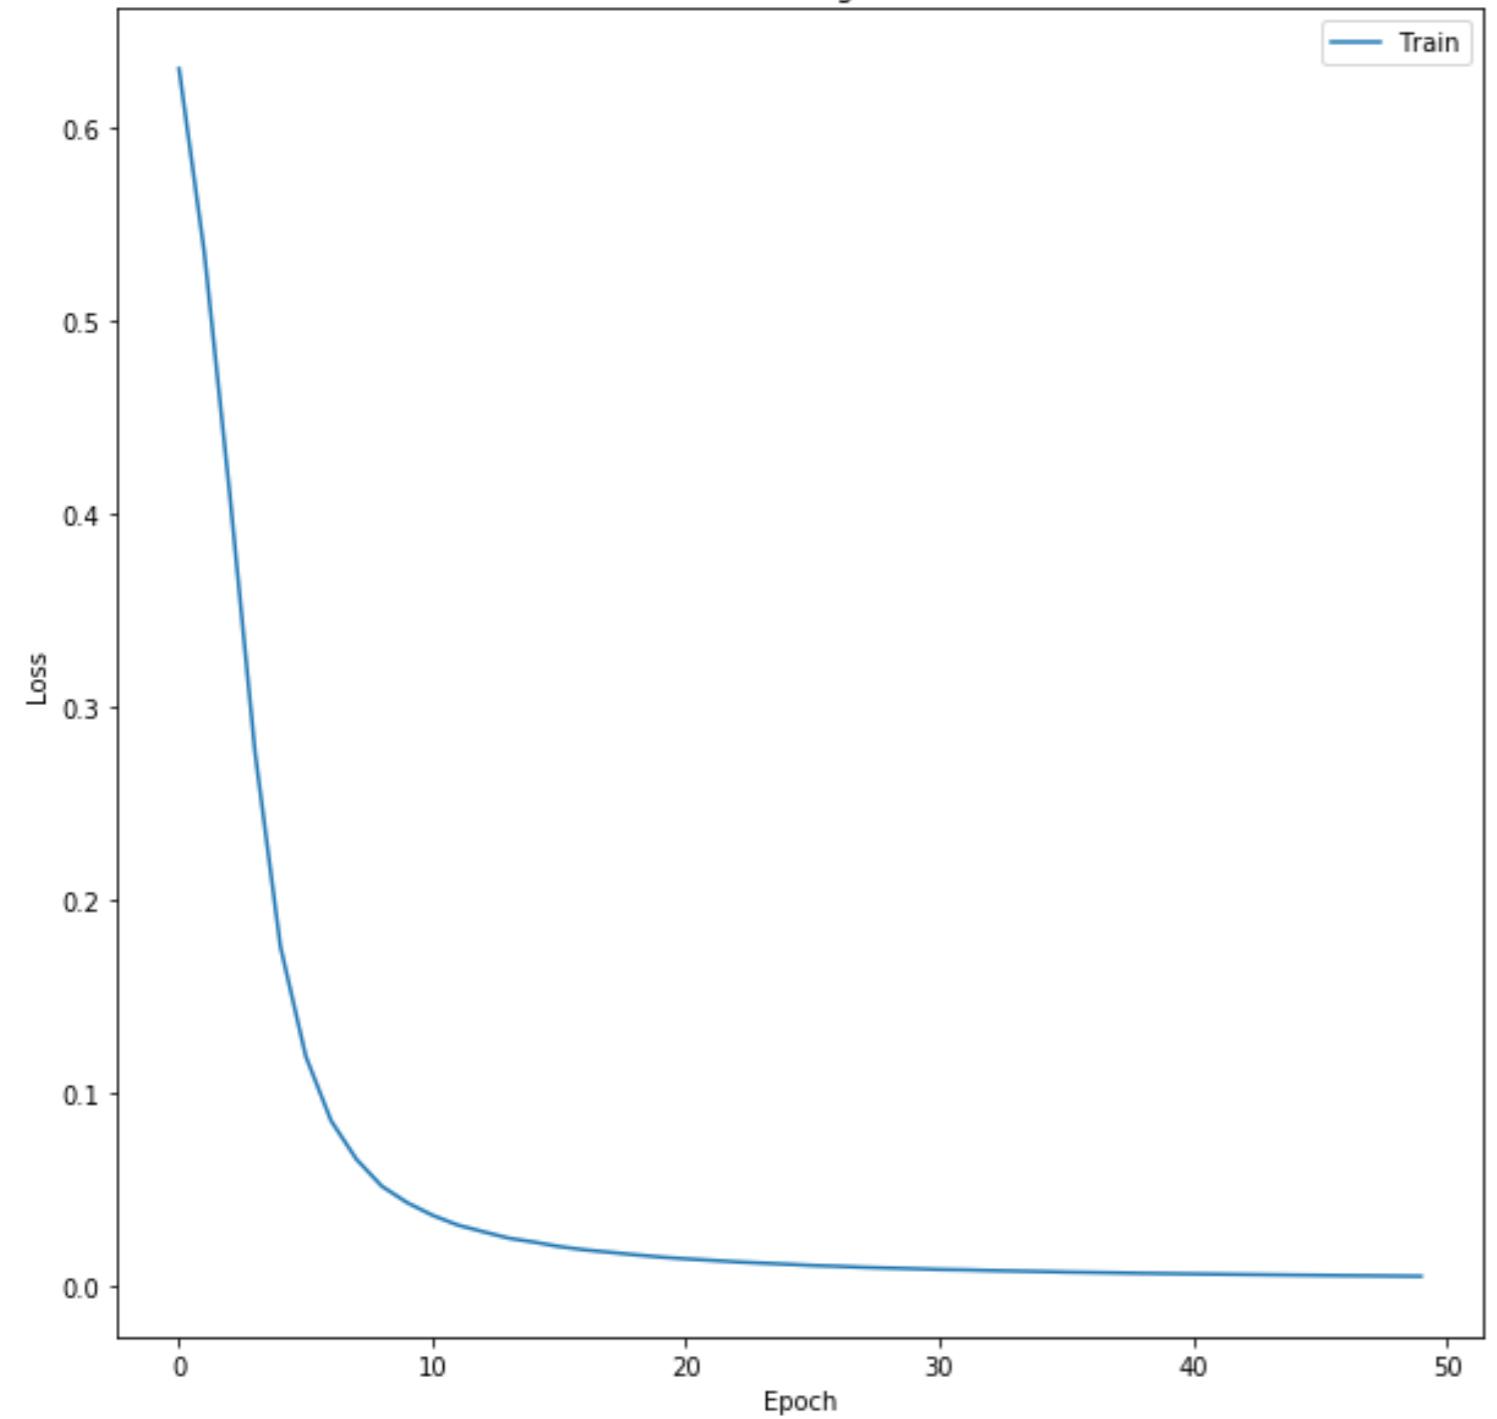
# Learning curves

INTRODUCTION TO DEEP LEARNING WITH KERAS

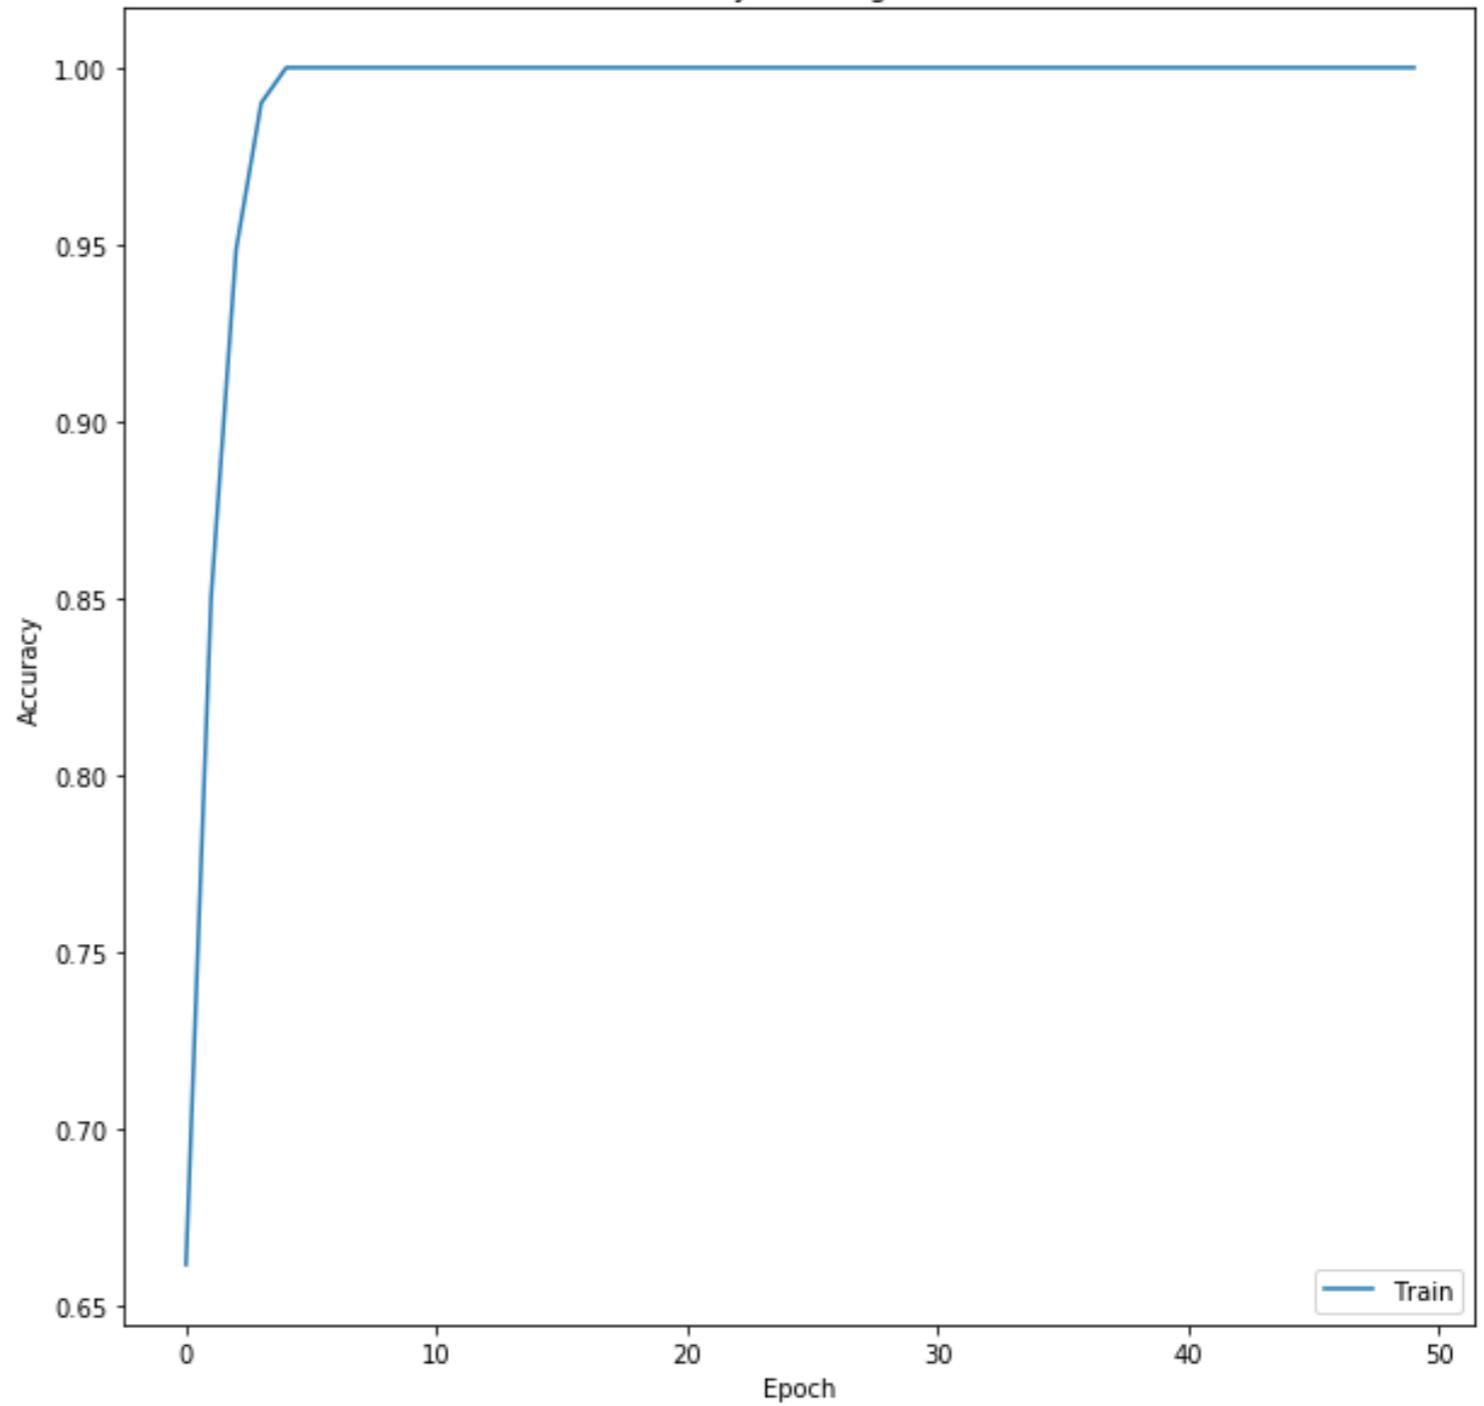


**Miguel Esteban**  
Data Scientist & Founder

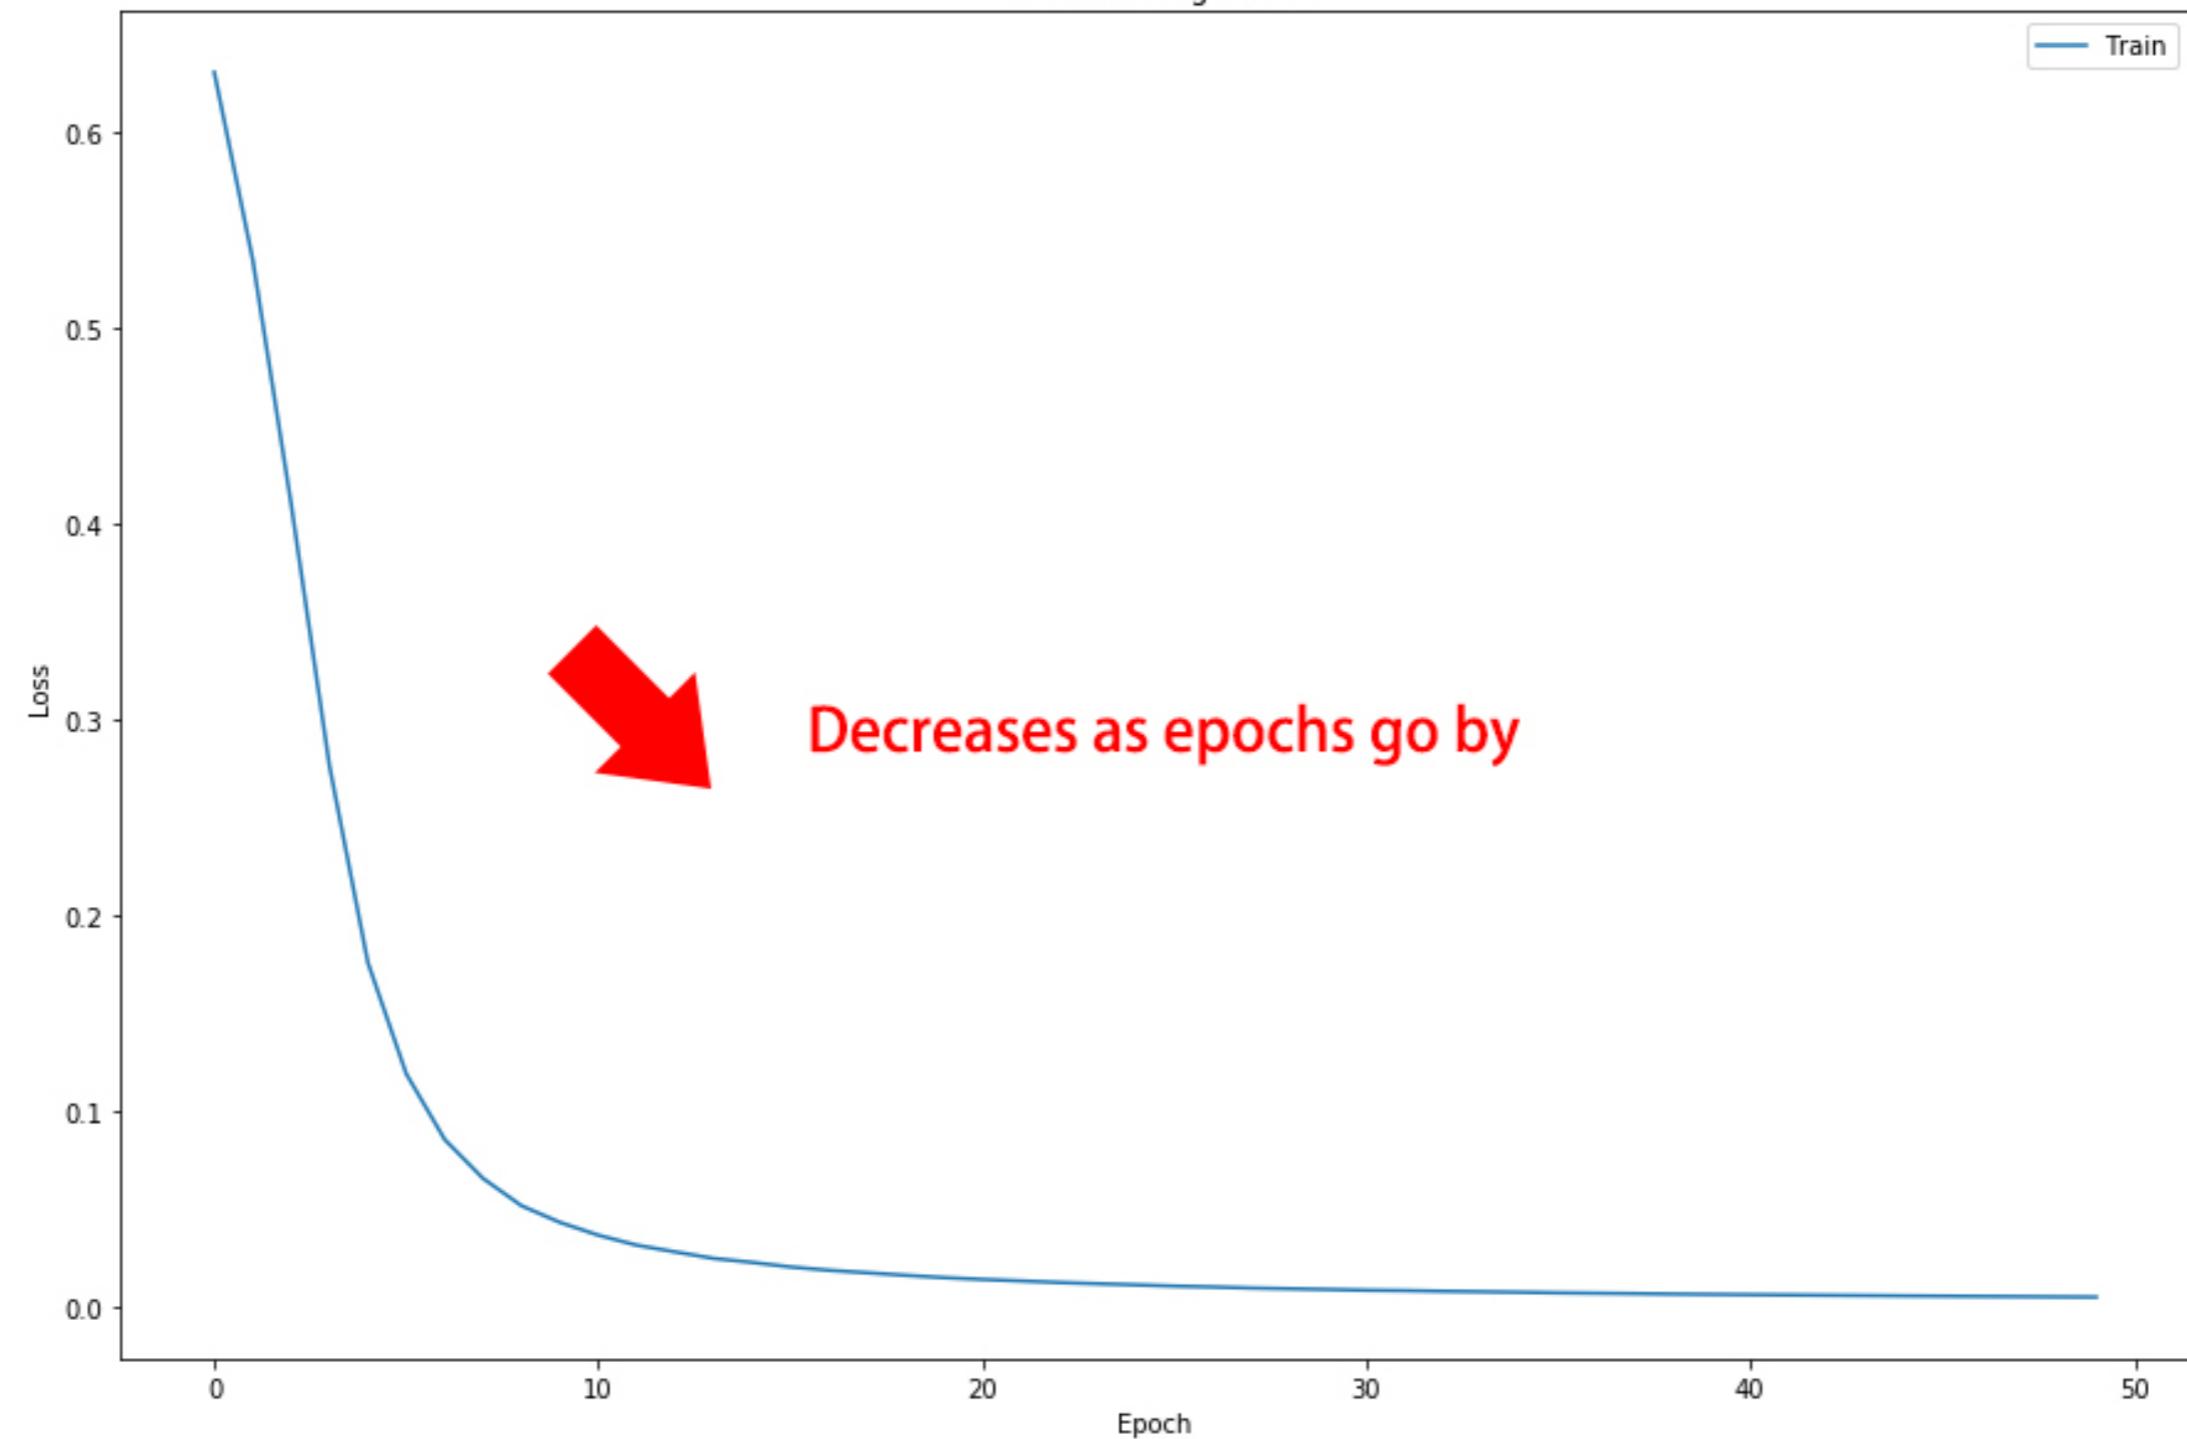
Loss Learning Curve

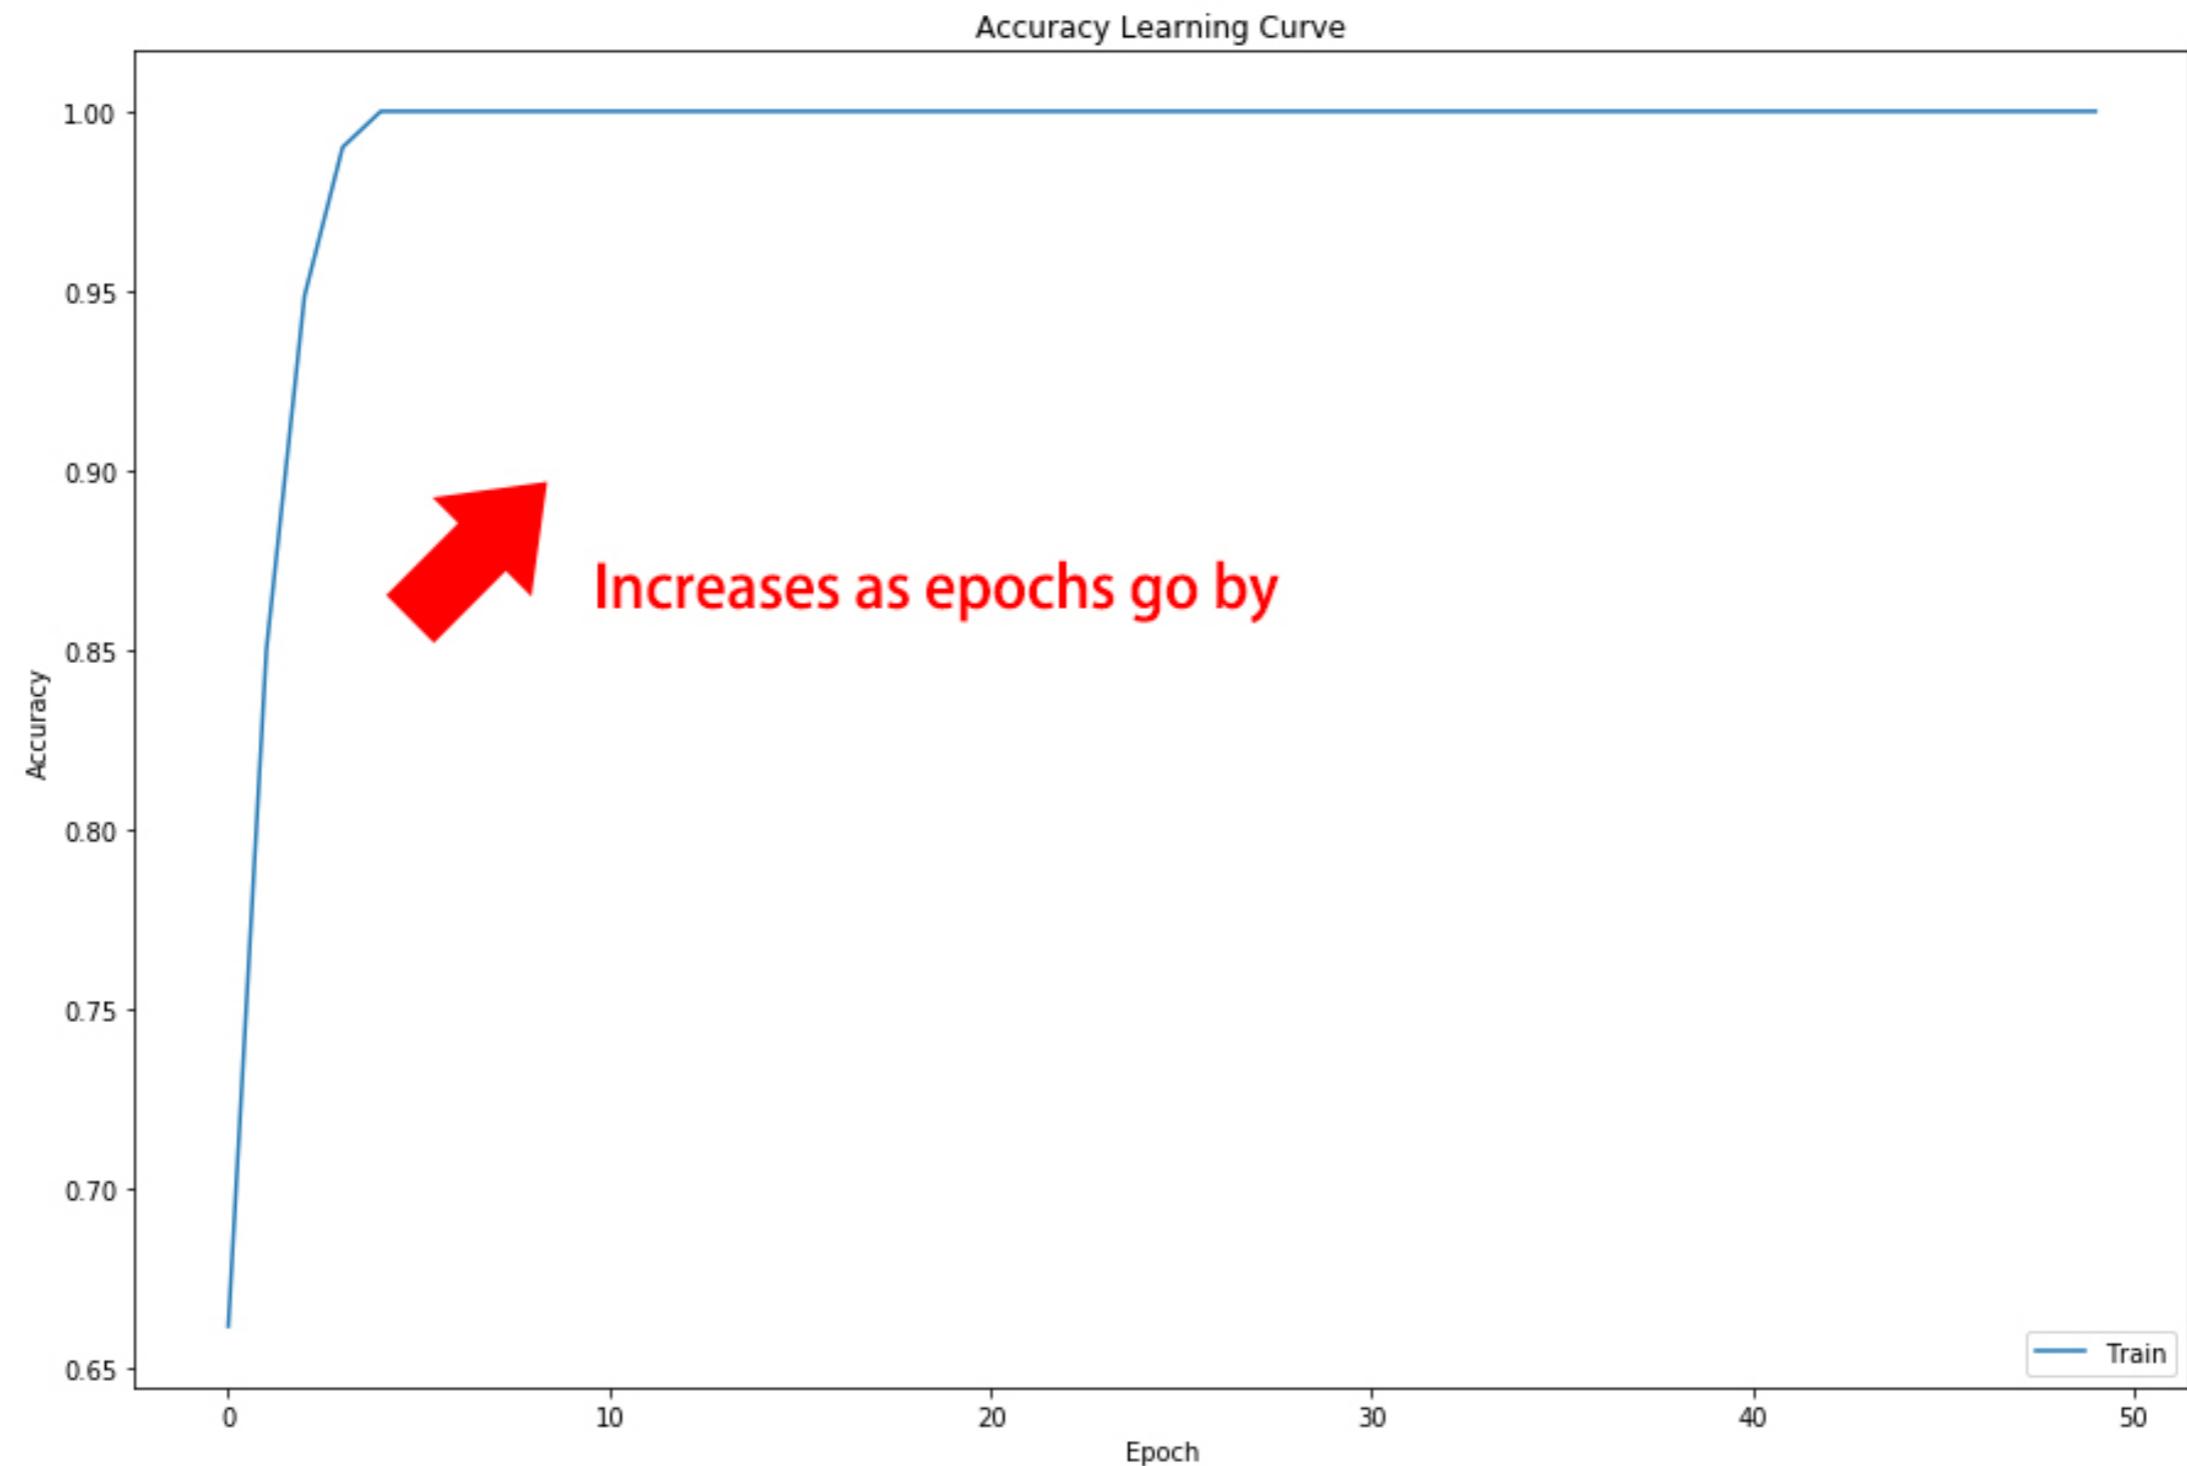


Accuracy Learning Curve

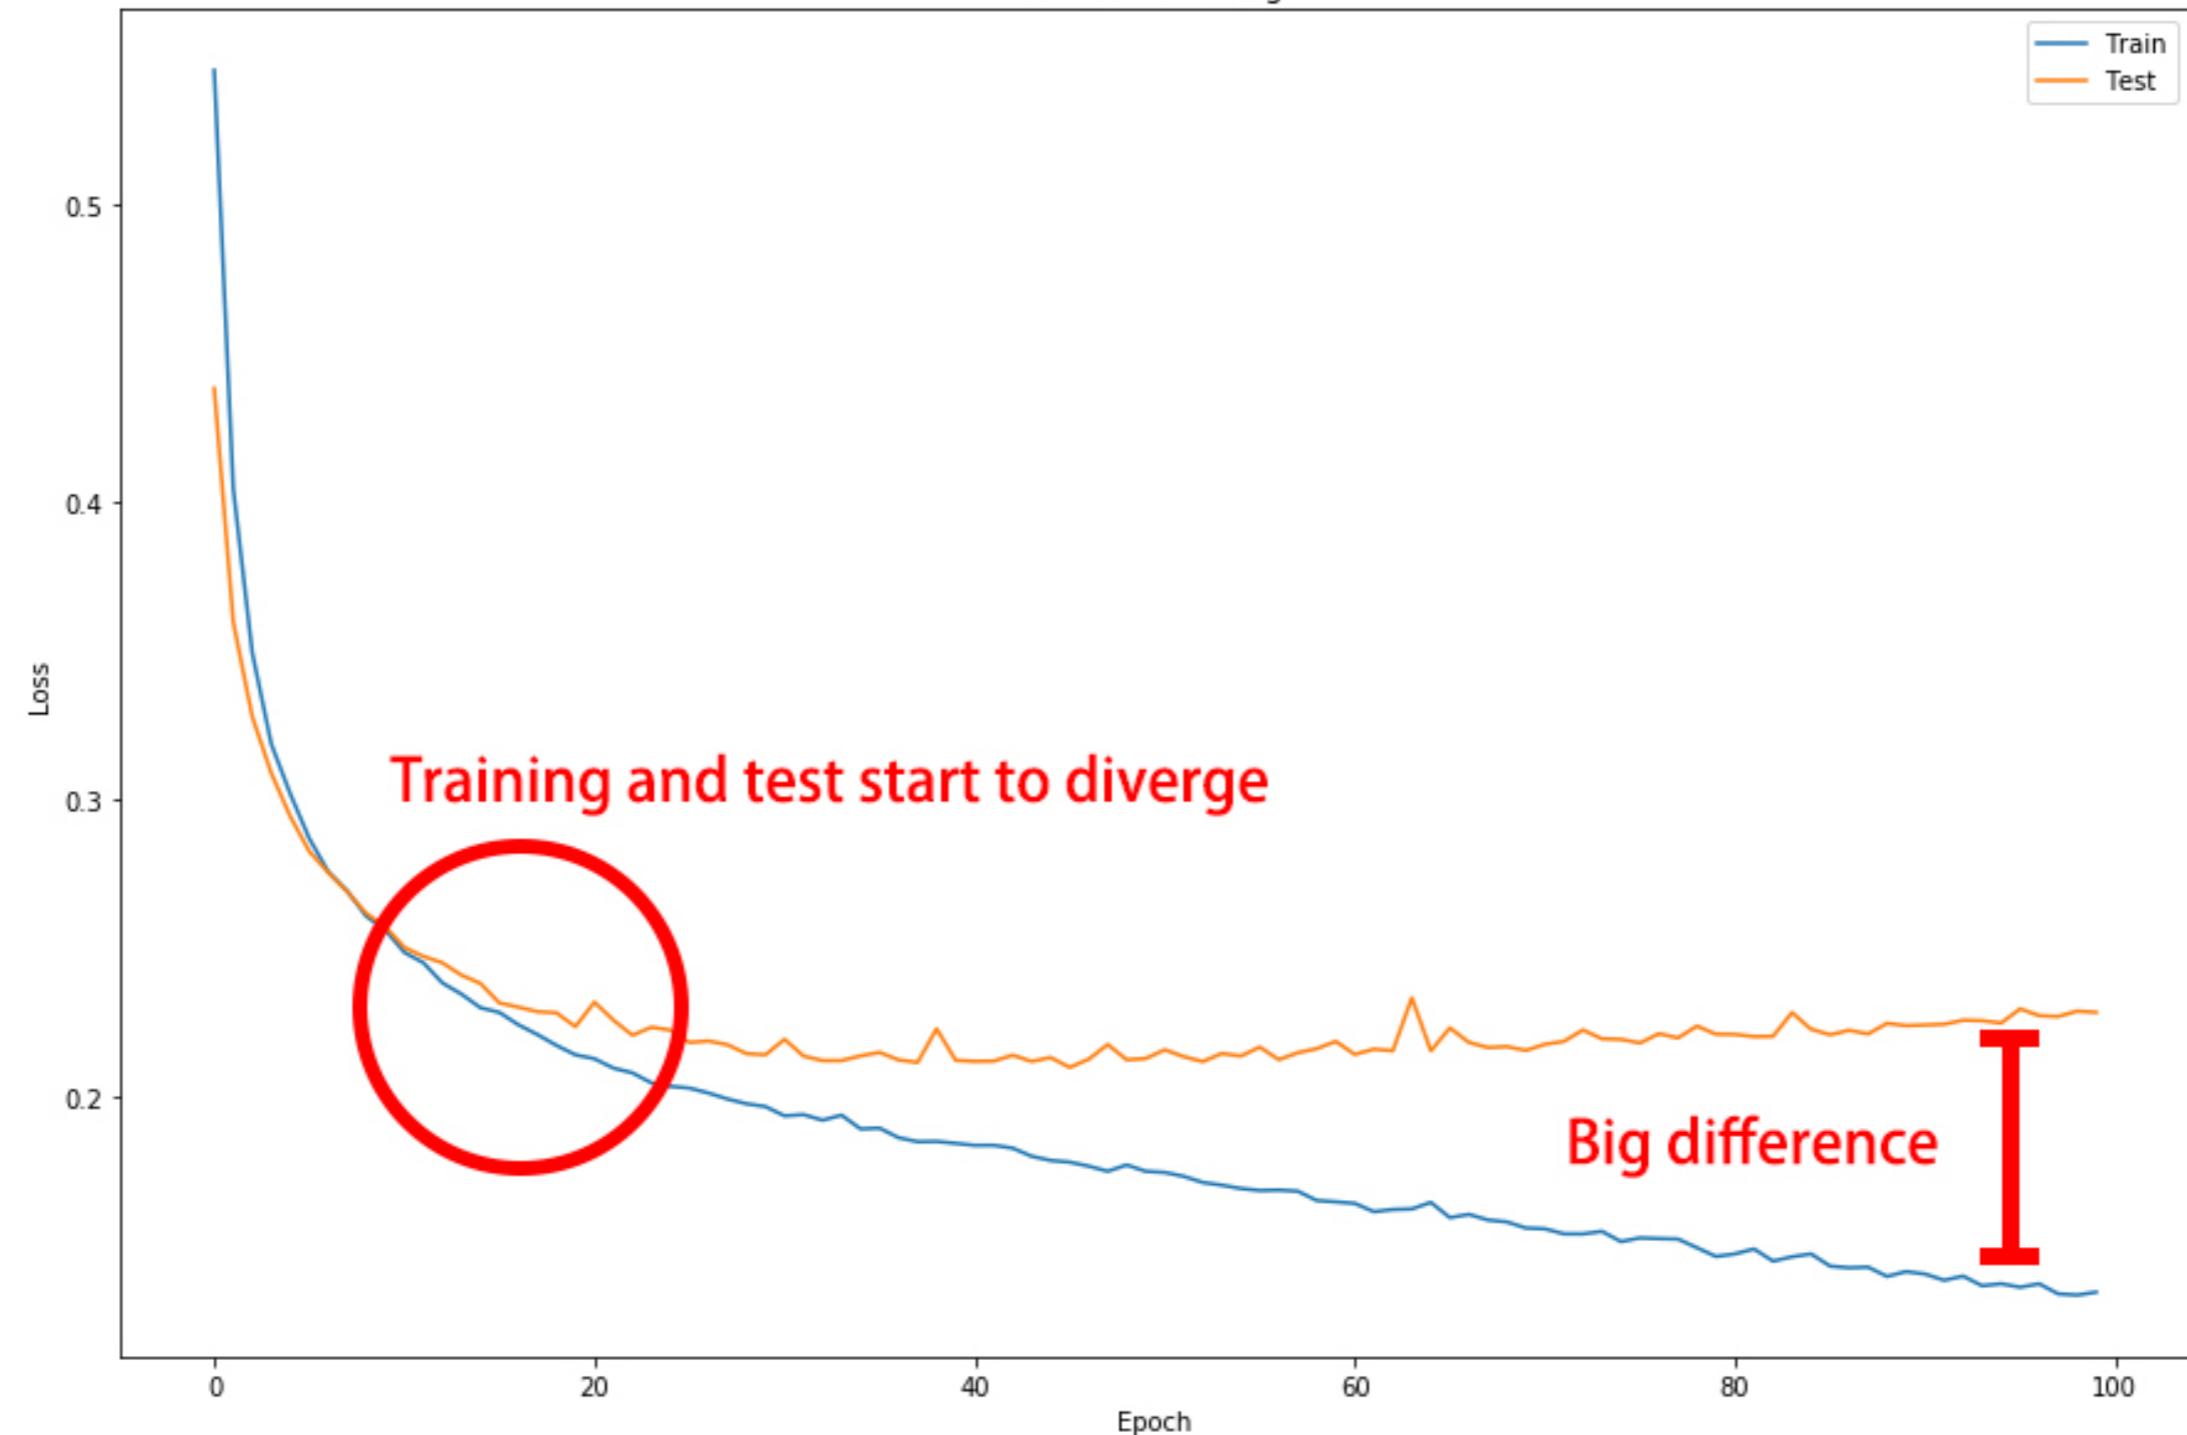


Loss Learning Curve

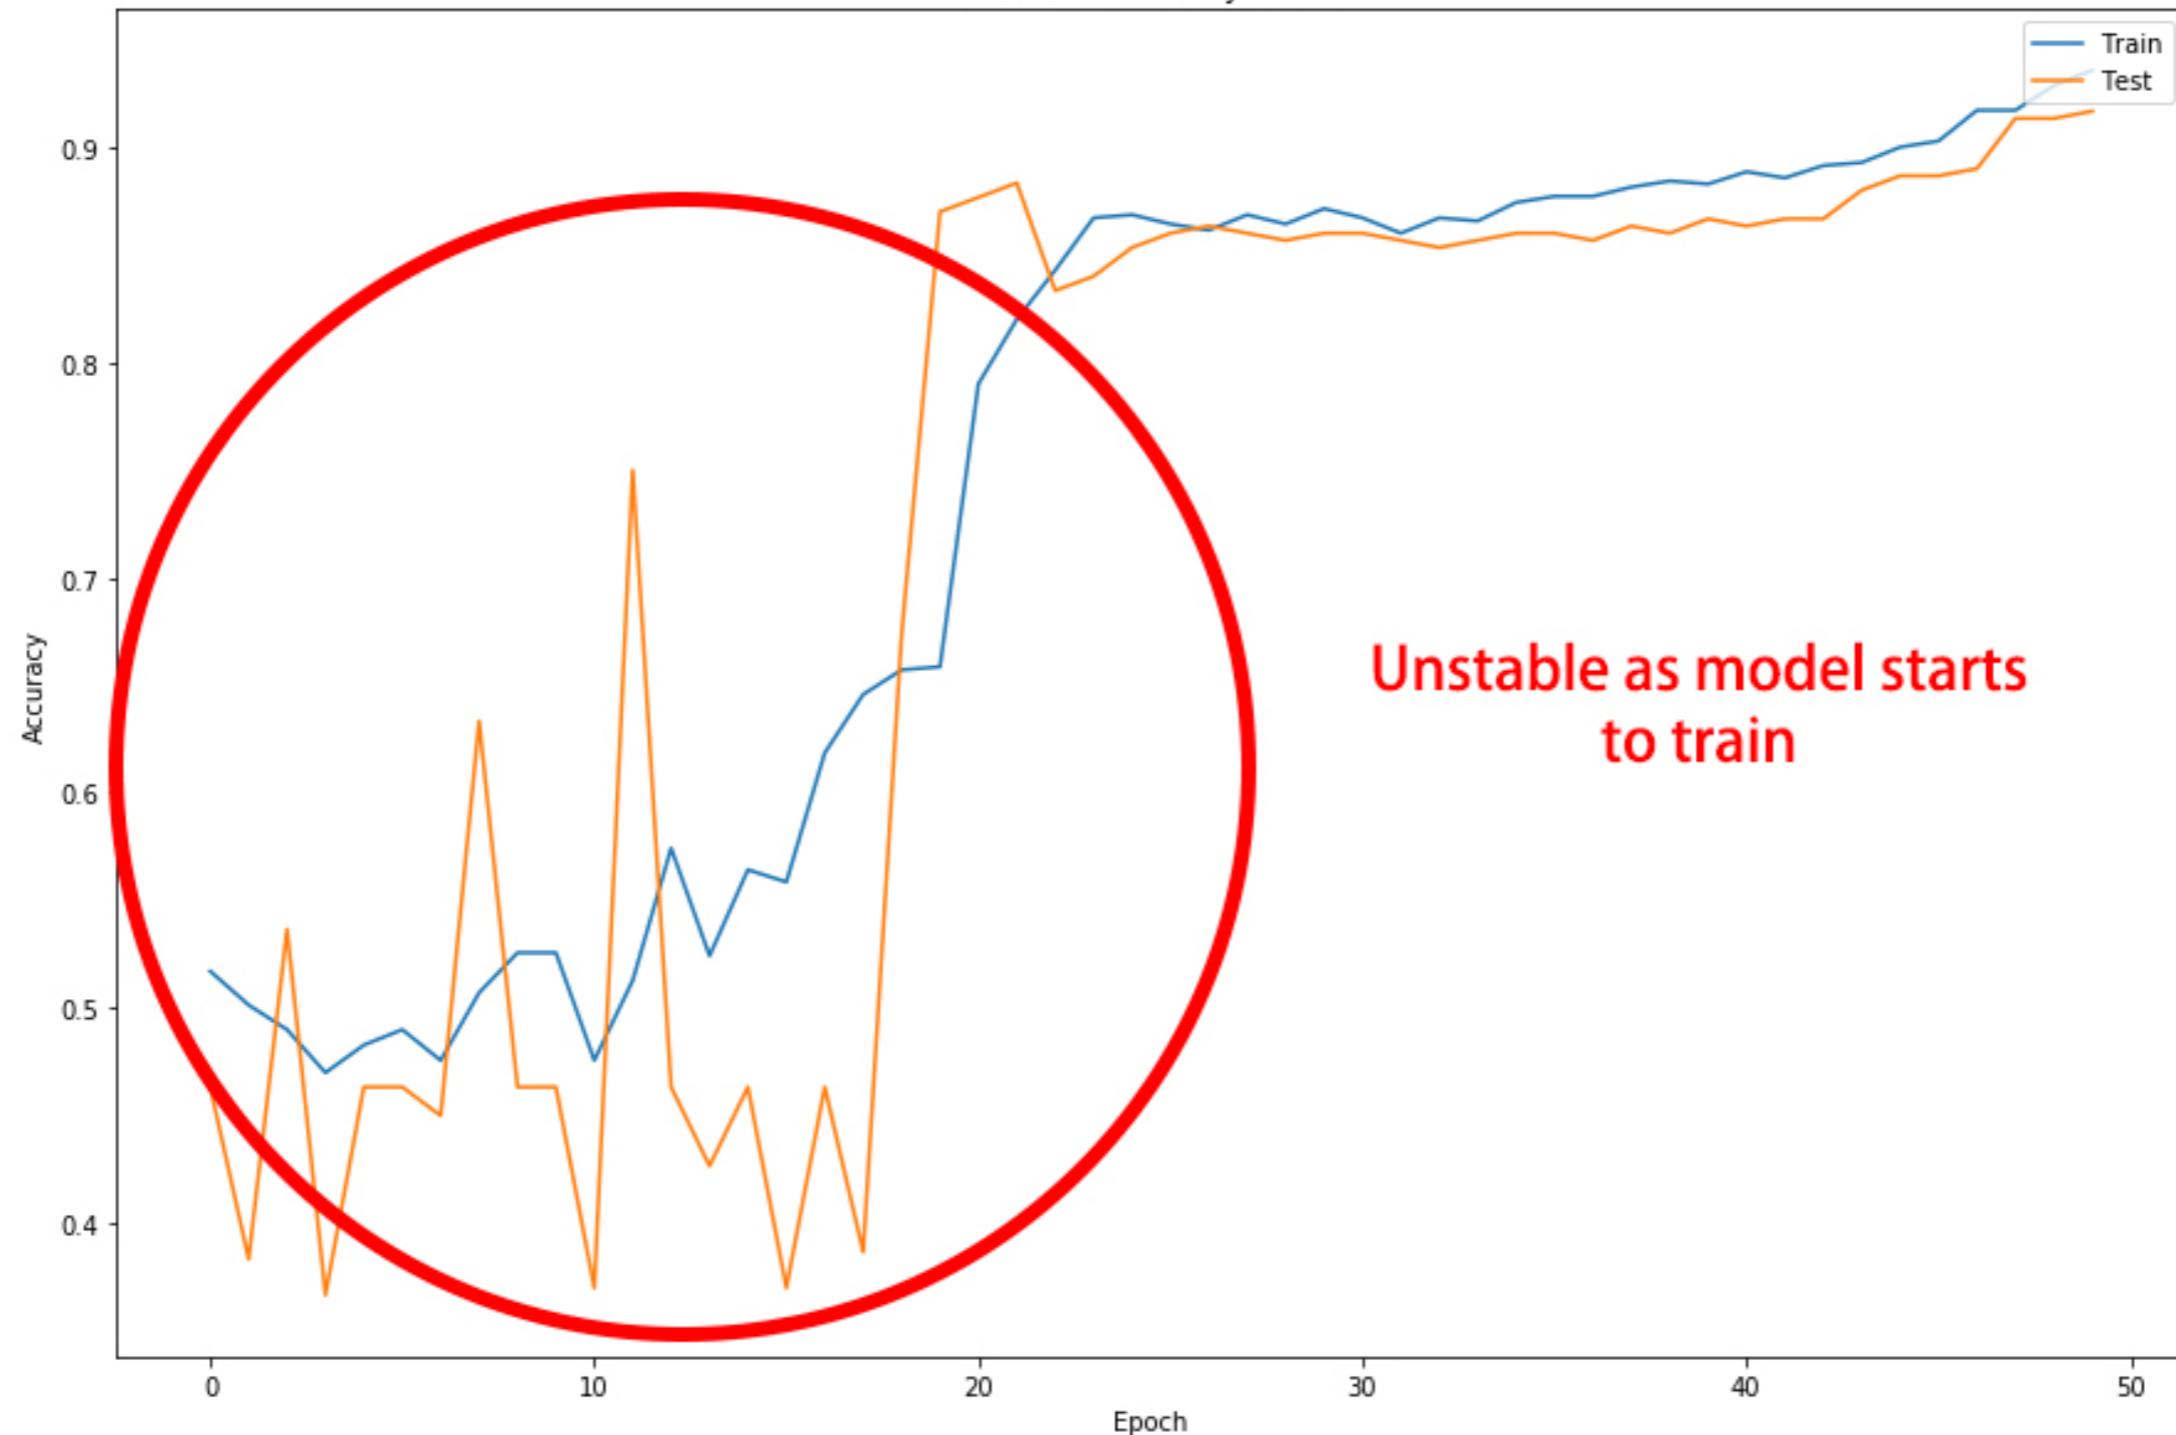


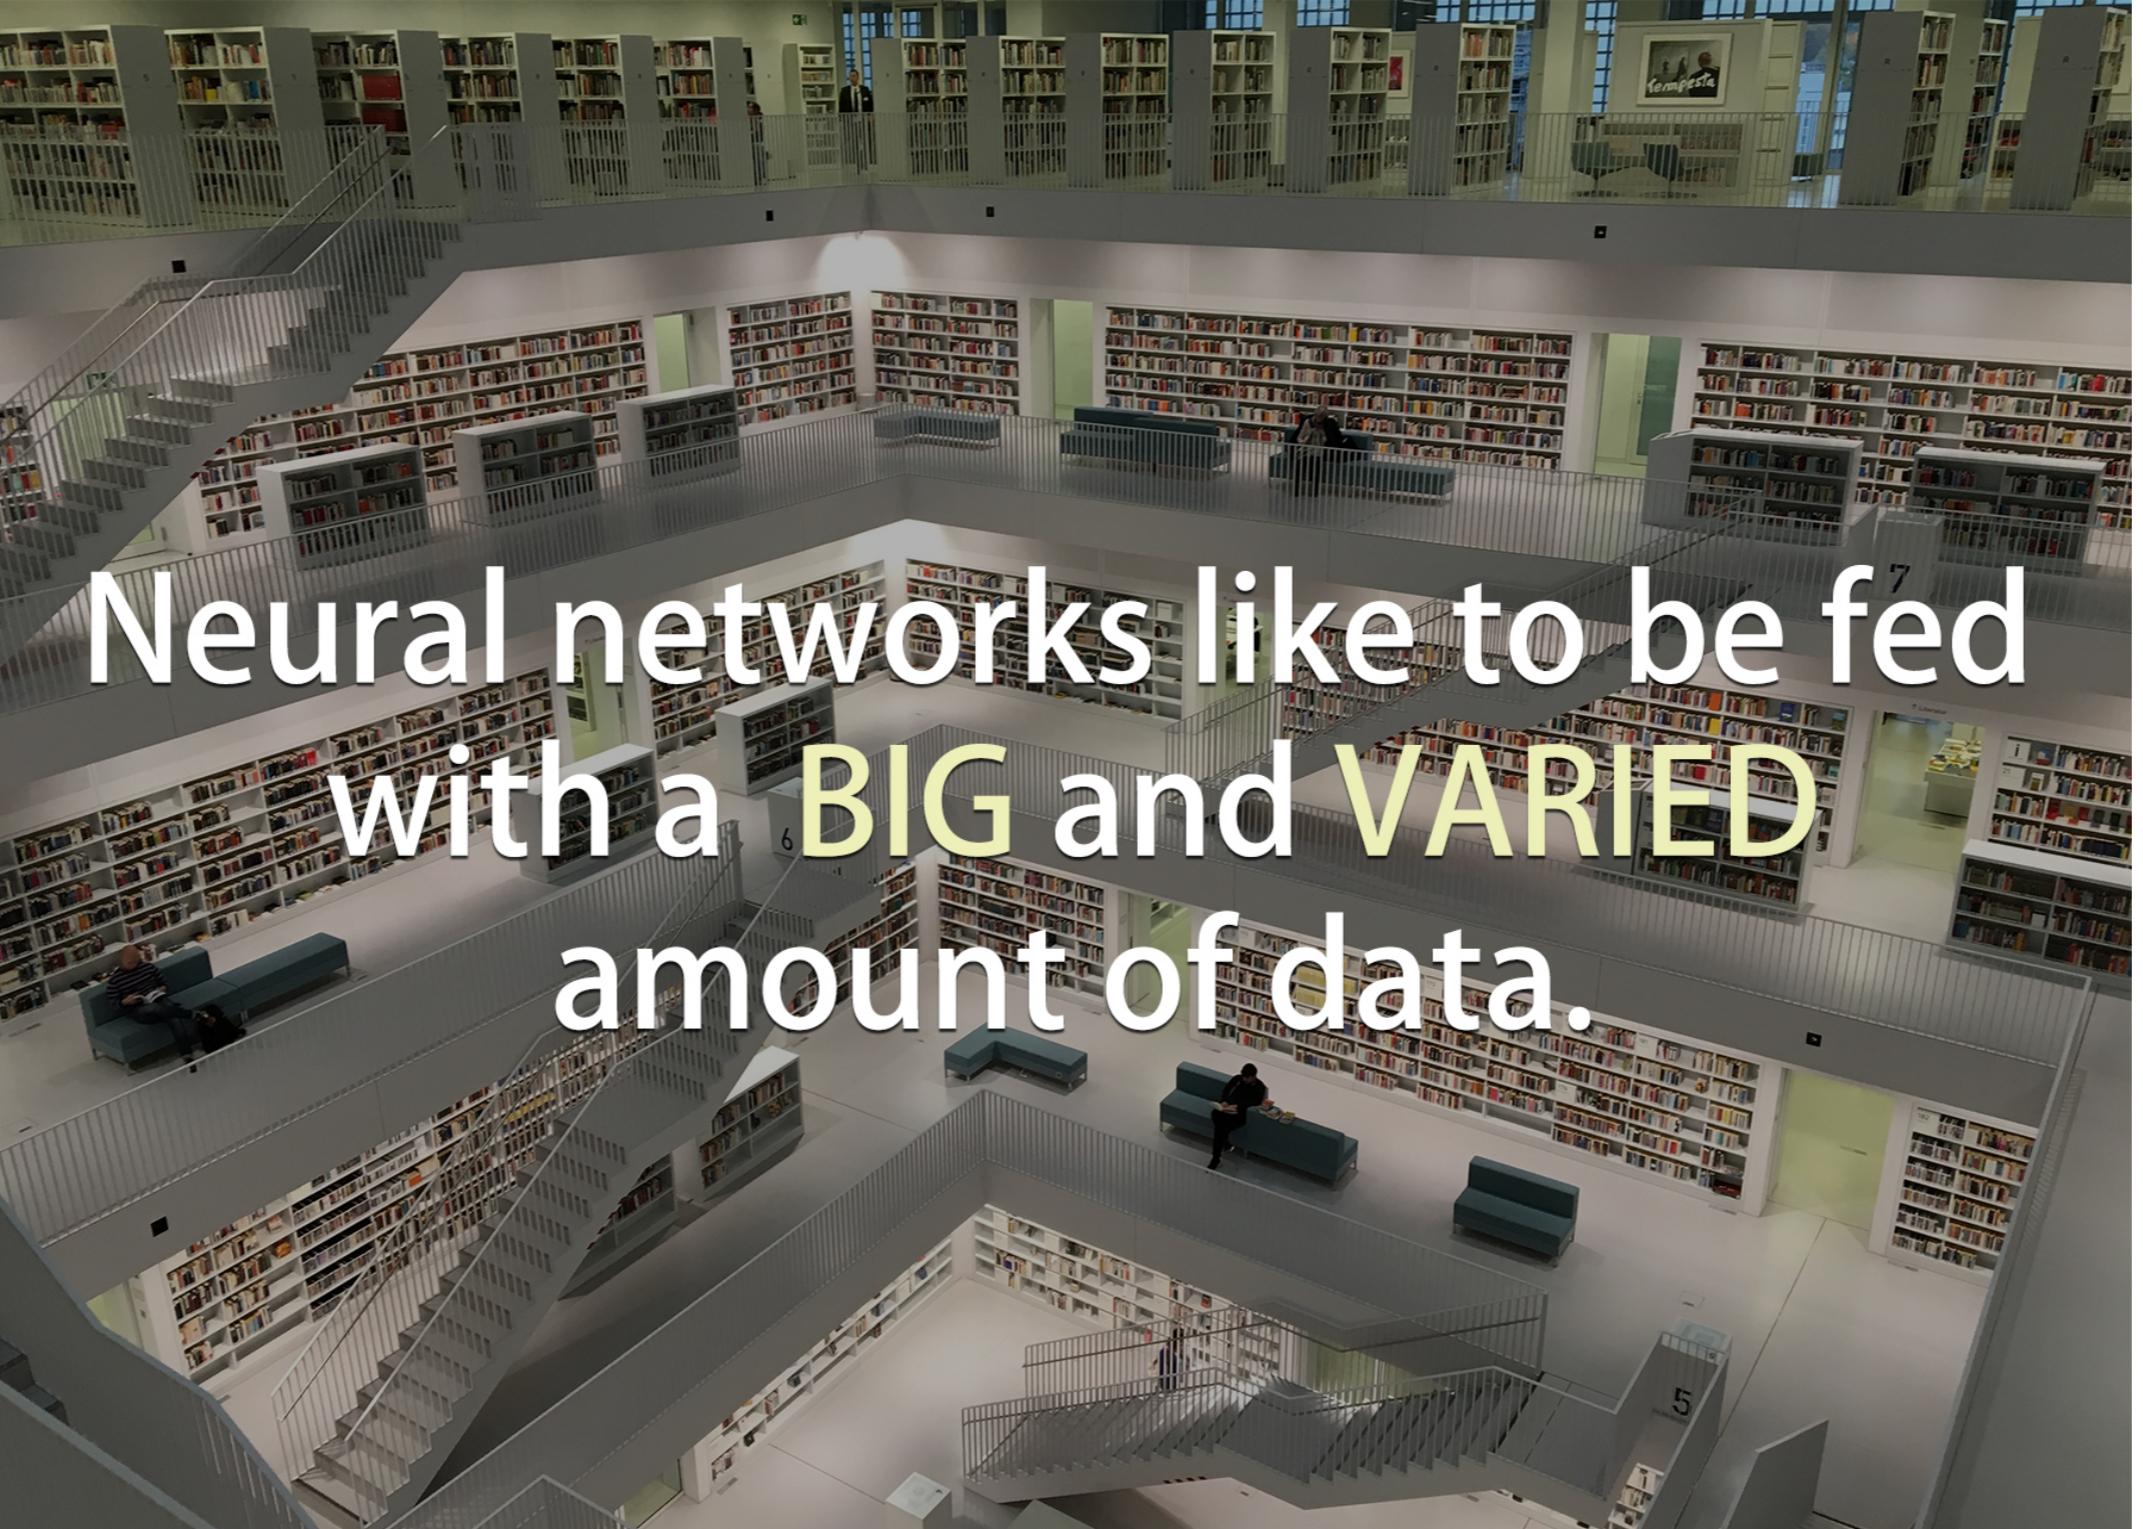


## Model Overfitting



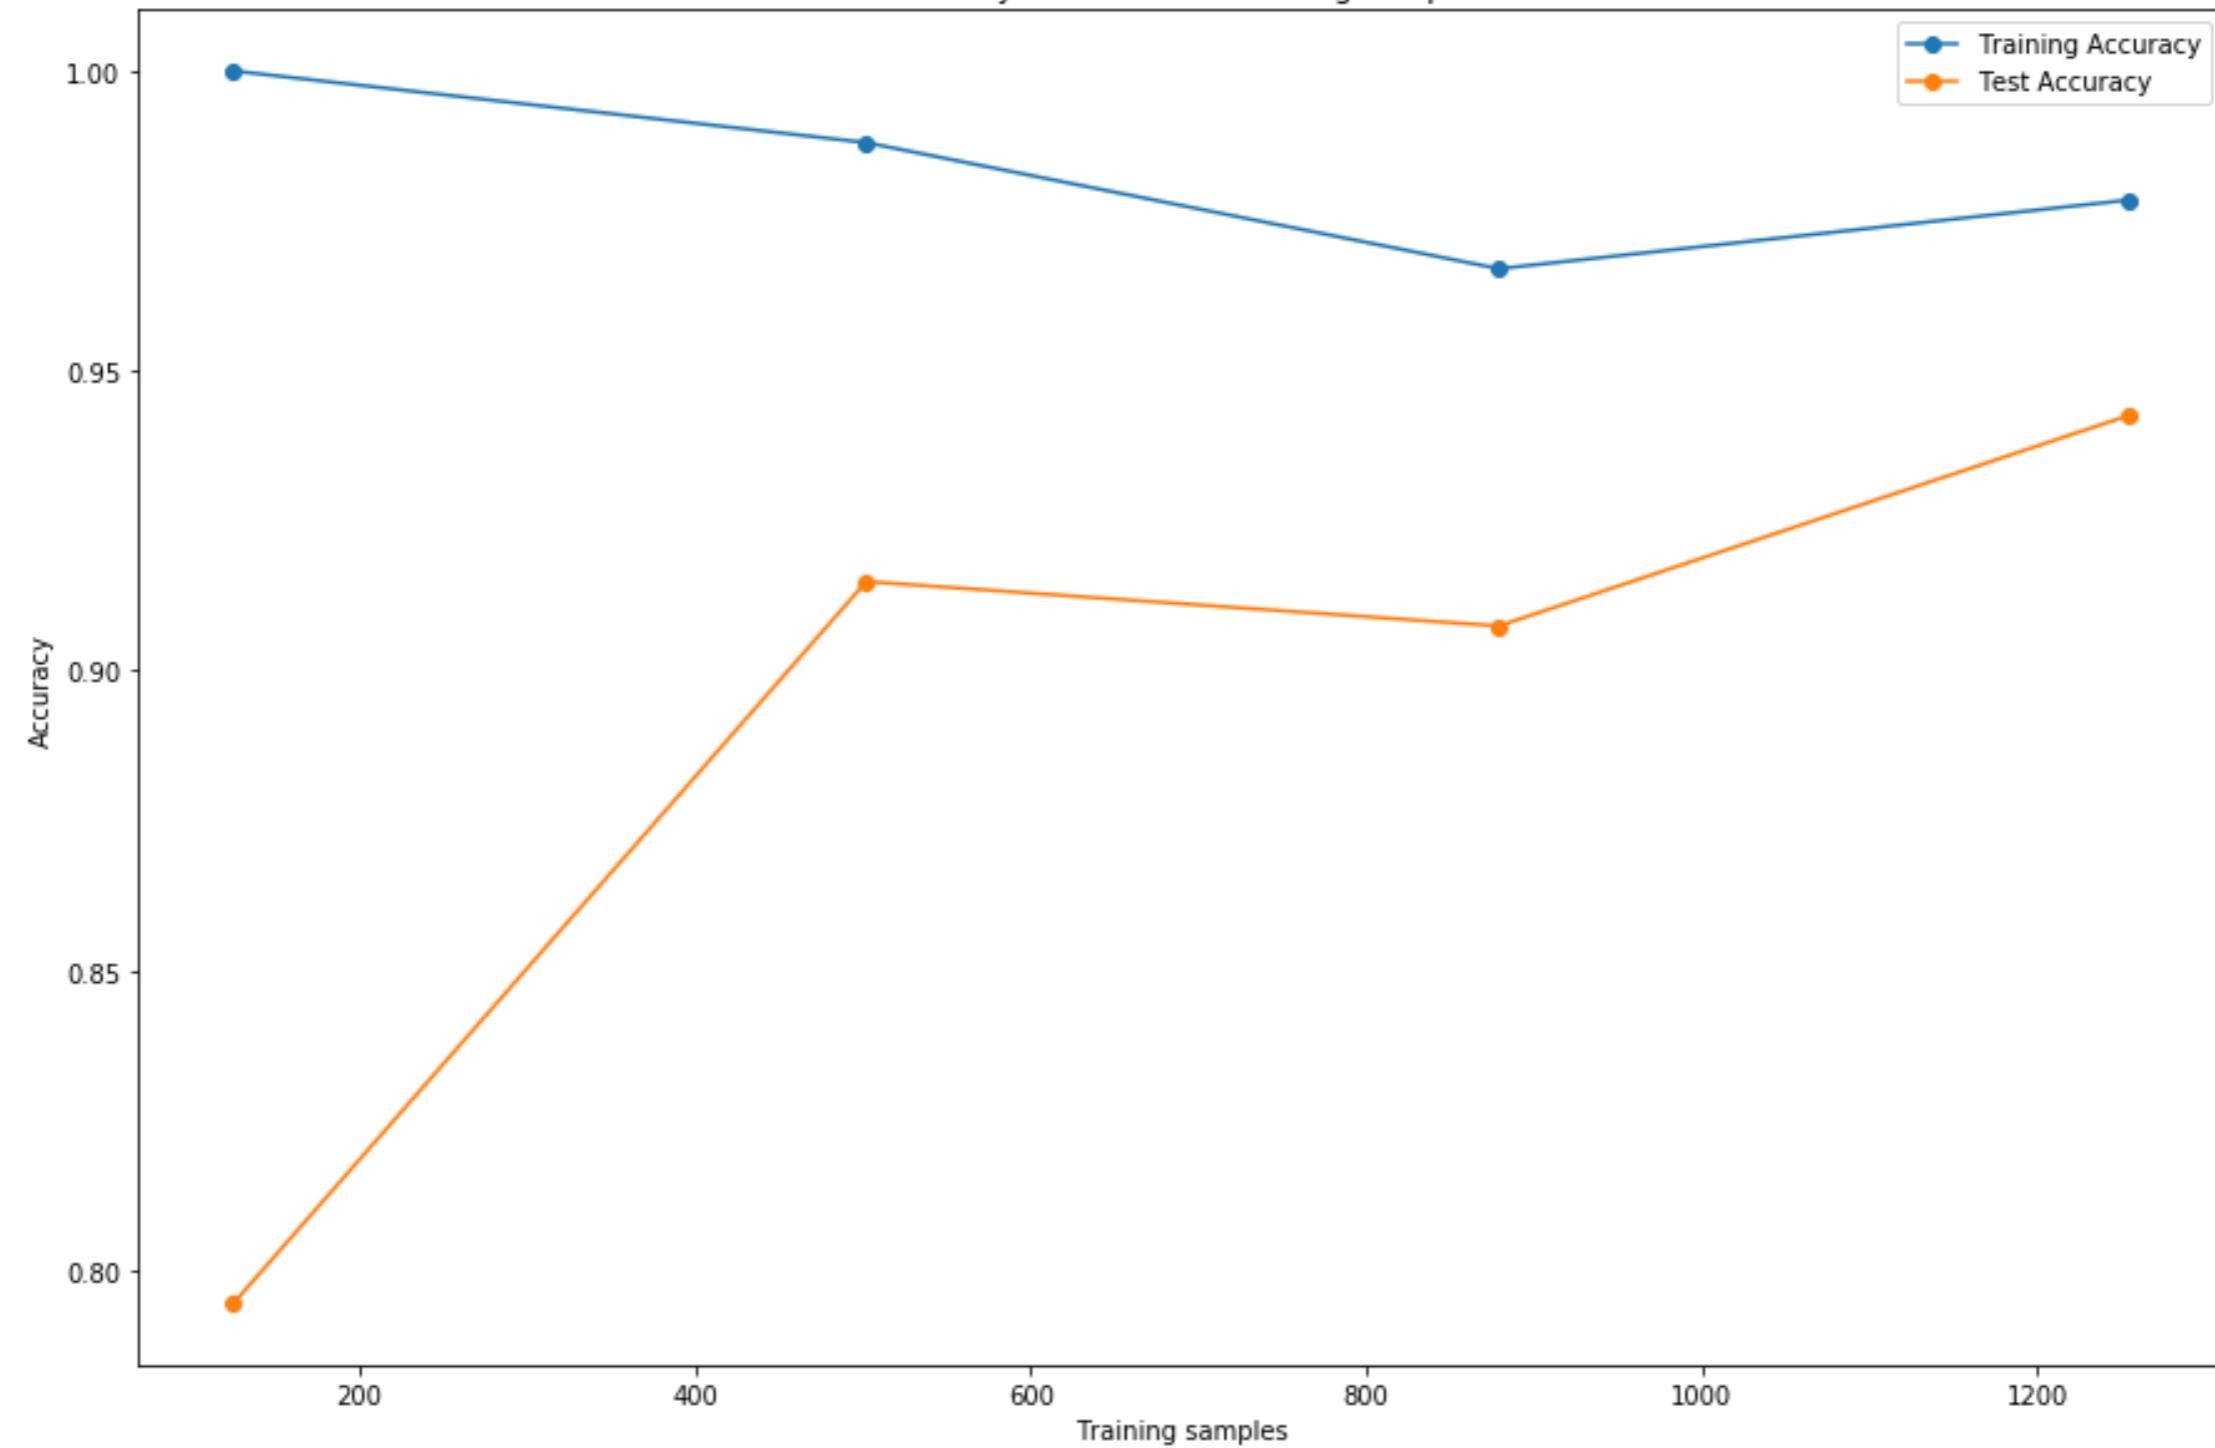
Unstable Accuracy Curve

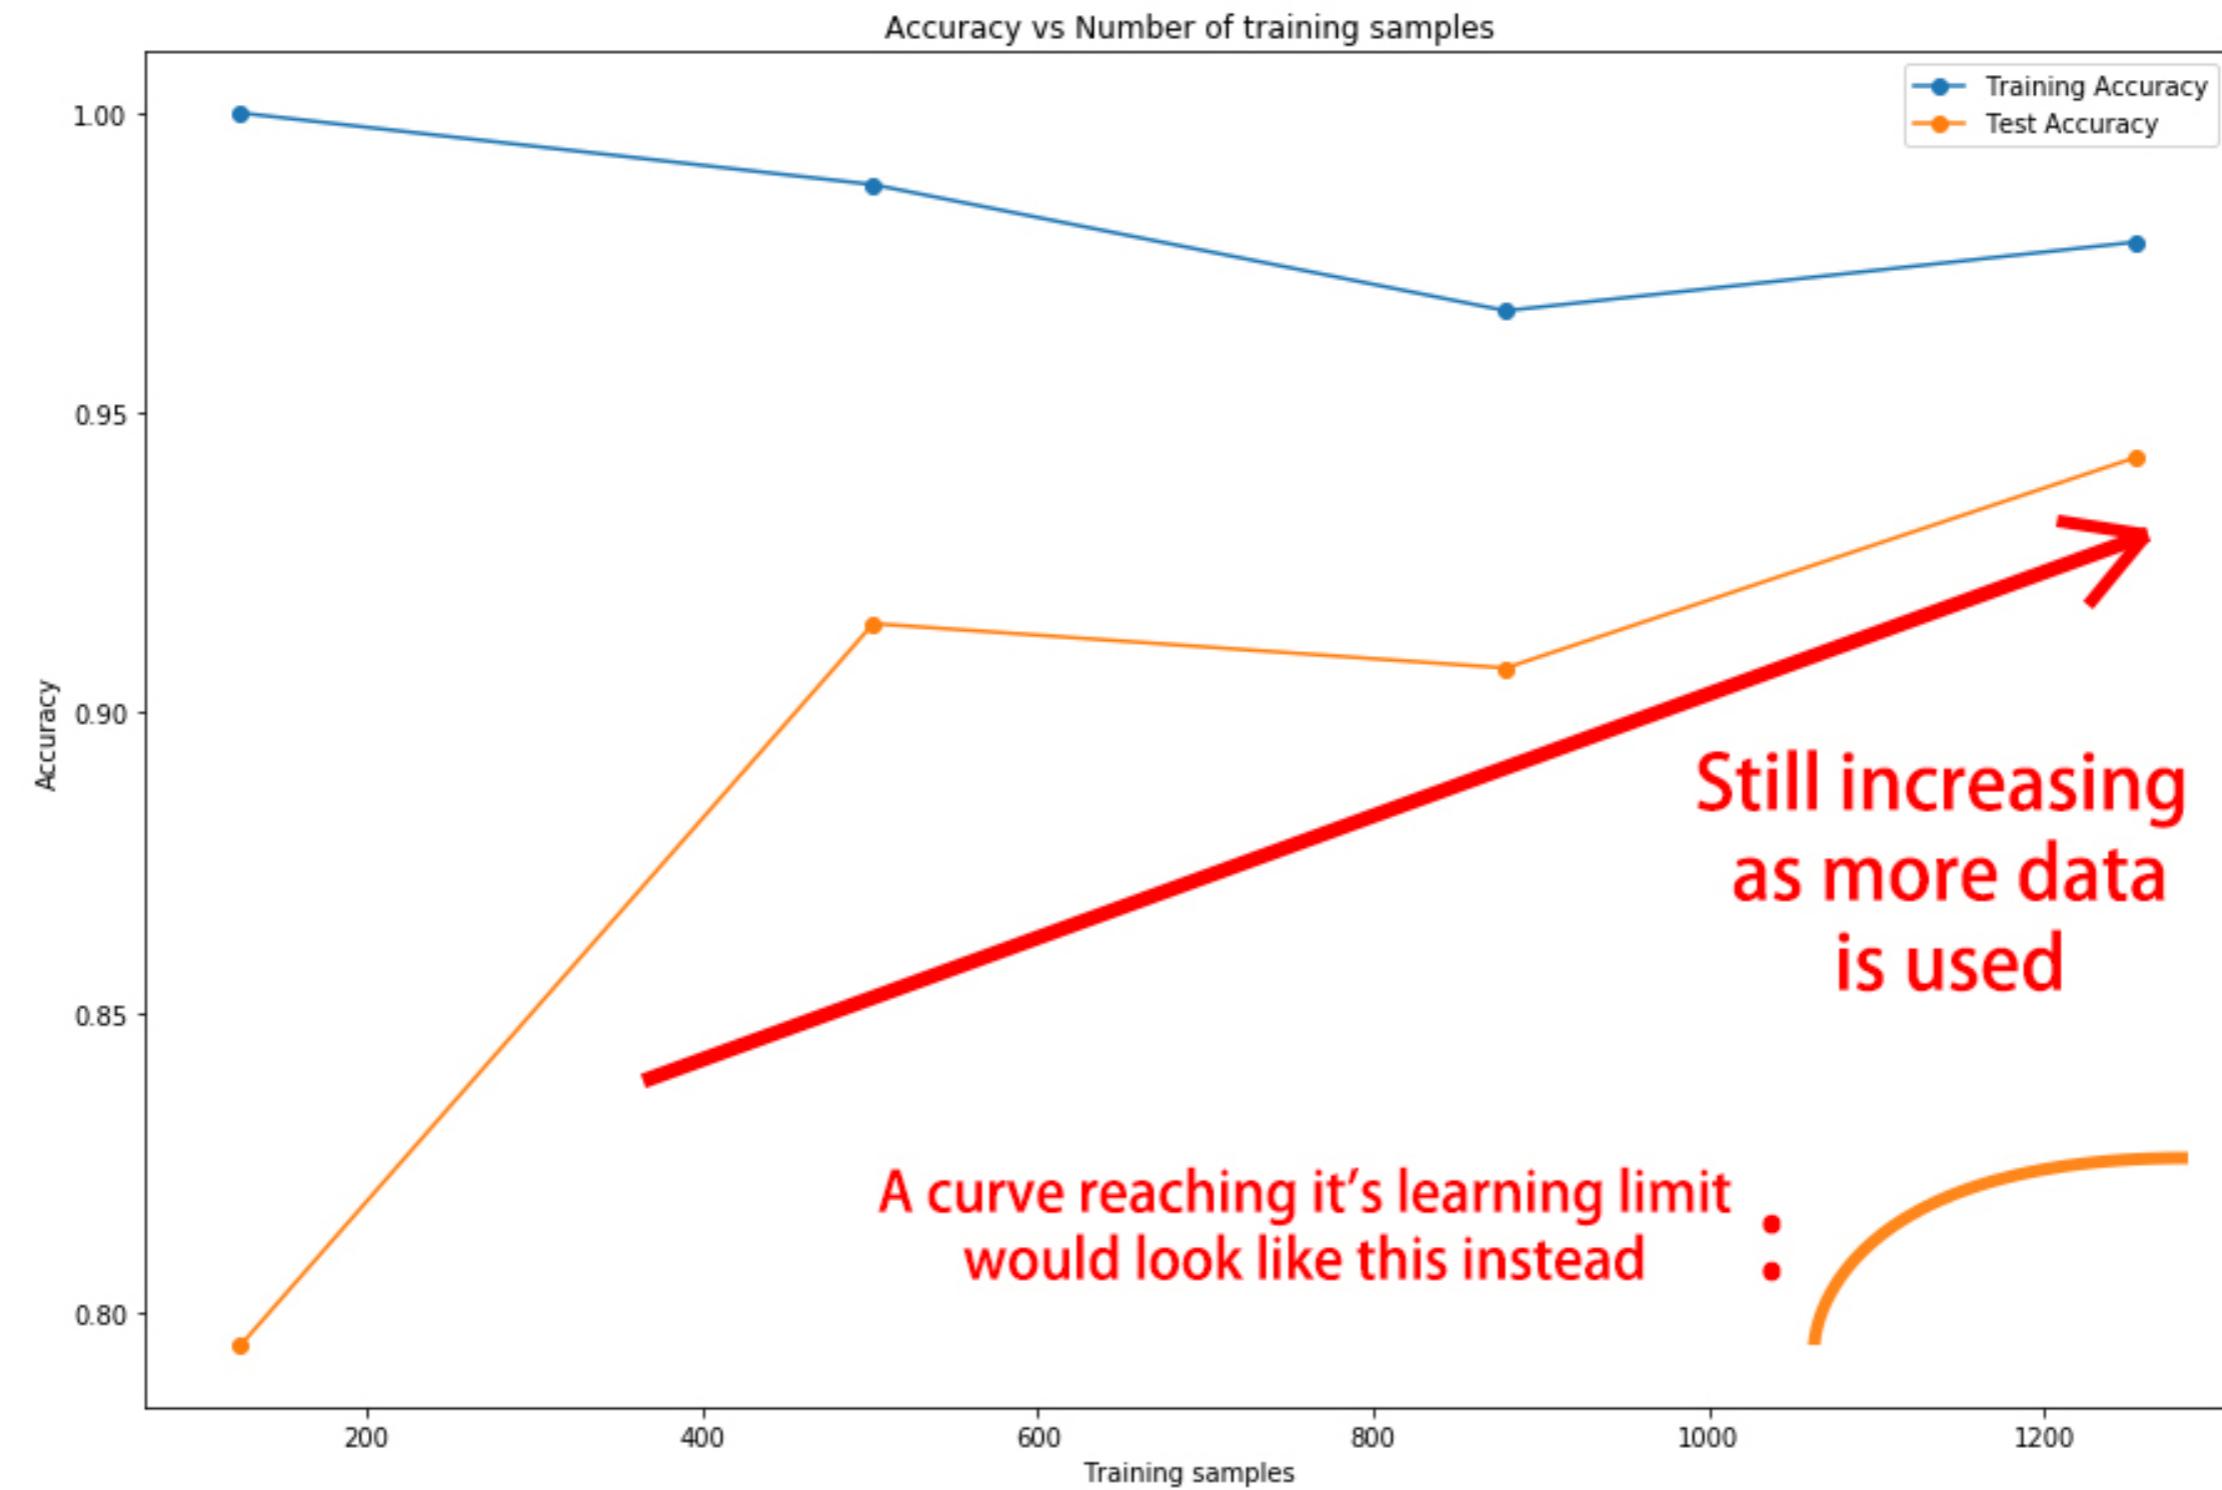




Neural networks like to be fed  
with a **BIG** and **VARIED**  
amount of data.

Accuracy vs Number of training samples





```
# Store initial model weights  
init_weights = model.get_weights()  
  
# Lists for storing accuracies  
train_accs = []  
tests_accs = []
```

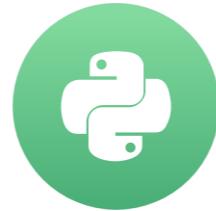
```
for train_size in train_sizes:  
    # Split a fraction according to train_size  
    X_train_frac, _, y_train_frac, _ =  
        train_test_split(X_train, y_train, train_size=train_size)  
    # Set model initial weights  
    model.set_weights(initial_weights)  
    # Fit model on the training set fraction  
    model.fit(X_train_frac, y_train_frac, epochs=100,  
              verbose=0,  
              callbacks=[EarlyStopping(monitor='loss', patience=1)])  
    # Get the accuracy for this training set fraction  
    train_acc = model.evaluate(X_train_frac, y_train_frac, verbose=0)[1]  
    train_accs.append(train_acc)  
    # Get the accuracy on the whole test set  
    test_acc = model.evaluate(X_test, y_test, verbose=0)[1]  
    test_accs.append(test_acc)  
    print("Done with size: ", train_size)
```

**Time to dominate all  
curves!**

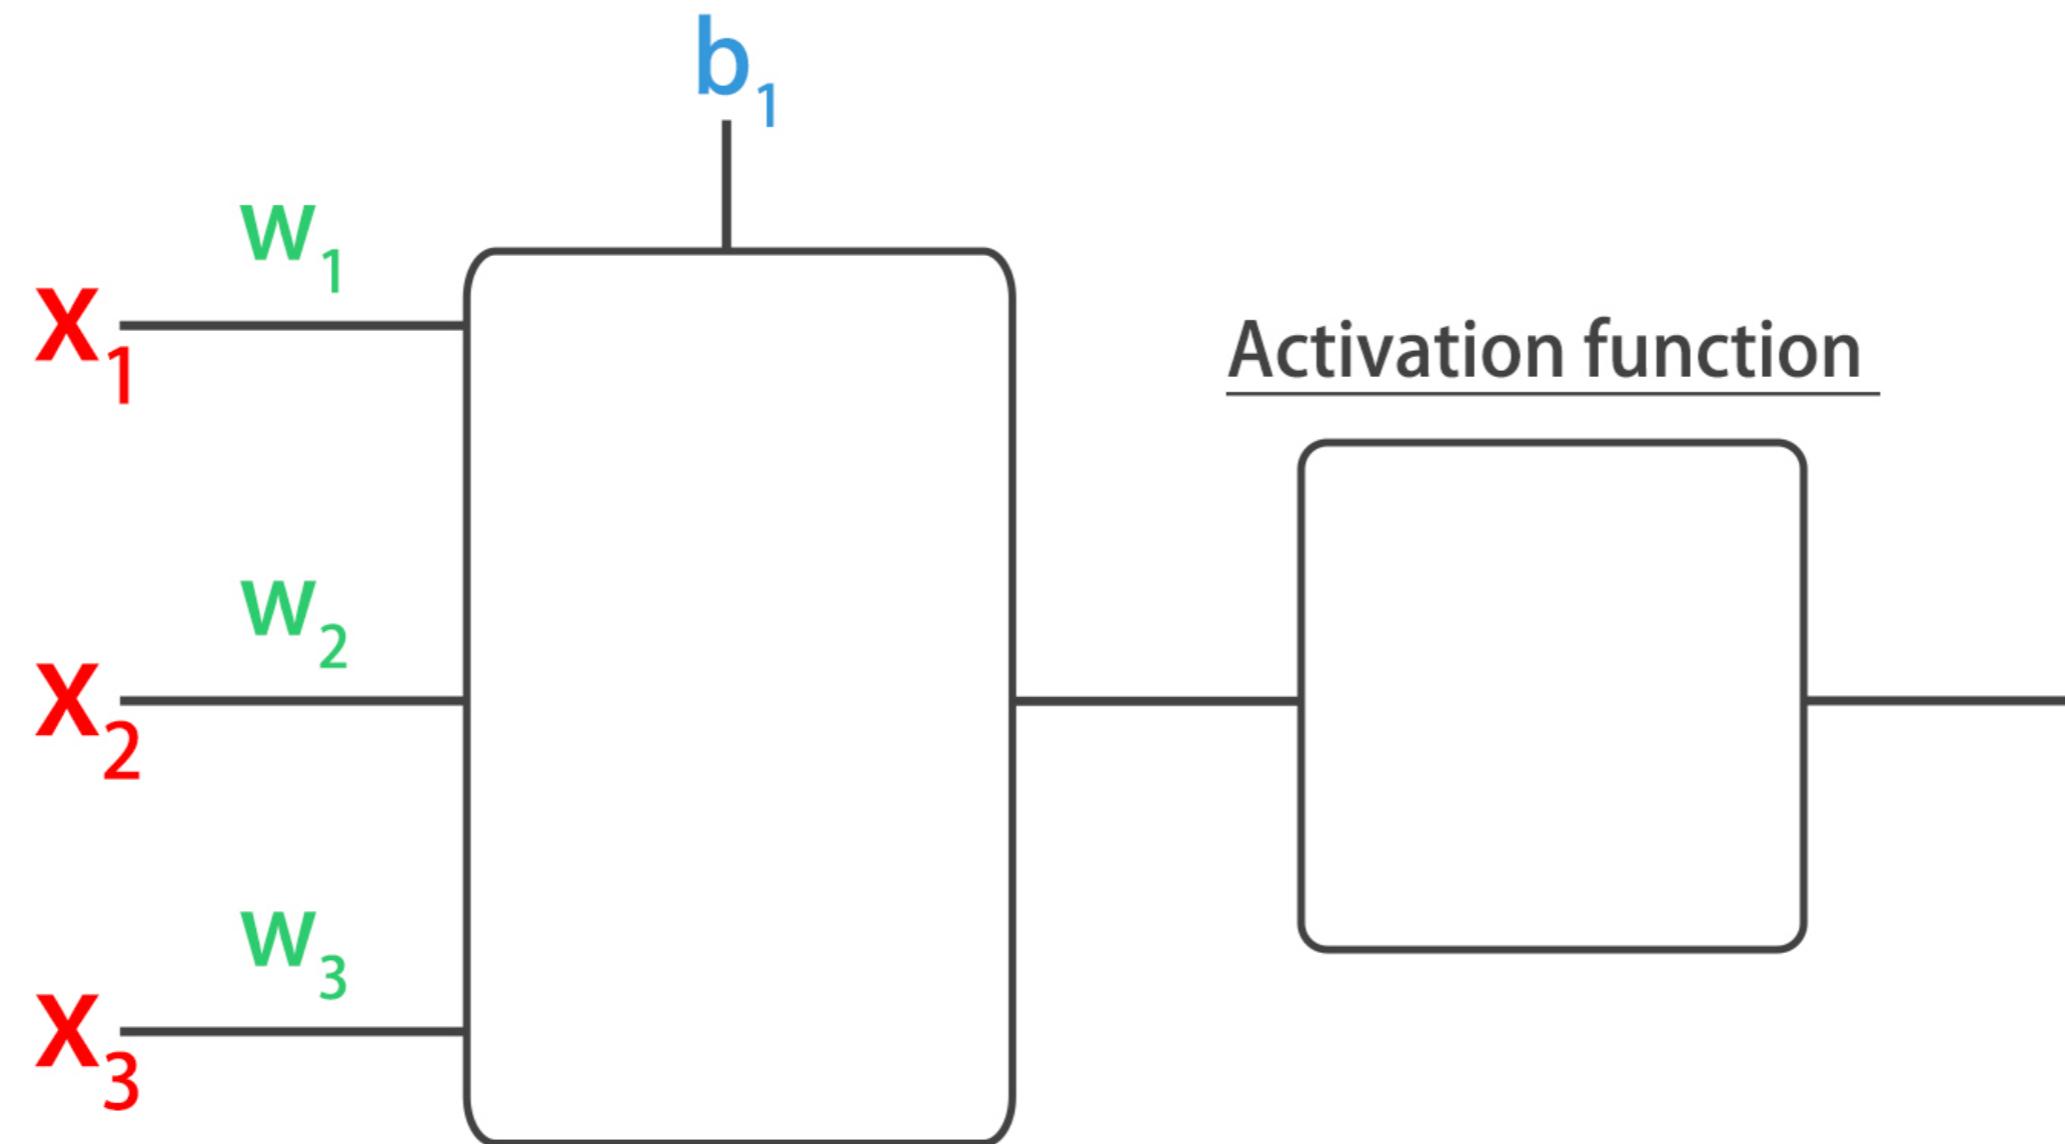
**INTRODUCTION TO DEEP LEARNING WITH KERAS**

# Activation functions

INTRODUCTION TO DEEP LEARNING WITH KERAS

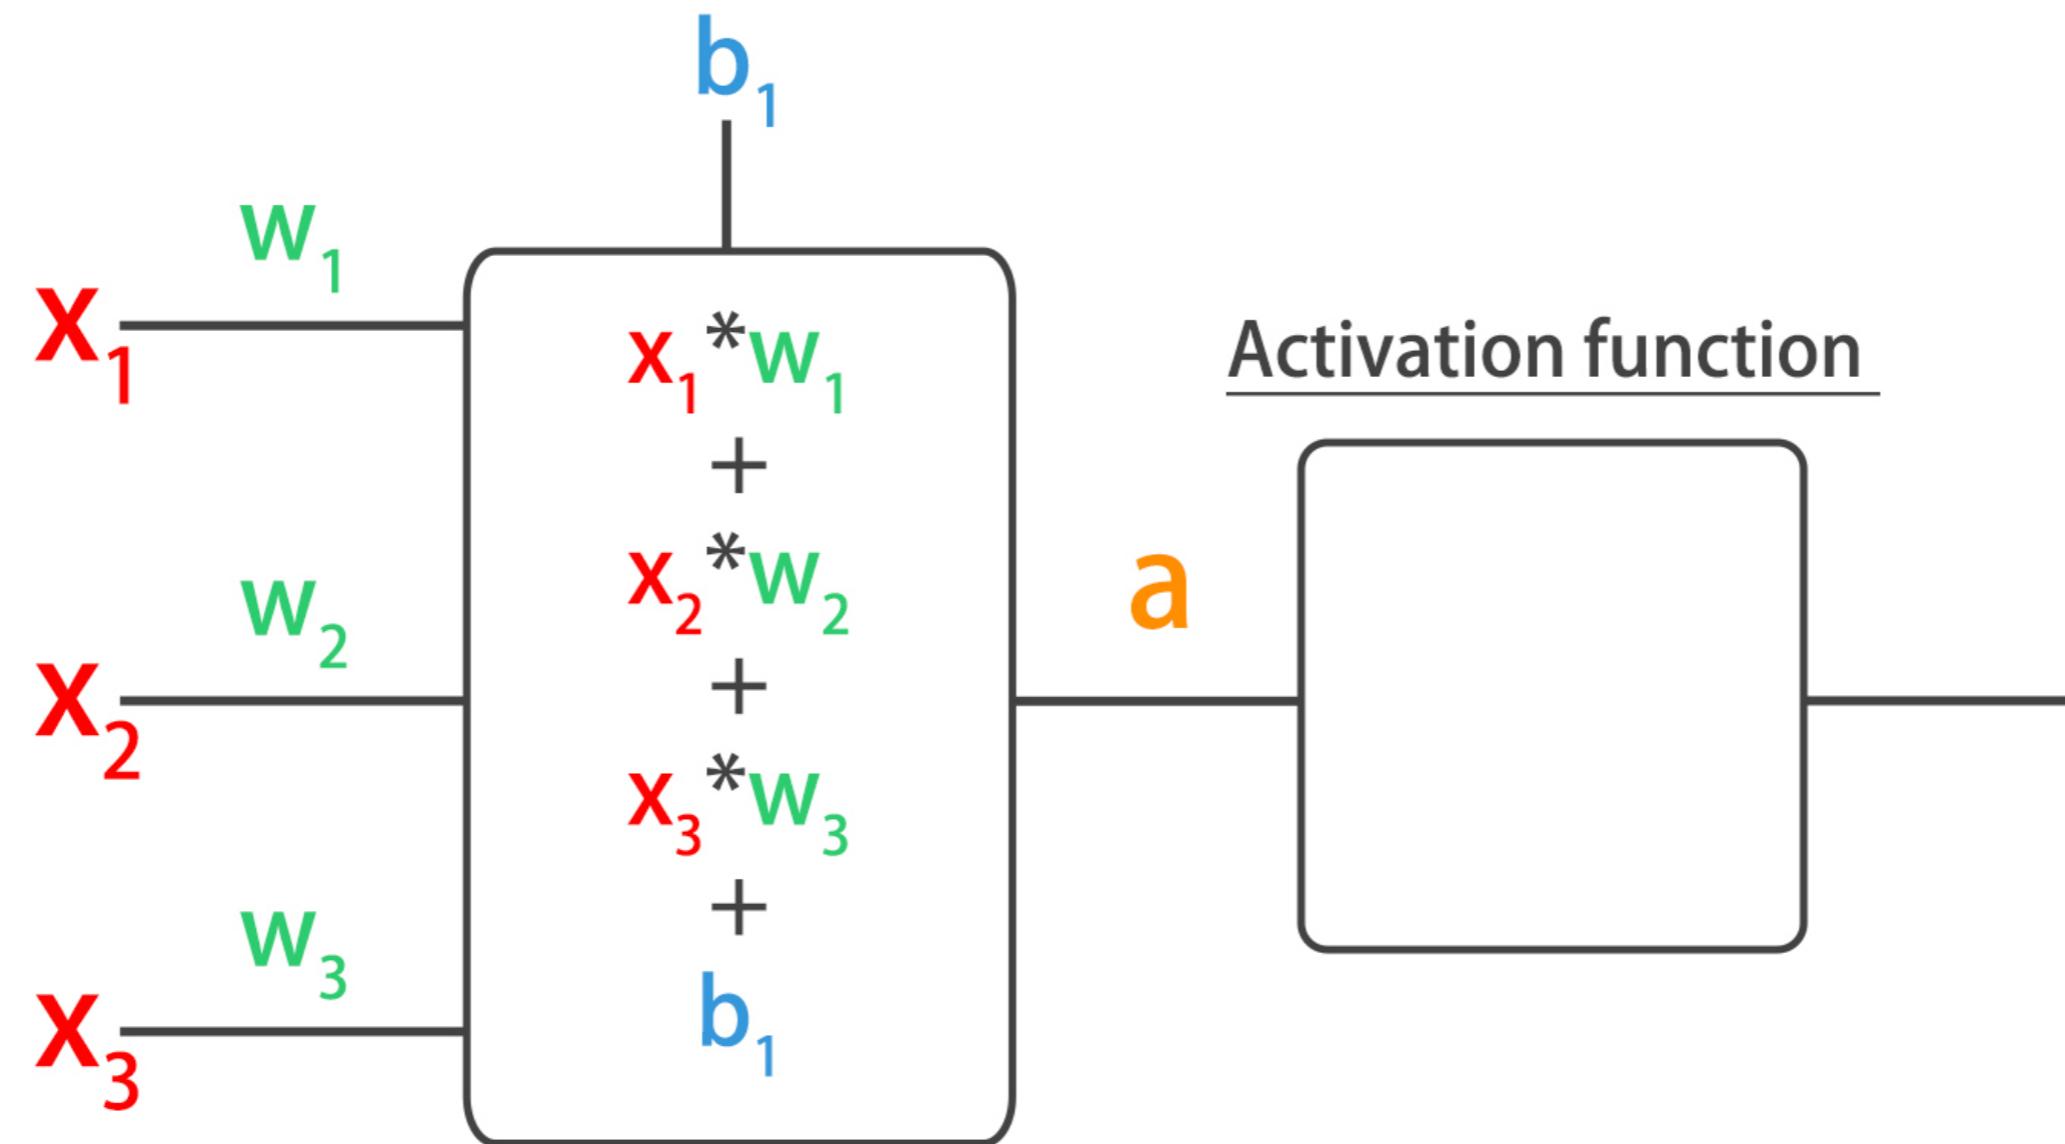


**Miguel Esteban**  
Data Scientist & Founder

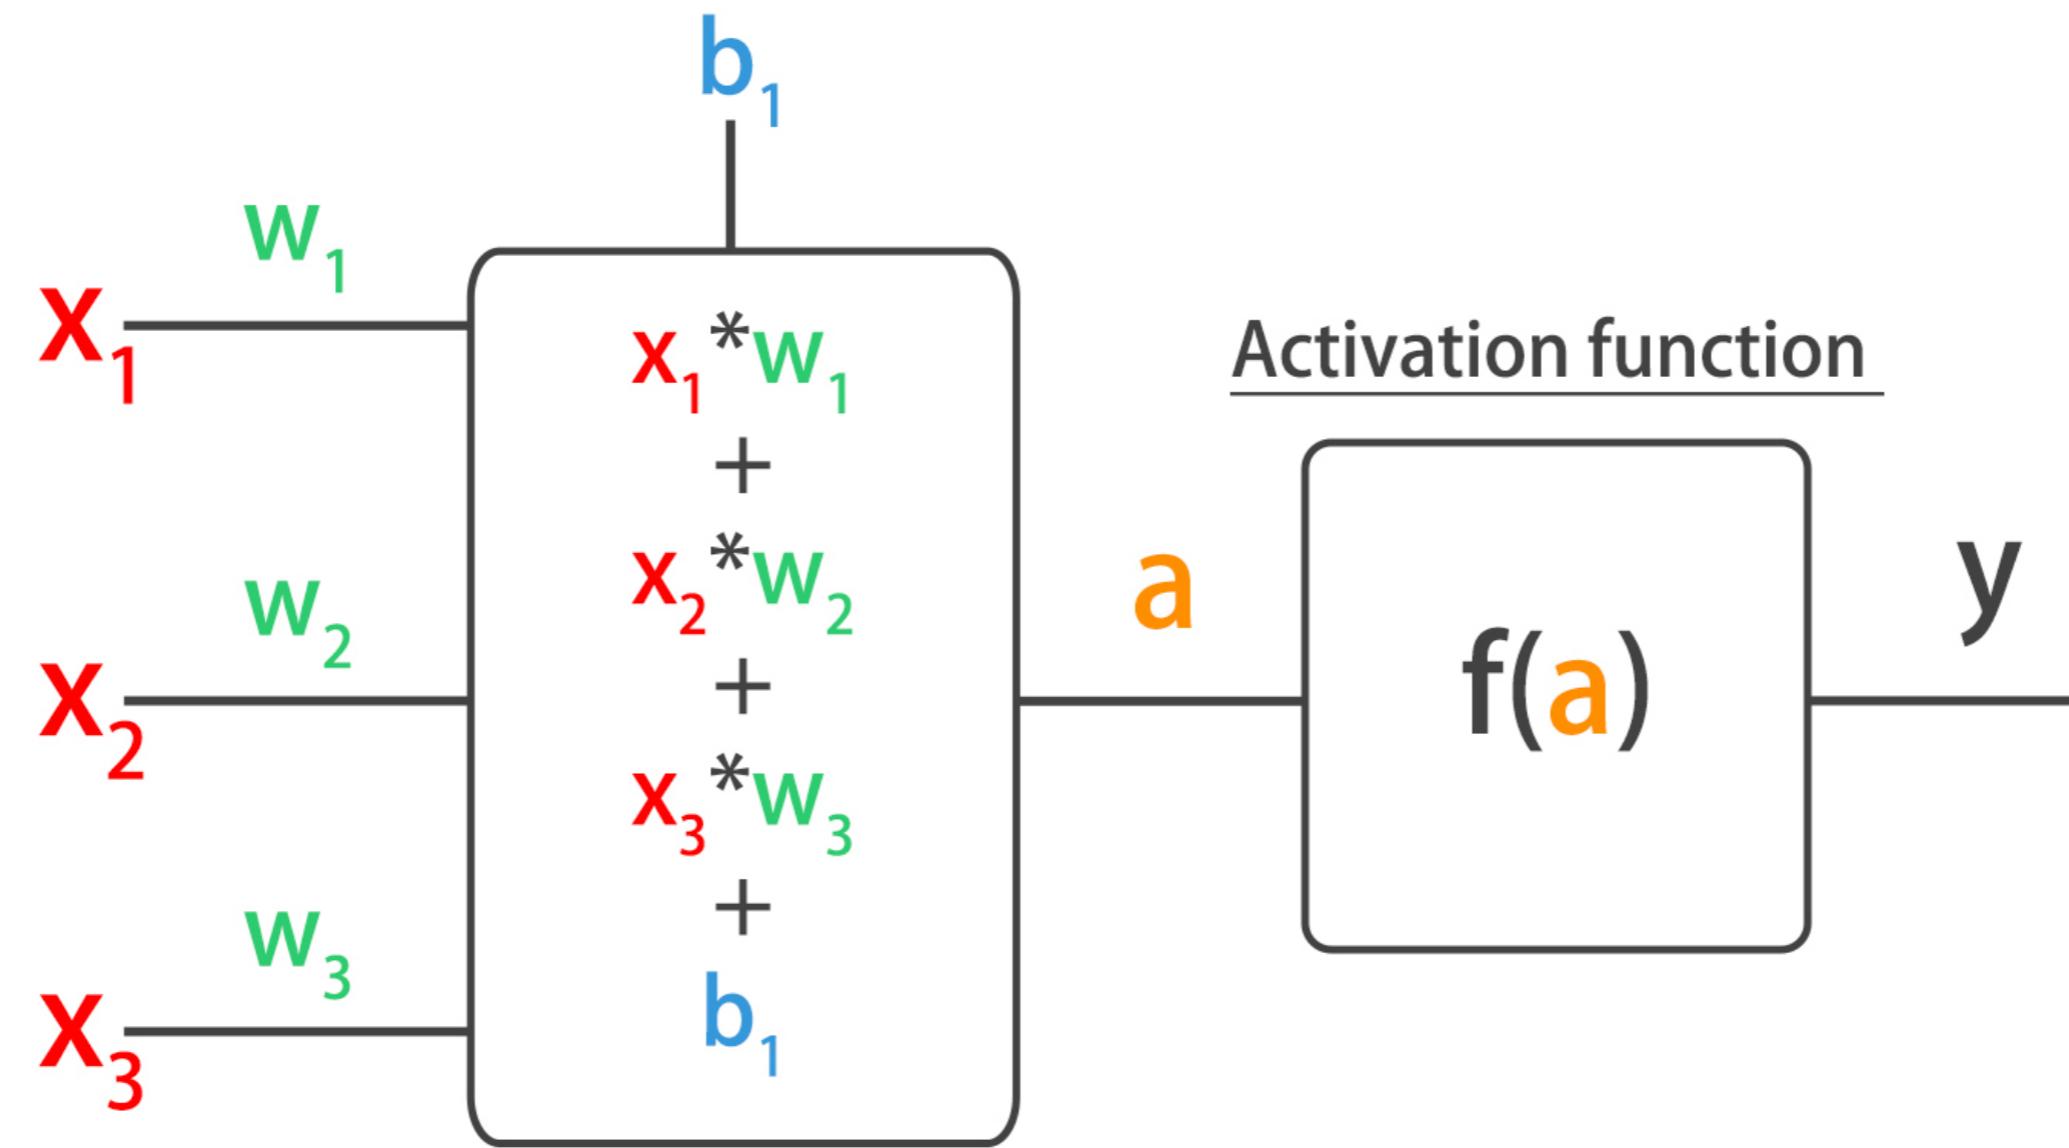


Neuron

$a = \text{sum of inputs} * \text{weights} + \text{bias}$

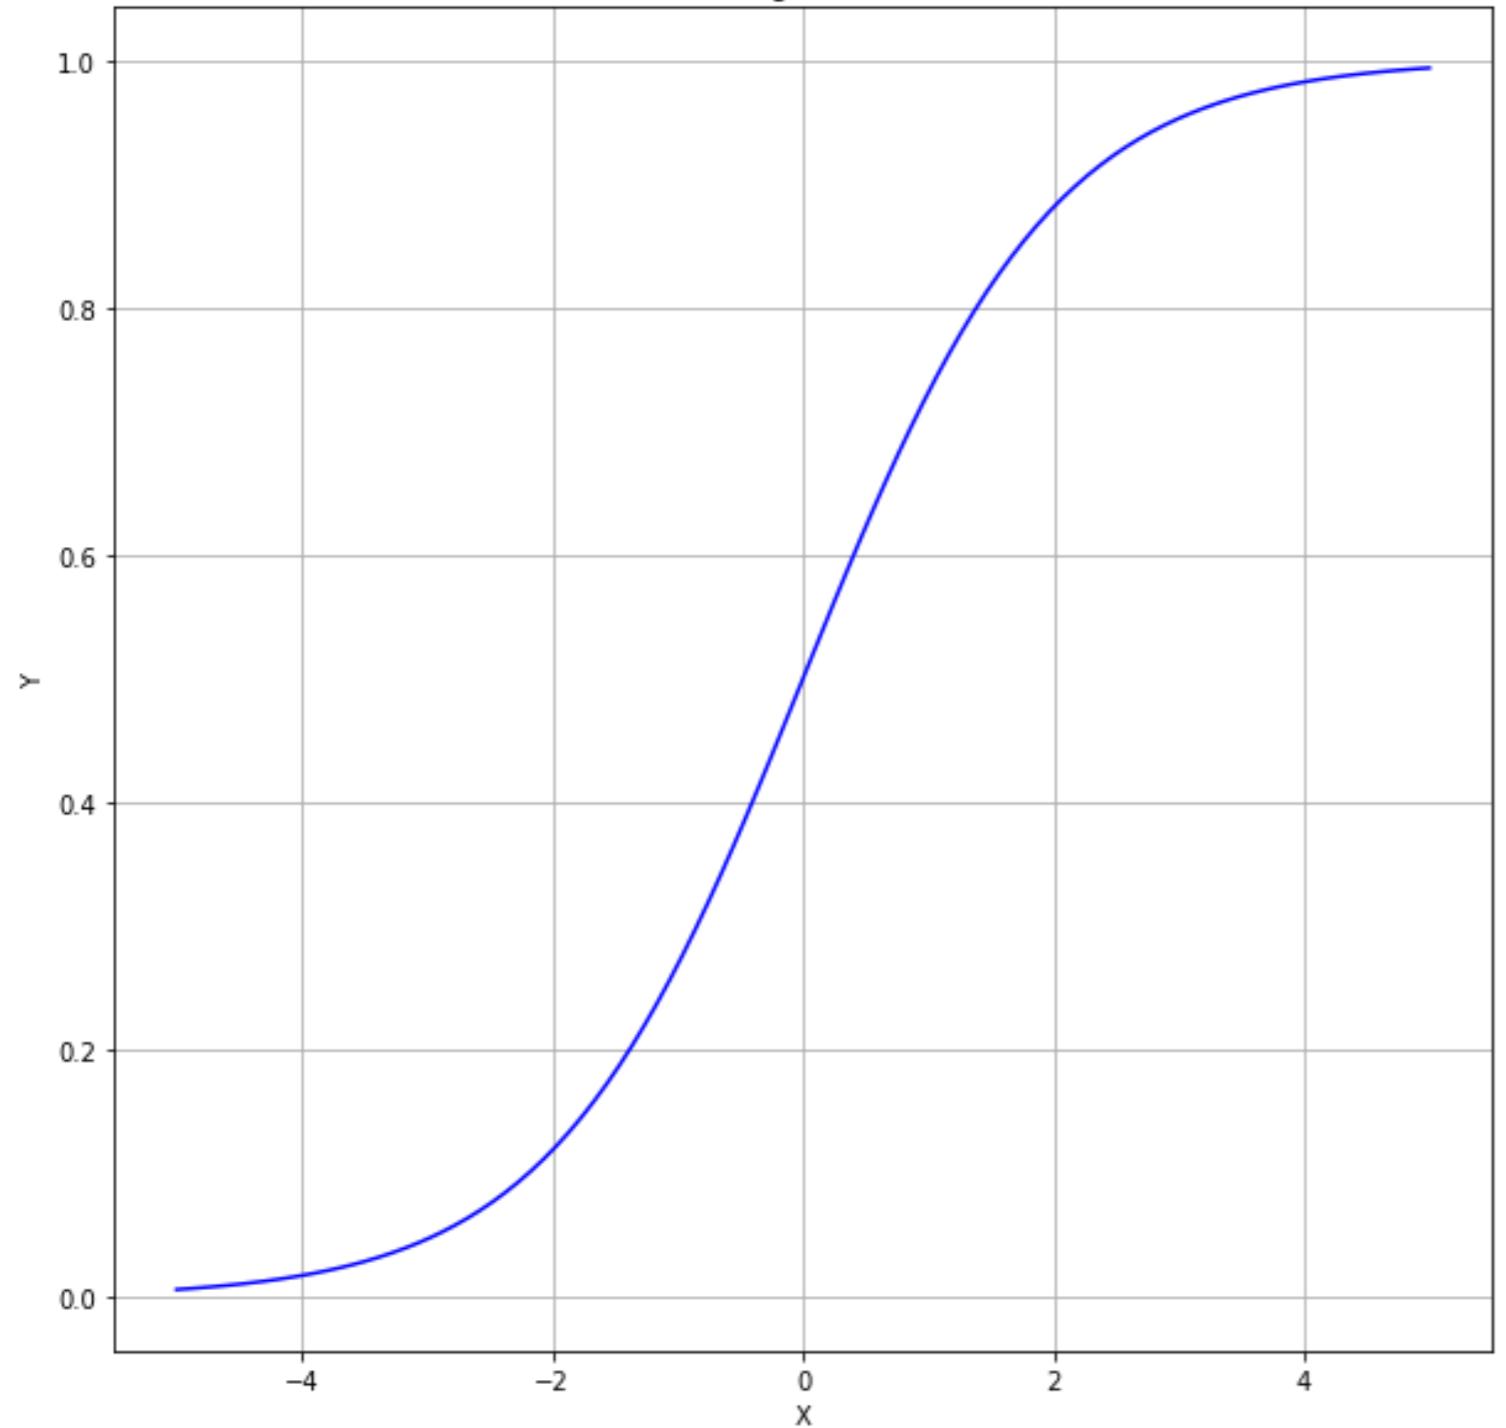


Activation function

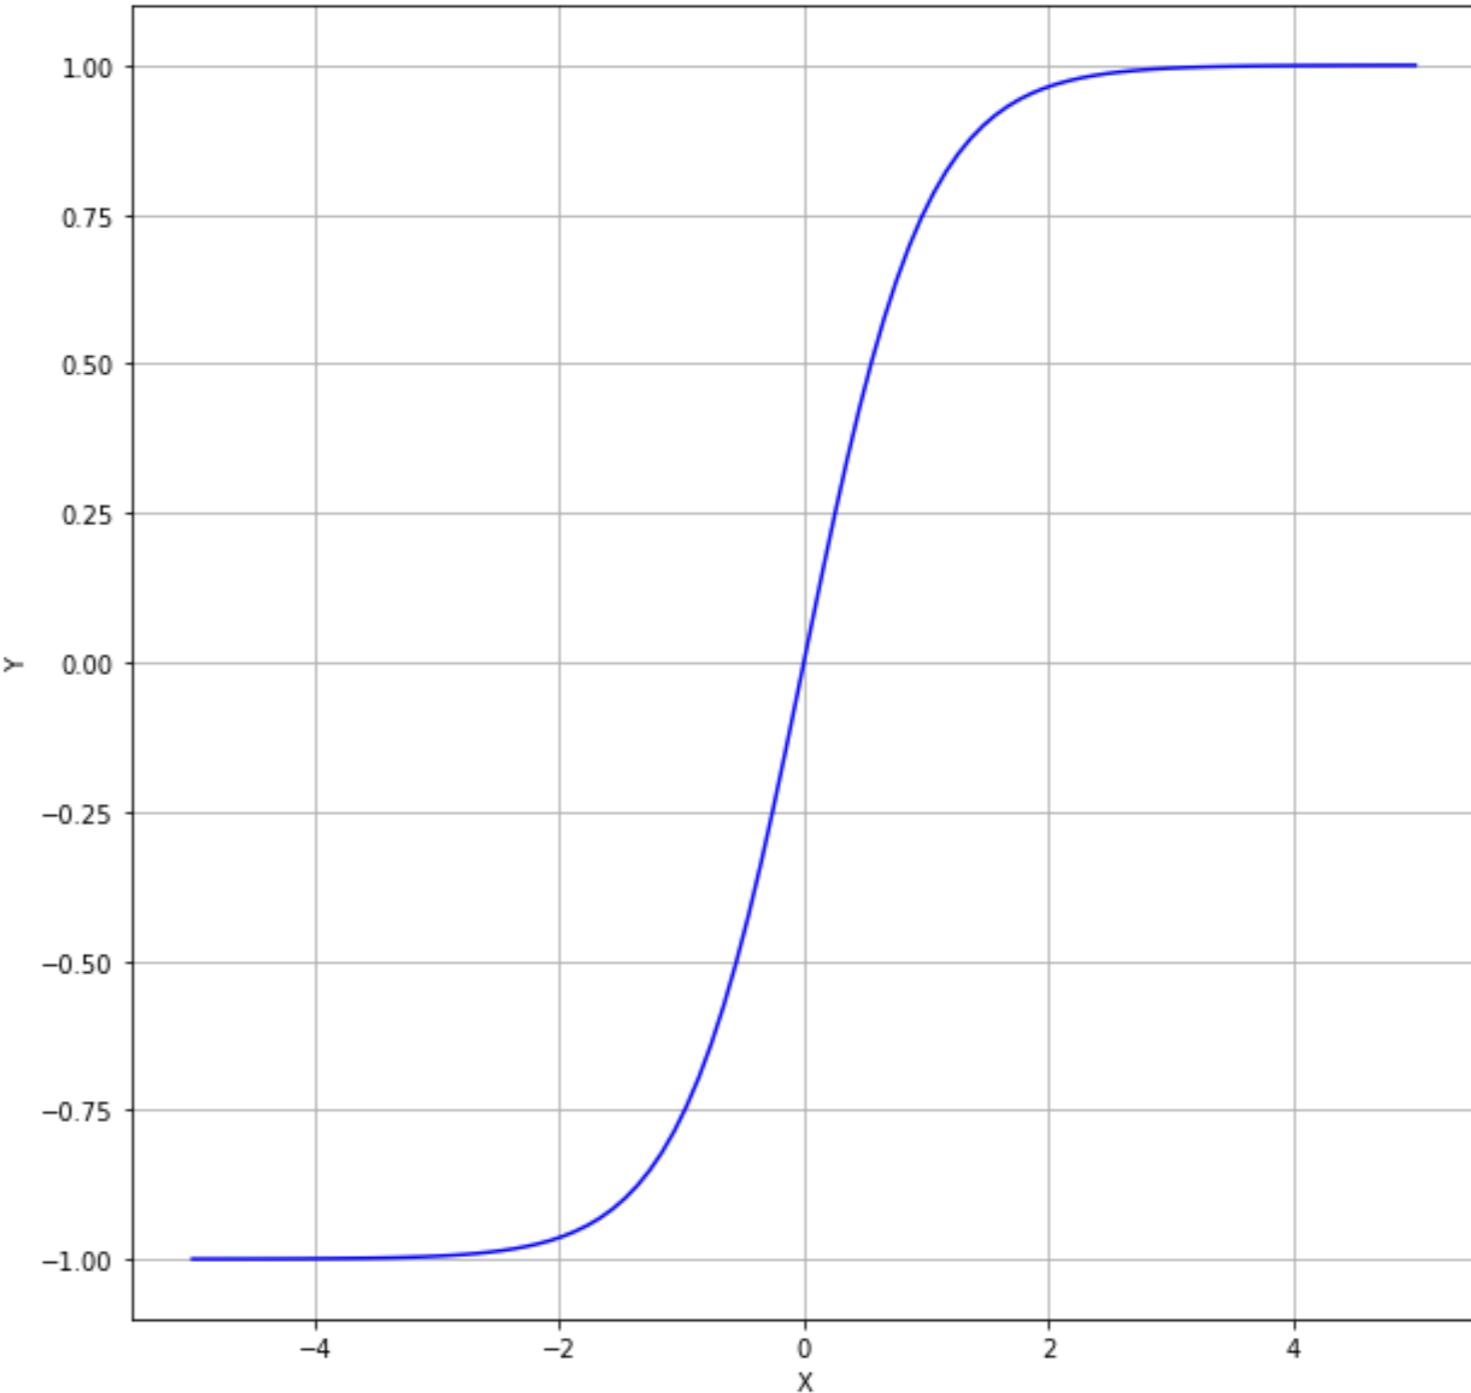


$$a = \text{sum of } x_i * w_i + b_1$$

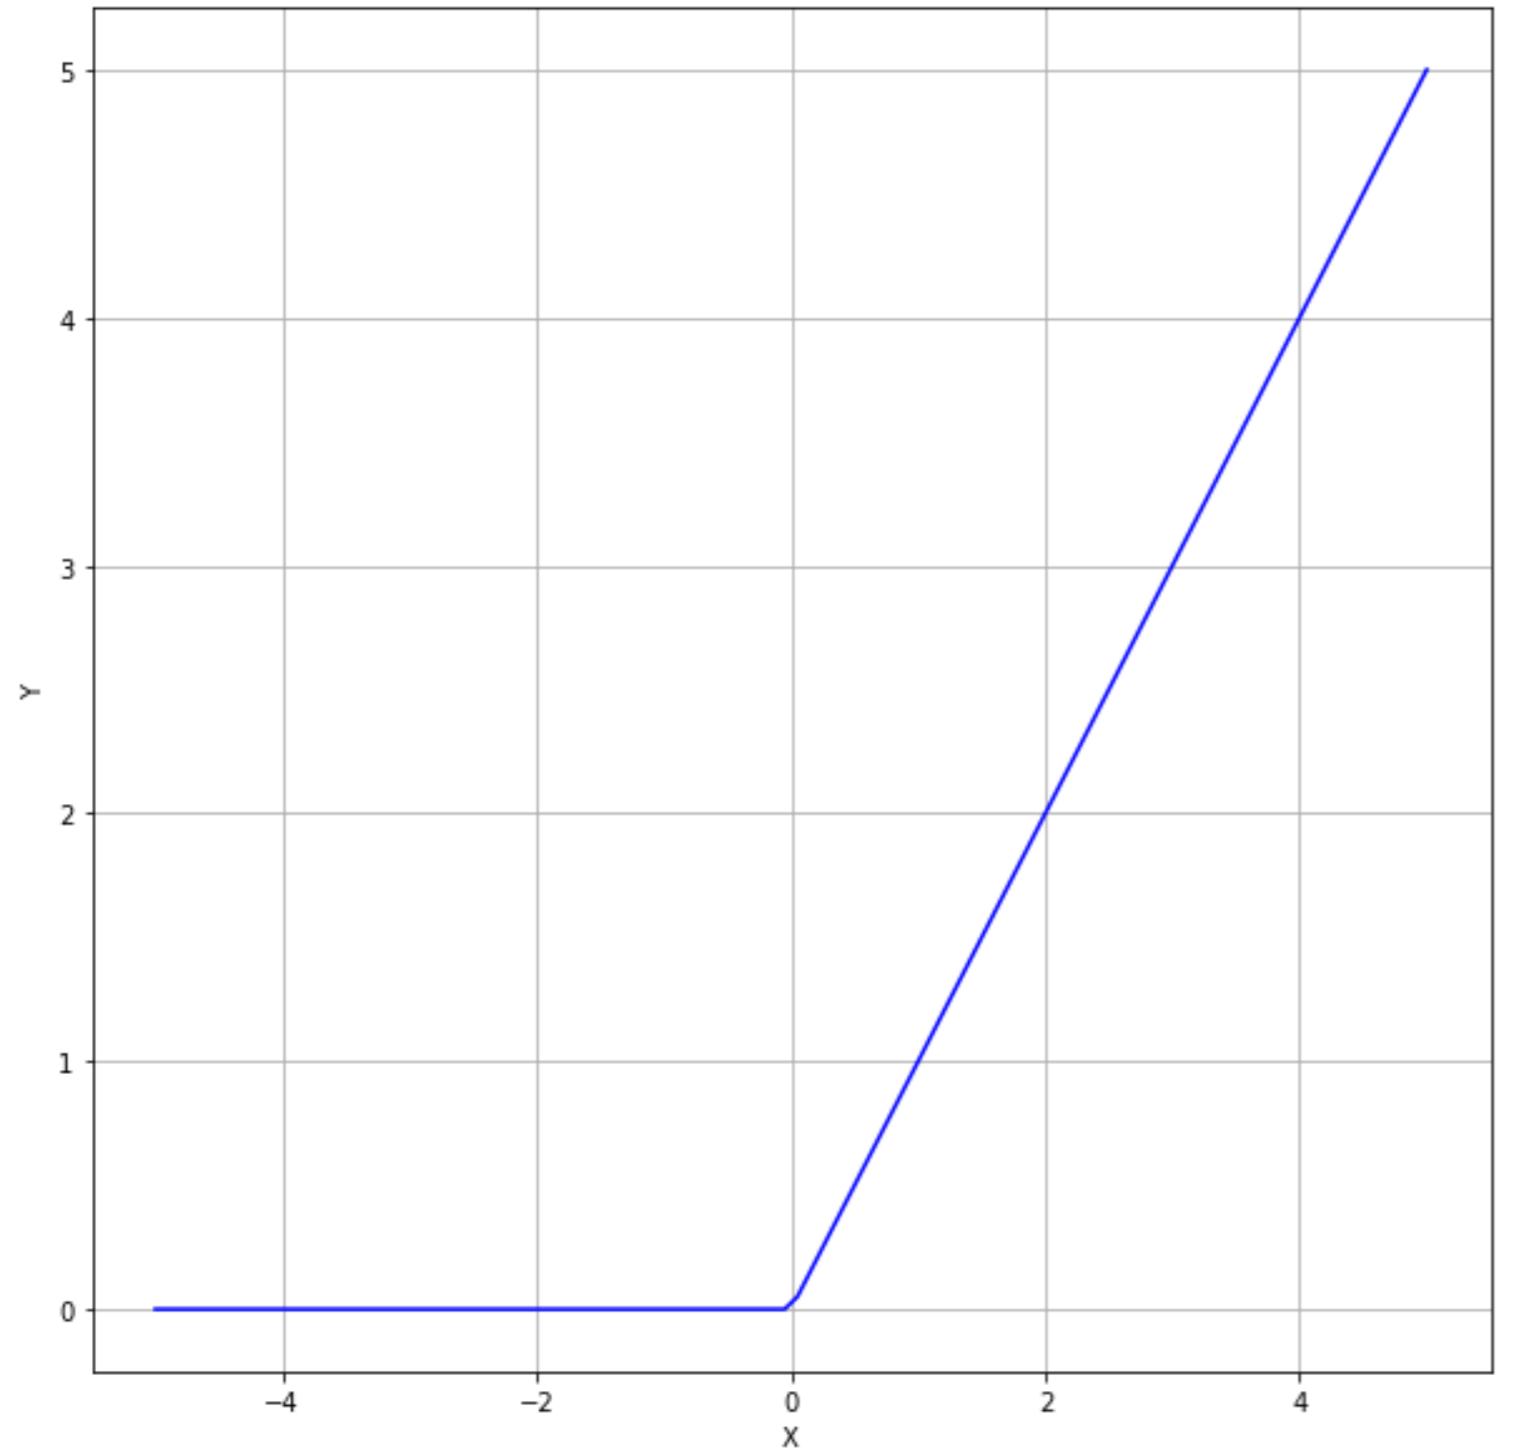
Sigmoid



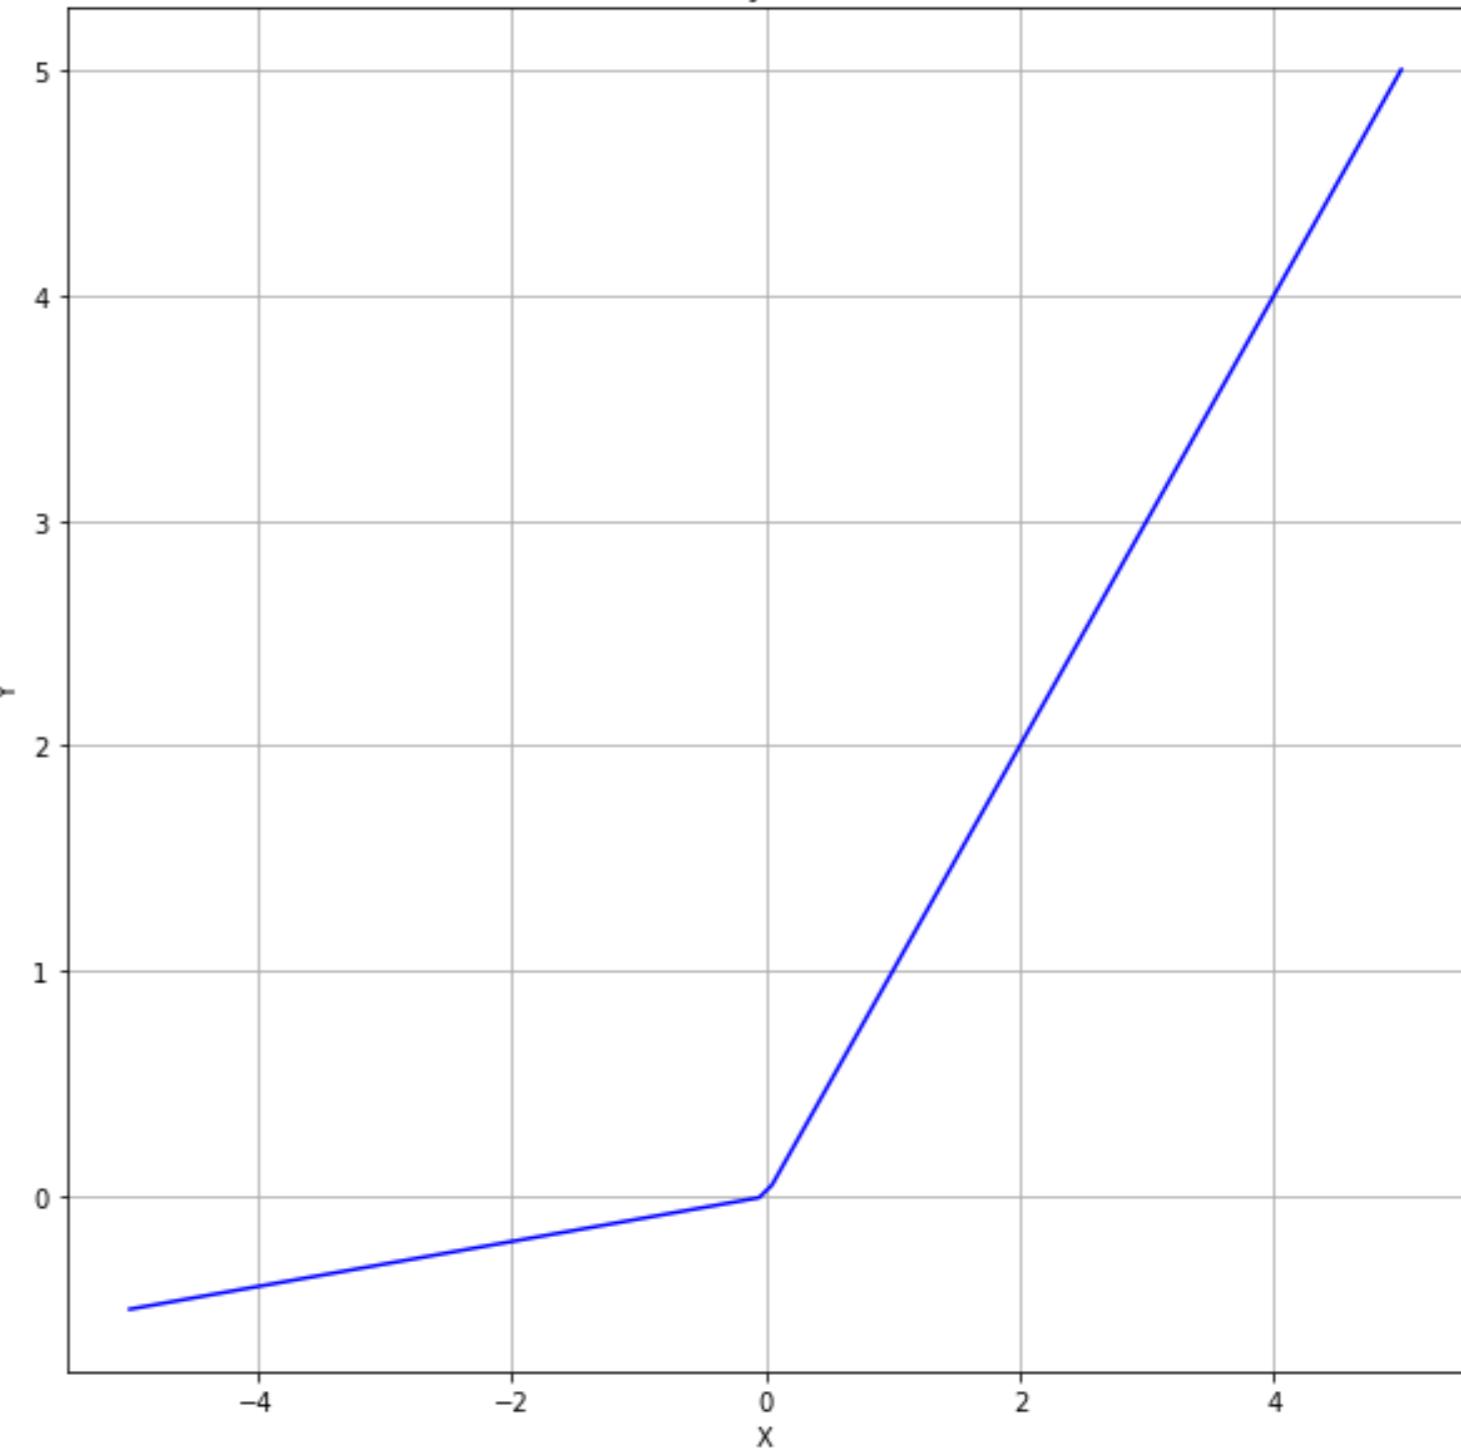
Tanh



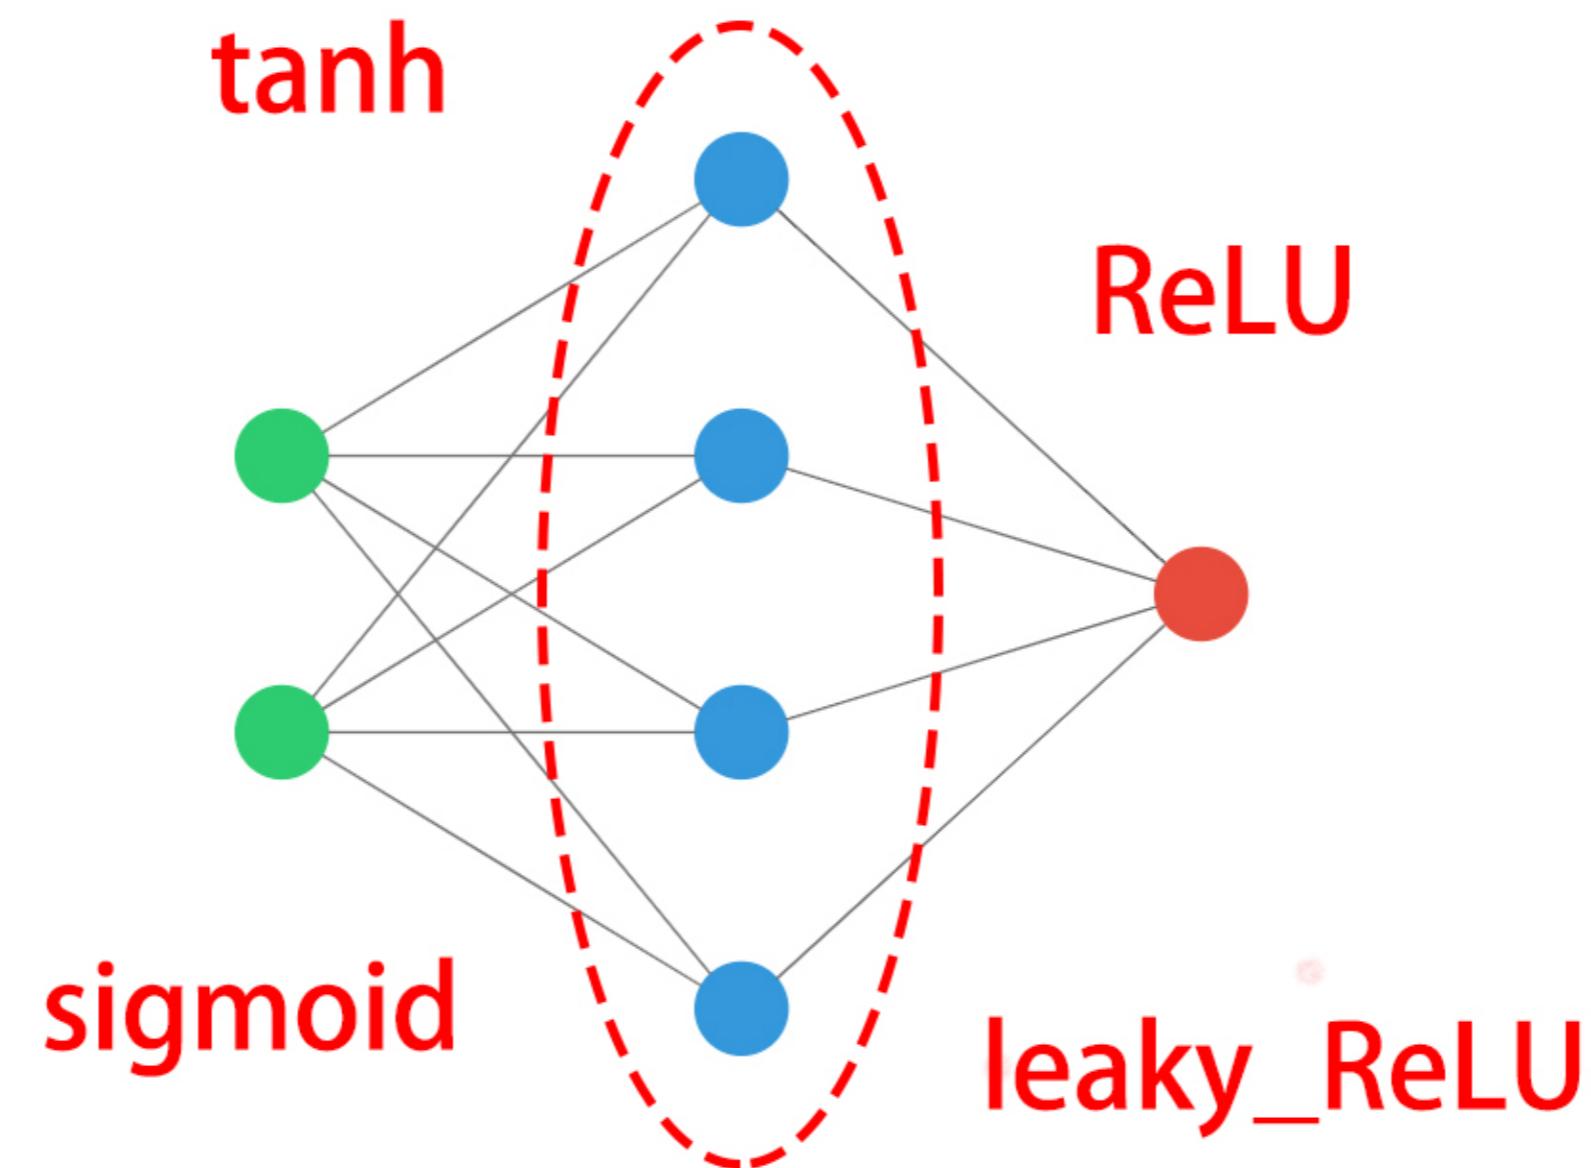
ReLU



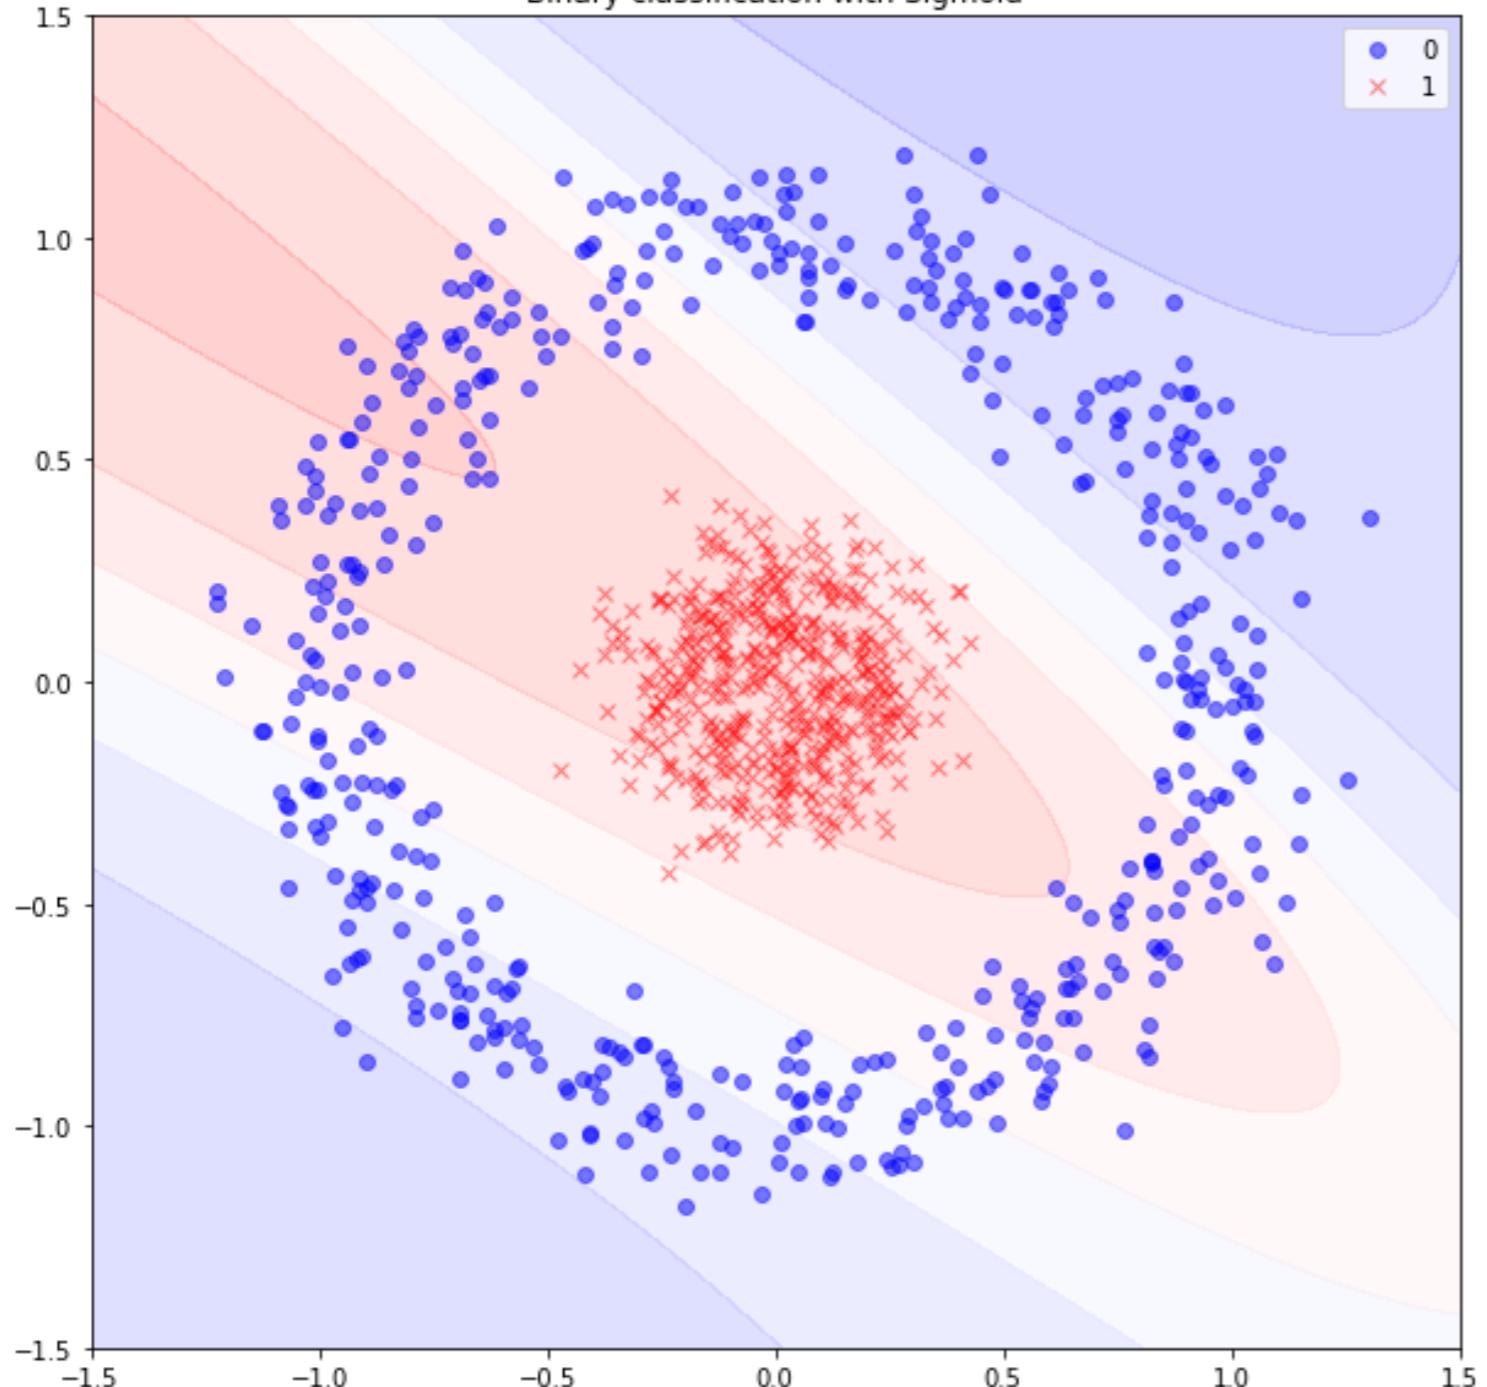
Leaky ReLU



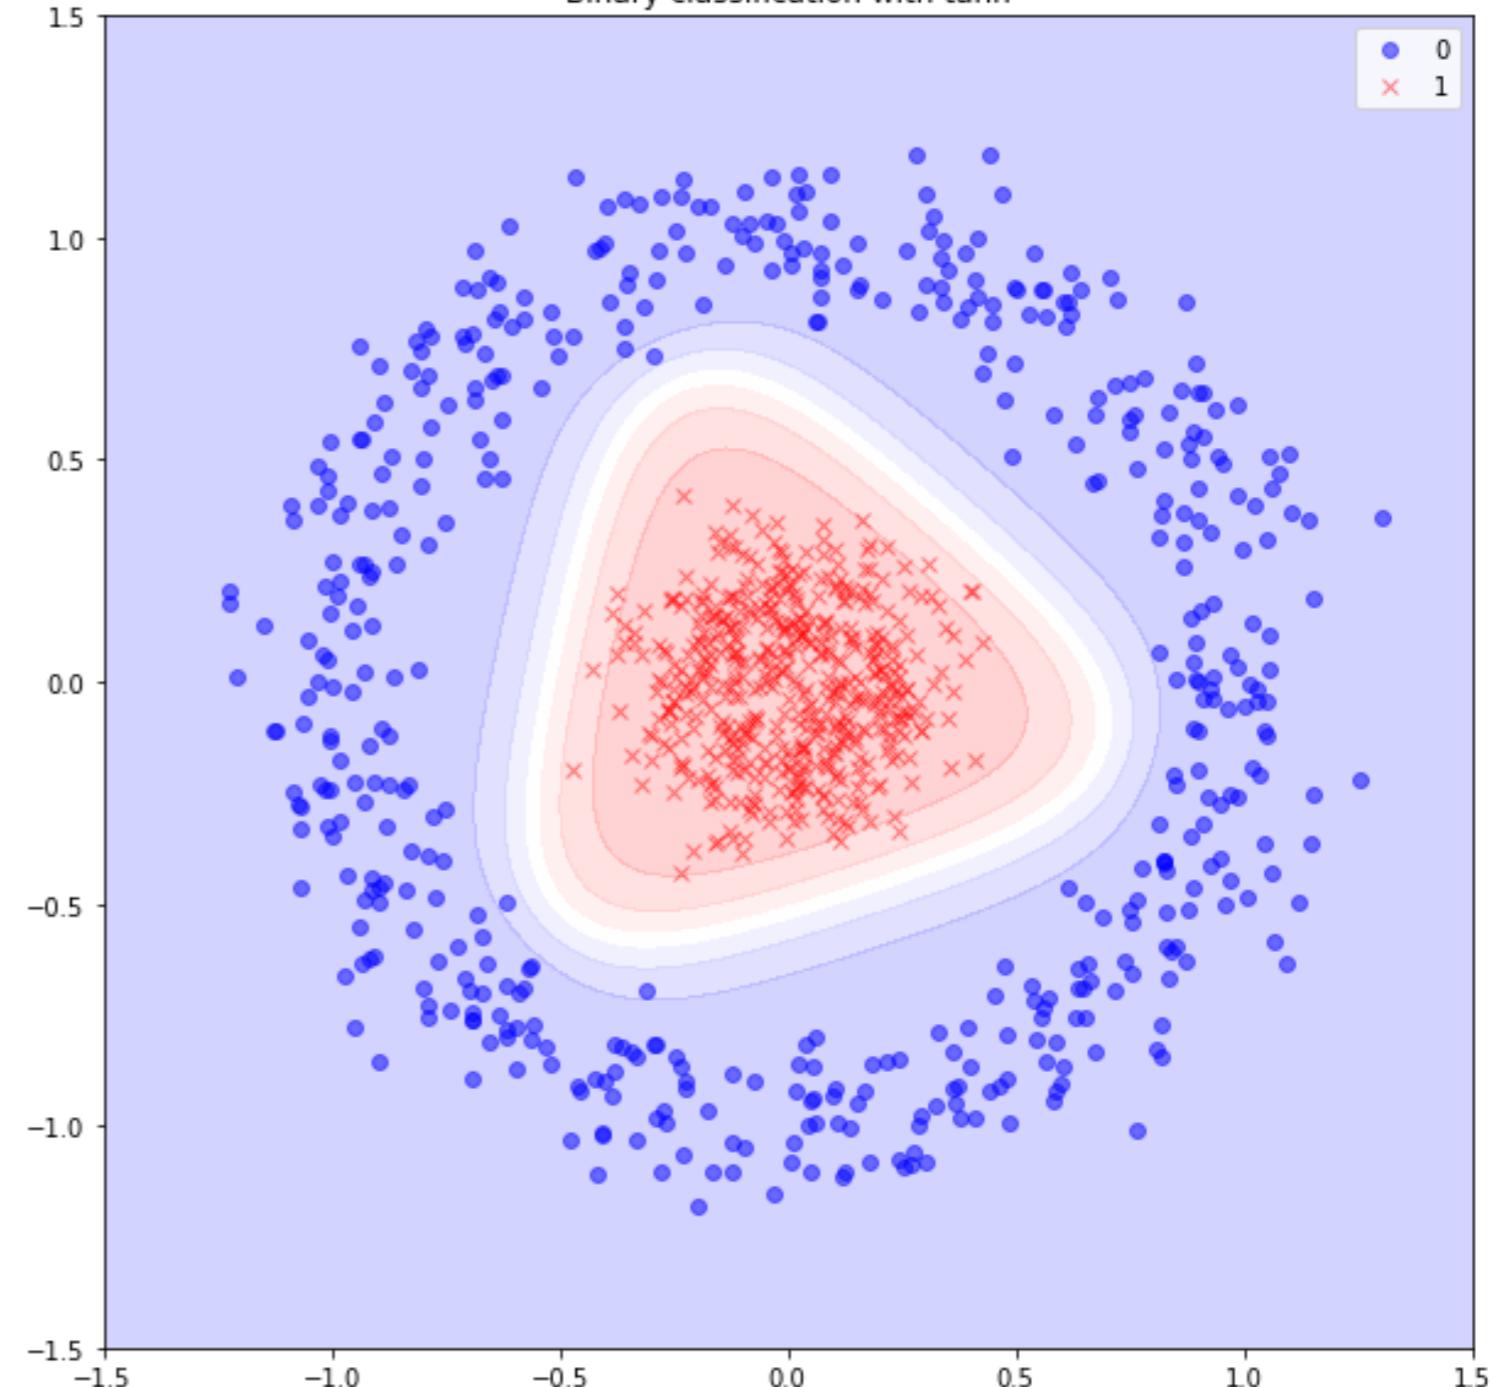
# Effects of activation functions



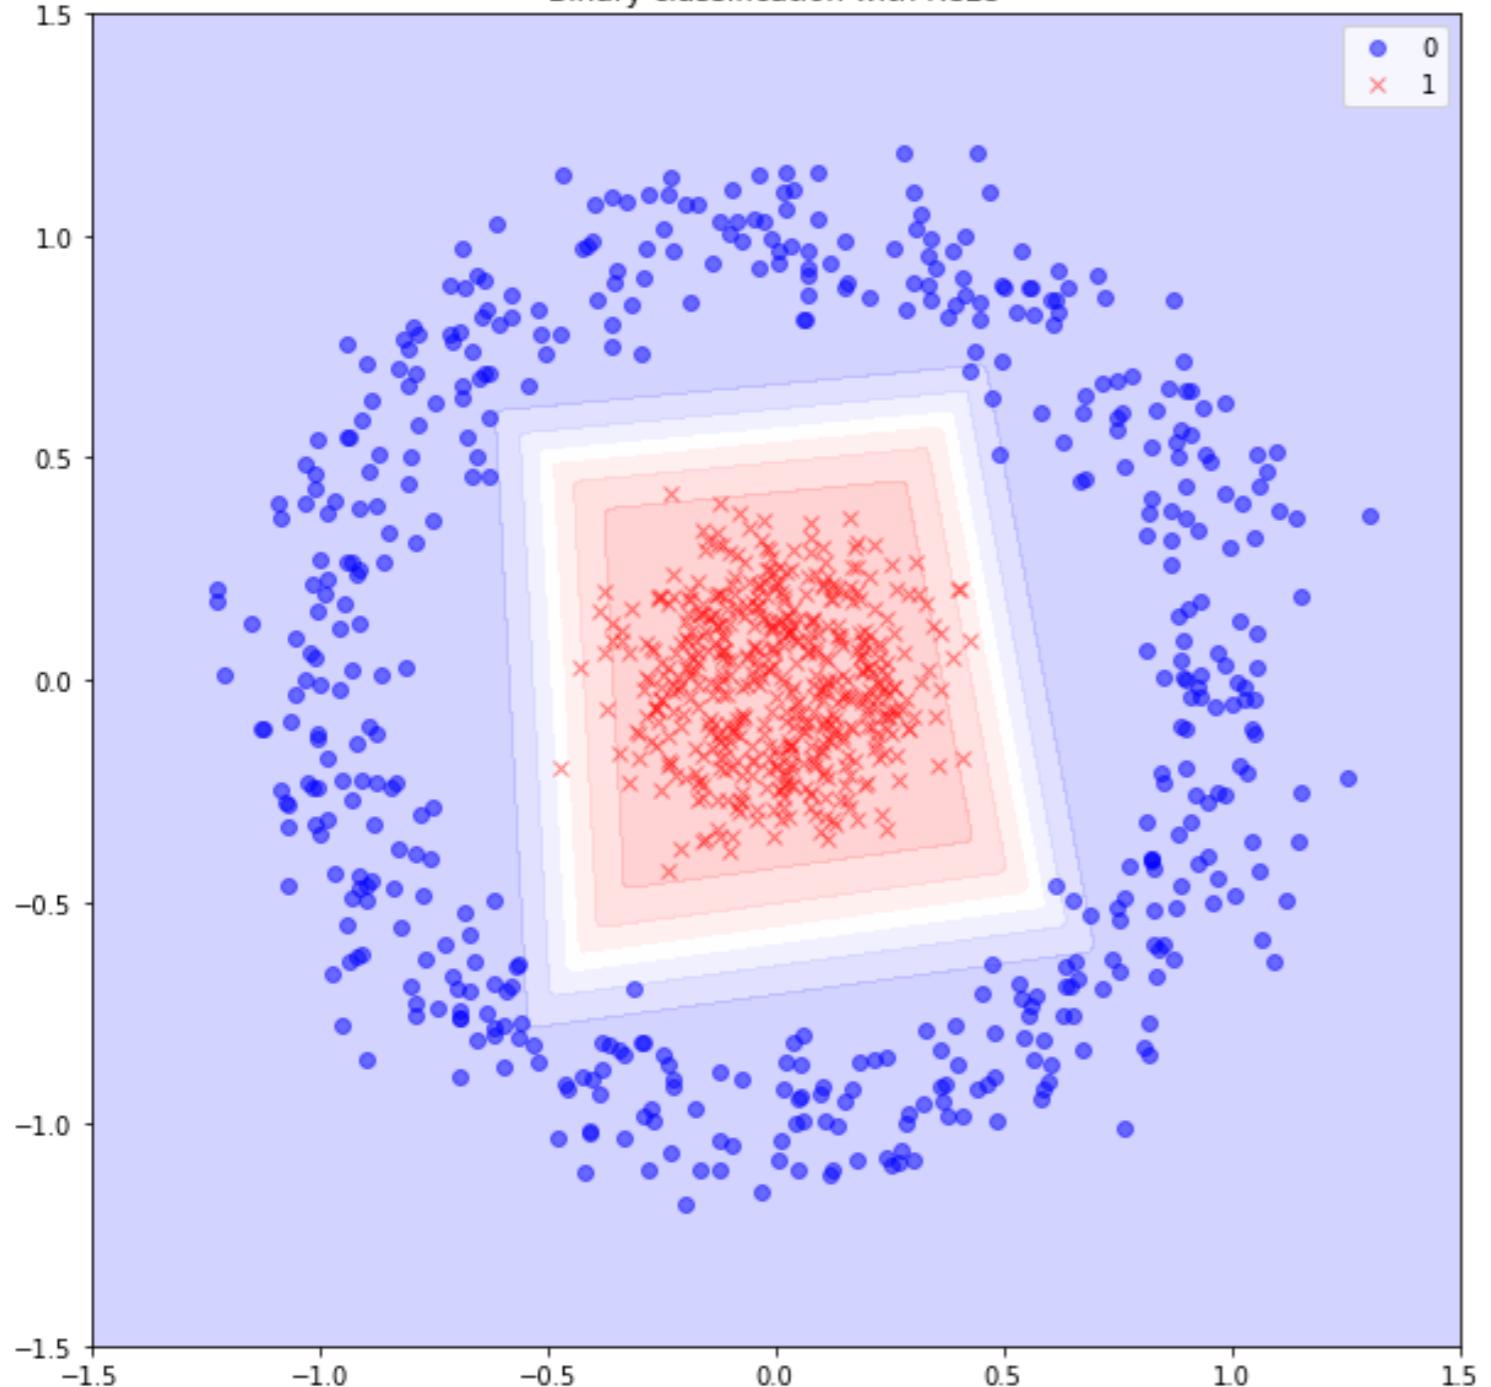
Binary classification with Sigmoid



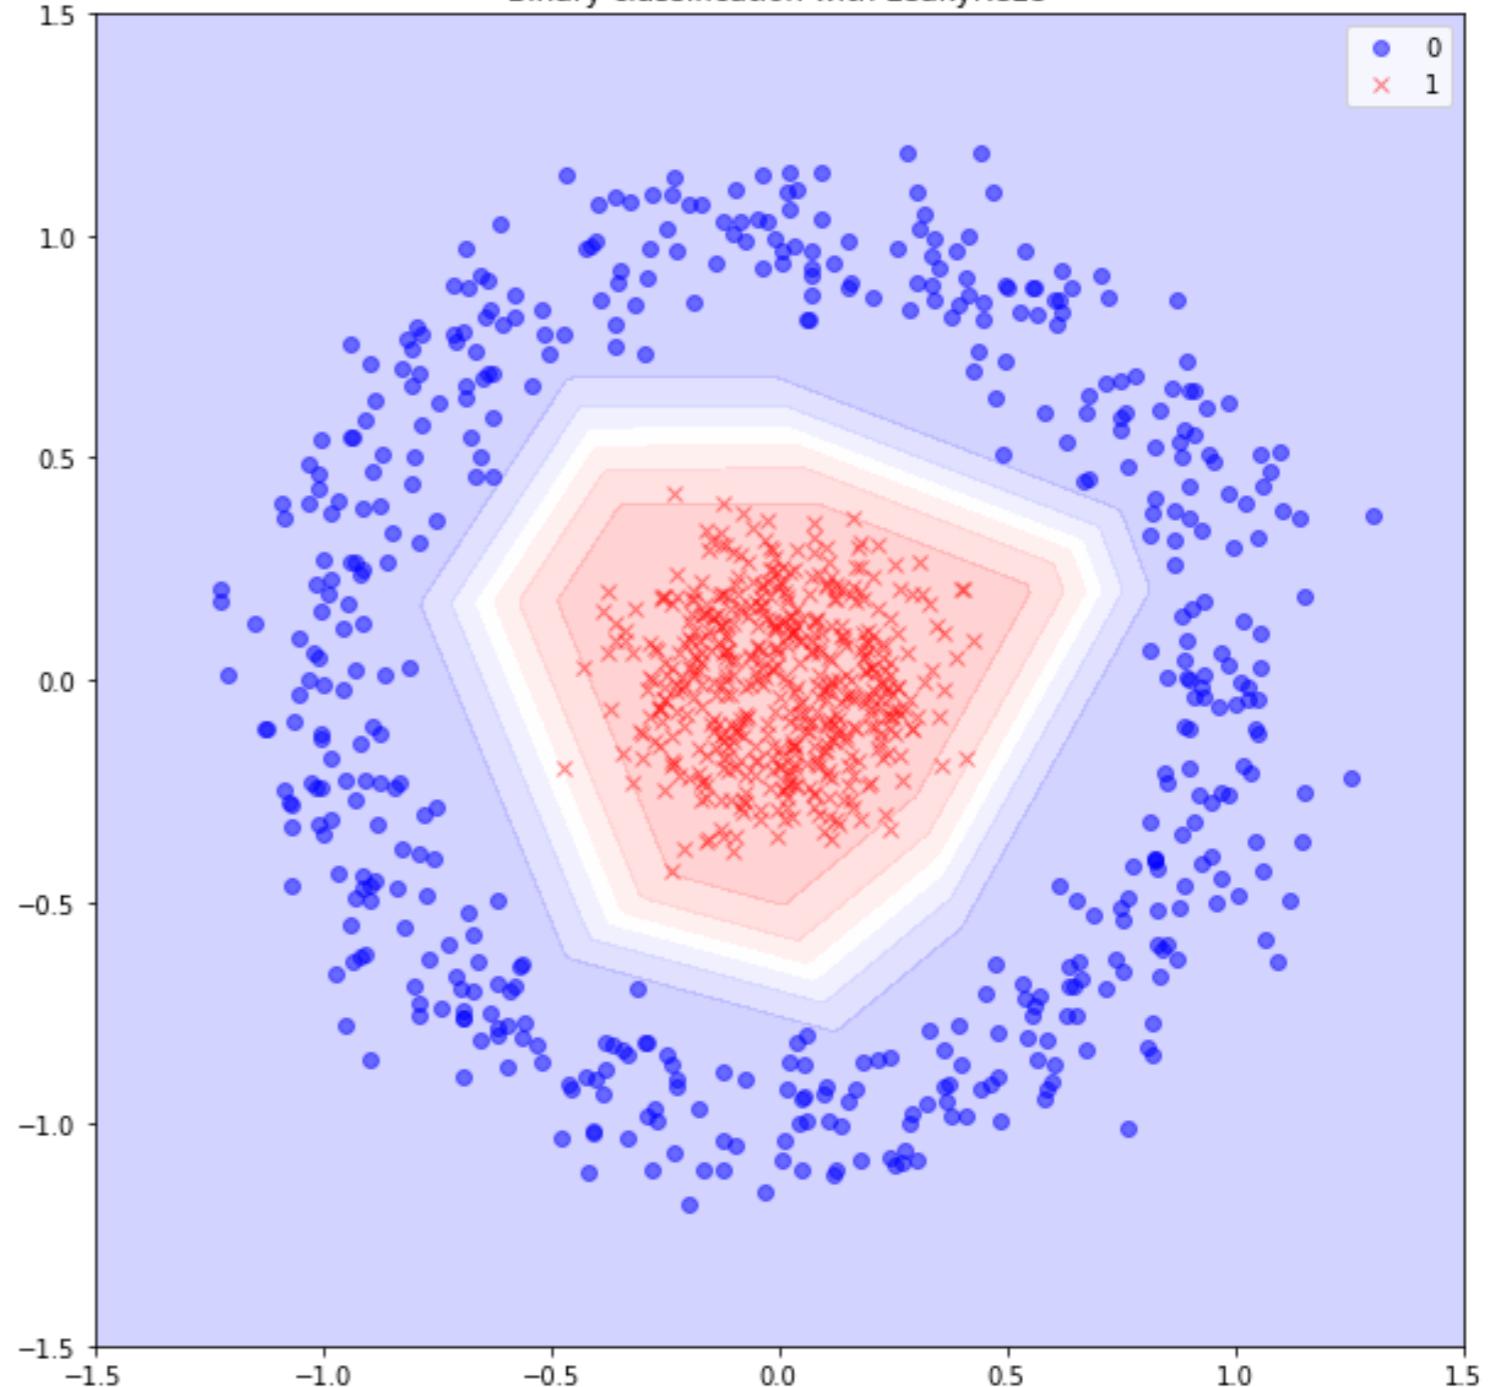
Binary classification with tanh



Binary classification with ReLU

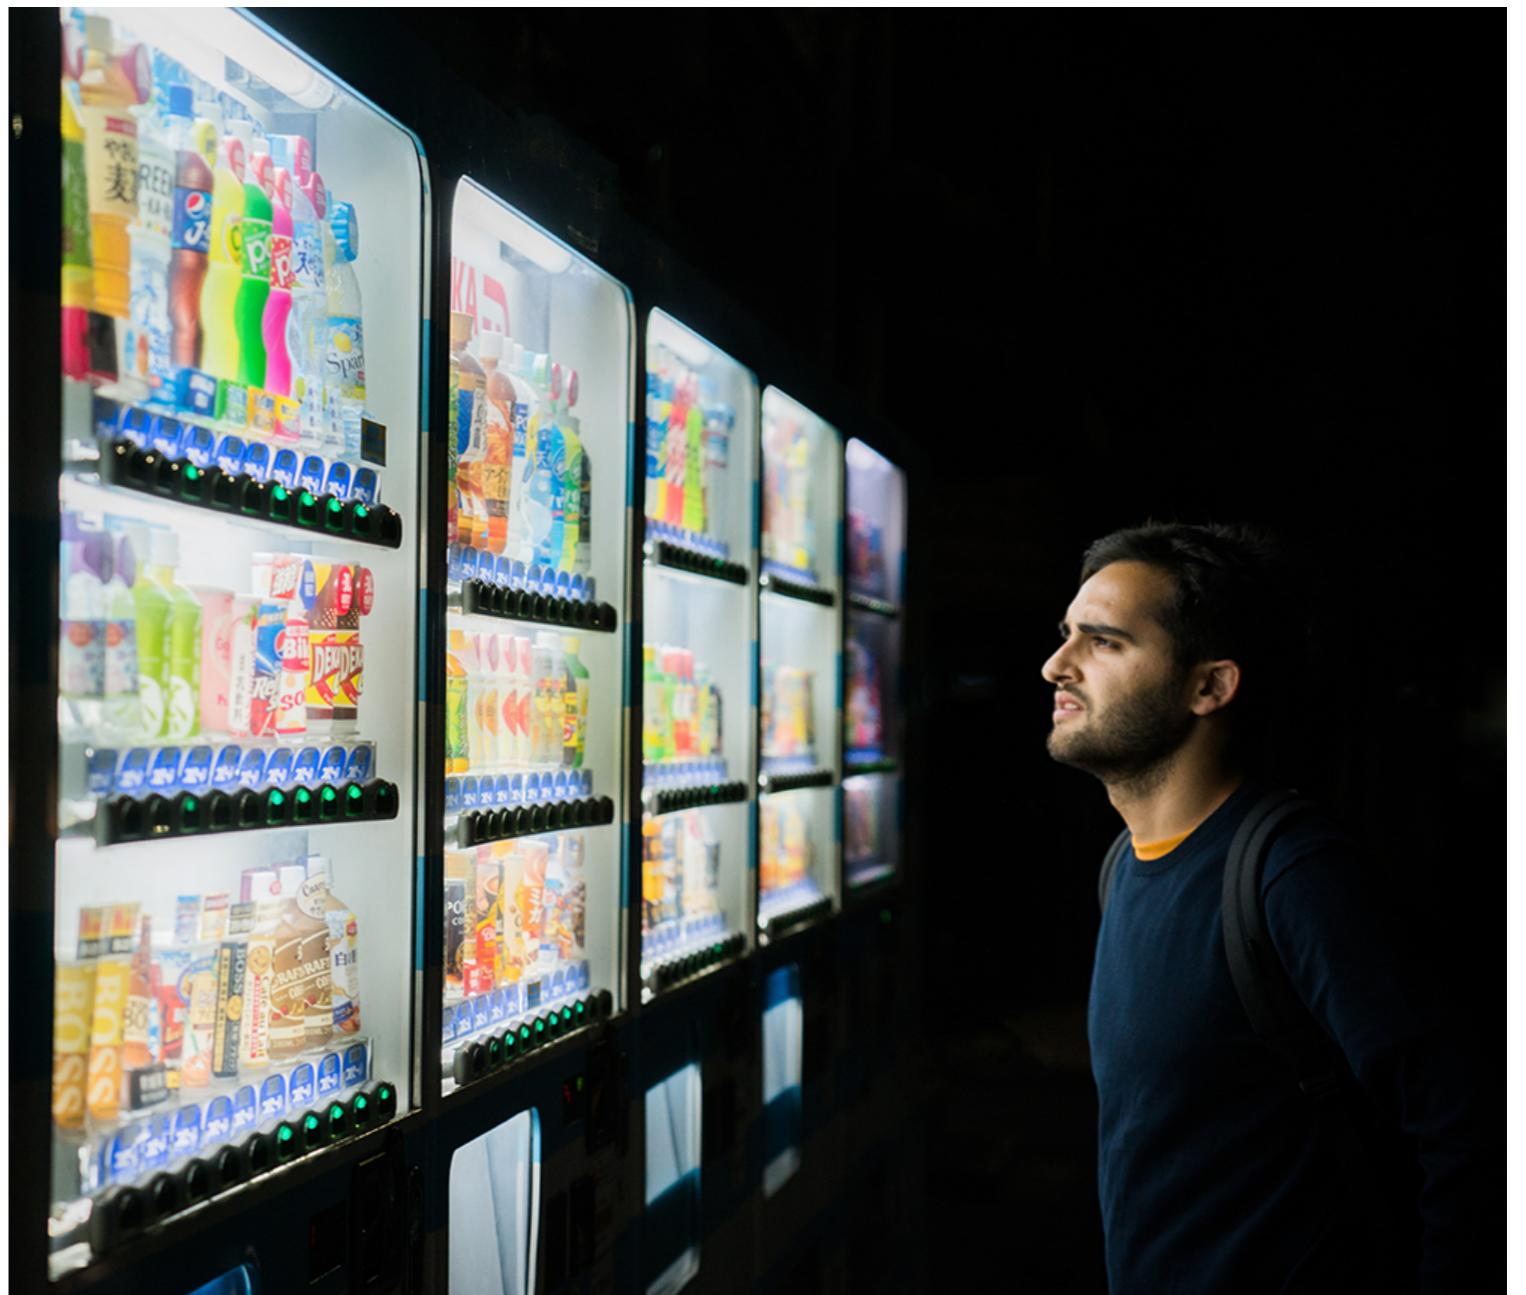


Binary classification with LeakyReLU



# Which activation function to use?

- No magic formula
- Different properties
- Depends on our problem
- Goal to achieve in a given layer
- ReLU are a good first choice
- Sigmoids not recommended for deep models
- Tune with experimentation



# Comparing activation functions

```
# Set a random seed
np.random.seed(1)

# Return a new model with the given activation
def get_model(act_function):

    model = Sequential()
    model.add(Dense(4, input_shape=(2,), activation=act_function))
    model.add(Dense(1, activation='sigmoid'))
    return model
```

# Comparing activation functions

```
# Activation functions to try out
activations = ['relu', 'sigmoid', 'tanh']

# Dictionary to store results
activation_results = {}

for funct in activations:
    model = get_model(act_function=funct)
    history = model.fit(X_train, y_train,
                         validation_data=(X_test, y_test),
                         epochs=100, verbose=0)
    activation_results[funct] = history
```

# Comparing activation functions

```
import pandas as pd

# Extract val_loss history of each activation function
val_loss_per_funct = {k:v.history['val_loss'] for k,v in activation_results.items()}

# Turn the dictionary into a pandas dataframe
val_loss_curves = pd.DataFrame(val_loss_per_funct)

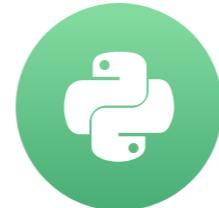
# Plot the curves
val_loss_curves.plot(title='Loss per Activation function')
```

# Let's practice!

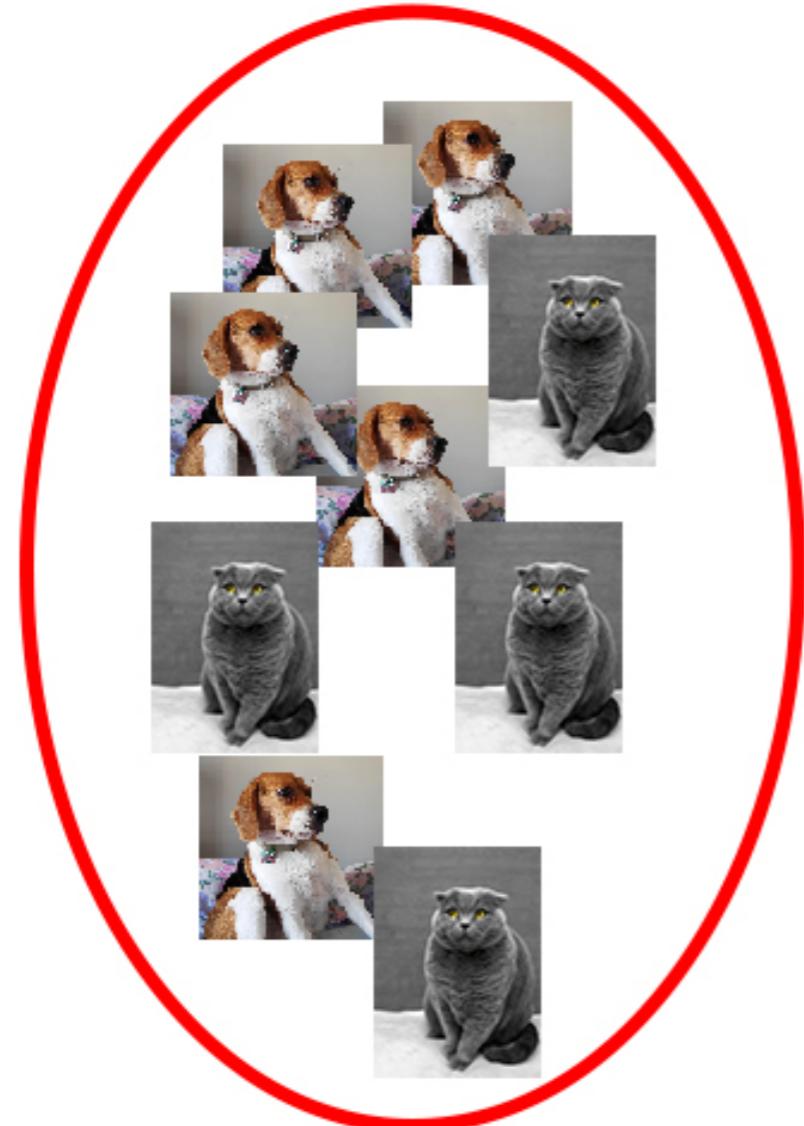
INTRODUCTION TO DEEP LEARNING WITH KERAS

# Batch size and batch normalization

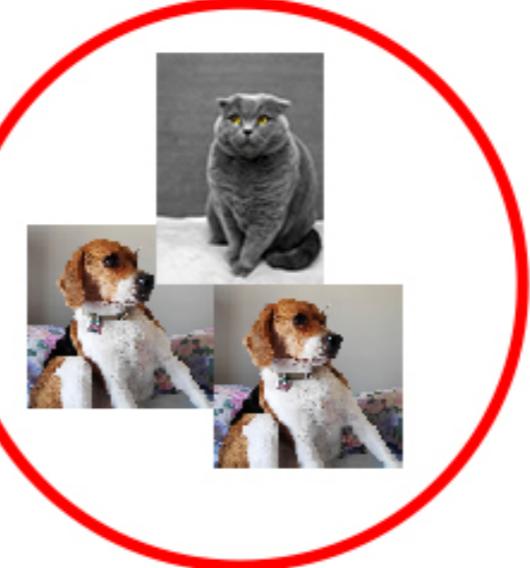
INTRODUCTION TO DEEP LEARNING WITH KERAS



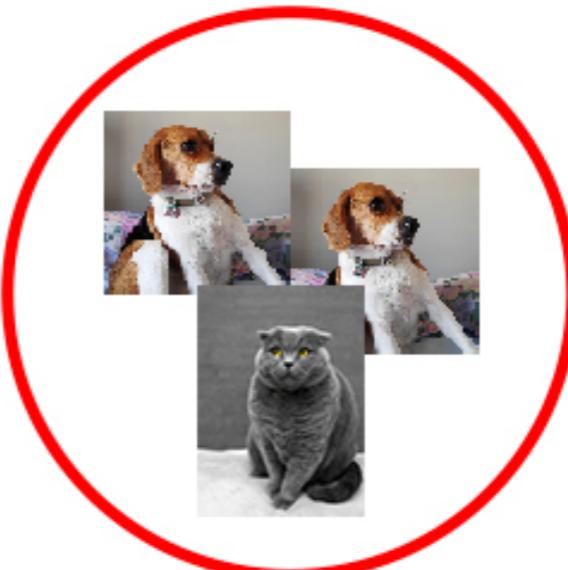
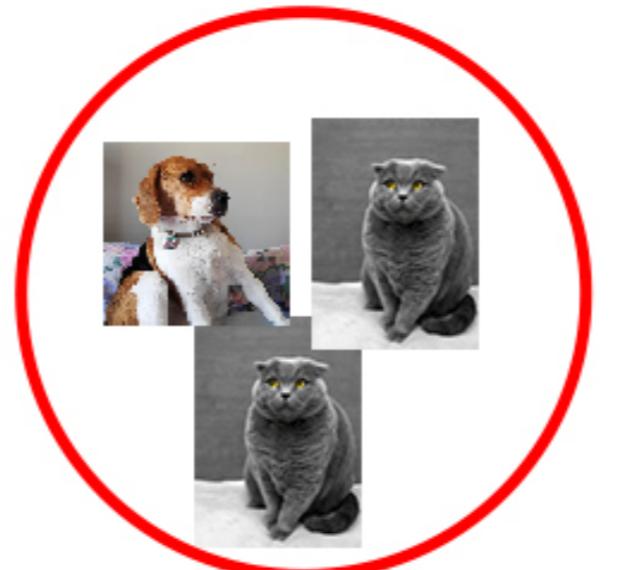
Miguel Esteban  
Data Scientist & Founder



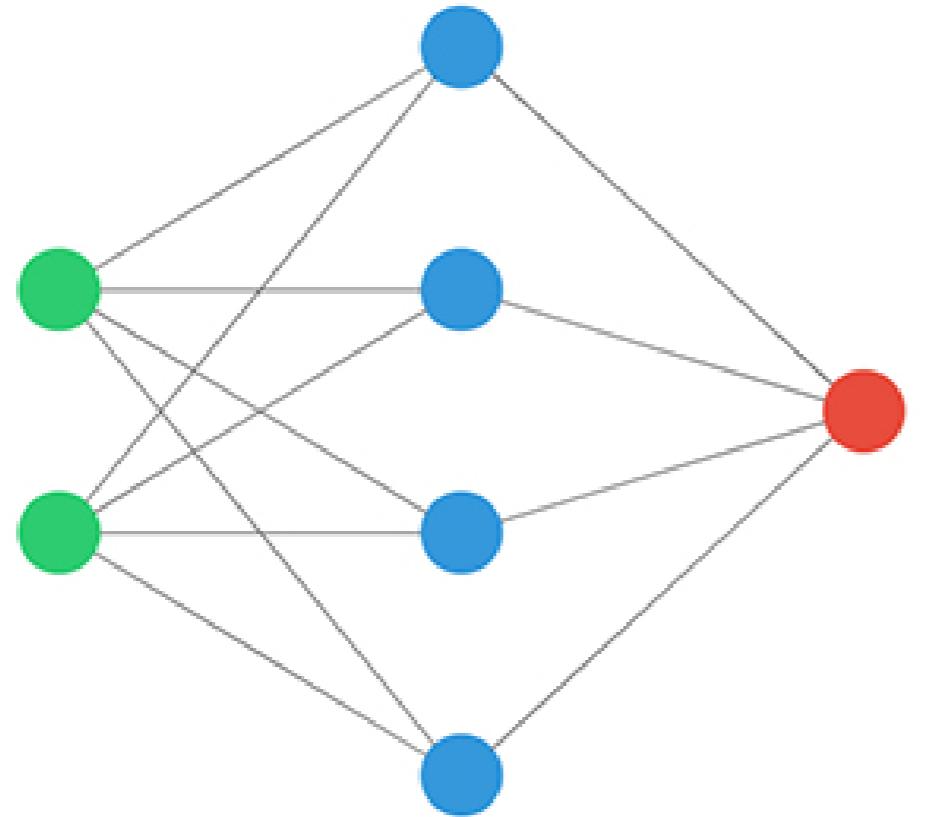
**Batch**



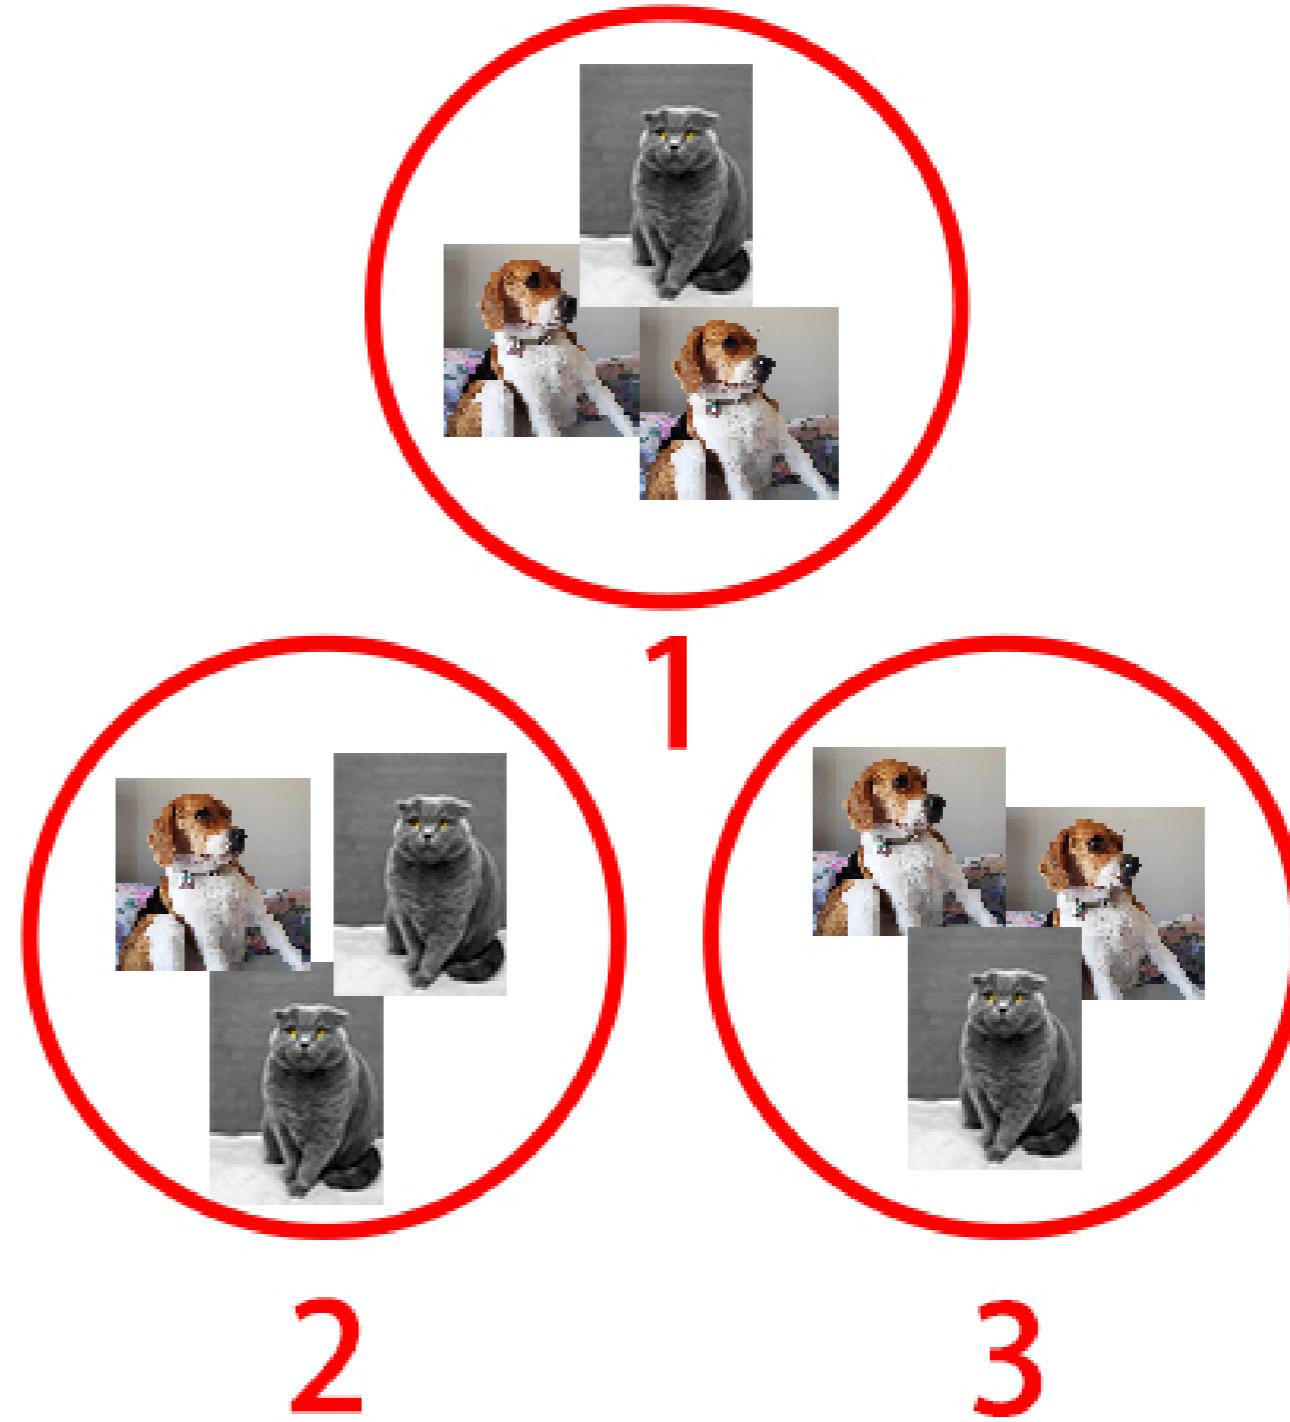
**Mini-batches**



The network is fed with 3 mini-batches



1 Epoch = 3 weight updates



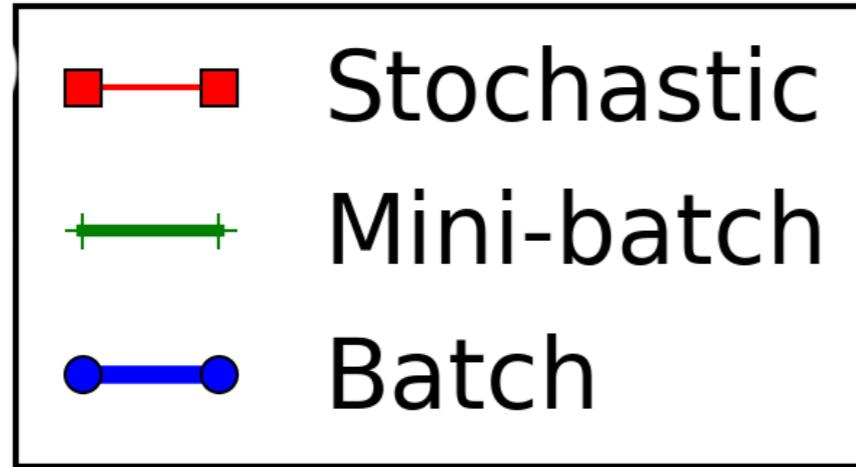
# Mini-batches

## Advantages

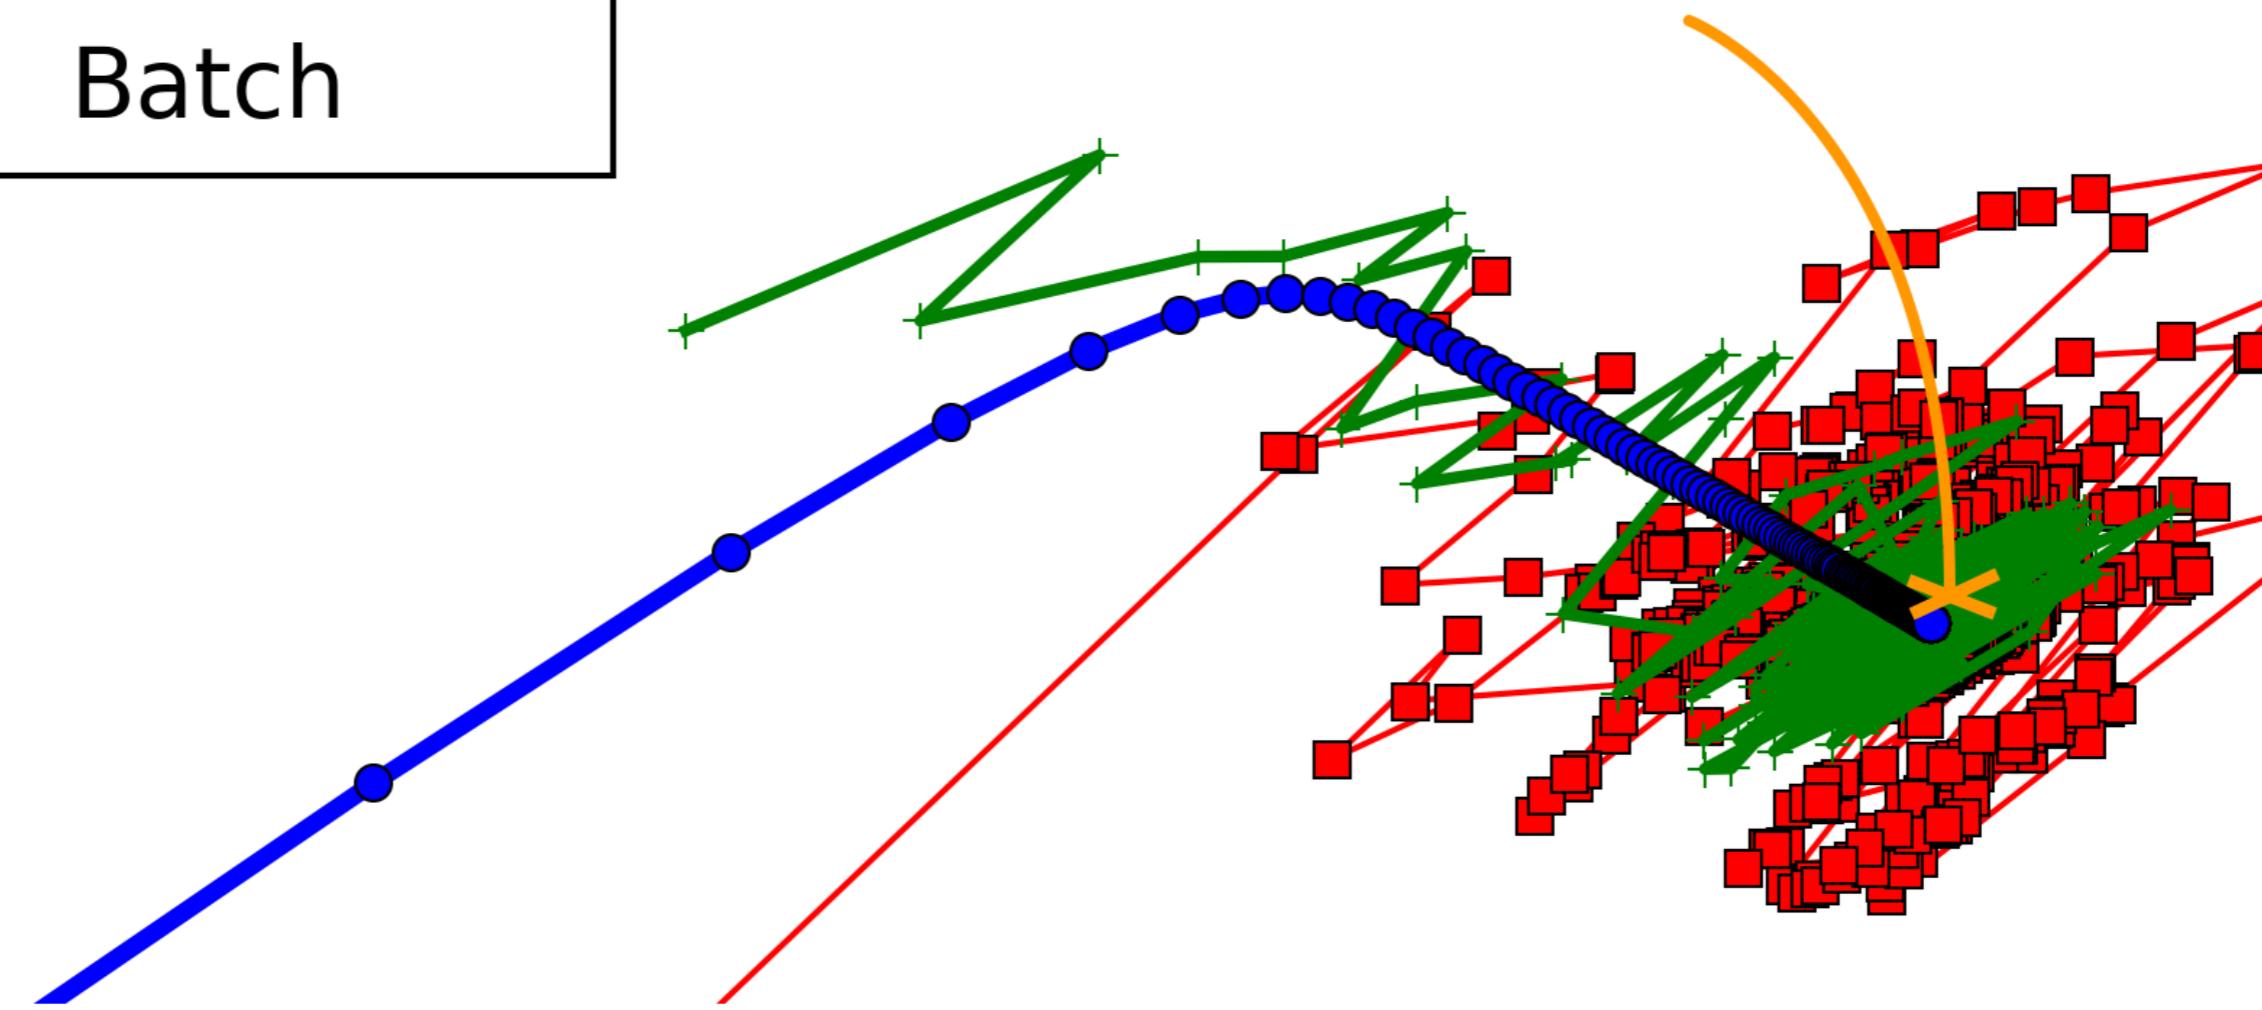
- Networks train faster (more weight updates in same amount of time)
- Less RAM memory required, can train on huge datasets
- Noise can help networks reach a lower error, escaping local minima

## Disadvantages

- More iterations need to be run
- Need to be adjusted, we need to find a good batch size



Best value for our weights  
(were loss is smaller)



<sup>1</sup> Stack Exchange

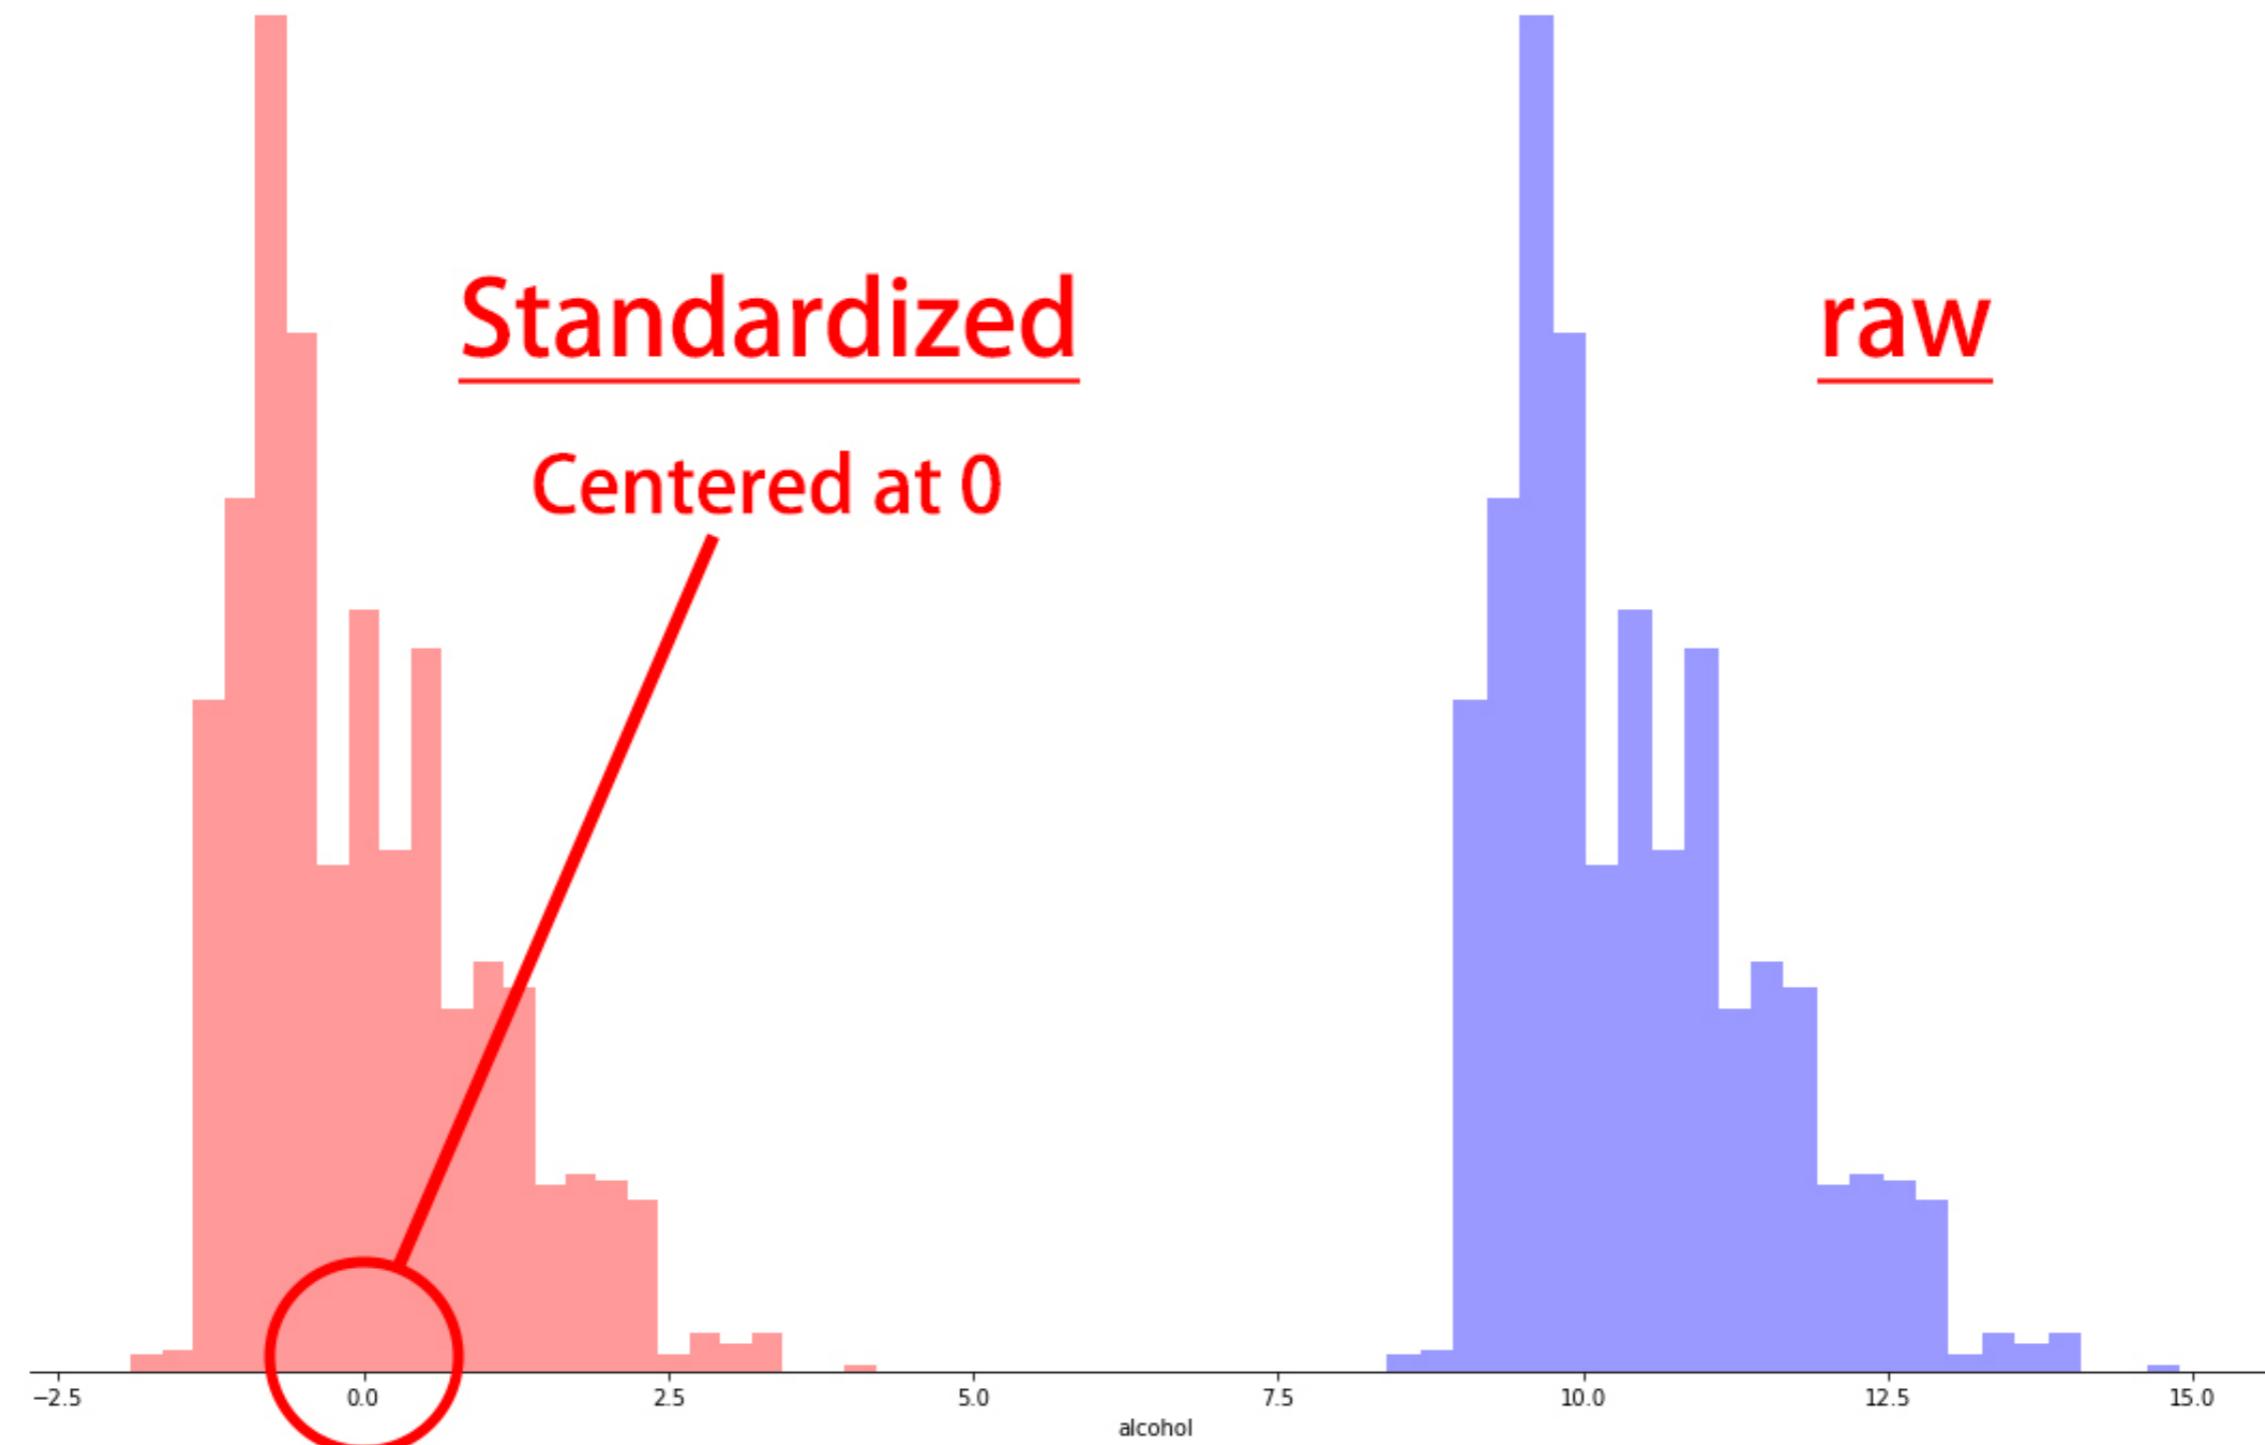
# Batch size in Keras

```
# Fitting an already built and compiled model  
model.fit(X_train, y_train, epochs=100, batch_size=128)  
^ ^ ^ ^ ^ ^ ^ ^ ^ ^ ^ ^
```

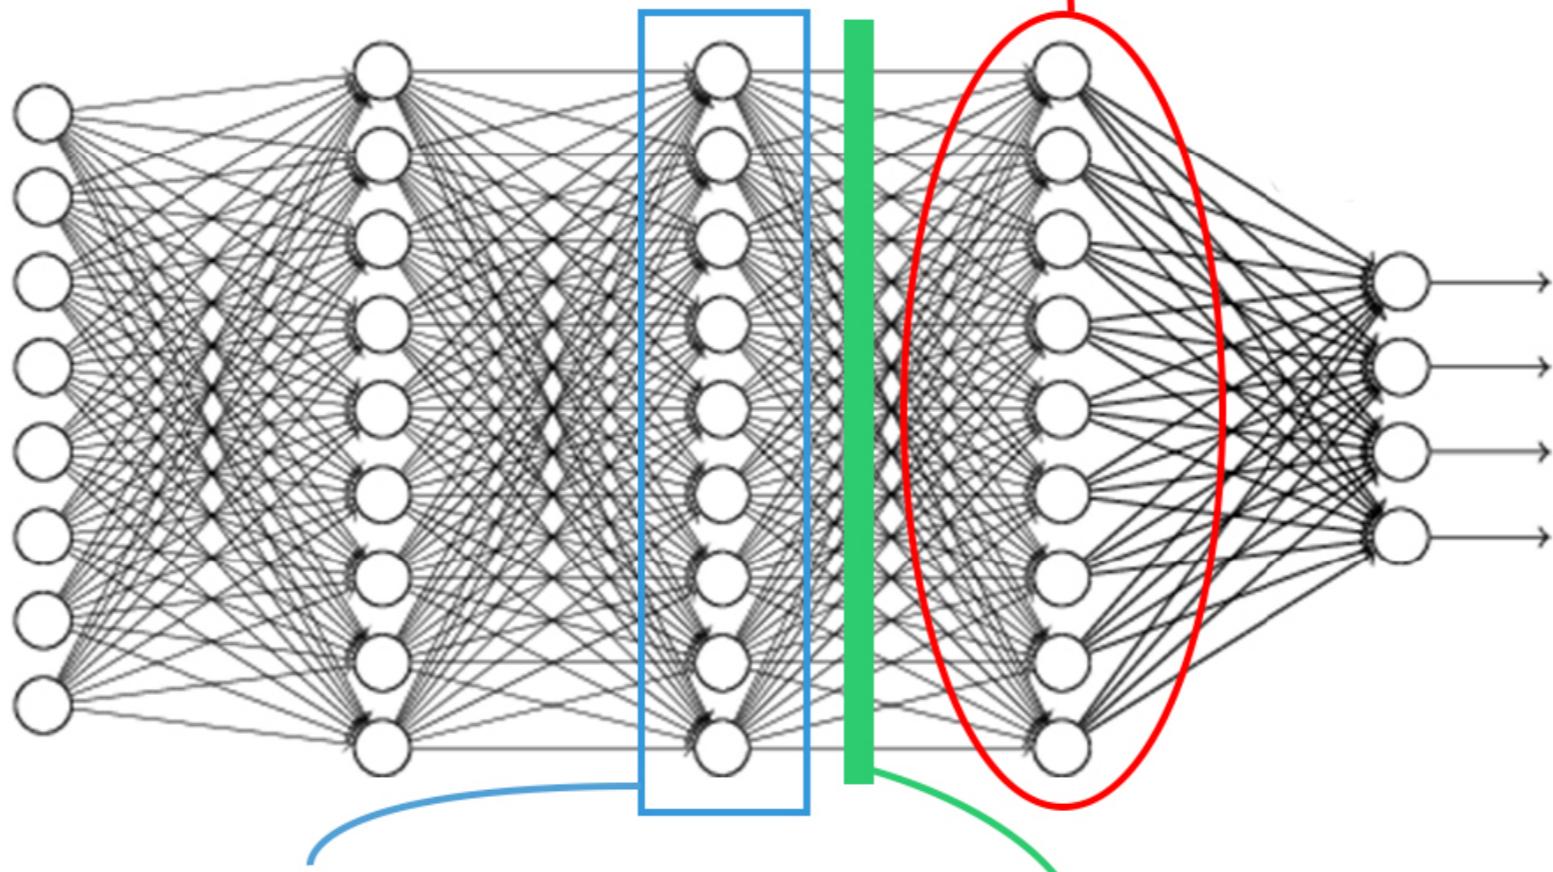
# Standardization (a normalization approach)

$\frac{\text{data} - \text{mean}}{\text{standard deviation}}$

---



1. This layer weights are trained based on the previous layer outputs it receives



2. But when this layer updates its weights via gradient descent, its outputs are also updated

3. Batch normalization makes sure that independently of the changes, the inputs to the next layer are normalized

# Batch normalization advantages

- Improves gradient flow
- Allows higher learning rates
- Reduces dependence on weight initializations
- Acts as an unintended form of regularization
- Limits internal covariate shift

# Batch normalization in Keras

```
# Import BatchNormalization from keras layers
from keras.layers import BatchNormalization

# Instantiate a Sequential model
model = Sequential()

# Add an input layer
model.add(Dense(3, input_shape=(2,), activation = 'relu'))

# Add batch normalization for the outputs of the layer above
model.add(BatchNormalization())

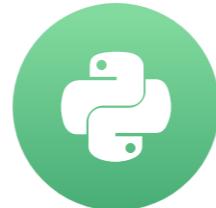
# Add an output layer
model.add(Dense(1, activation='sigmoid'))
```

# Let's practice!

INTRODUCTION TO DEEP LEARNING WITH KERAS

# Hyperparameter tuning

INTRODUCTION TO DEEP LEARNING WITH KERAS



Miguel Esteban  
Data Scientist & Founder

# Neural network hyperparameters

- Number of layers
- Number of neurons per layer
- Layer order
- Layer activations
- Batch sizes
- Learning rates
- Optimizers
- ...

# Sklearn recap

```
# Import RandomizedSearchCV
from sklearn.model_selection import RandomizedSearchCV

# Instantiate your classifier
tree = DecisionTreeClassifier()

# Define a series of parameters to look over
params = {'max_depth':[3,None], "max_features":range(1,4), 'min_samples_leaf': range(1,4)}

# Perform random search with cross validation
tree_cv = RandomizedSearchCV(tree, params, cv=5)

tree_cv.fit(X,y)

# Print the best parameters
print(tree_cv.best_params_)
```

```
{'min_samples_leaf': 1, 'max_features': 3, 'max_depth': 3}
```

# Turn a Keras model into a Sklearn estimator

```
# Function that creates our Keras model
def create_model(optimizer='adam', activation='relu'):

    model = Sequential()
    model.add(Dense(16, input_shape=(2, ), activation=activation))
    model.add(Dense(1, activation='sigmoid'))
    model.compile(optimizer=optimizer, loss='binary_crossentropy')
    return model

# Import sklearn wrapper from keras
from keras.wrappers.scikit_learn import KerasClassifier

# Create a model as a sklearn estimator
model = KerasClassifier(build_fn=create_model, epochs=6, batch_size=16)
```

# Cross-validation

```
# Import cross_val_score
from sklearn.model_selection import cross_val_score

# Check how your keras model performs with 5 fold crossvalidation
kfold = cross_val_score(model, X, y, cv=5)

# Print the mean accuracy per fold
kfold.mean()
```

0.913333

```
# Print the standard deviation per fold
kfold.std()
```

0.110754

# Tips for neural networks hyperparameter tuning

- Random search is preferred over grid search
- Don't use many epochs
- Use a smaller sample of your dataset
- Play with batch sizes, activations, optimizers and learning rates

# Random search on Keras models

```
# Define a series of parameters
params = dict(optimizer=['sgd', 'adam'], epochs=3,
               batch_size=[5, 10, 20], activation=['relu', 'tanh'])

# Create a random search cv object and fit it to the data
random_search = RandomizedSearchCV(model, params_dist=params, cv=3)

random_search_results = random_search.fit(X, y)

# Print results
print("Best: %f using %s".format(random_search_results.best_score_,
random_search_results.best_params_))
```

```
Best: 0.94 using {'optimizer': 'adam', 'epochs': 3, 'batch_size': 10, 'activation': 'rel
```

# Tuning other hyperparameters

```
def create_model(nl=1,nn=256):  
    model = Sequential()  
    model.add(Dense(16, input_shape=(2,), activation='relu'))  
    # Add as many hidden layers as specified in nl  
    for i in range(nl):  
        # Layers have nn neurons  
        model.add(Dense(nn, activation='relu'))  
    # End defining and compiling your model...
```

# Tuning other hyperparameters

```
# Define parameters, named just like in create_model()  
params = dict(nl=[1, 2, 9], nn=[128,256,1000])  
  
# Repeat the random search...  
  
# Print results...
```

```
Best: 0.87 using {'nl': 2, 'nn': 128}
```

# Let's tune some networks!

INTRODUCTION TO DEEP LEARNING WITH KERAS

# Tensors, layers and autoencoders

INTRODUCTION TO DEEP LEARNING WITH KERAS



Miguel Esteban  
Data Scientist & Founder

# Accessing Keras layers

```
# Accessing the first layer of a Keras model  
first_layer = model.layers[0]  
  
# Printing the layer, and its input, output and weights  
print(first_layer.input)  
print(first_layer.output)  
print(first_layer.weights)
```

```
<tf.Tensor 'dense_1_input:0' shape=(?, 3) dtype=float32>
```

```
<tf.Tensor 'dense_1/Relu:0' shape=(?, 2) dtype=float32>
```

```
[<tf.Variable 'dense_1/kernel:0' shape=(3, 2) dtype=float32_ref>,  
<tf.Variable 'dense_1/bias:0' shape=(2,) dtype=float32_ref>]
```

# What are tensors?

```
# Defining a rank 2 tensor (2 dimensions)
T2 = [[1,2,3],
      [4,5,6],
      [7,8,9]]

# Defining a rank 3 tensor (3 dimensions)
T3 = [[1,2,3],
      [4,5,6],
      [7,8,9],  
  
      [10,11,12],
      [13,14,15],
      [16,17,18],  
  
      [19,20,21],
      [22,23,24],
      [25,26,27]]
```

```
# Import Keras backend
import keras.backend as K

# Get the input and output tensors of a model layer
inp = model.layers[0].input
out = model.layers[0].output

# Function that maps layer inputs to outputs
inp_to_out = K.function([inp], [out])

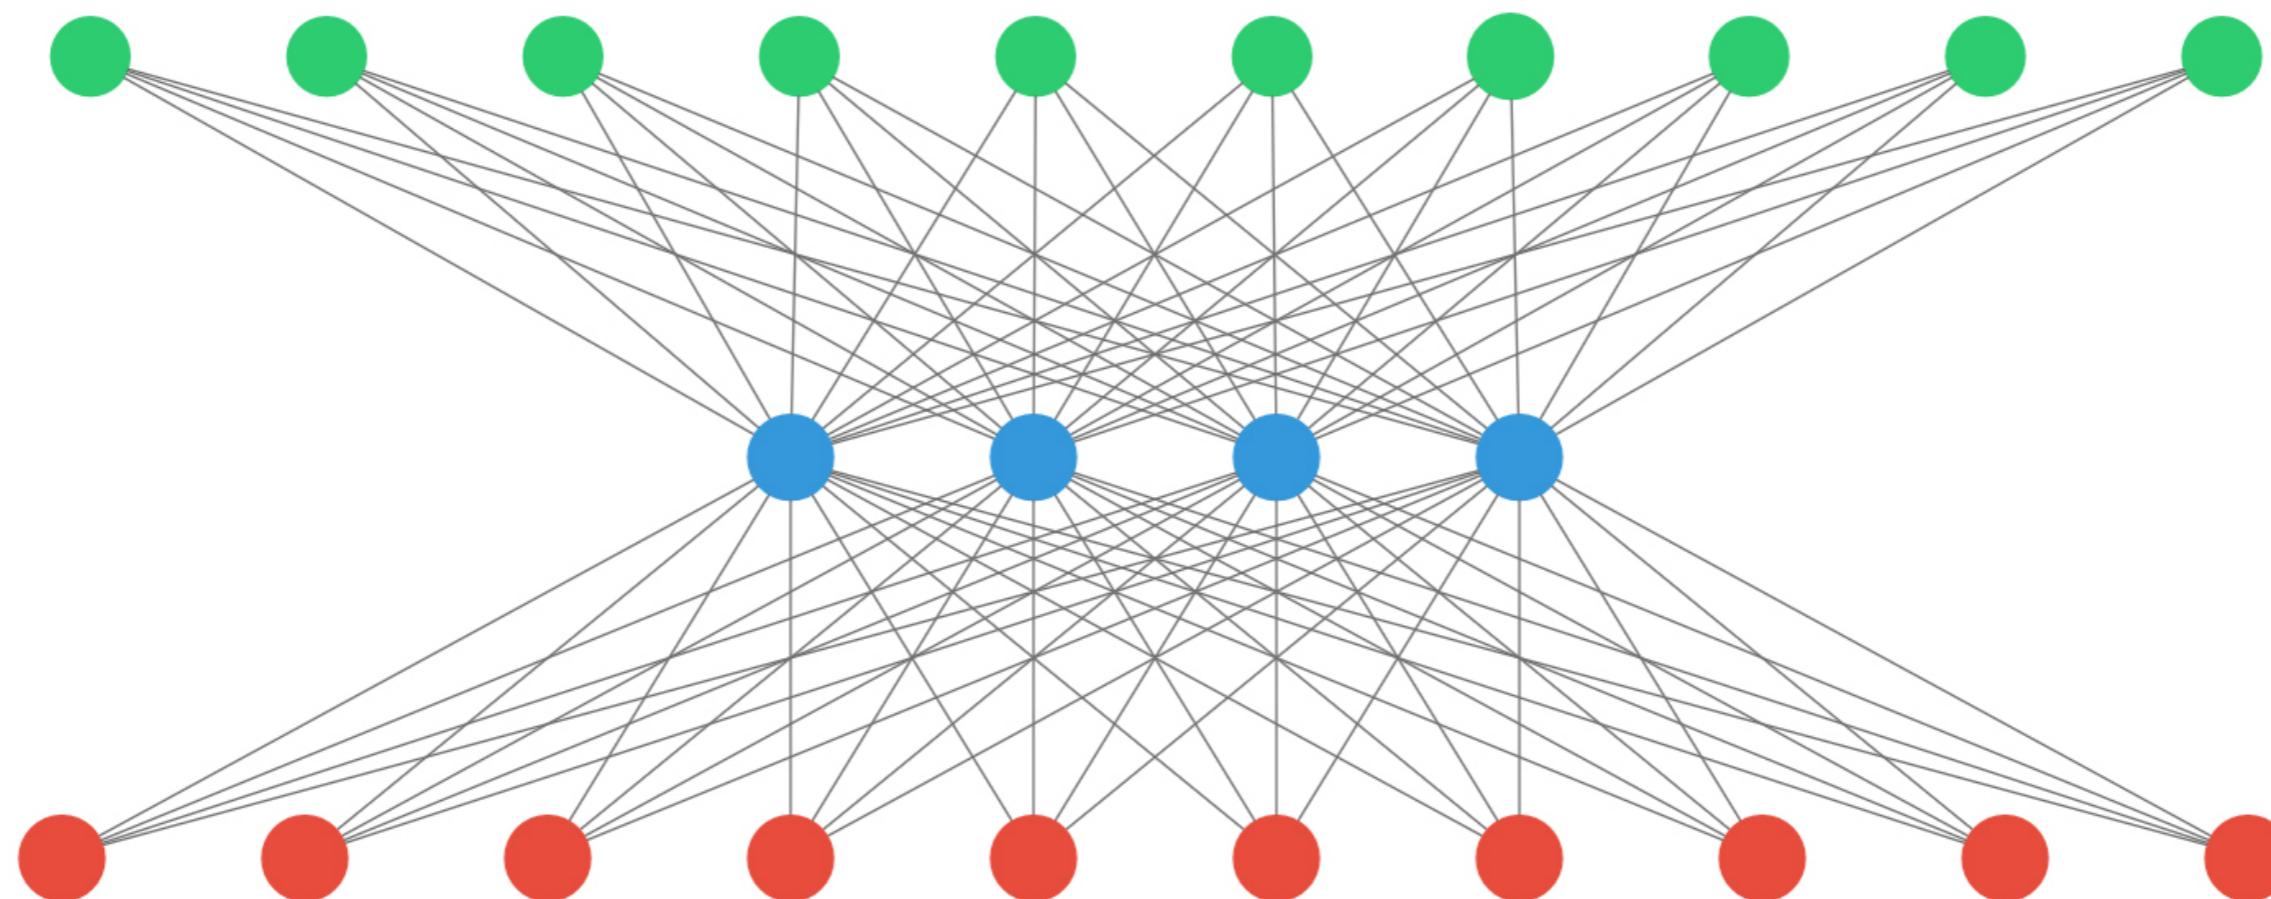
# We pass an input and get the output we'd get in that first layer
print(inp_to_out([X_train]))
```

```
# Outputs of the first layer per sample in X_train
[array([[0.7, 0], ..., [0.1, 0.3]])]
```

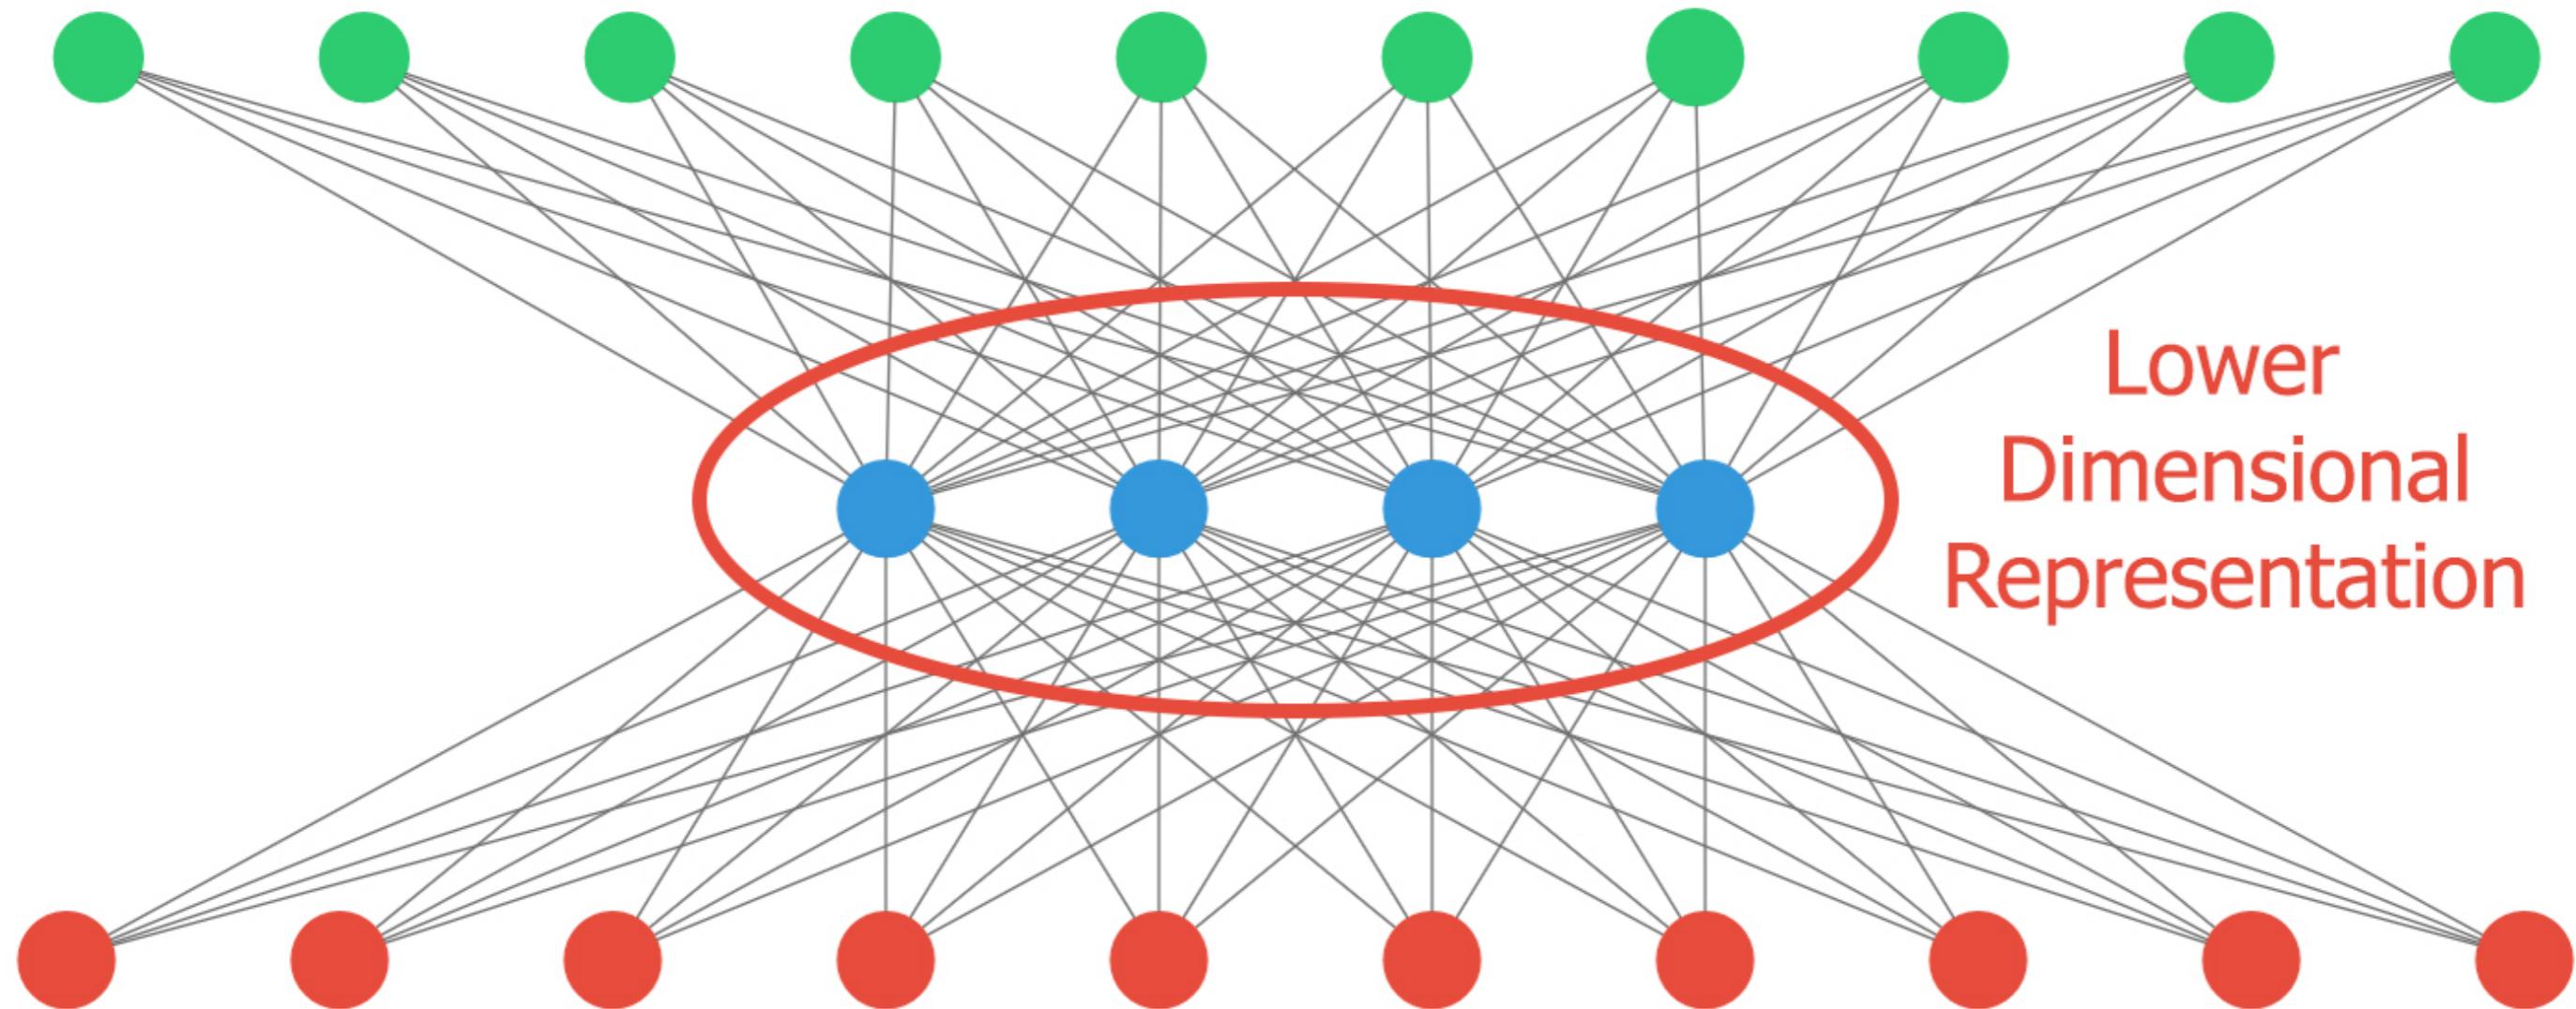
it's time to introduce a new architecture...

# Autoencoders!

INPUTS == OUTPUTS



# Autoencoders!



# Autoencoder use cases

- Dimensionality reduction:
  - Smaller dimensional space representation of our inputs.
- De-noising data:
  - If trained with clean data, irrelevant noise will be filtered out during reconstruction.
- Anomaly detection:
  - A poor reconstruction will result when the model is fed with unseen inputs.
- ...



# Building a simple autoencoder

```
# Instantiate a sequential model
autoencoder = Sequential()

# Add a hidden layer of 4 neurons and an input layer of 100
autoencoder.add(Dense(4, input_shape=(100, ), activation='relu'))

# Add an output layer of 100 neurons
autoencoder.add(Dense(100, activation='sigmoid'))

# Compile your model with the appropiate loss
autoencoder.compile(optimizer='adam', loss='binary_crossentropy')
```

# Breaking it into an encoder

```
# Building a separate model to encode inputs  
encoder = Sequential()  
encoder.add(autoencoder.layers[0])  
  
# Predicting returns the four hidden layer neuron outputs  
encoder.predict(X_test)
```

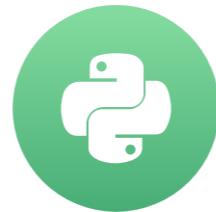
```
# Four numbers for each observation in X_test  
array([10.0234375, 5.833543, 18.90444, 9.20348],...)
```

# Let's experiment!

INTRODUCTION TO DEEP LEARNING WITH KERAS

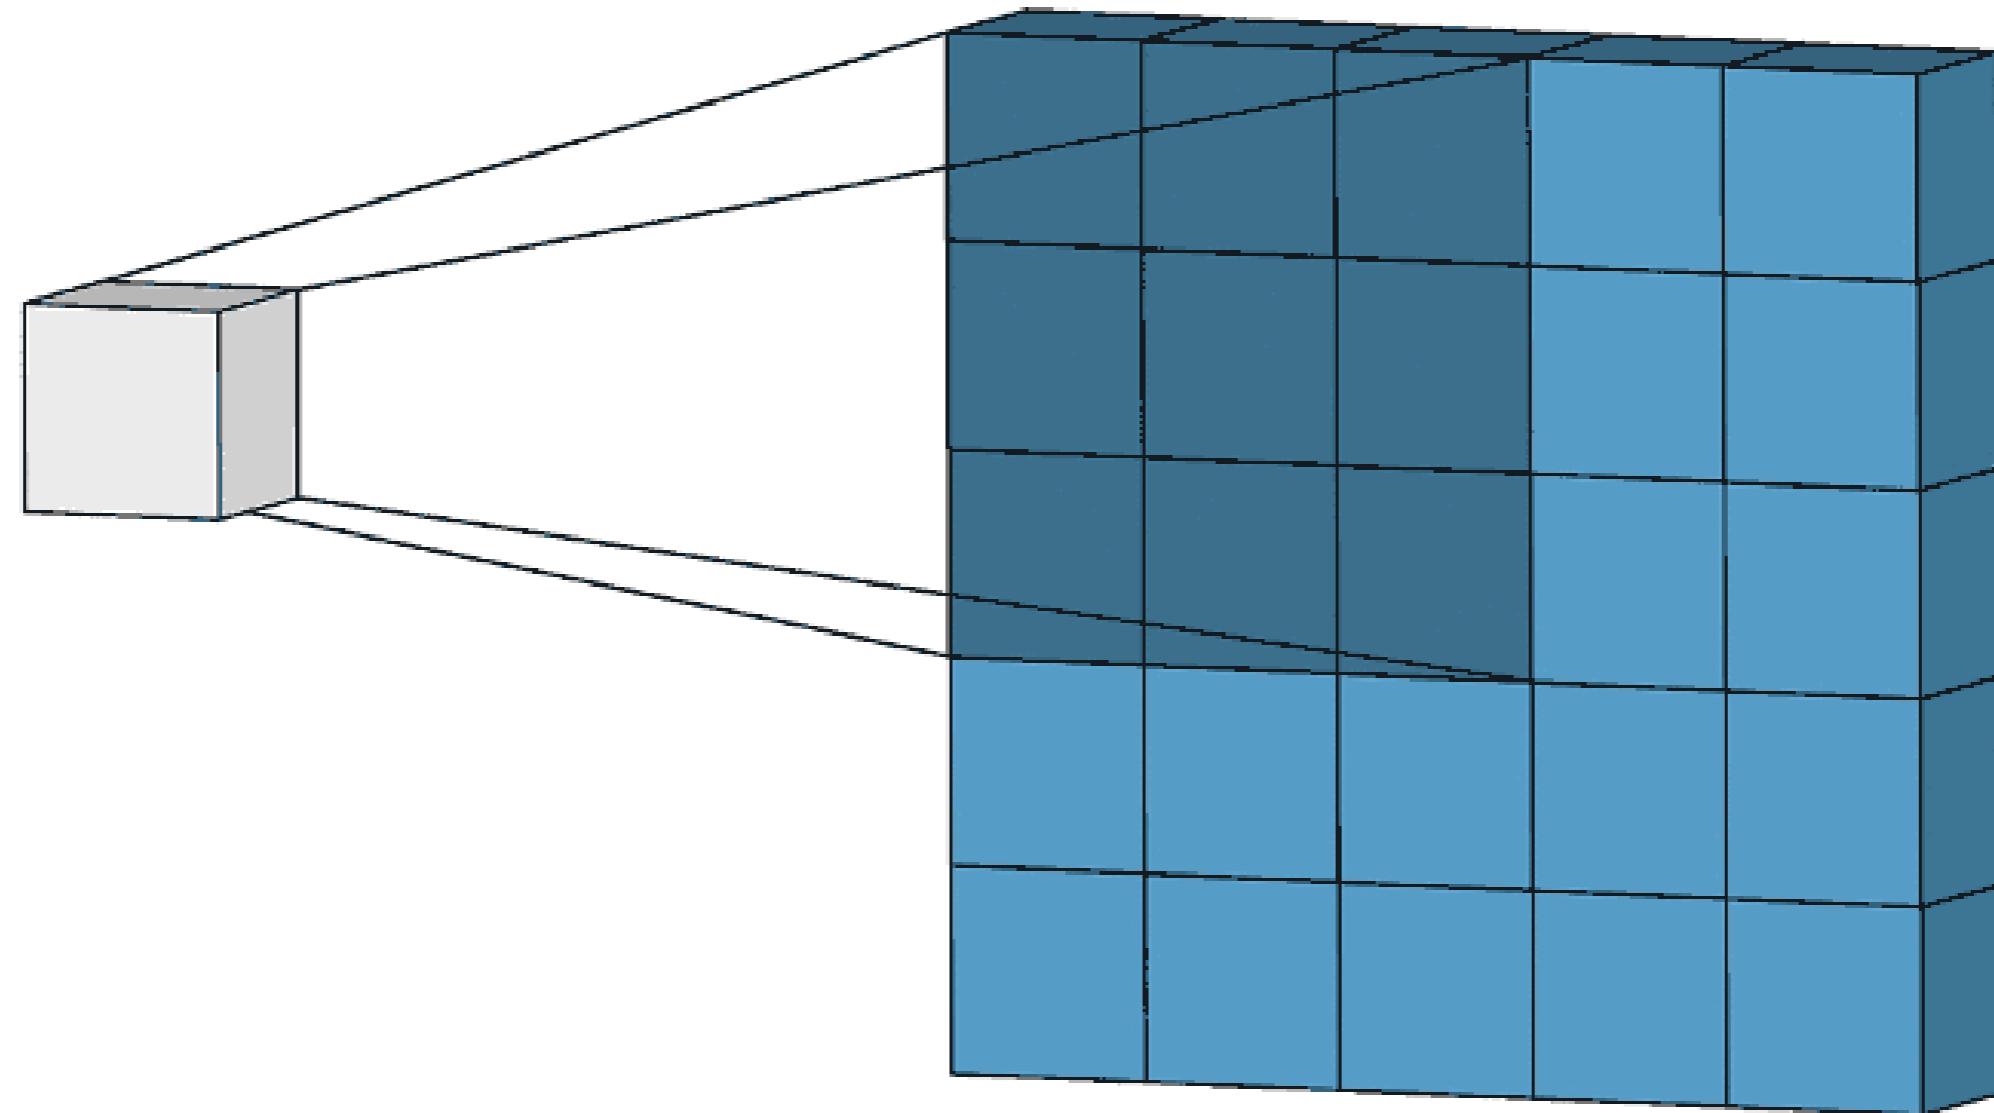
# Intro to CNNs

INTRODUCTION TO DEEP LEARNING WITH KERAS



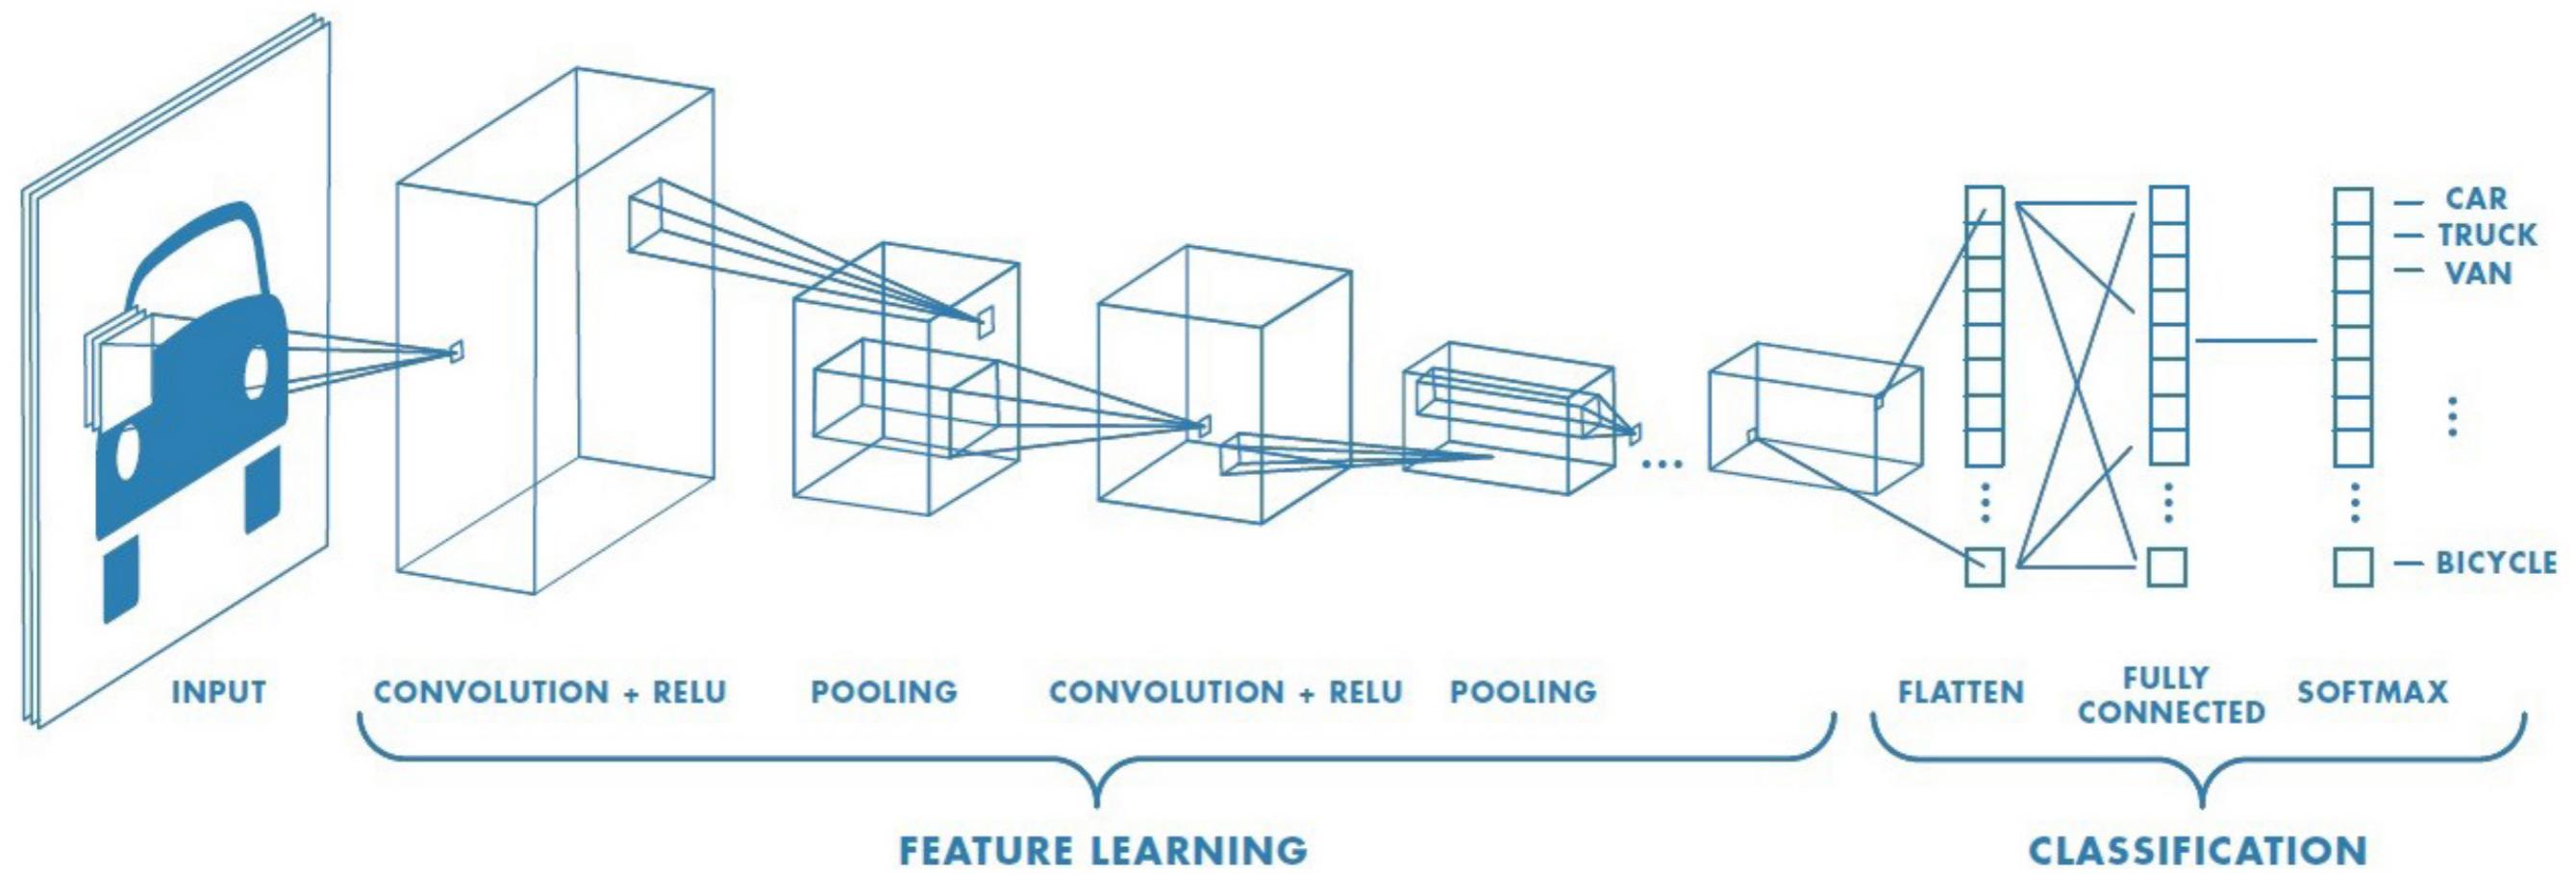
Miguel Esteban

Data Scientist & Founder



3 0	3 1	2 2	1	0
0 2	0 2	1 0	3	1
3 0	1 1	2 2	2	3
2	0	0	2	2
2	0	0	0	1

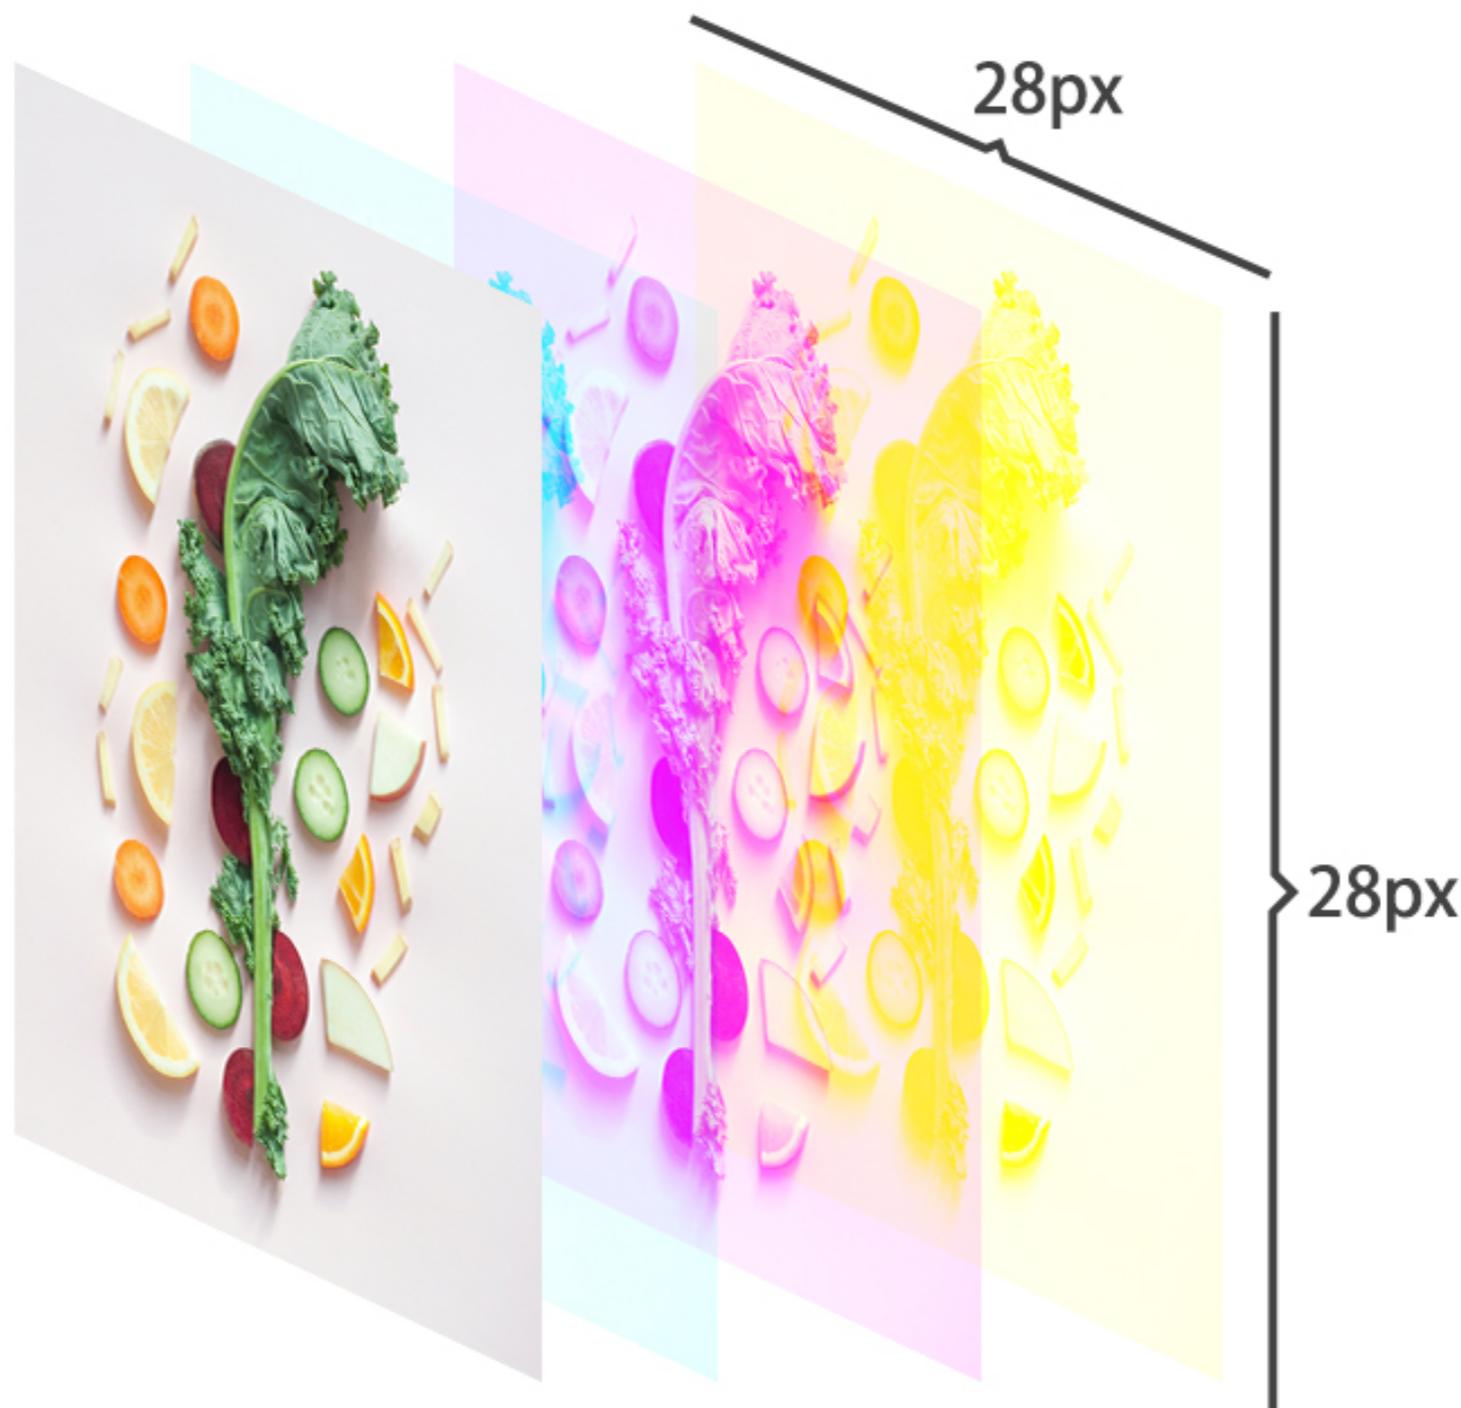
12.0	12.0	17.0
10.0	17.0	19.0
9.0	6.0	14.0



`input_shape  
(WIDTH,HEIGHT,CHANNELS)`

`input_shape = (28,28,3)`

Color Images have 3 channels  
**RGB** (Red Green Blue)



```
# Import Conv2D layer and Flatten from keras layers
from keras.layers import Dense, Conv2D, Flatten

# Instantiate your model as usual
model = Sequential()

# Add a convolutional layer with 32 filters of size 3x3
model.add(Conv2D(filters=32,
                 kernel_size=3,
                 input_shape=(28, 28, 1),
                 activation='relu'))

# Add another convolutional layer
model.add(Conv2D(8, kernel_size=3, activation='relu'))

# Flatten the output of the previous layer
model.add(Flatten())

# End this multiclass model with 3 outputs and softmax
model.add(Dense(3, activation='softmax'))
```

1



2



3

ResNet50  
(50 Layers)

4

Predicted  
("cheeseburger", 0.99),  
('bagel', 0.01),  
('bakery', 0.00)

# Pre-processing images for ResNet50

```
# Import image from keras preprocessing
from keras.preprocessing import image

# Import preprocess_input from keras applications resnet50
from keras.applications.resnet50 import preprocess_input

# Load the image with the right target size for your model
img = image.load_img(img_path, target_size=(224, 224))

# Turn it into an array
img = image.img_to_array(img)

# Expand the dimensions so that it's understood by our network:
# img.shape turns from (224, 224, 3) into (1, 224, 224, 3)
img = np.expand_dims(img, axis=0)

# Pre-process the img in the same way training images were
img = preprocess_input(img)
```

# Using the ResNet50 model in Keras

```
# Import ResNet50 and decode_predictions from keras.applications.resnet50
from keras.applications.resnet50 import ResNet50, decode_predictions

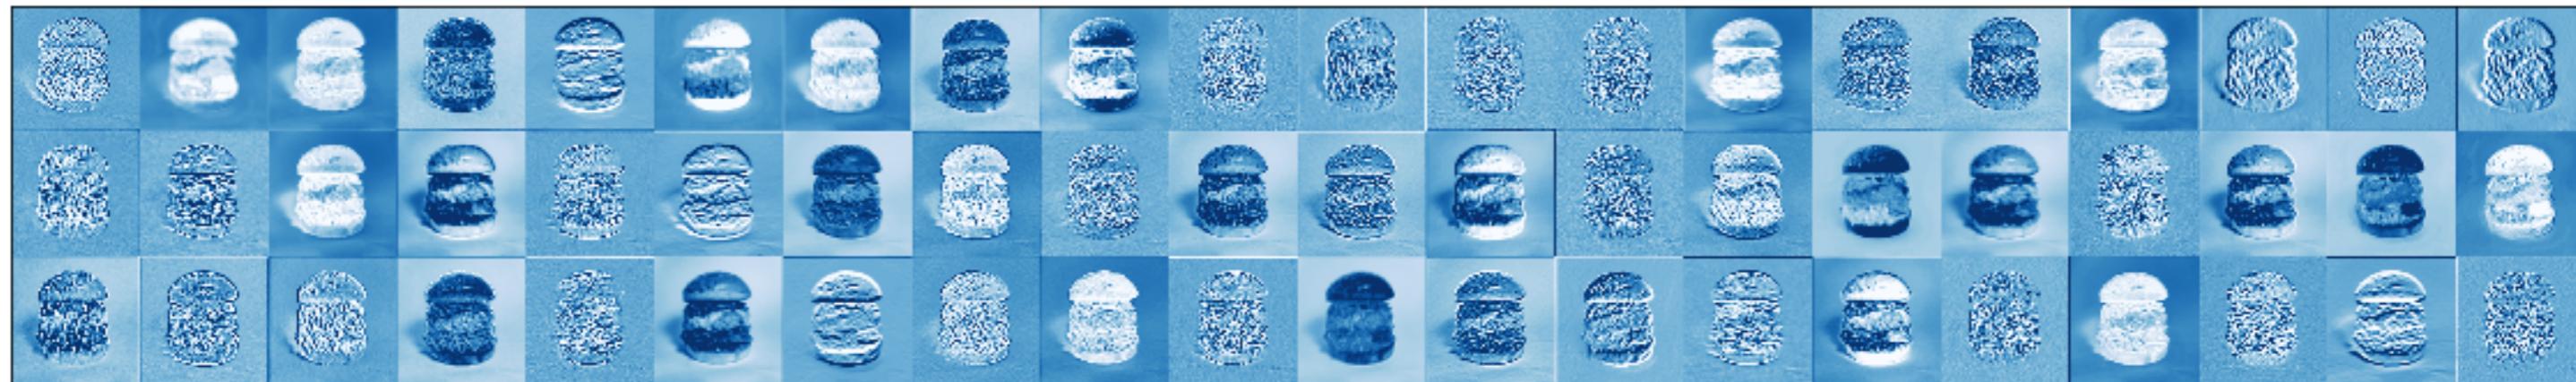
# Instantiate a ResNet50 model with imagenet weights
model = ResNet50(weights='imagenet')

# Predict with ResNet50 on our img
preds = model.predict(img)

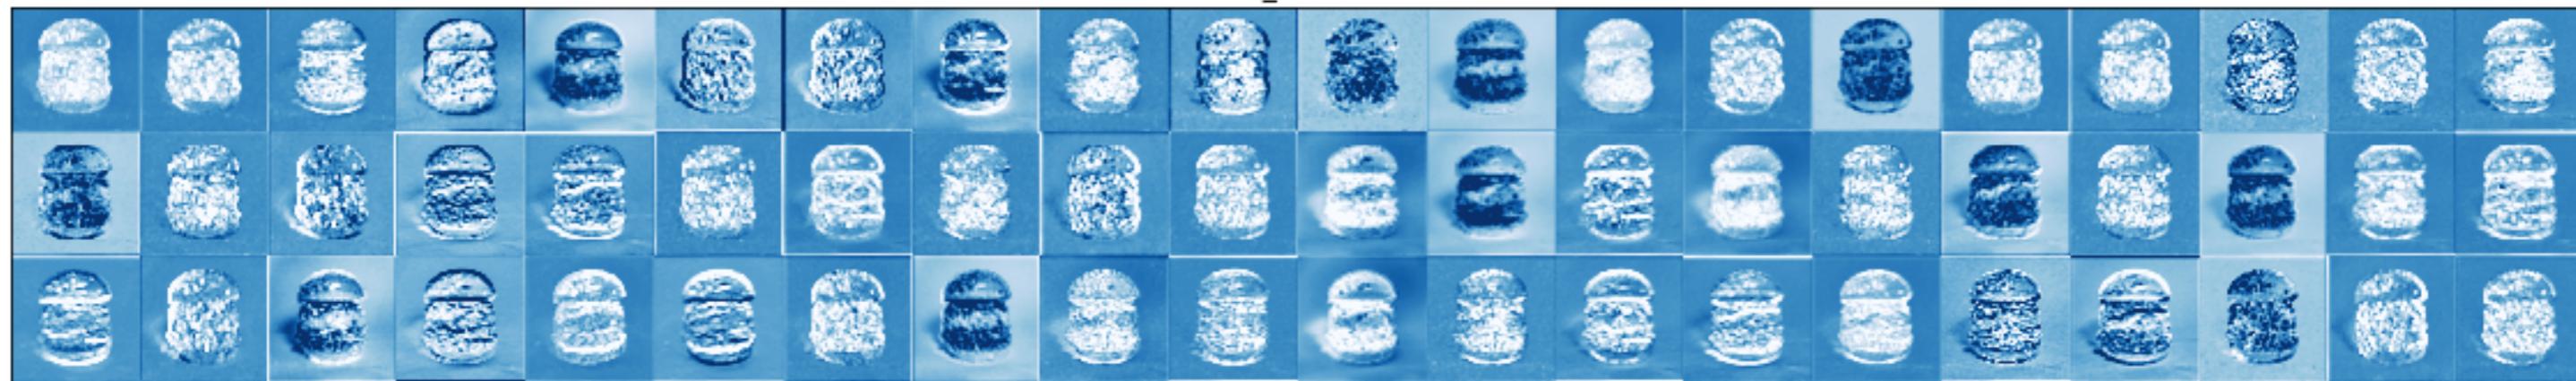
# Decode predictions and print it
print('Predicted:', decode_predictions(preds, top=1)[0])
```

```
Predicted: [(['n07697313', 'cheeseburger', 0.9868016])]
```

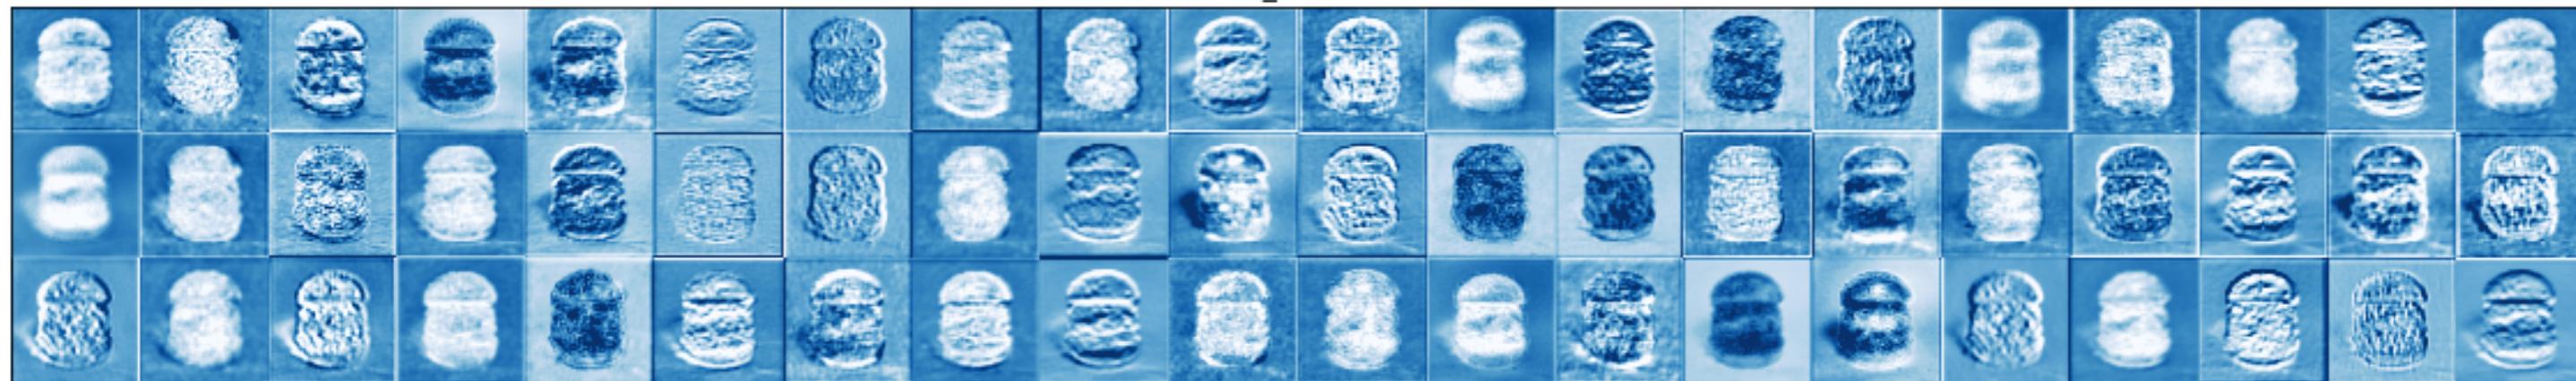
conv1



res2a\_branch2a



res2a\_branch2b



# Let's experiment!

INTRODUCTION TO DEEP LEARNING WITH KERAS

# Intro to LSTMs

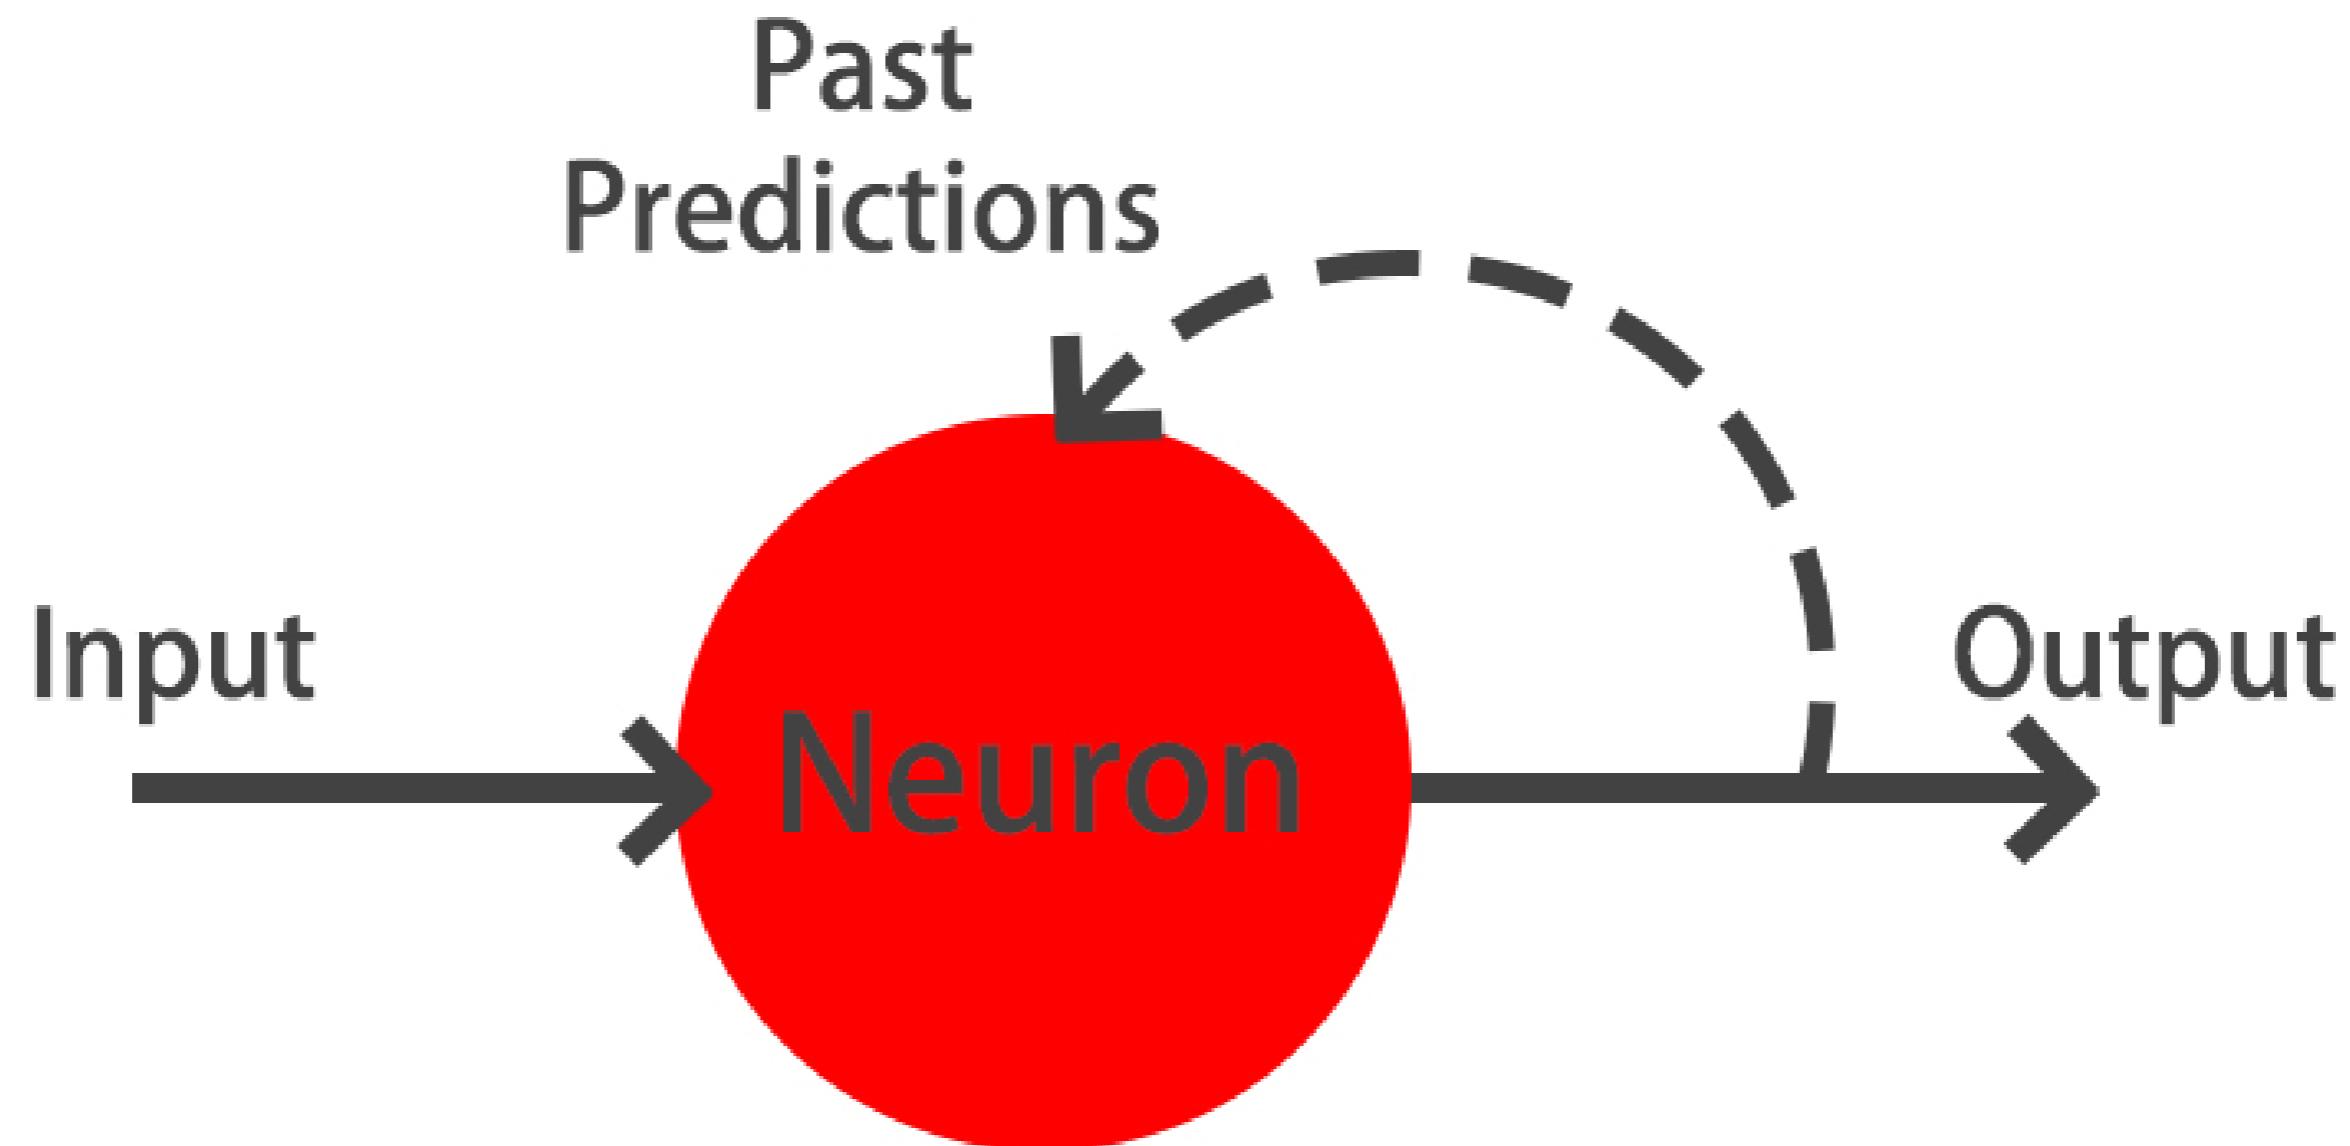
INTRODUCTION TO DEEP LEARNING WITH KERAS

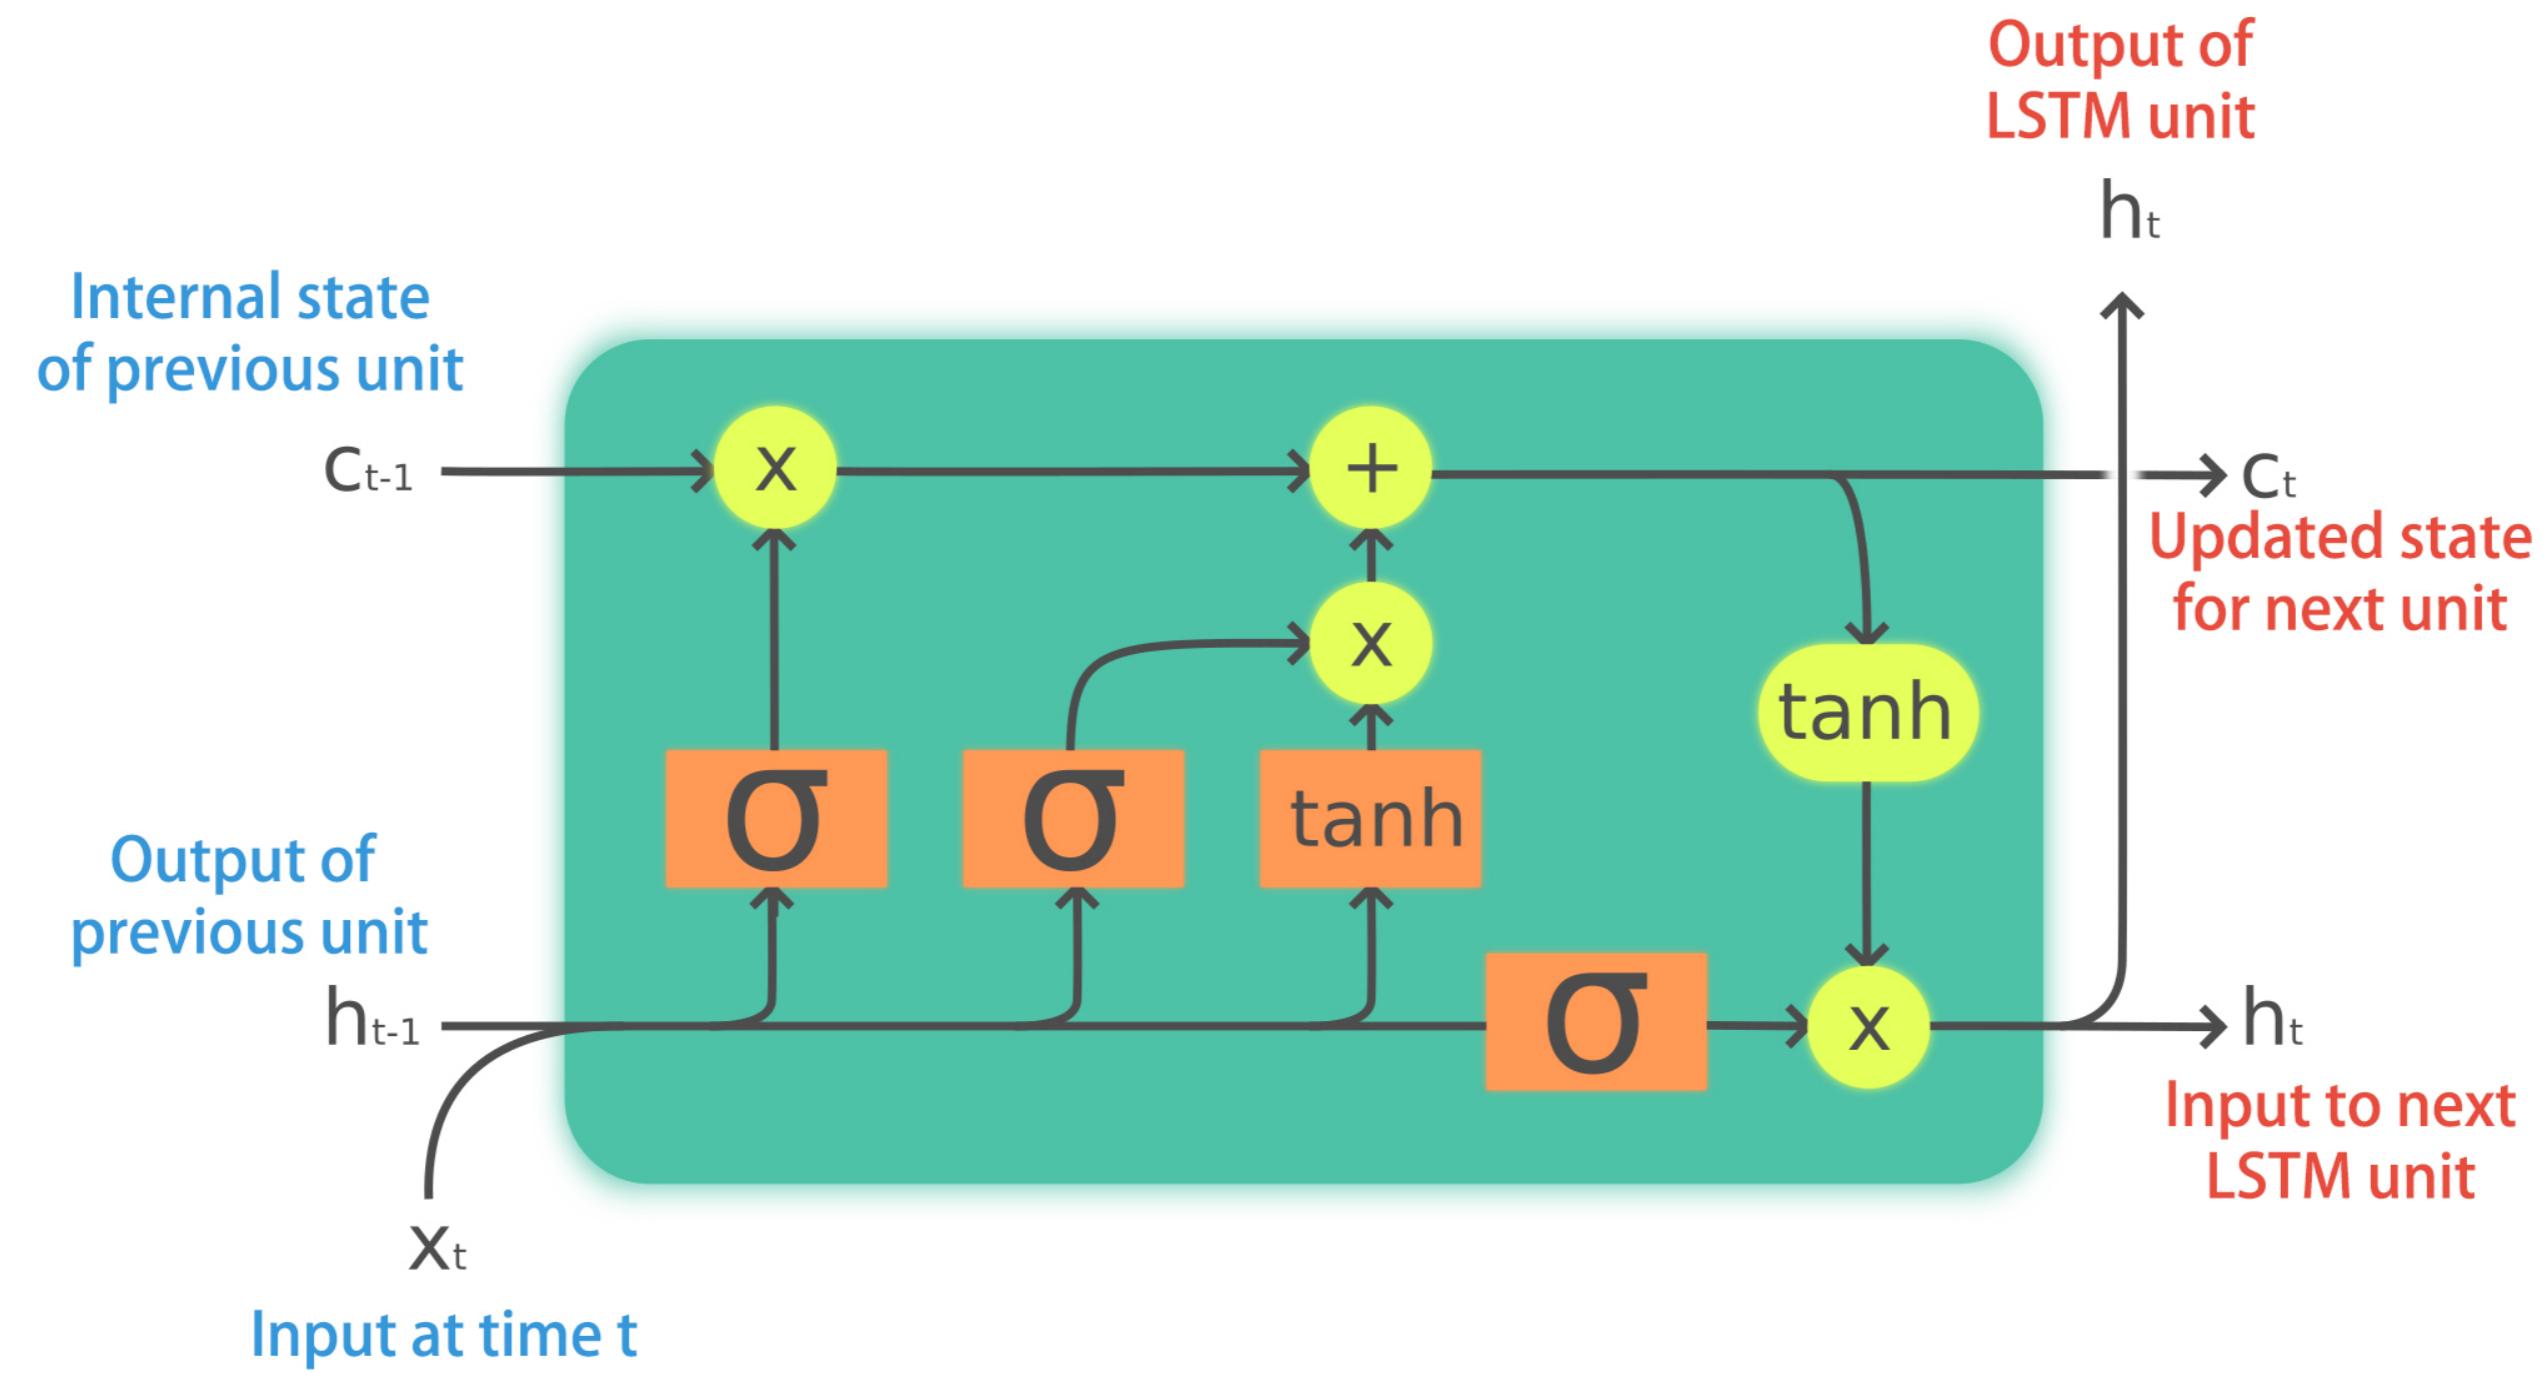


Miguel Esteban

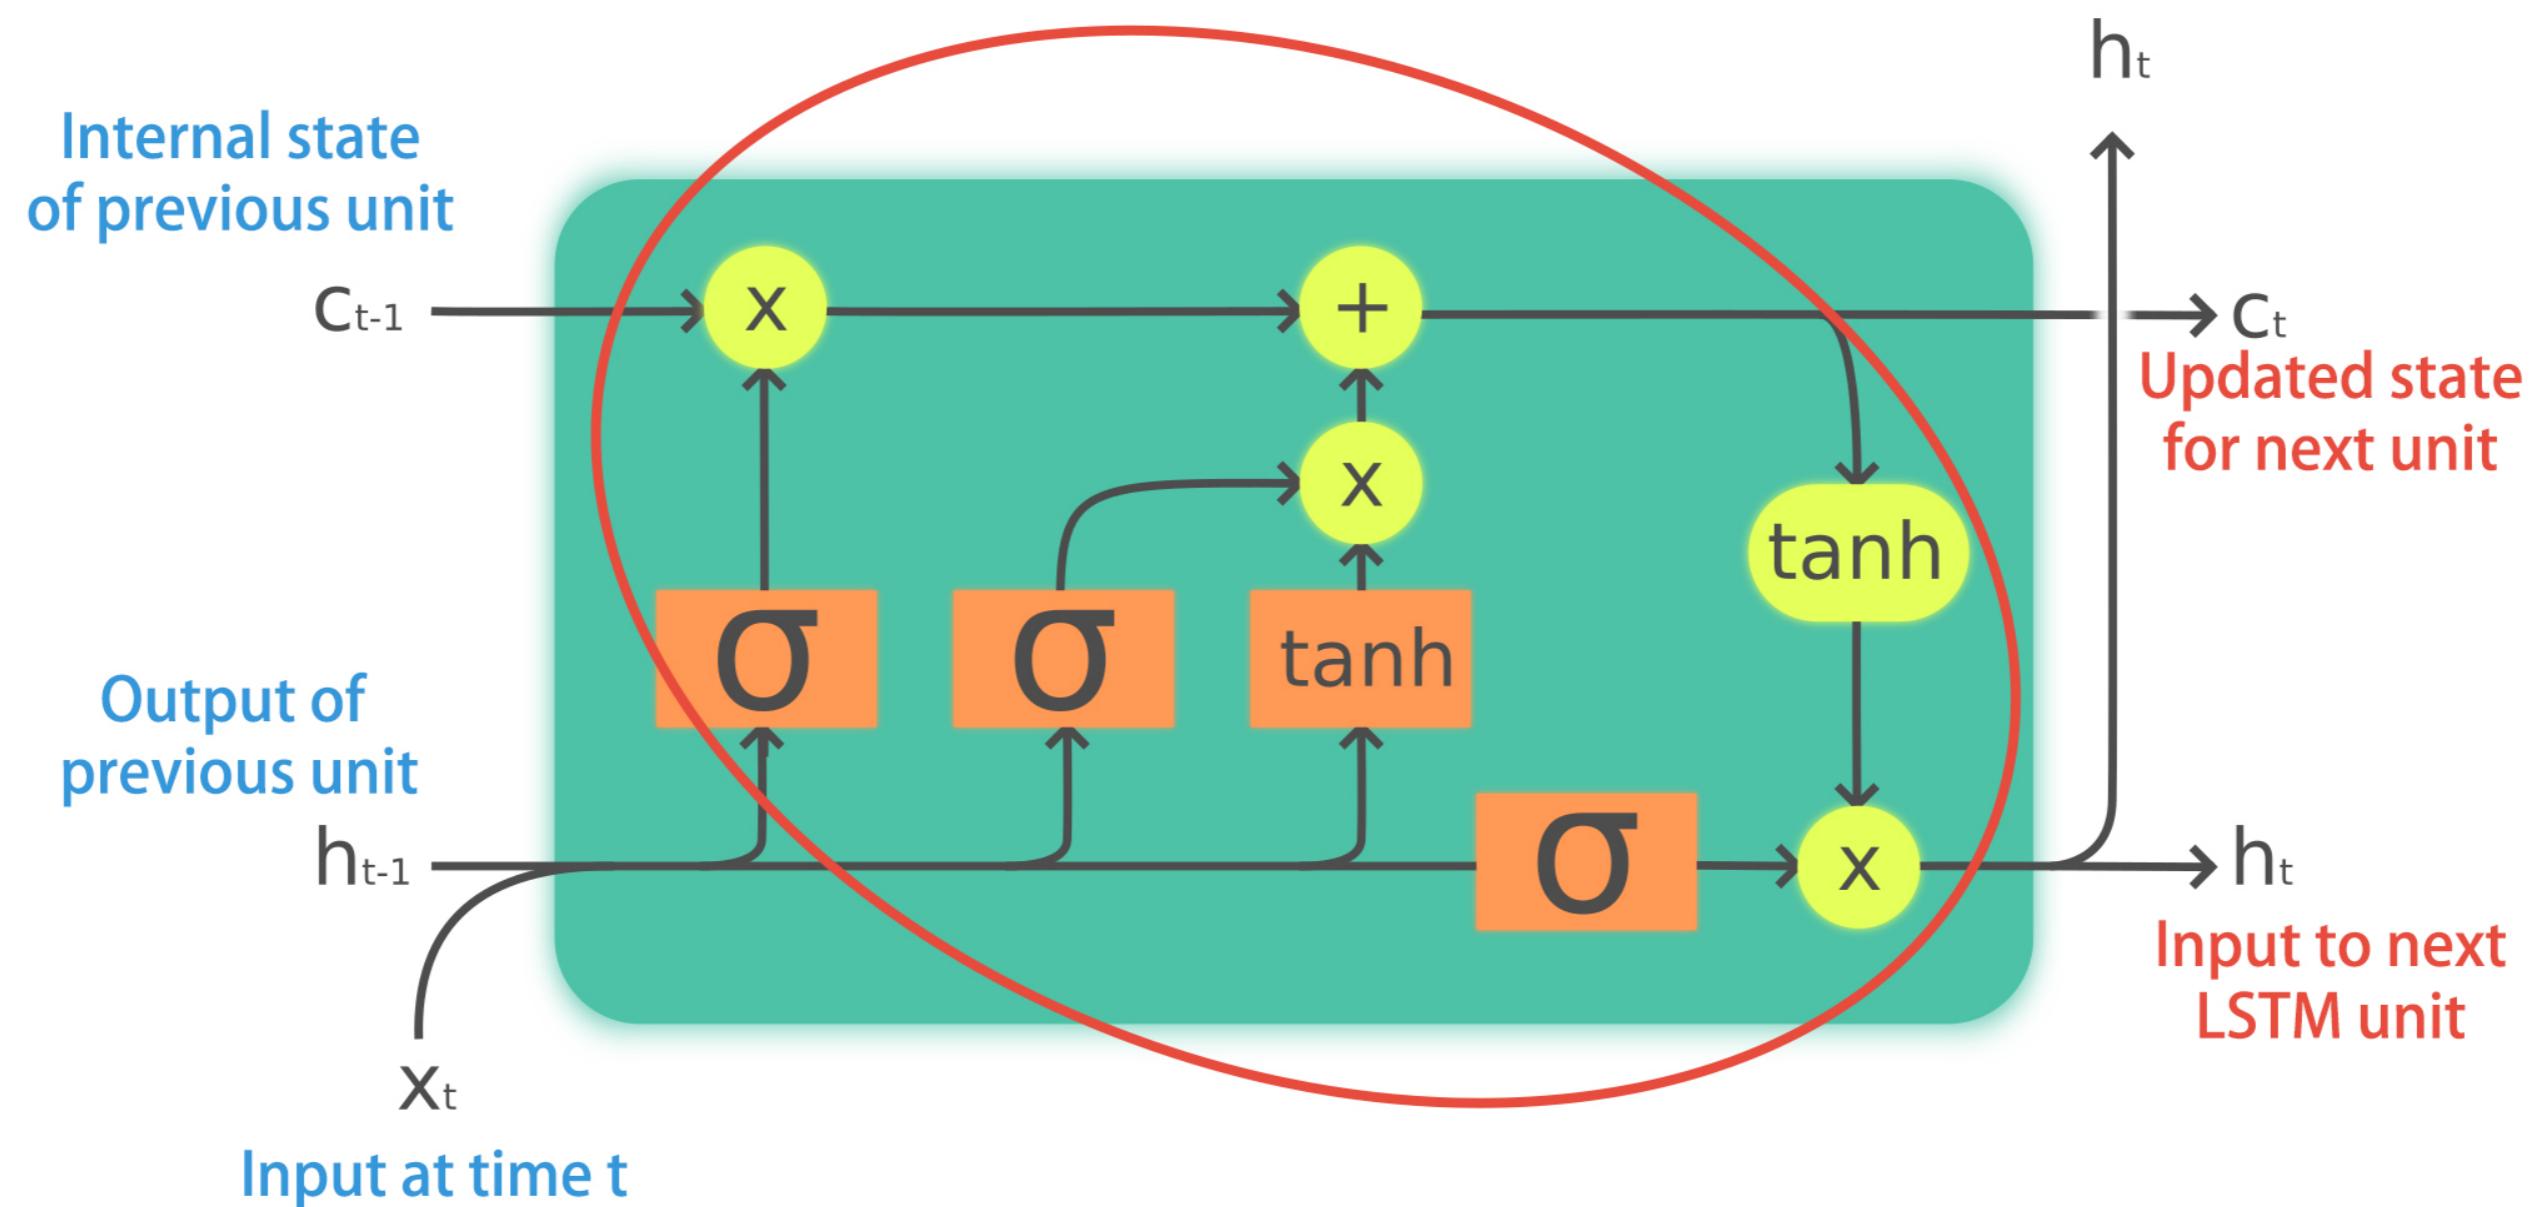
Data Scientist & Founder

# What are RNNs?





They learn what to ignore, what to keep  
and to select the most important pieces  
of past information.



# When to use LSTMs?

- Image captioning
- Speech to text
- Text translation
- Document summarization
- Text generation
- Musical composition
- ...

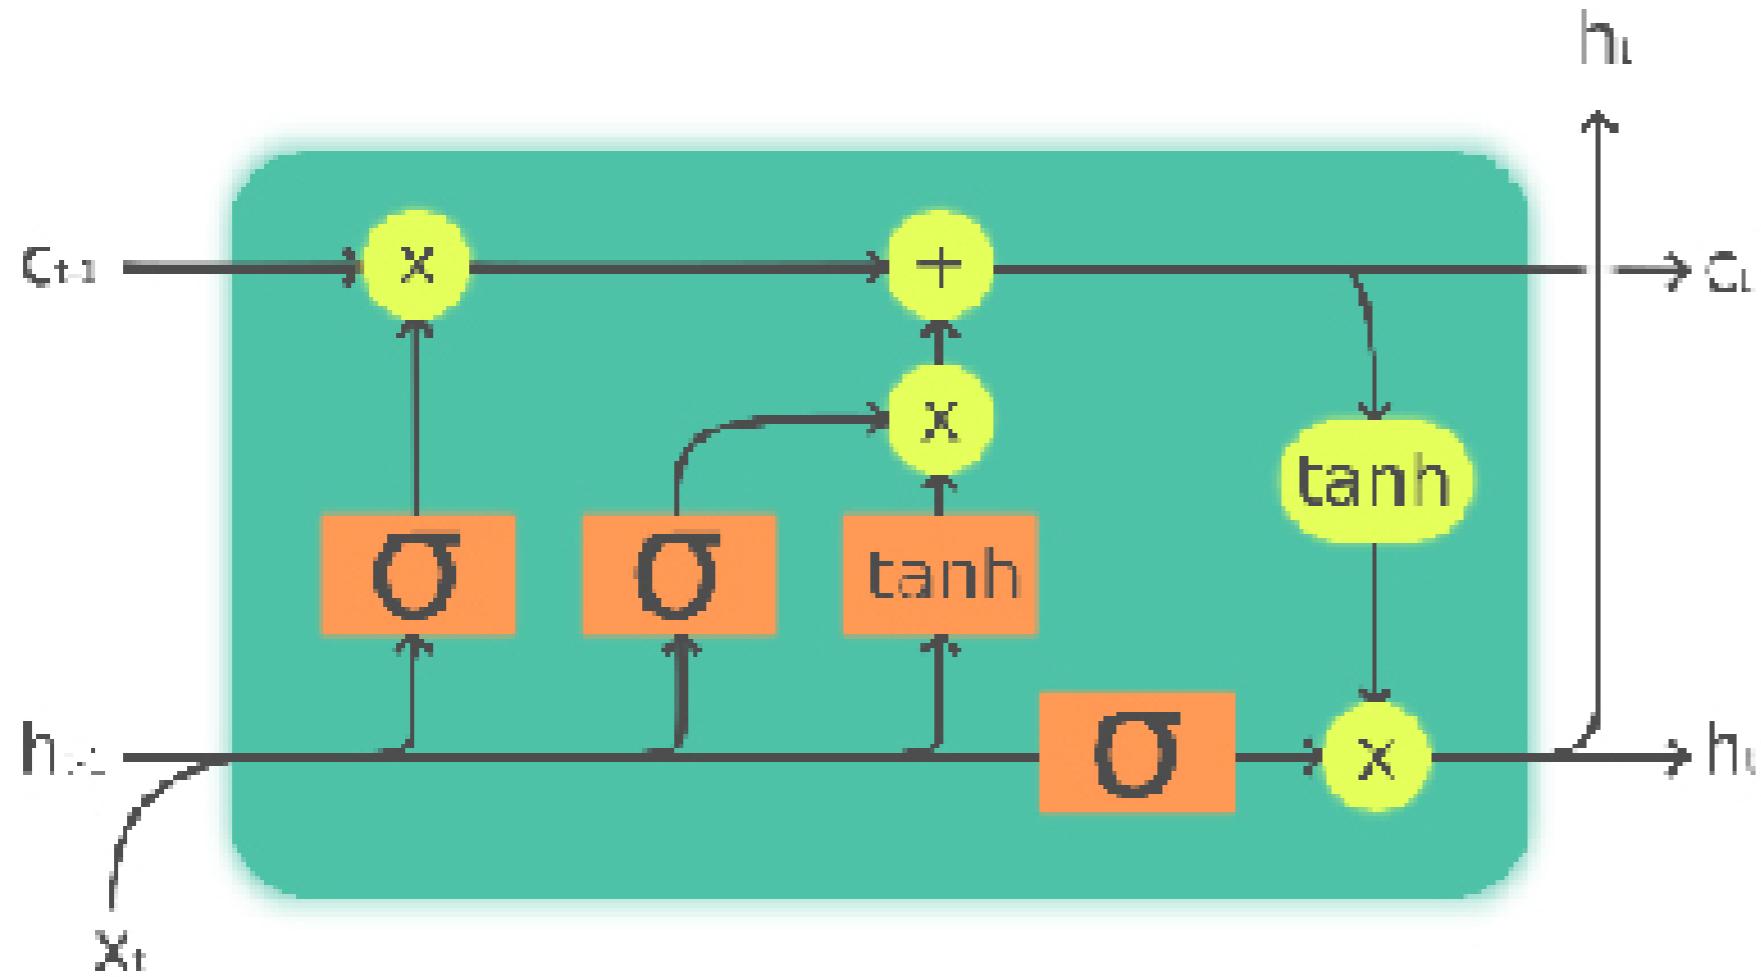


"man in black shirt is playing guitar."



"construction worker in orange safety vest is working on road."

<sup>1</sup> Karpathy, A., & Fei <sup>2</sup> Fei, L. (2015). Deep visual <sup>3</sup> semantic alignments for generating image descriptions.

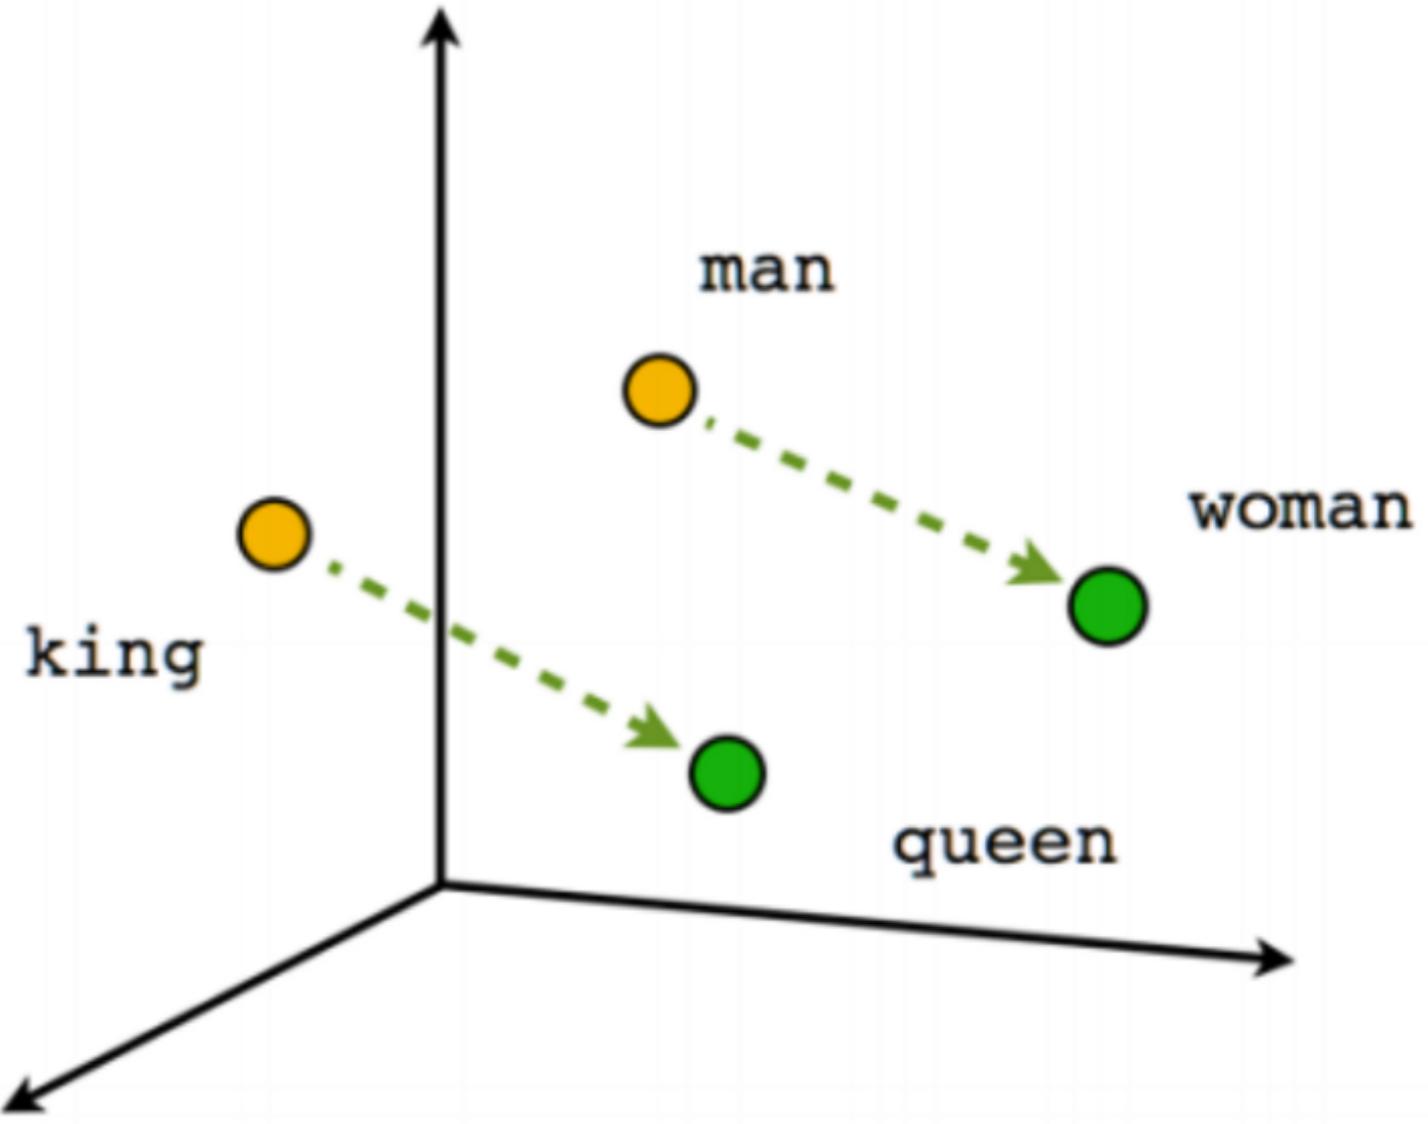
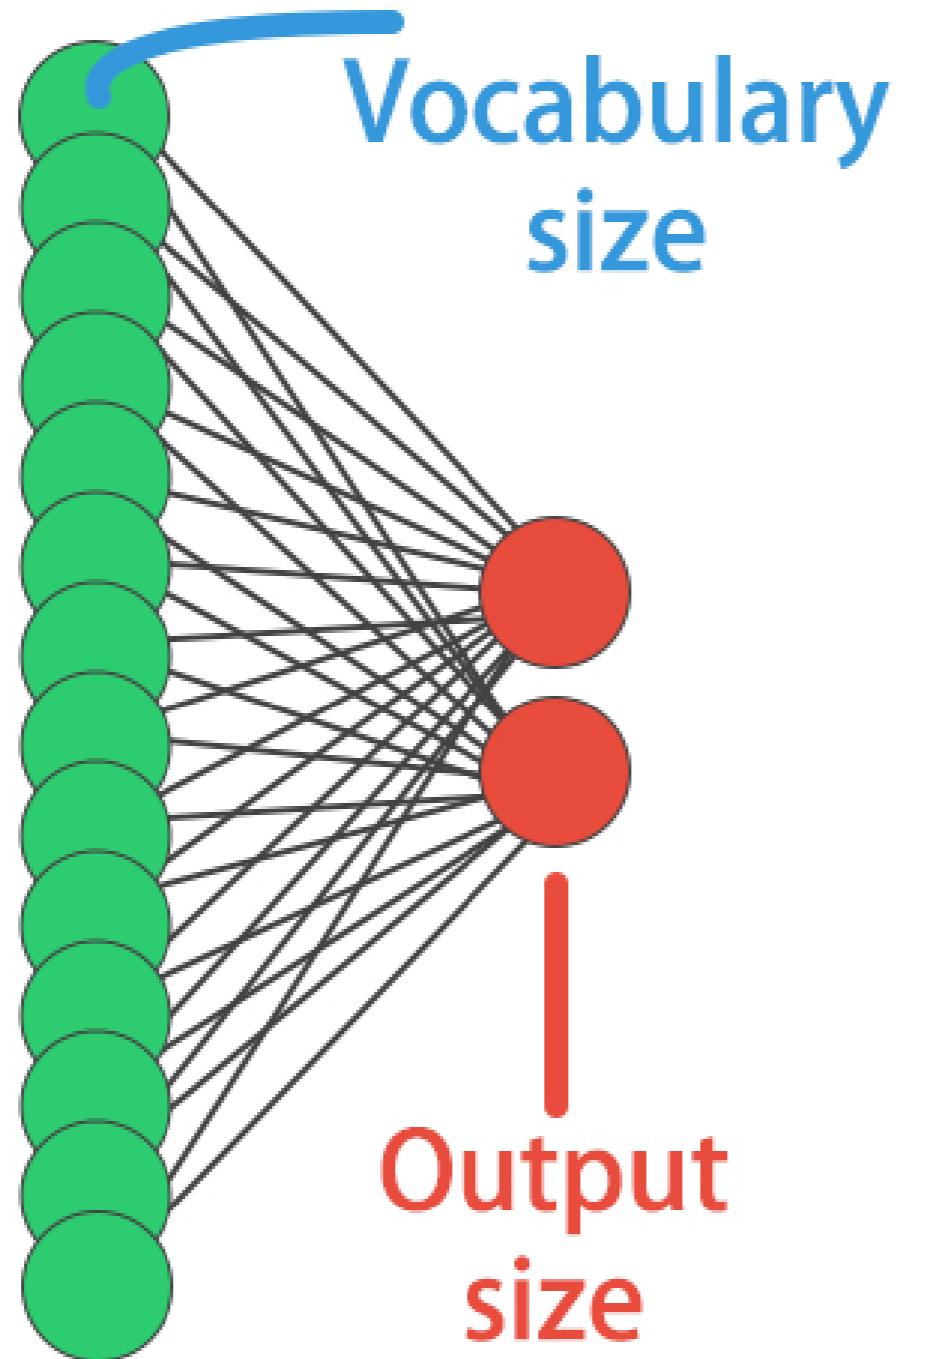


+ Text

this is a sentence



42 11 23 1



```
text = 'Hi this is a small sentence'

# We choose a sequence length
seq_len = 3

# Split text into a list of words
words = text.split()
```

```
['Hi', 'this', 'is', 'a', 'small', 'sentence']
```

```
# Make lines
lines = []
for i in range(seq_len, len(words) + 1):
    line = ' '.join(words[i-seq_len:i])
    lines.append(line)
```

```
['Hi this is', 'this is a', 'is a small', 'a small sentence']
```

```
# Import Tokenizer from keras preprocessing text
from keras.preprocessing.text import Tokenizer

# Instantiate Tokenizer
tokenizer = Tokenizer()

# Fit it on the previous lines
tokenizer.fit_on_texts(lines)

# Turn the lines into numeric sequences
sequences = tokenizer.texts_to_sequences(lines)
```

```
array([[5, 3, 1], [3, 1, 2], [1, 2, 4], [2, 4, 6]])
```

```
print(tokenizer.index_word)
```

```
{1: 'is', 2: 'a', 3: 'this', 4: 'small', 5: 'hi', 6: 'sentence'}
```

```
# Import Dense, LSTM and Embedding layers
from keras.layers import Dense, LSTM, Embedding
model = Sequential()

# Vocabulary size
vocab_size = len(tokenizer.index_word) + 1

# Starting with an embedding layer
model.add(Embedding(input_dim=vocab_size, output_dim=8, input_length=2))

# Adding an LSTM layer
model.add(LSTM(8))

# Adding a Dense hidden layer
model.add(Dense(8, activation='relu'))

# Adding an output layer with softmax
model.add(Dense(vocab_size, activation='softmax'))
```

# Let's do it!

## INTRODUCTION TO DEEP LEARNING WITH KERAS

# You're done!

INTRODUCTION TO DEEP LEARNING WITH KERAS

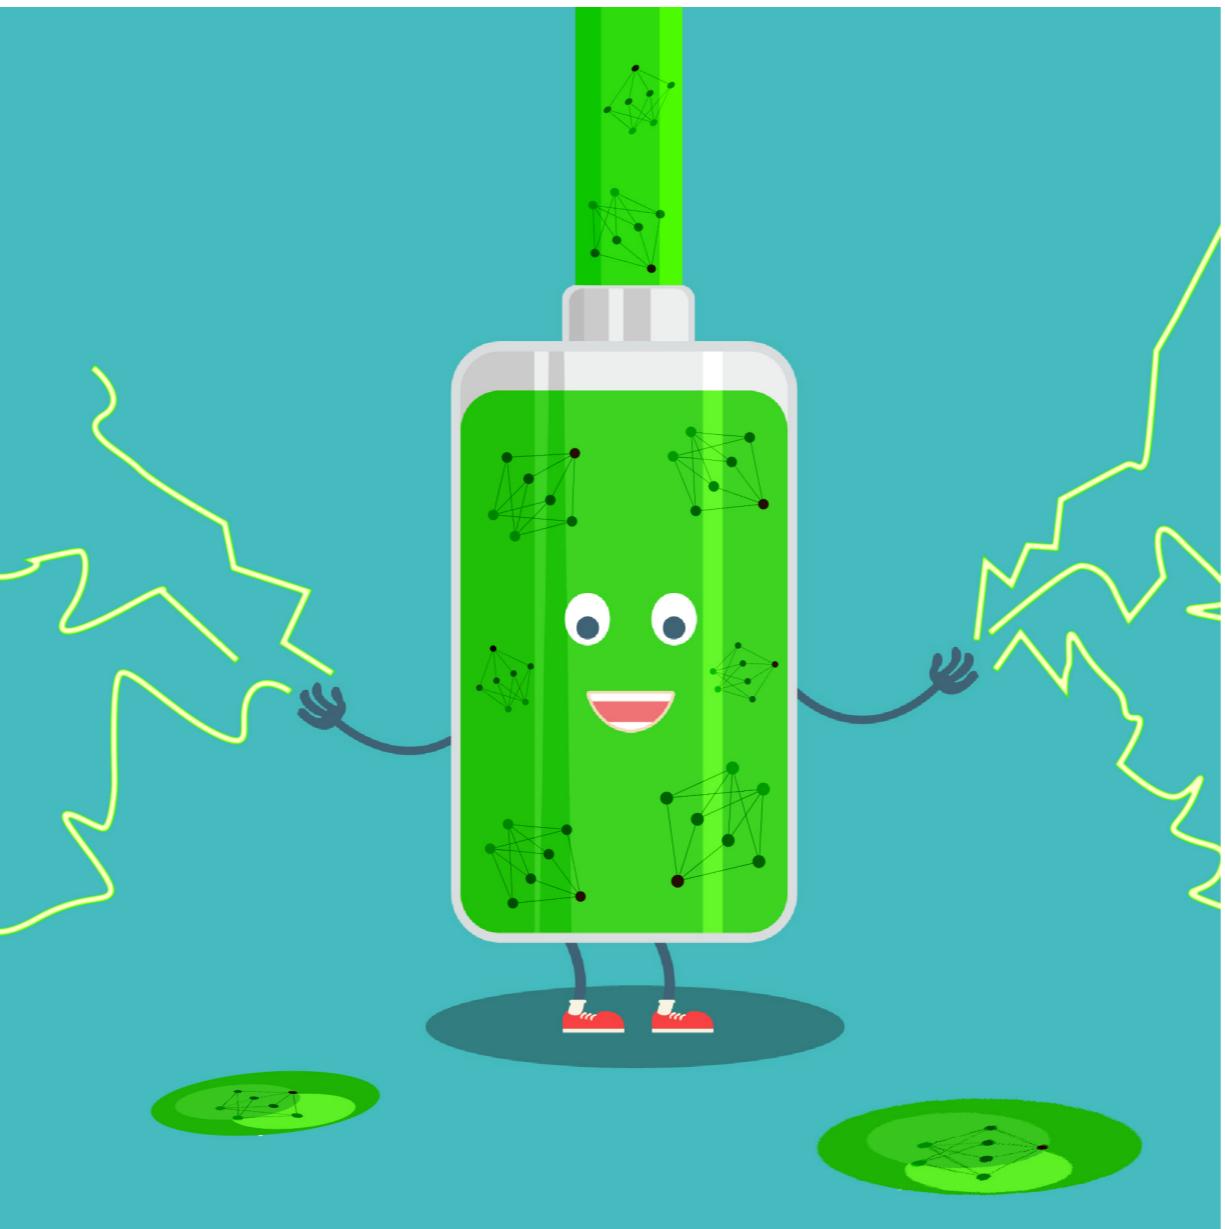


**Miguel Esteban**  
Data Scientist & Founder

# Congratulations!



# You've learned a lot



# What you've learned

- Basics of neural networks
- Building sequential models
- Building models for regression
- Approaching binary classification, multi-class and multi-label problems with neural networks
- Activation functions
- Hyperparameter optimization
- Autoencoders
- De-noising images
- CNN concepts
- Use pre-trained models
- Visualize convolutions
- LSTMs concepts
- Work with LSTMs and text
- All this by using many different datasets and learning a lot of Keras utility functions.

# What could you learn next?

- Go deeper into CNNs
- Go deeper into LSTMs
- Keras Functional API
- Models that share layers, models with several branches
- GANs: Generative Adversarial Networks
- Deeplearning projects



It's sad to say goodbye  
my friend

# Have a good one!

INTRODUCTION TO DEEP LEARNING WITH KERAS